

# CSE261: Computer Architecture

## 12. Processor (4): Pipelining #1

Seongil Wi

# HW2 will be Released Soon!



Implementing single-cycle datapath and control

- Due: Nov. 21, 11:49PM

# Summary: Single-cycle and Multicycle CPU

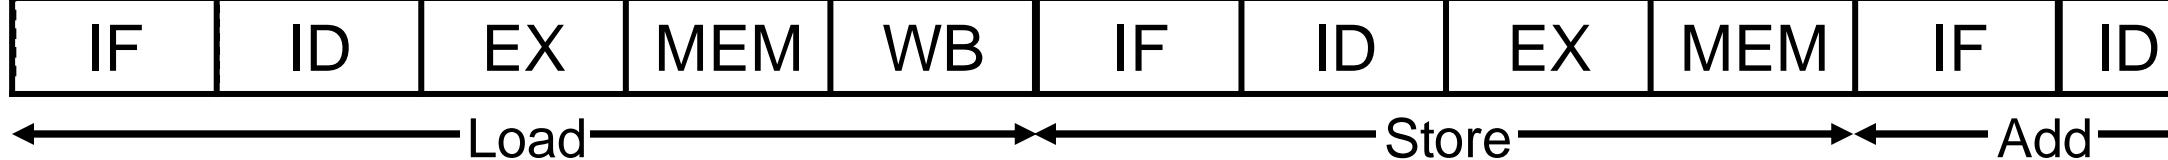
Clocks

Single-cycle Implementation



Clocks

Multicycle Implementation



*Is there a way to increase the speed further?*

# Pipelining Concept

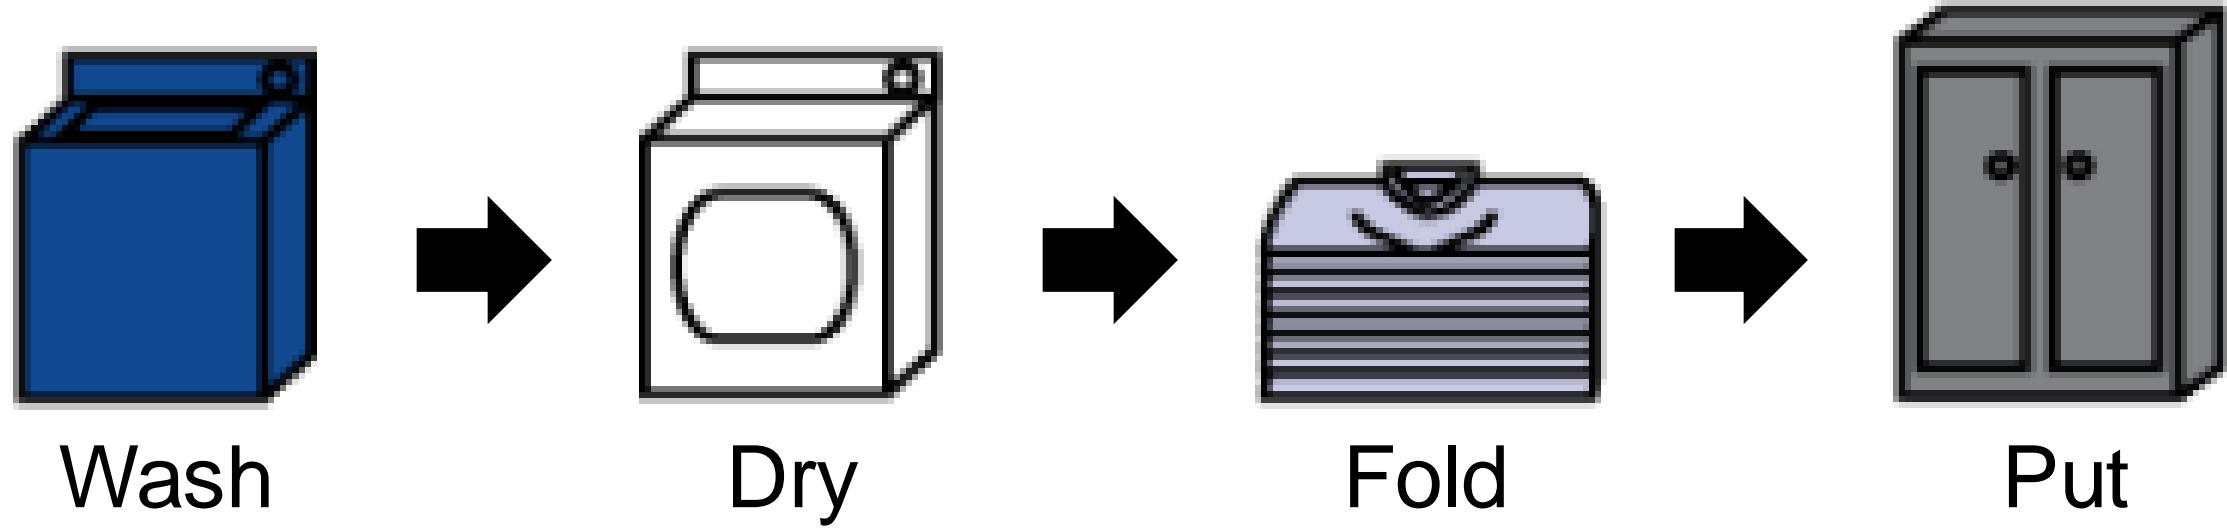
# Pipelining



Several instruction are executed at the same time to enhance the performance

# Motivation Example: Sequential Laundry Processing<sup>6</sup>

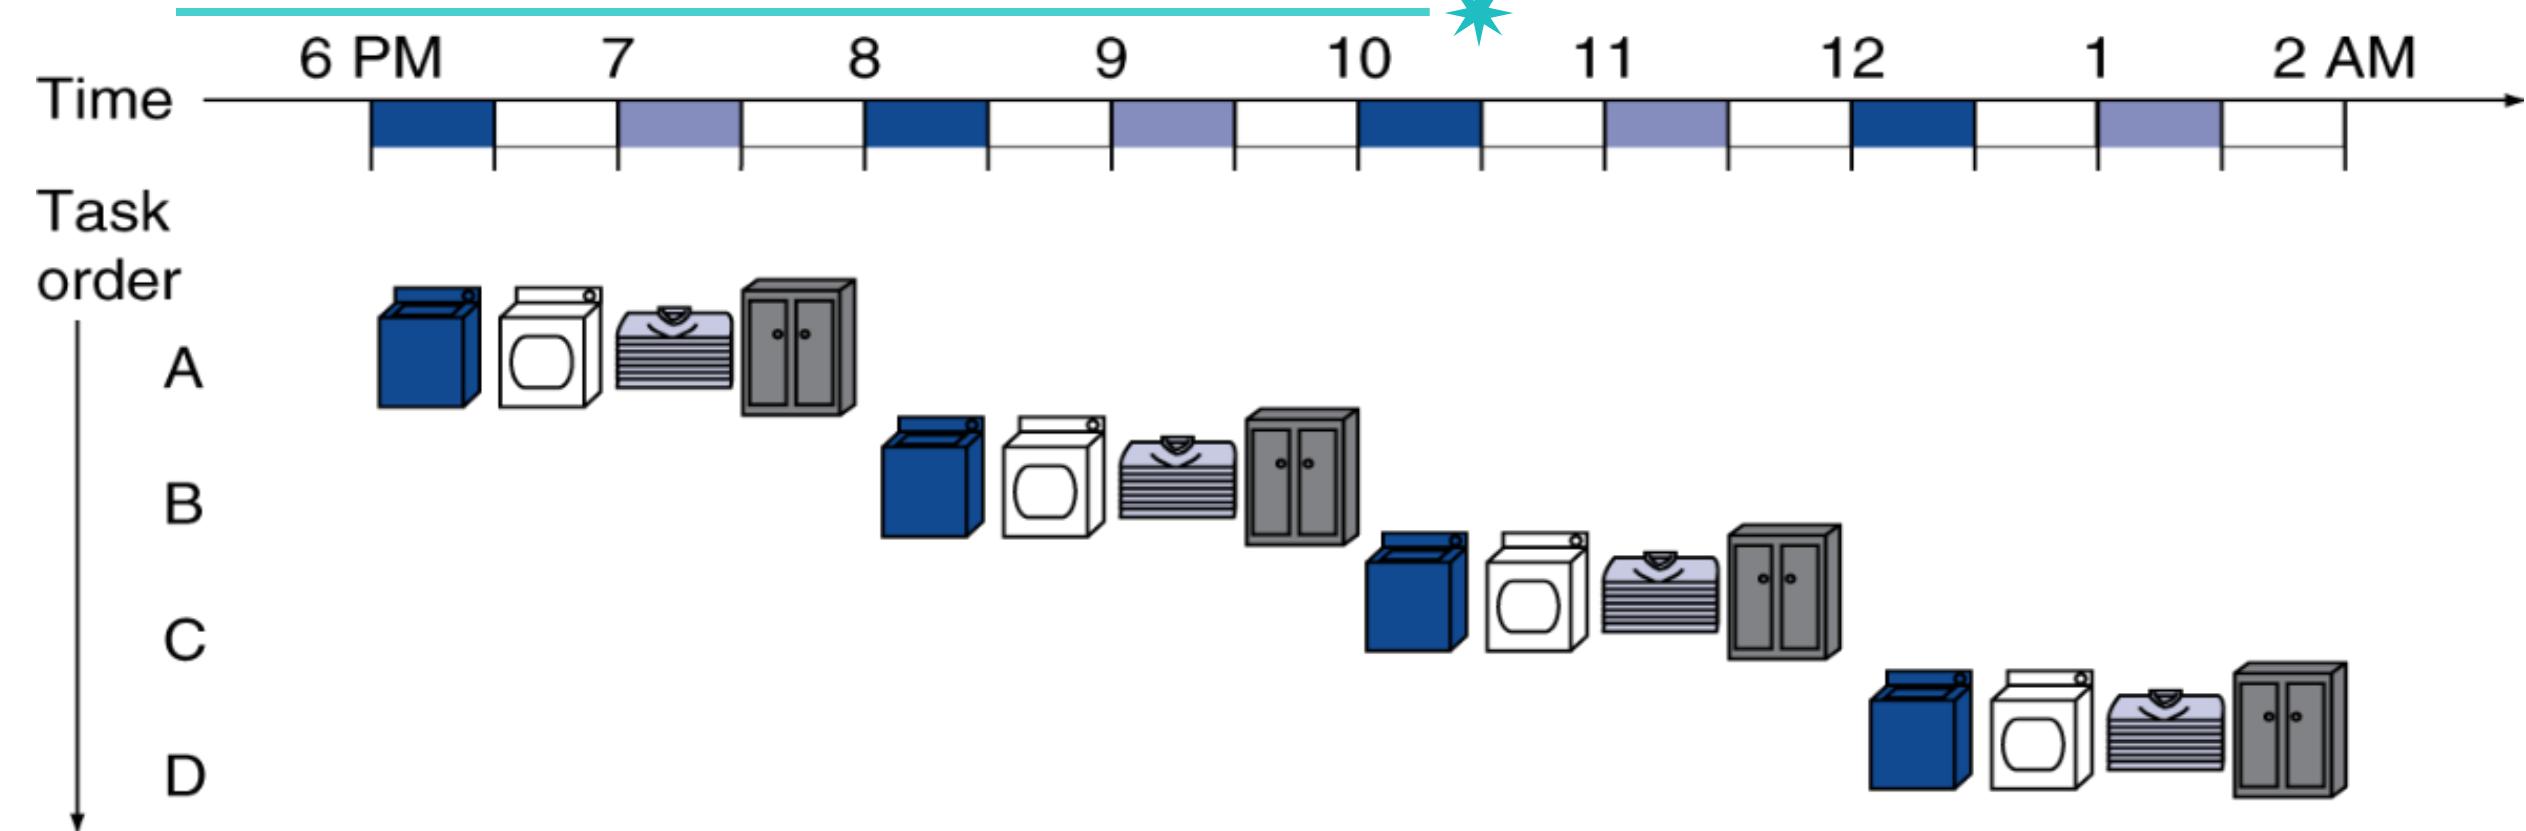
---



**Assume you have a set of  
clothes to wash: SET A, B, C, D**

# Motivation Example: Sequential Laundry Processing

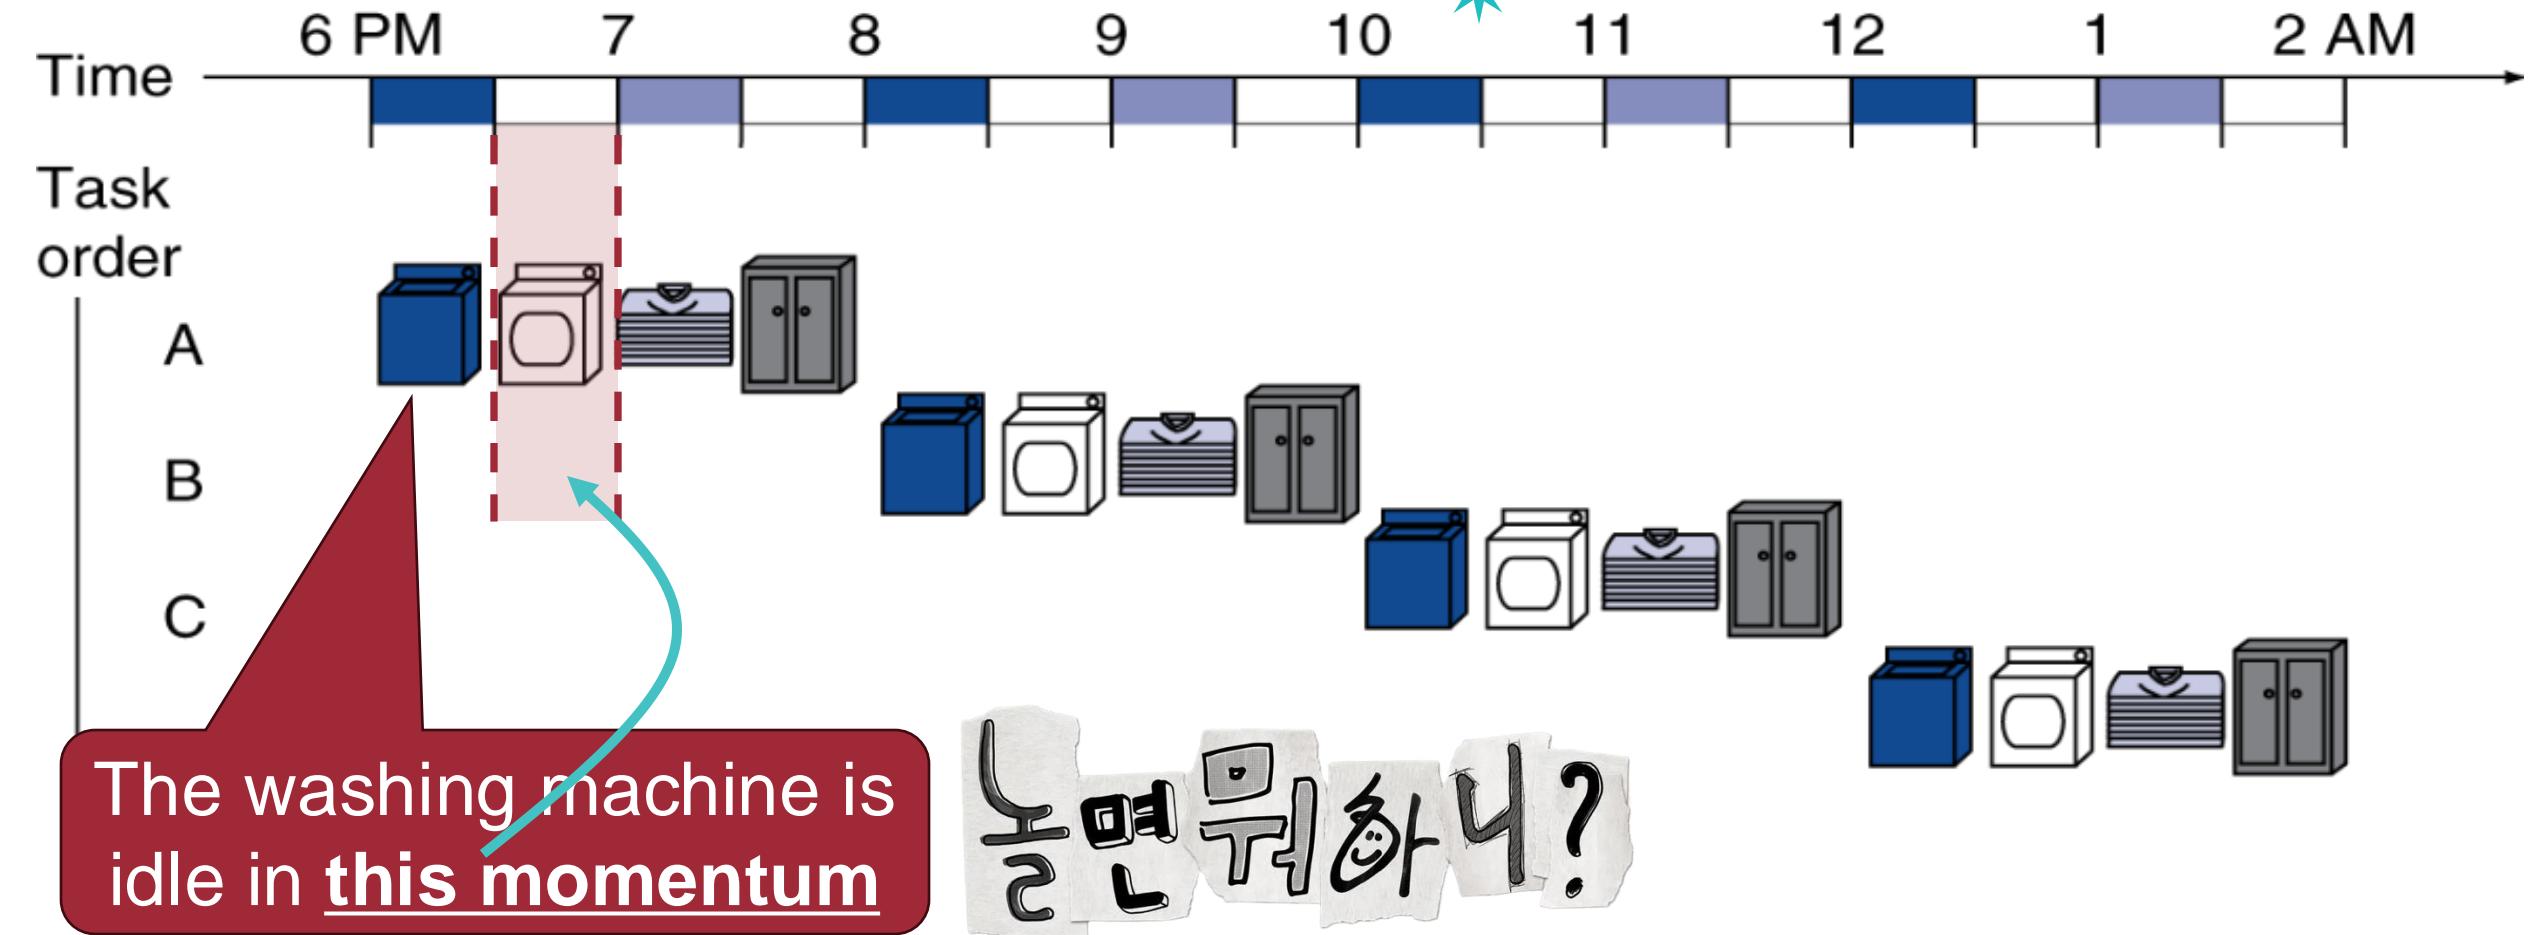
7



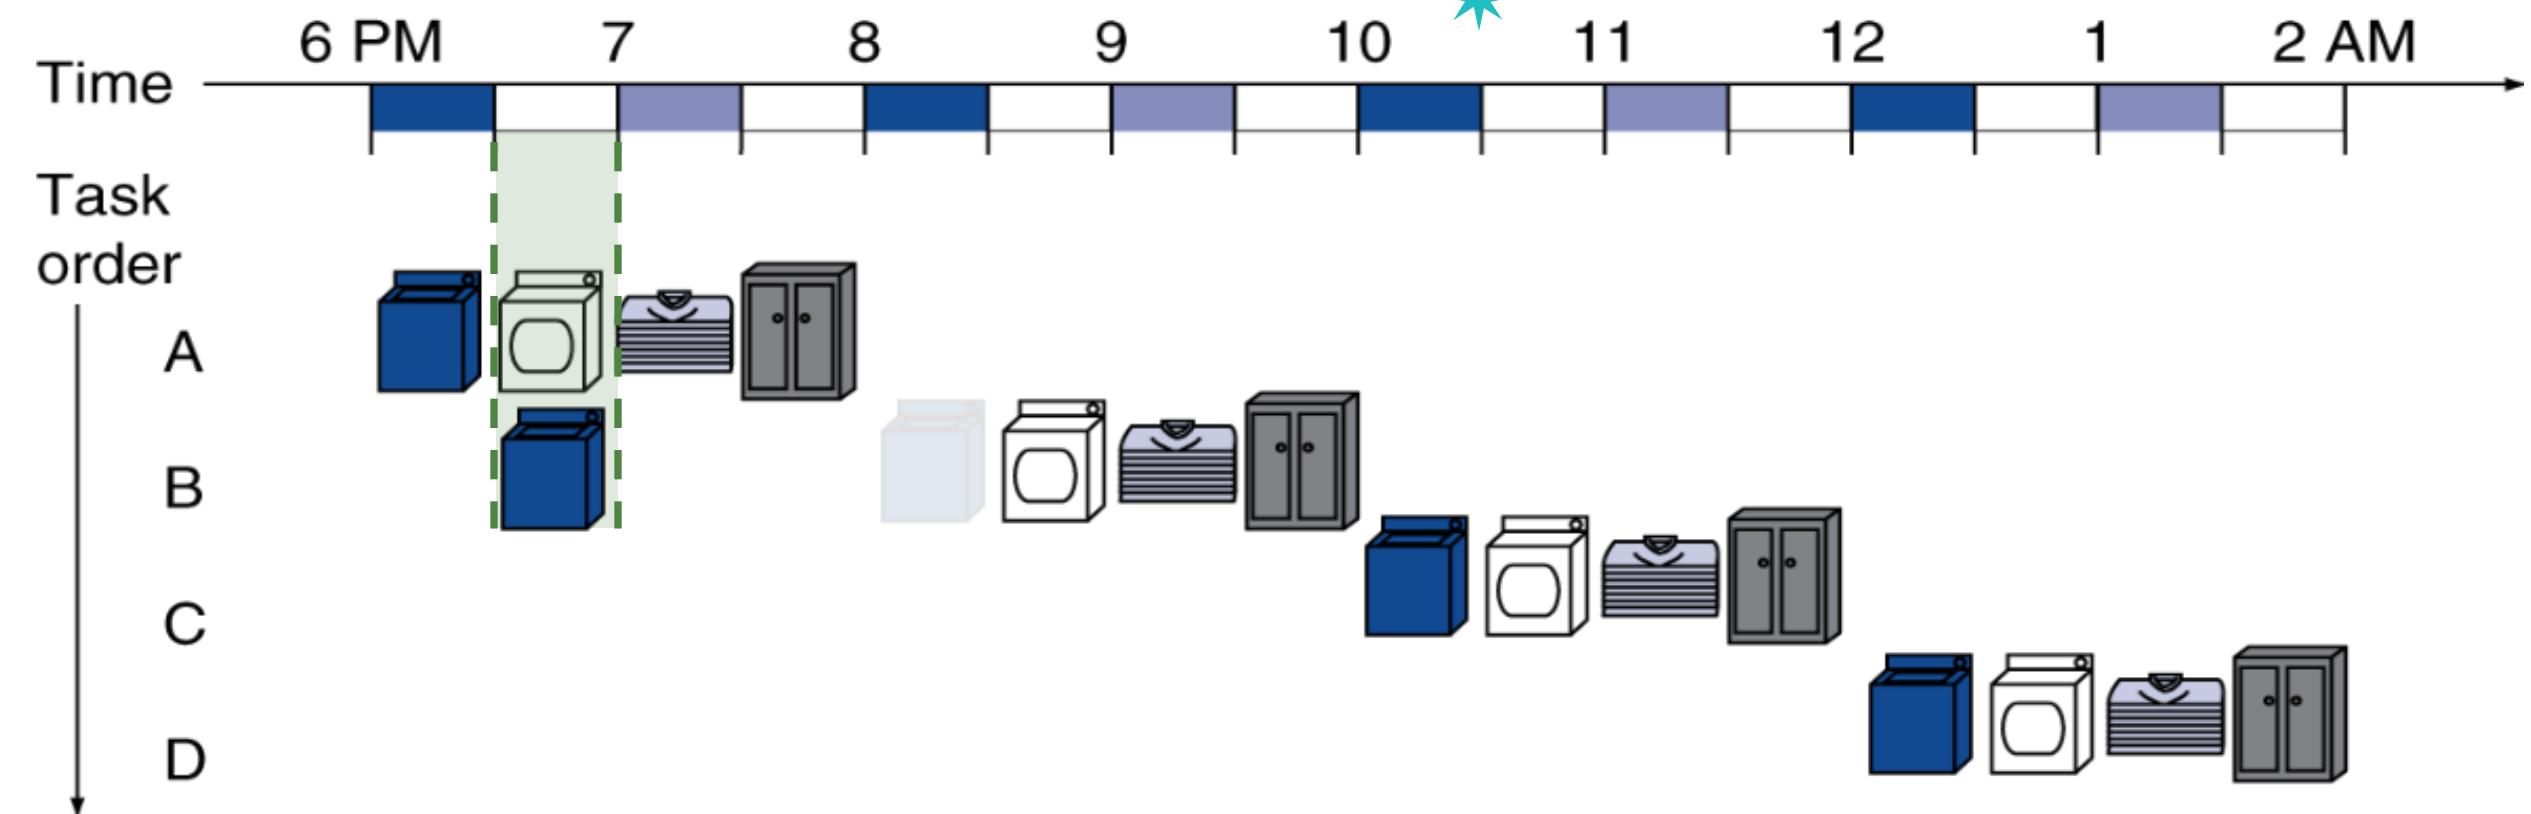
**Single- and multi-cycle implementation start the next task only after the previous task has been completed**

**Execution time:**  
 $2 + 2 + 2 + 2$   
= 8 hours

# Let's Think about it in a More Natural Way!



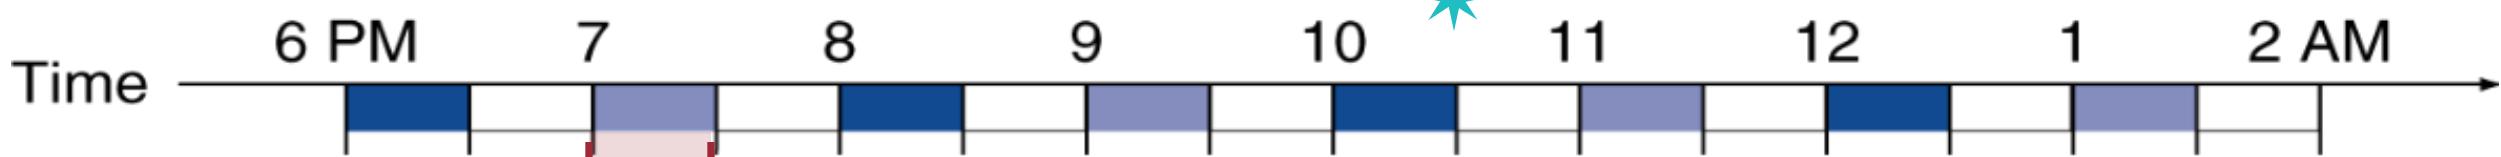
# Solution: Execute Washing at the Same Time



# Problem Again...



10



Task  
order

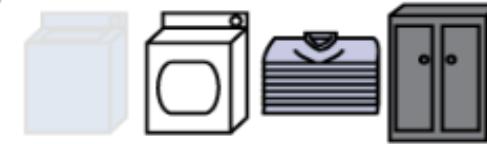
A



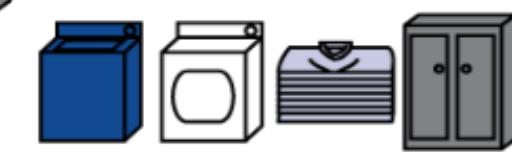
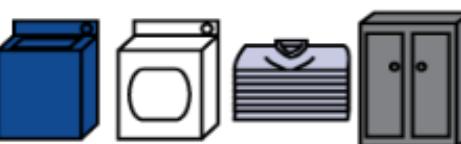
B



C

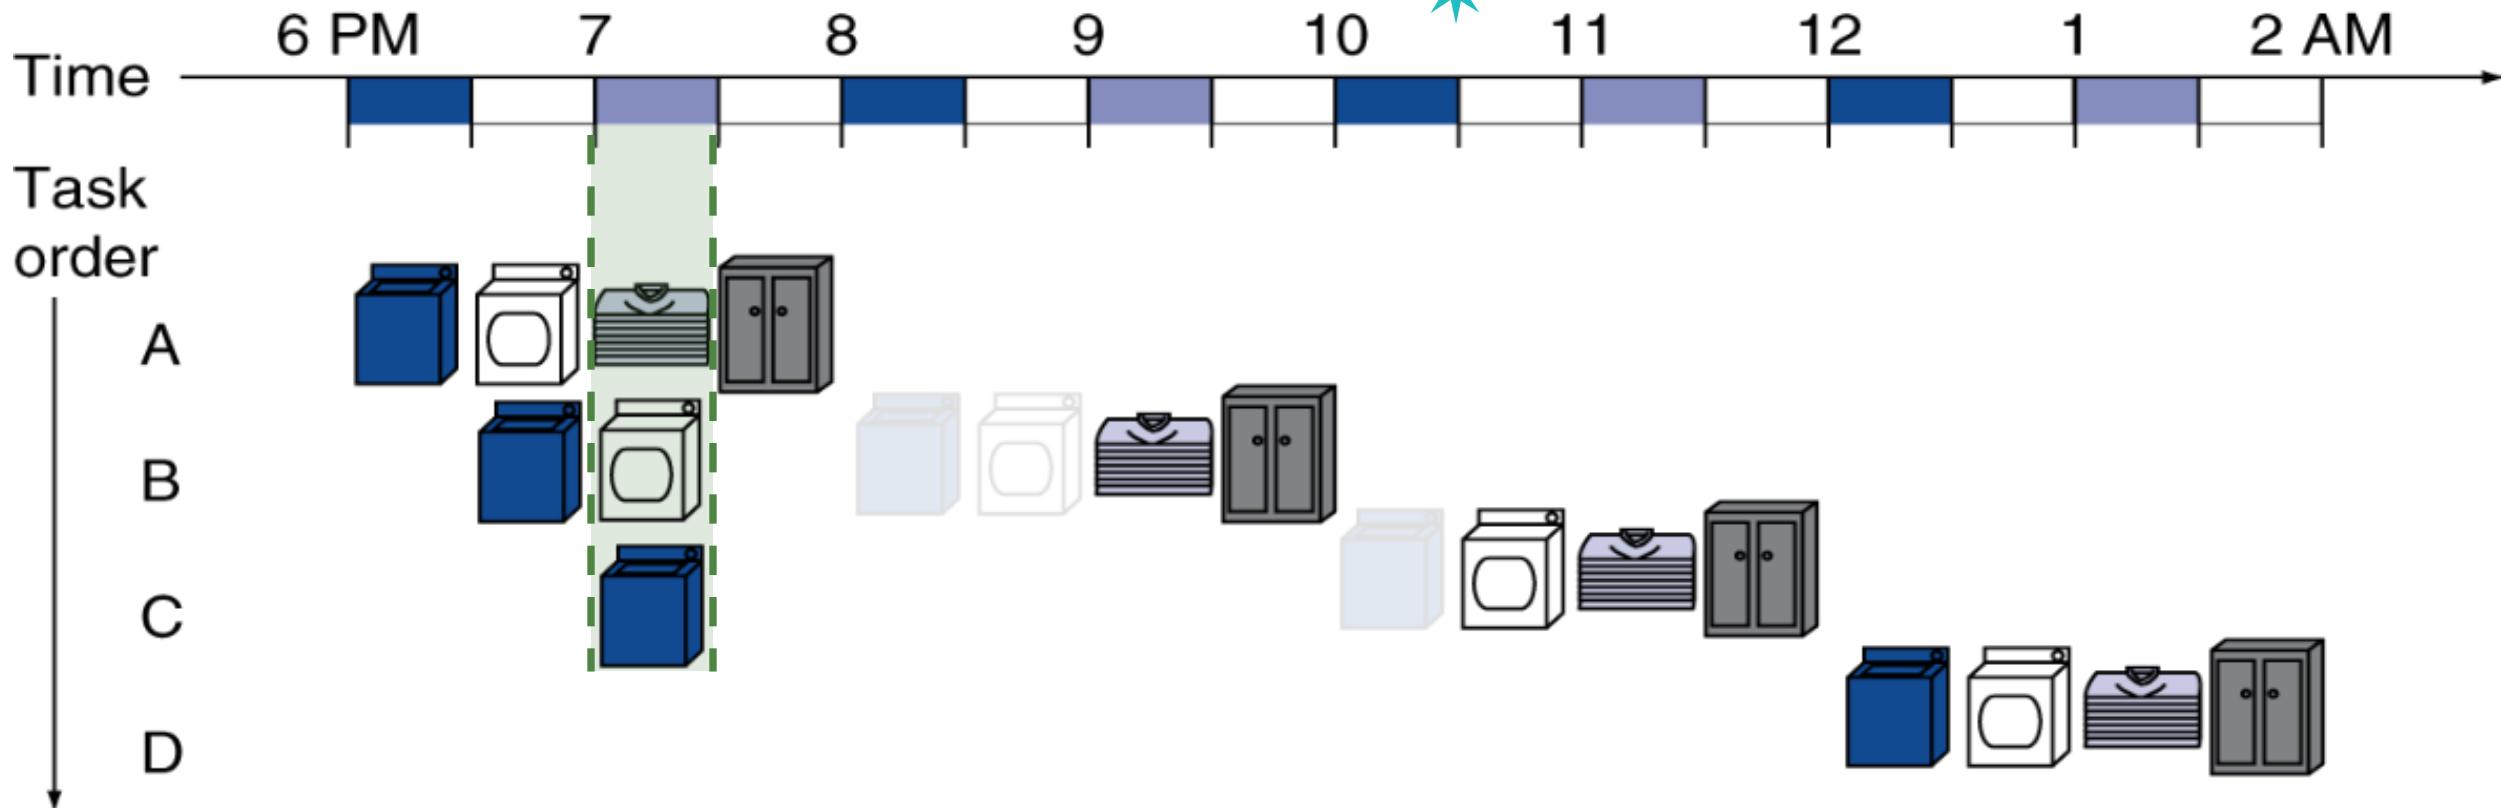


D



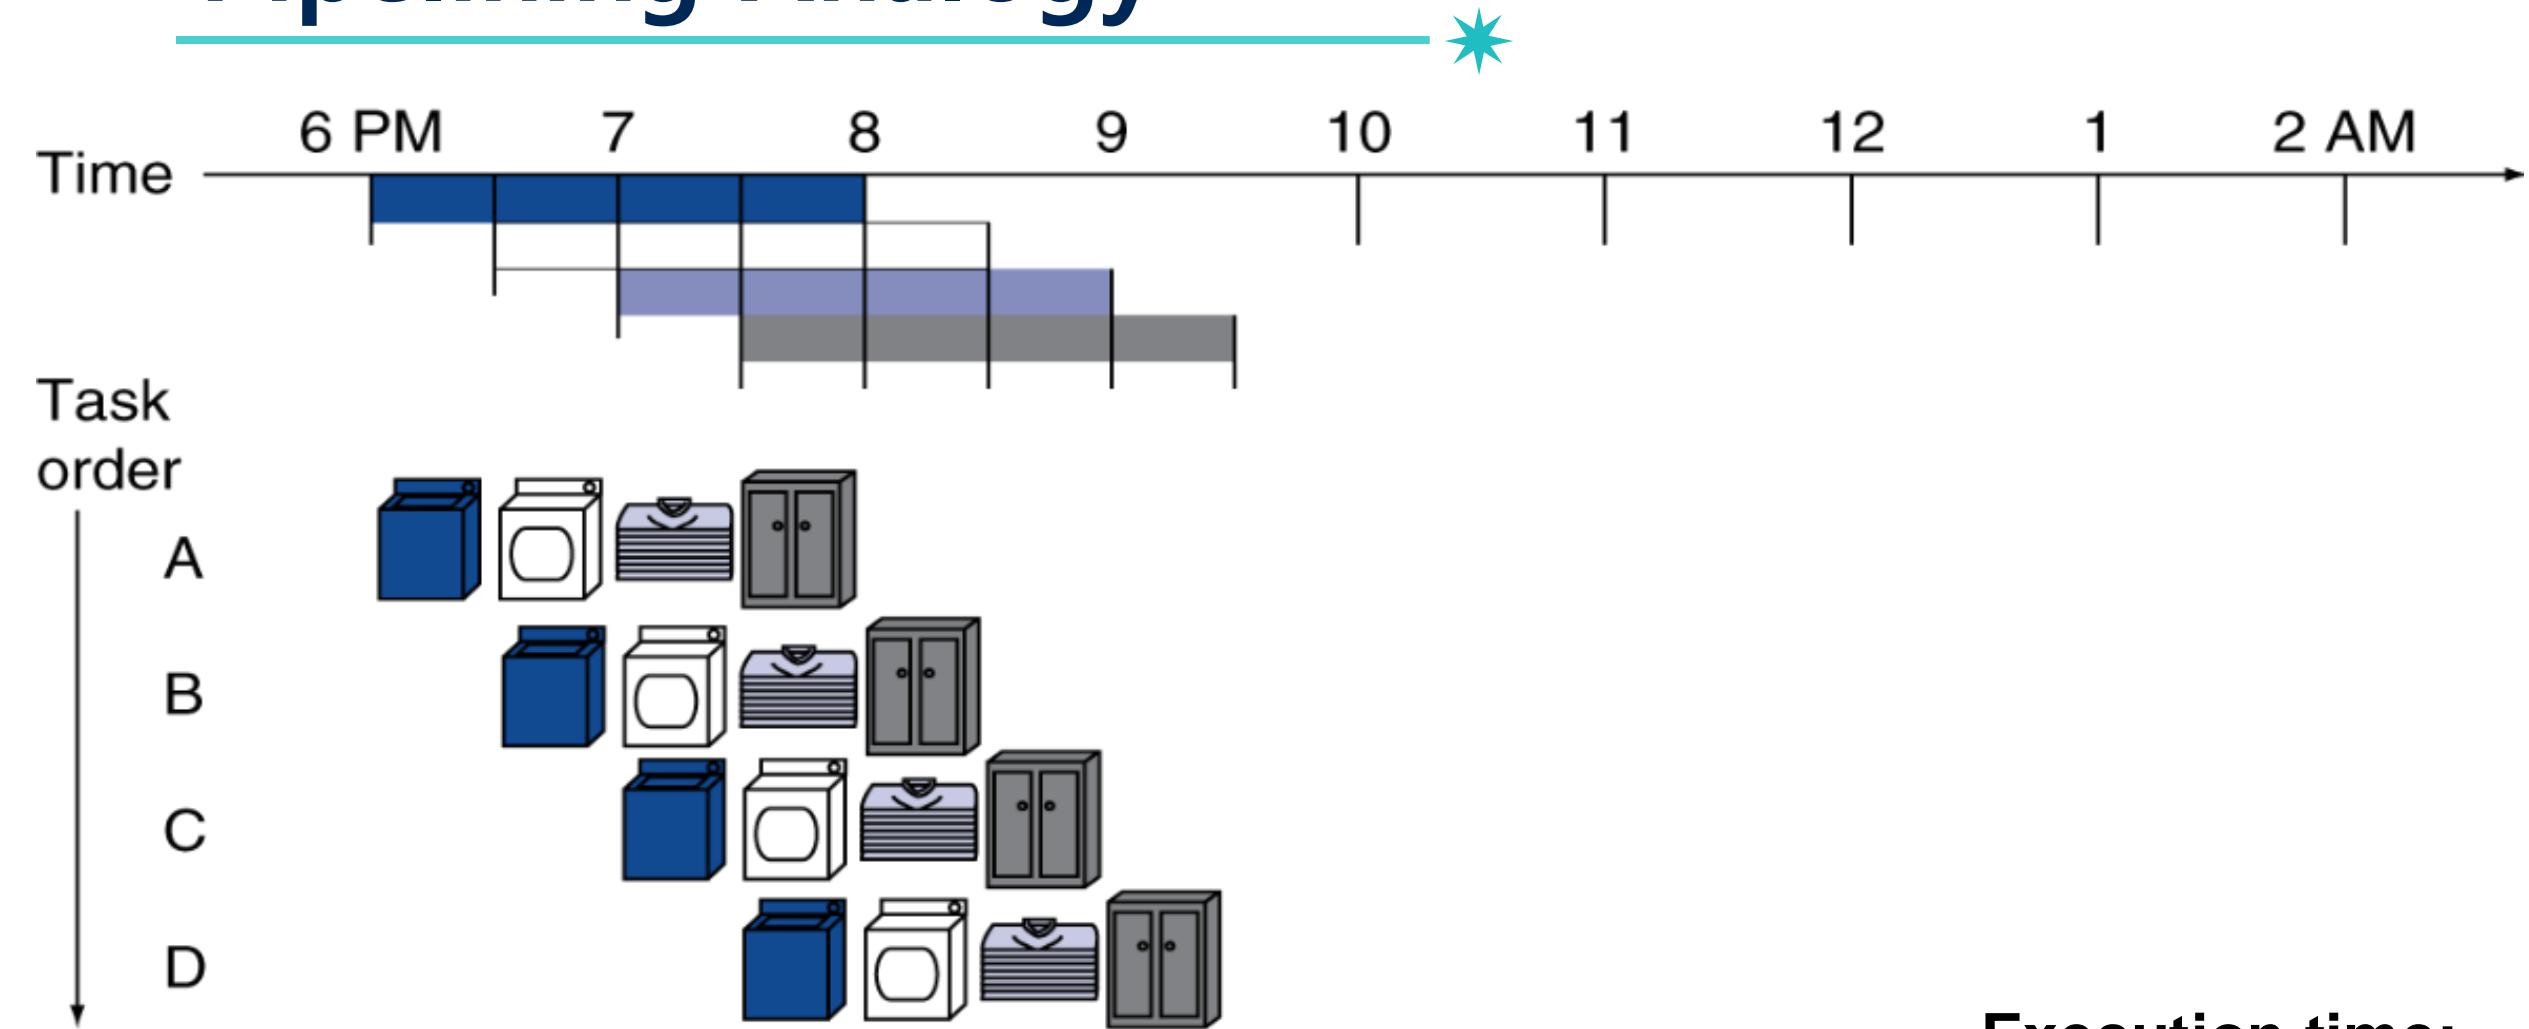
The washing machine and dry  
machine is idle in this momentum

# Solution: Execute Washing/drying at the Same Time



# Pipelining Analogy

12

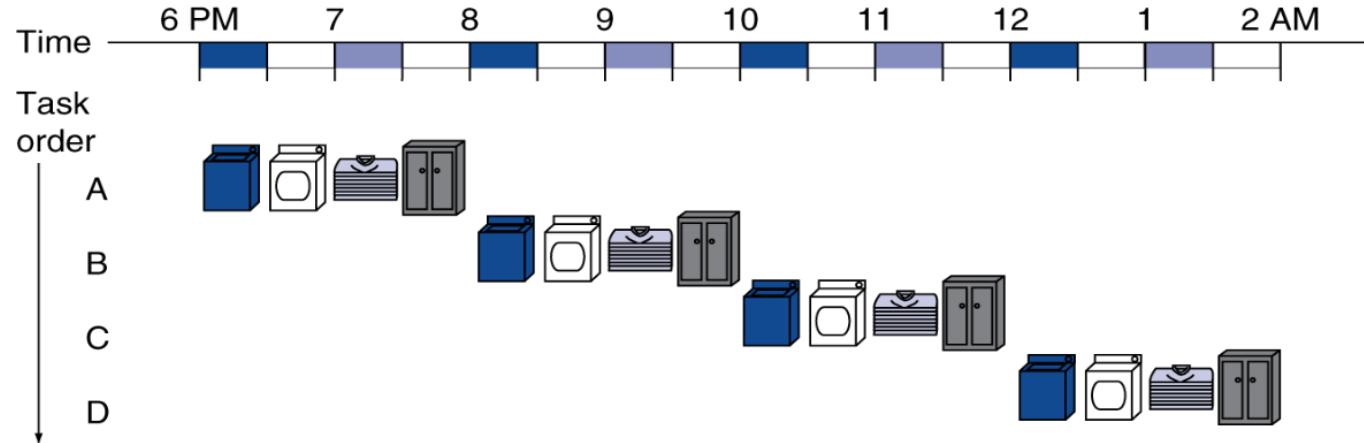


Pipelining overlaps execution

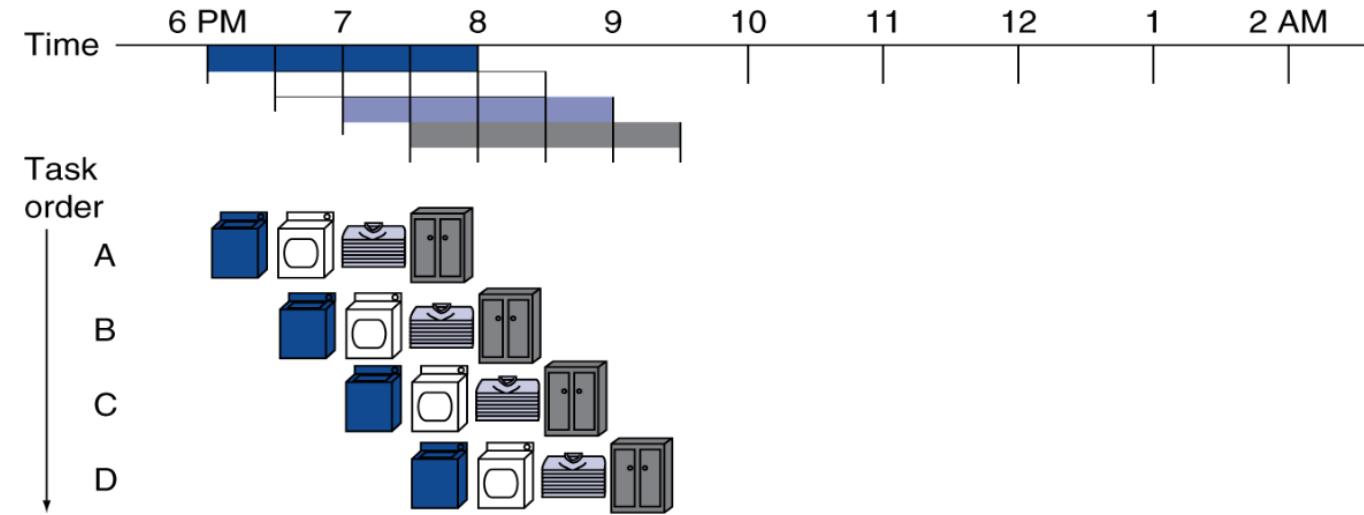
**Execution time:**  
 $2 + 0.5 + 0.5 + 0.5 = 3.5 \text{ hours}$

# Performance Comparison: Speedup

13



**Without pipelining**  
**execution time:**  
 $2 + 2 + 2 + 2$   
= 8 hours



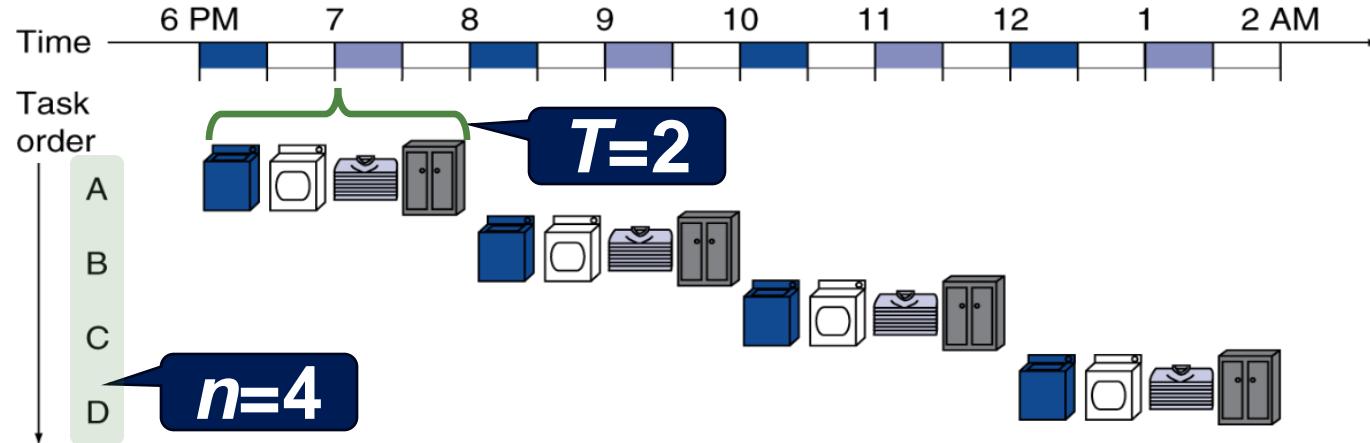
**Pipelining**  
**execution time:**  
 $2 + 0.5 + 0.5 + 0.5$   
= 3.5 hours

Speedup?  
 $8/3.5 = 2.3$

# Let's Think about Speedup in General Terms

14

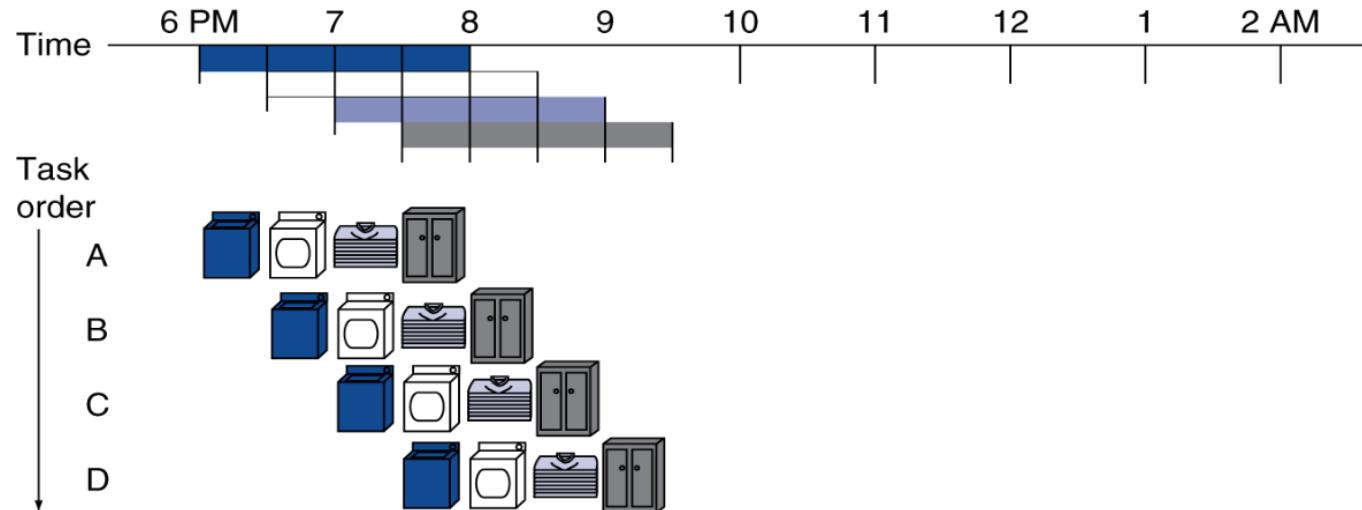
Suppose there are  $n$  jobs to execute and each job takes  $T$  hours



Without pipelining  
execution time:

$$2 + 2 + 2 + 2  
= 8 \text{ hours}$$

$n \cdot T$



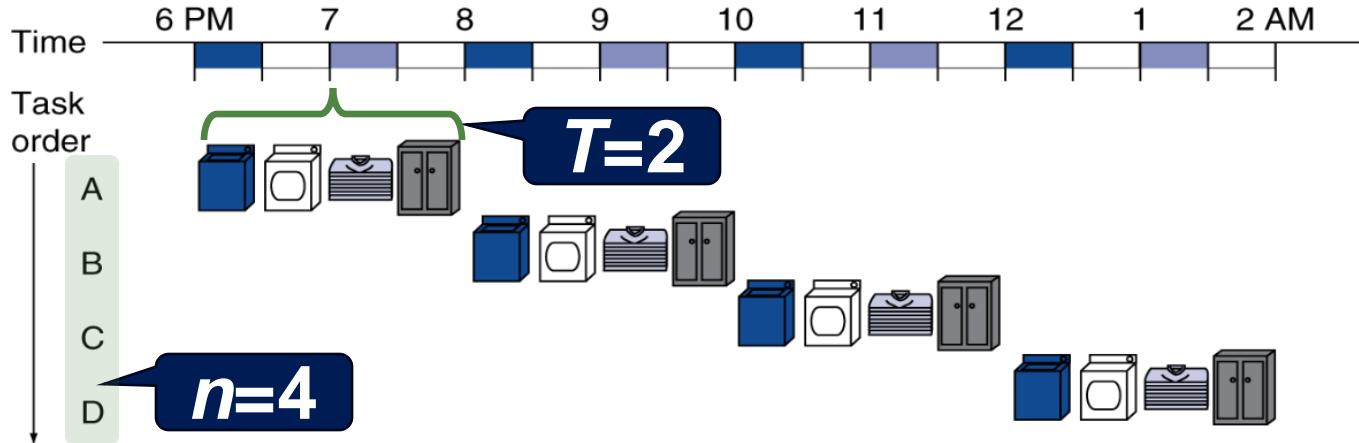
Pipelining  
execution time:

$$2 + 0.5 + 0.5 + 0.5  
= 3.5 \text{ hours}$$

# Let's Think about Speedup in General Terms

15

Suppose there are  $n$  jobs to execute and each job takes  $T$  hours

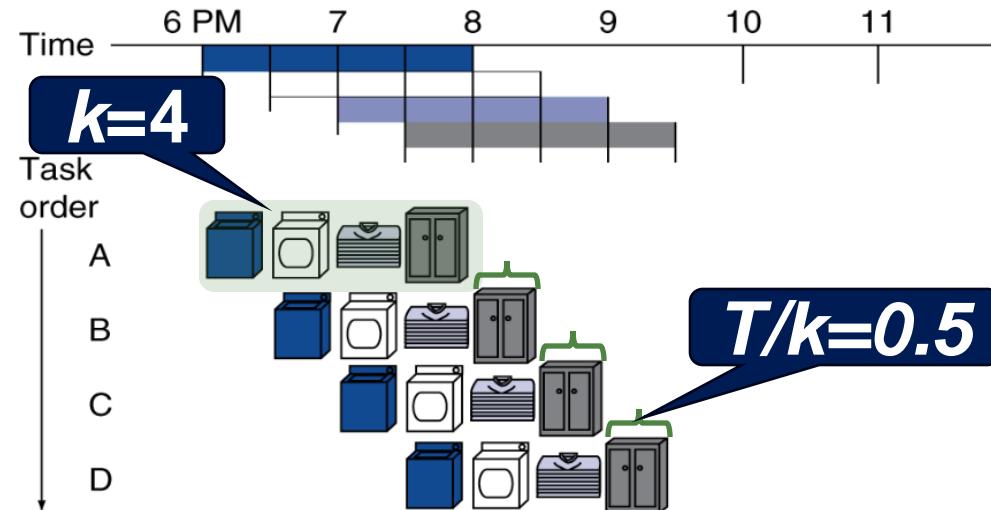


Without pipelining  
execution time:

$$2 + 2 + 2 + 2  
= 8 \text{ hours}$$

$n*T$

Suppose we have  $k$  stage



Pipelining  
execution time:

$$2 + 0.5 + 0.5 + 0.5  
= 3.5 \text{ hours}$$

$(n+k-1) * T/K$

# Pipelining Speedup

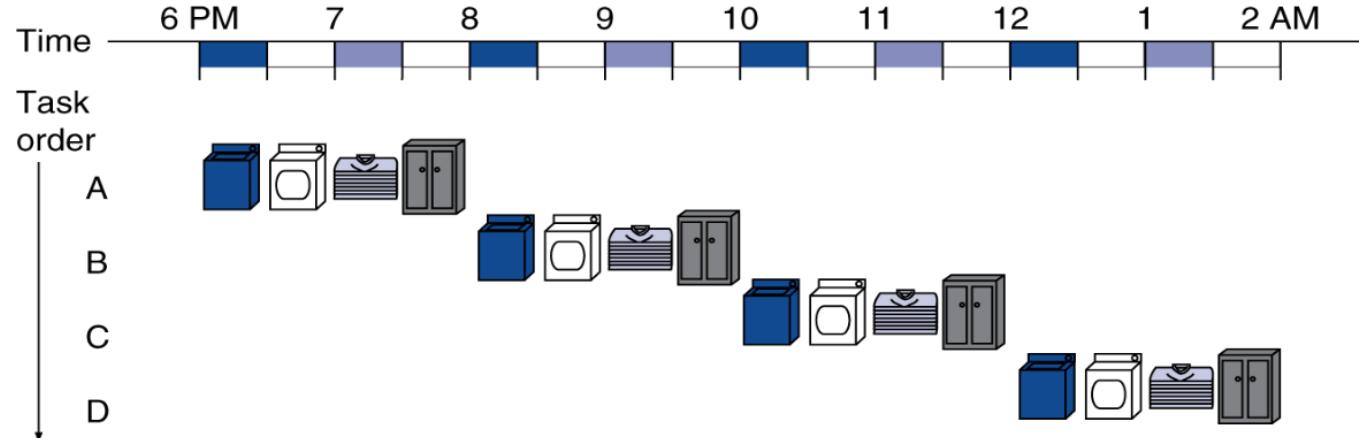


- Execution time *without pipelining*:  $n*T$
- Execution time *with pipelining*:  $(n+k-1)*T/k$
- **Pipelining speedup:**  $\frac{n*T}{(n+k-1)*T/k} = \frac{n*k}{n+k-1}$
- **If  $n$  goes infinity, pipelining speedup becomes  $k$**

$$\lim_{n \rightarrow \infty} \frac{n*k}{n+k-1} = k$$

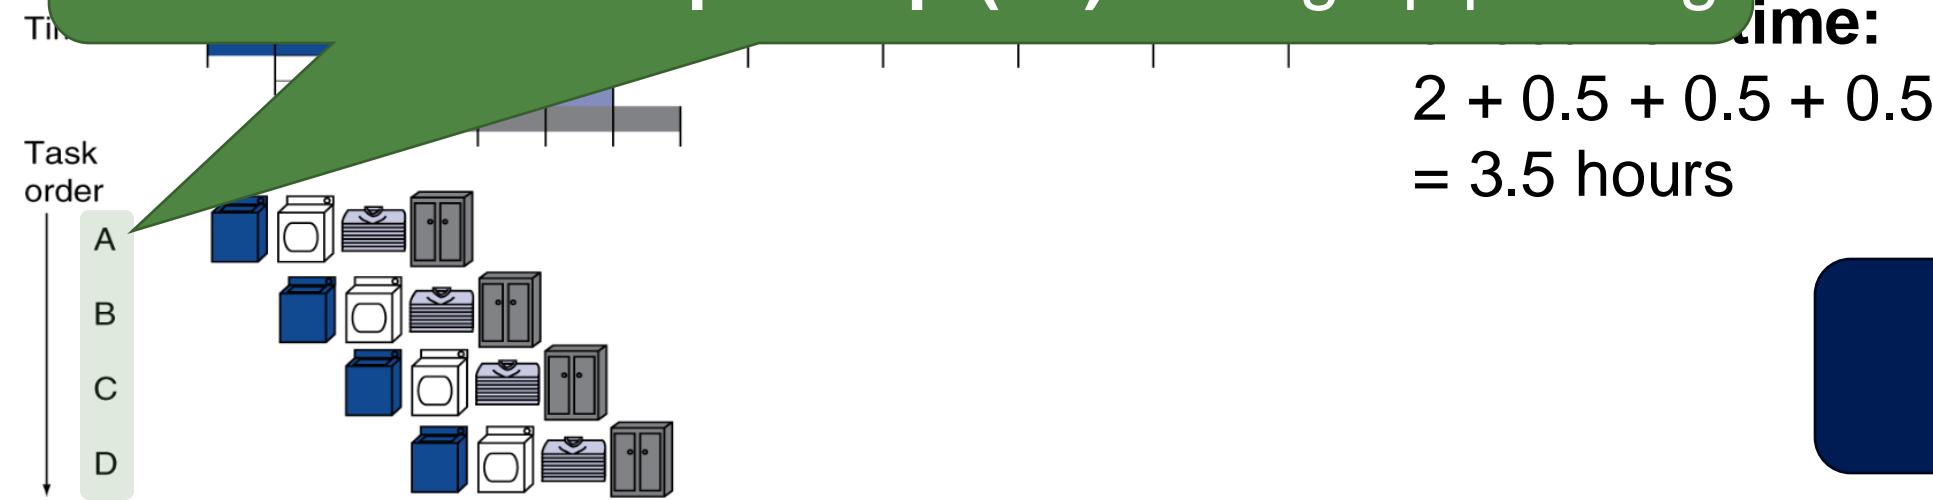
If we have a large number of jobs, we can achieve a **speedup of about  $k$  times** through pipelining

# Performance Comparison: Speedup



**Without pipelining execution time:**  
 $2 + 2 + 2 + 2 = 8 \text{ hours}$

If we have a large number of jobs, we can achieve **about 4 speedup (x4)** through pipelining



**With pipelining execution time:**  
 $2 + 0.5 + 0.5 + 0.5 = 3.5 \text{ hours}$

Speedup?  
 $8/3.5 = 2.3$

# Pipeline Lessons

---

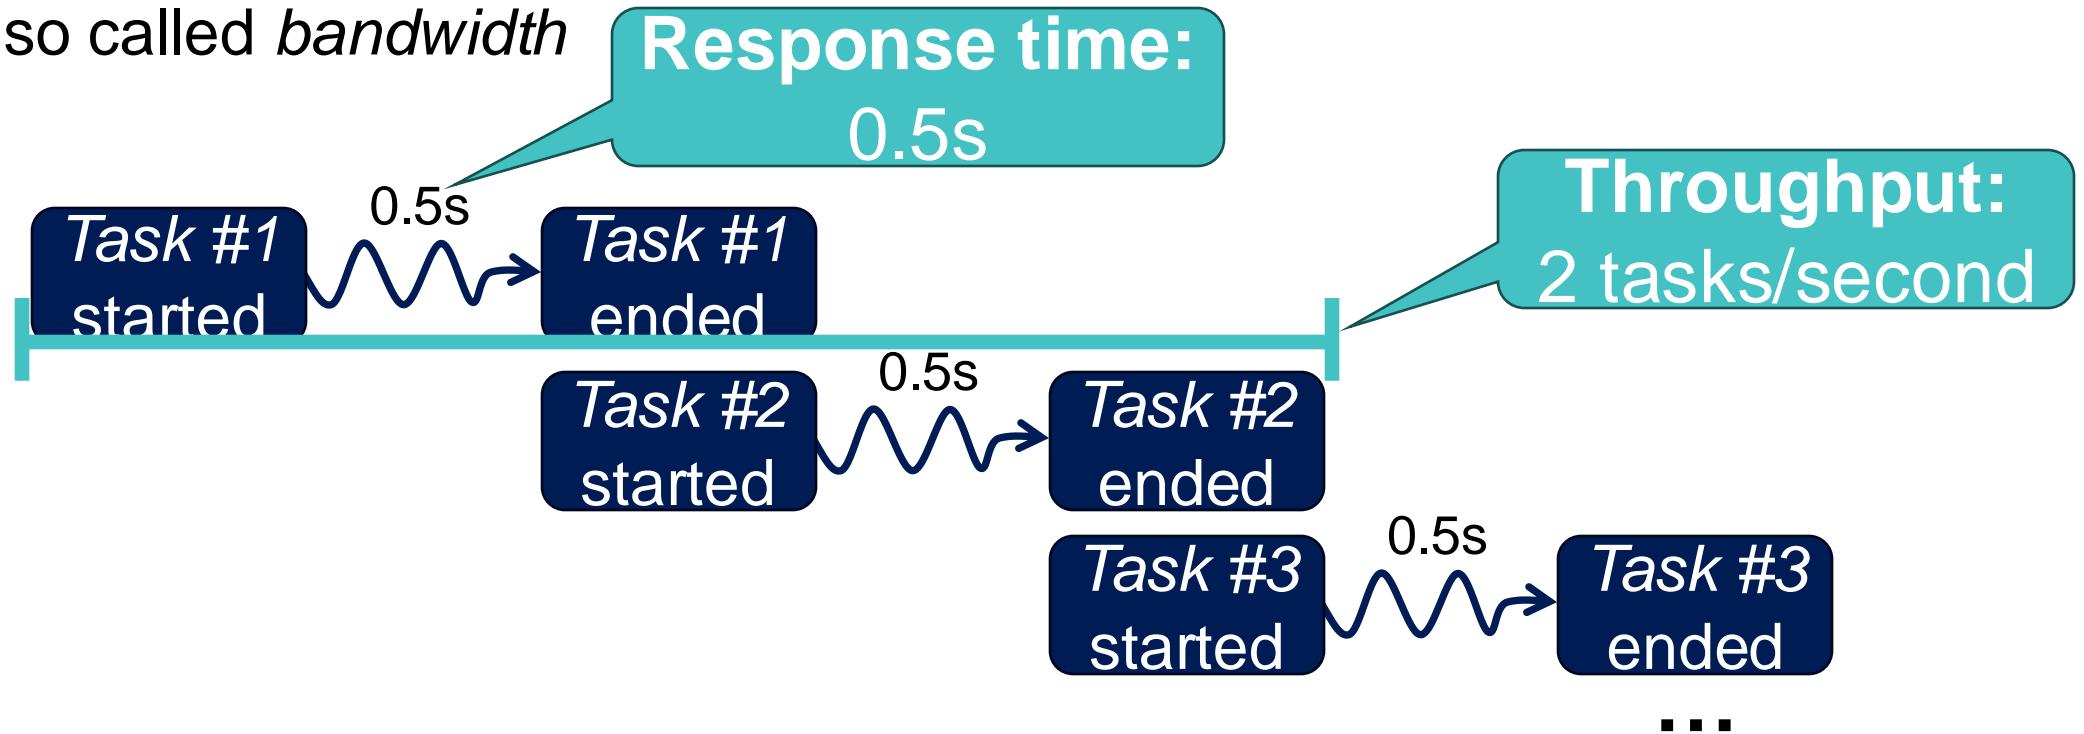


- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload

# Performance Metrics



- **Response time:** the time between the start and completion of a task
  - Also called *execution time, latency*
- **Throughput:** total work done per unit time
  - E.g., # of tasks/transactions/... per hour
  - Also called *bandwidth*



# Pipeline Lessons

---

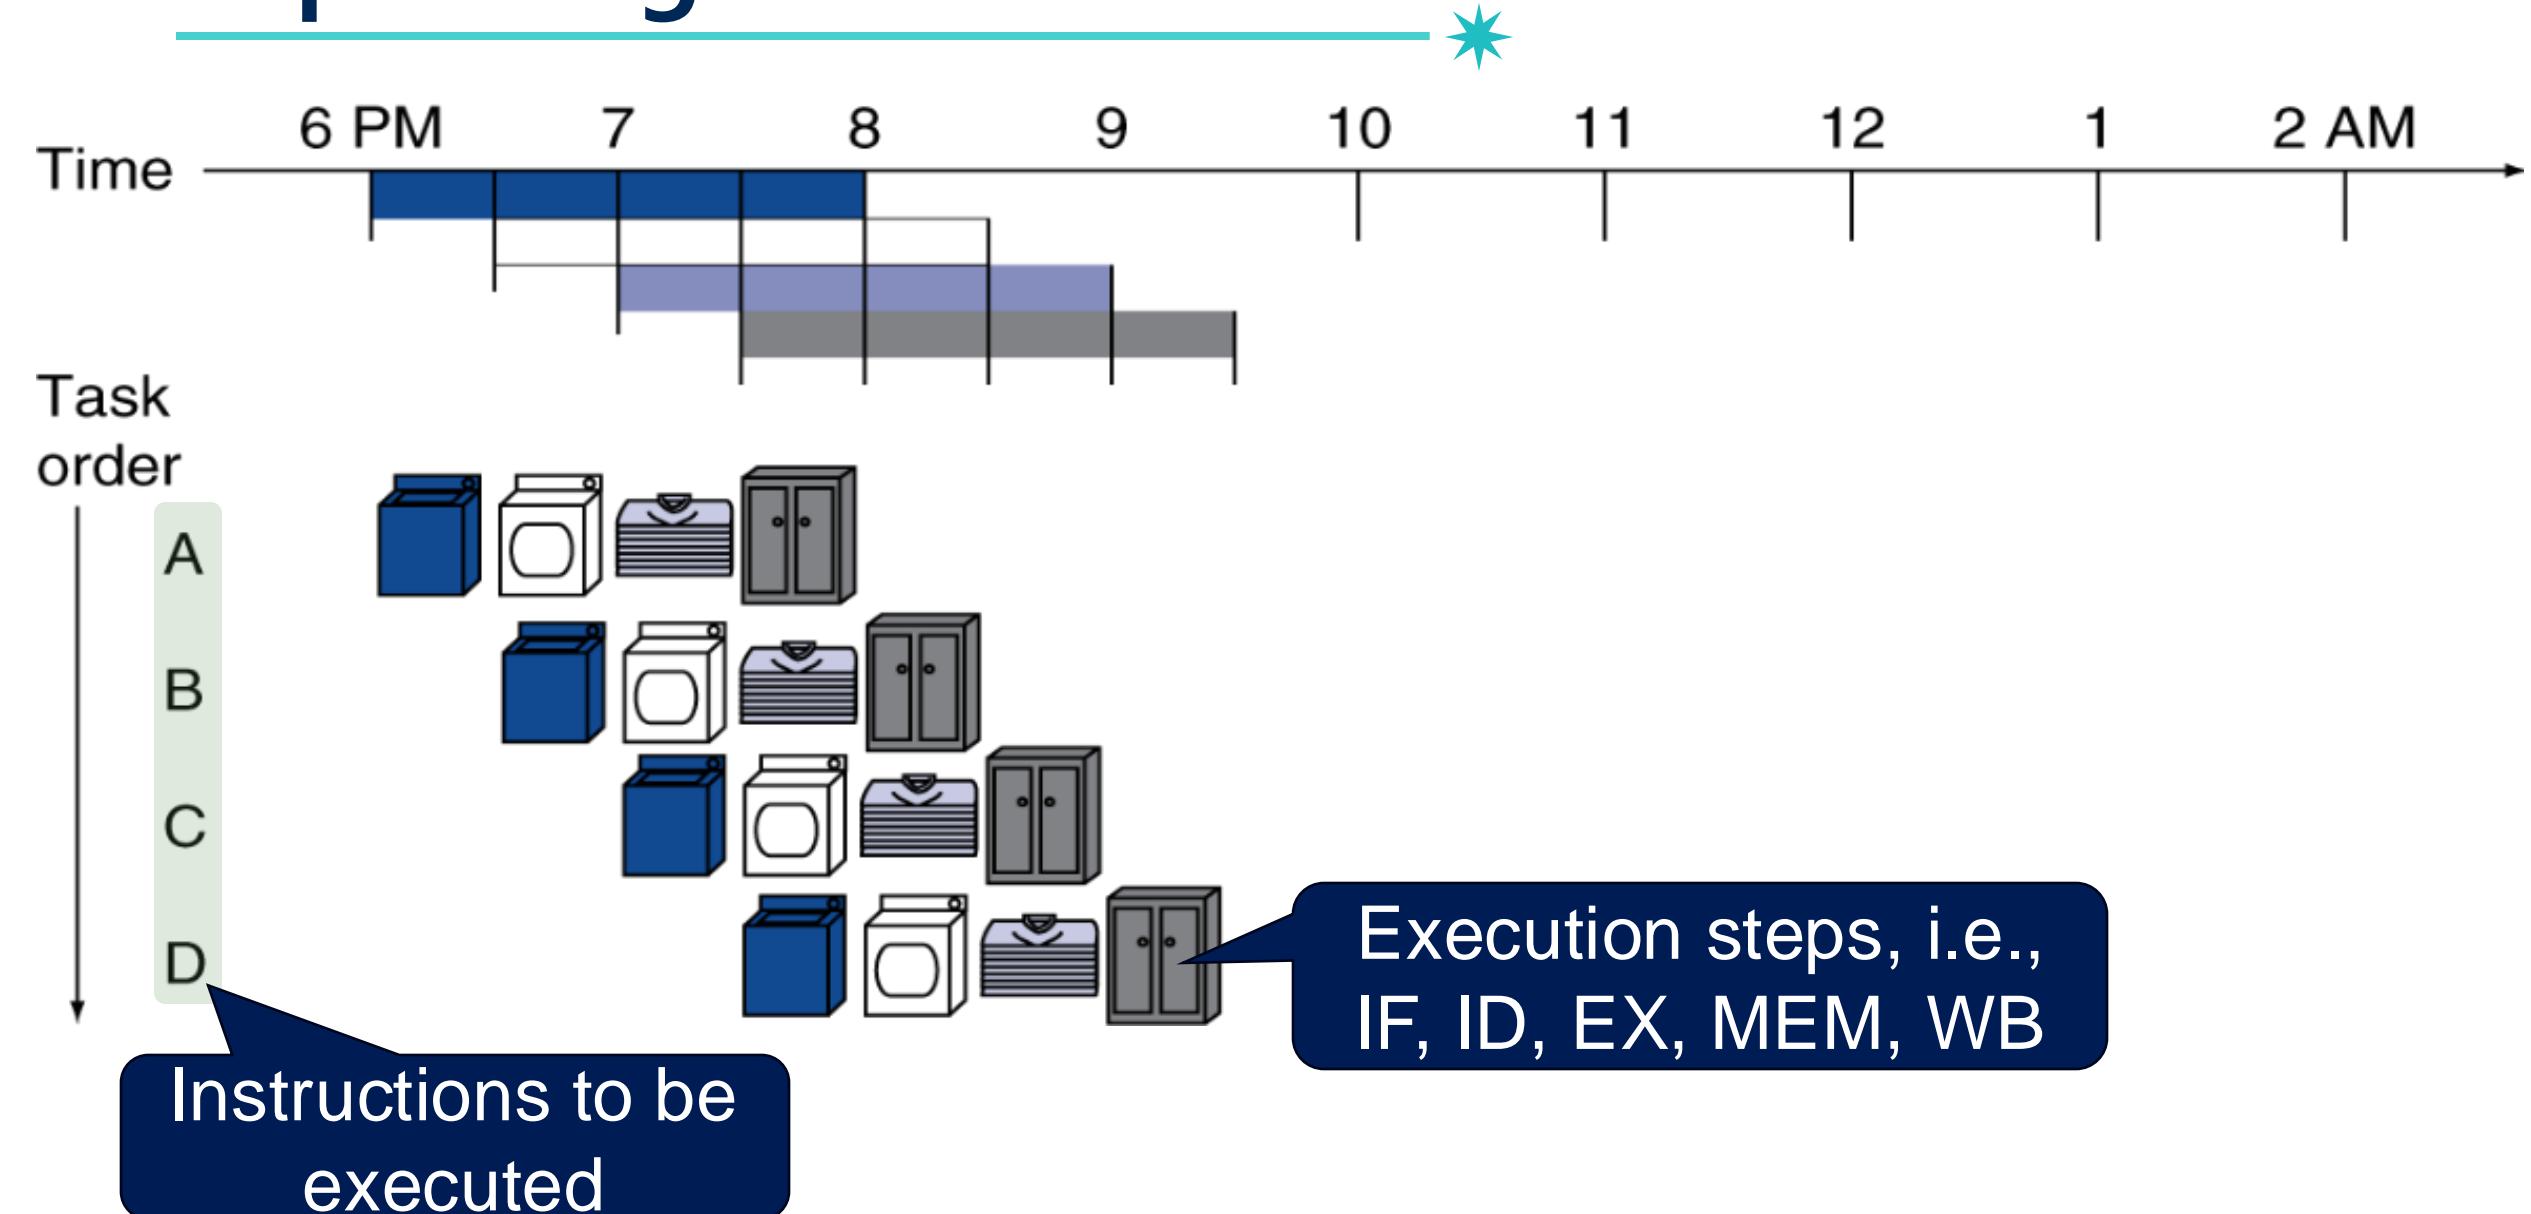


- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- Multiple tasks operating simultaneously using different resources
- Potential speedup = # pipe stages

# Pipelining Overview

# Pipelining Overview

22



# MIPS Pipelining

---



- Five stages ( $k=5$ ), one step per stage
  1. **IF**: Instruction fetch from memory
  2. **ID**: Instruction decode & register read
  3. **EX**: Execute operation or calculate address
  4. **MEM**: Access memory operand
  5. **WB**: Write result back to register

# MIPS Pipelining



- **Sequential execution**



add \$s0, \$t1, \$t2



lw \$1, 100(\$s0)



and \$t2, \$s0, \$s1

- **Pipelined execution**

add \$s0, \$t1, \$t2



lw \$1, 100(\$s0)



and \$t2, \$s0, \$s1



...

# Single-cycle vs. Multicycle vs. Pipeline

25

Clocks

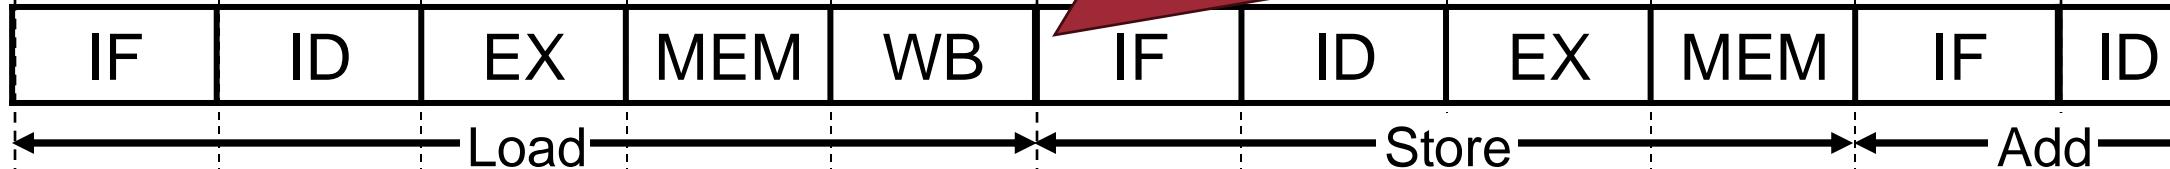
Single-cycle Implementation



Waste time

Clocks

Multicycle Implementation



The execution is still sequential

# Single-cycle vs. Multicycle vs. Pipeline

26

Clocks

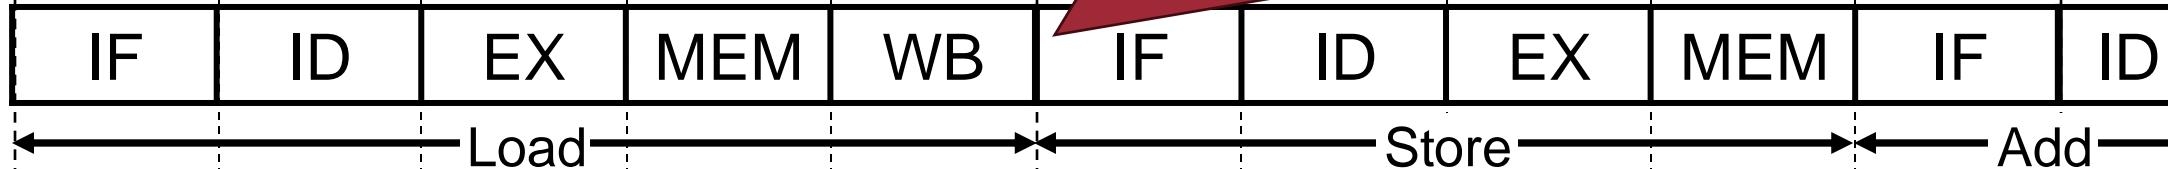
Single-cycle Implementation



Waste time

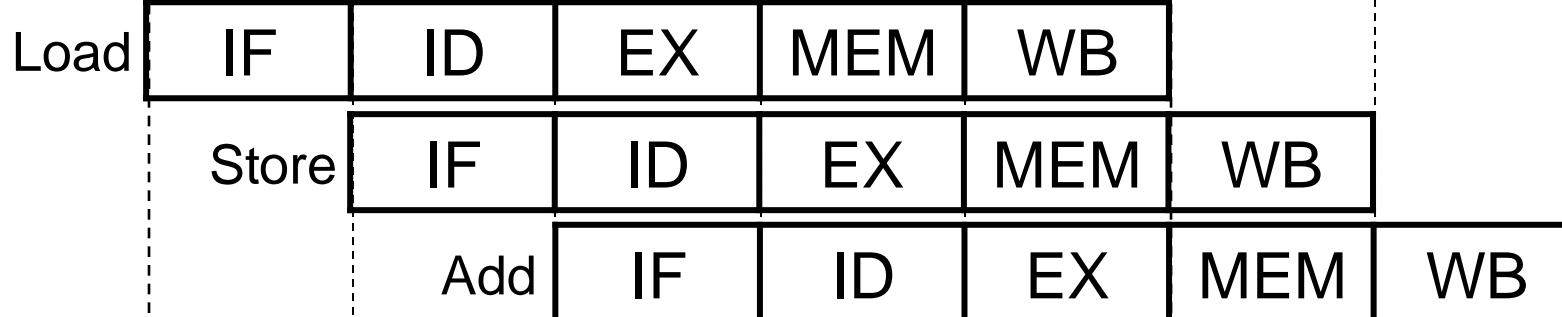
Clocks

Multicycle Implementation



The execution is still sequential

Pipeline Implementation



# Pipeline Performance

27

- Calculate cycle time assuming negligible delays except:
  - Memory access (200ps), ALU (200ps), Register access (100ps)

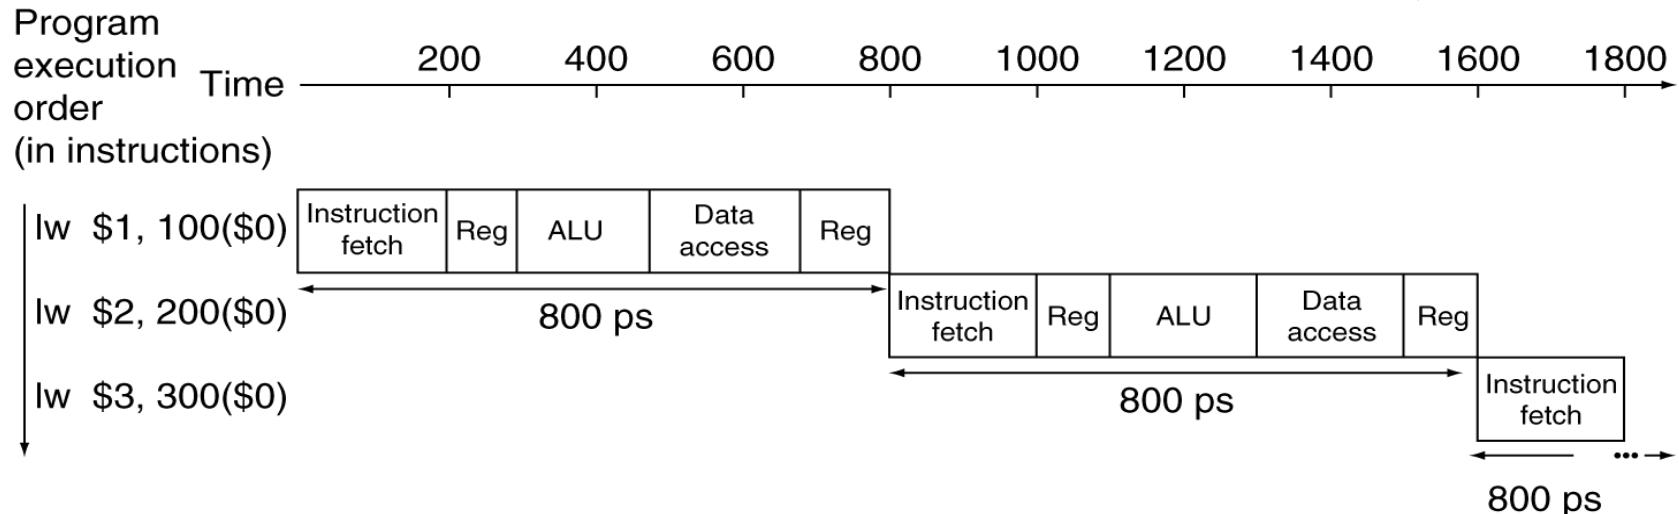
Instruction	Functional units used by the instruction class				
	Instruction fetch	Instruction decode	Execution	Memory Access	Register Write-back
R-type	200ps	100ps	200ps		100ps
Load word	200ps	100ps	200ps	200ps	100ps
Store word	200ps	100ps	200ps	200ps	
Conditional Branch	200ps	100ps	200ps		
Jump	200ps				

The diagram shows five rows corresponding to the instruction types. Each row contains a black arrow pointing to the right, followed by a blue rectangular box containing the total pipeline cycle time. The cycle times are: R-type (600ps), Load word (800ps), Store word (700ps), Conditional Branch (500ps), and Jump (200ps).

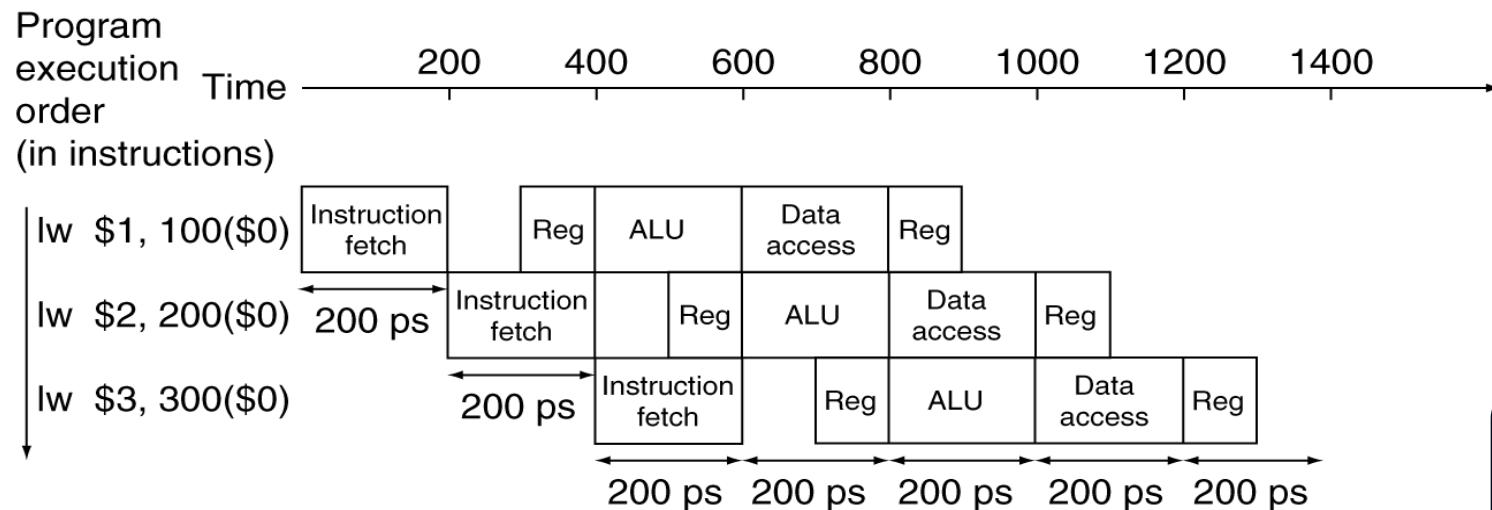
600ps  
800ps  
700ps  
500ps  
200ps

# Pipeline Performance: Speed Up?

28



**Without pipelining execution time:**  
 $800 + 800 + 800 = 2400 \text{ ps}$



**Pipelining execution time:**  
 $800 + 200 + 200 + 200 = 1400 \text{ ps}$

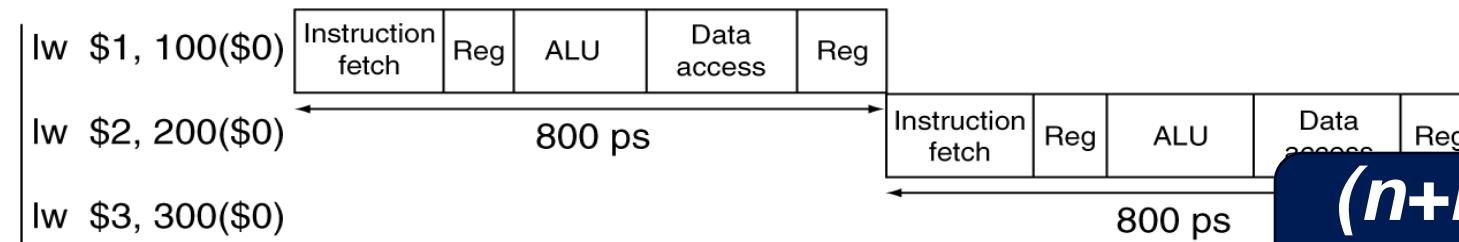
Speedup?  
 $2400/1400 \approx 1.7$

# What if We have 1,000,003 Instructions?

29

Program  
execution  
order  
(in instructions)

Time 200 400 600 800 1000 1200 1400 1600 1800



Without pipelining  
execution time:

$$n*T = n*800\text{ps}$$

$$(n+k-1) * T/K = \\ (n+5-1) * 200\text{ps}$$

Program  
execution  
order  
(in instructions)

Time 200 400 600 800 1000 1200 1400



Pipelining  
execution time:

$$\text{Speedup} = \frac{n*800\text{ps}}{(n+4)*200\text{ps}} \quad (n=1,000,003)$$

$$= 800,002,400 \text{ ps} / 200,01,400 \text{ ps} \cong 4.00$$

# MIPS ISA is Designed to Work Well with Pipelining

---



- **All instructions are 32-bits**
  - Consistent and short instruction fetch/decode times
- **Few and regular instruction formats**
  - Consistent and short instruction decode time
- **Load/store addressing**
  - Fixed sequence: calculate address in 3rd stage, access memory in 4th stage

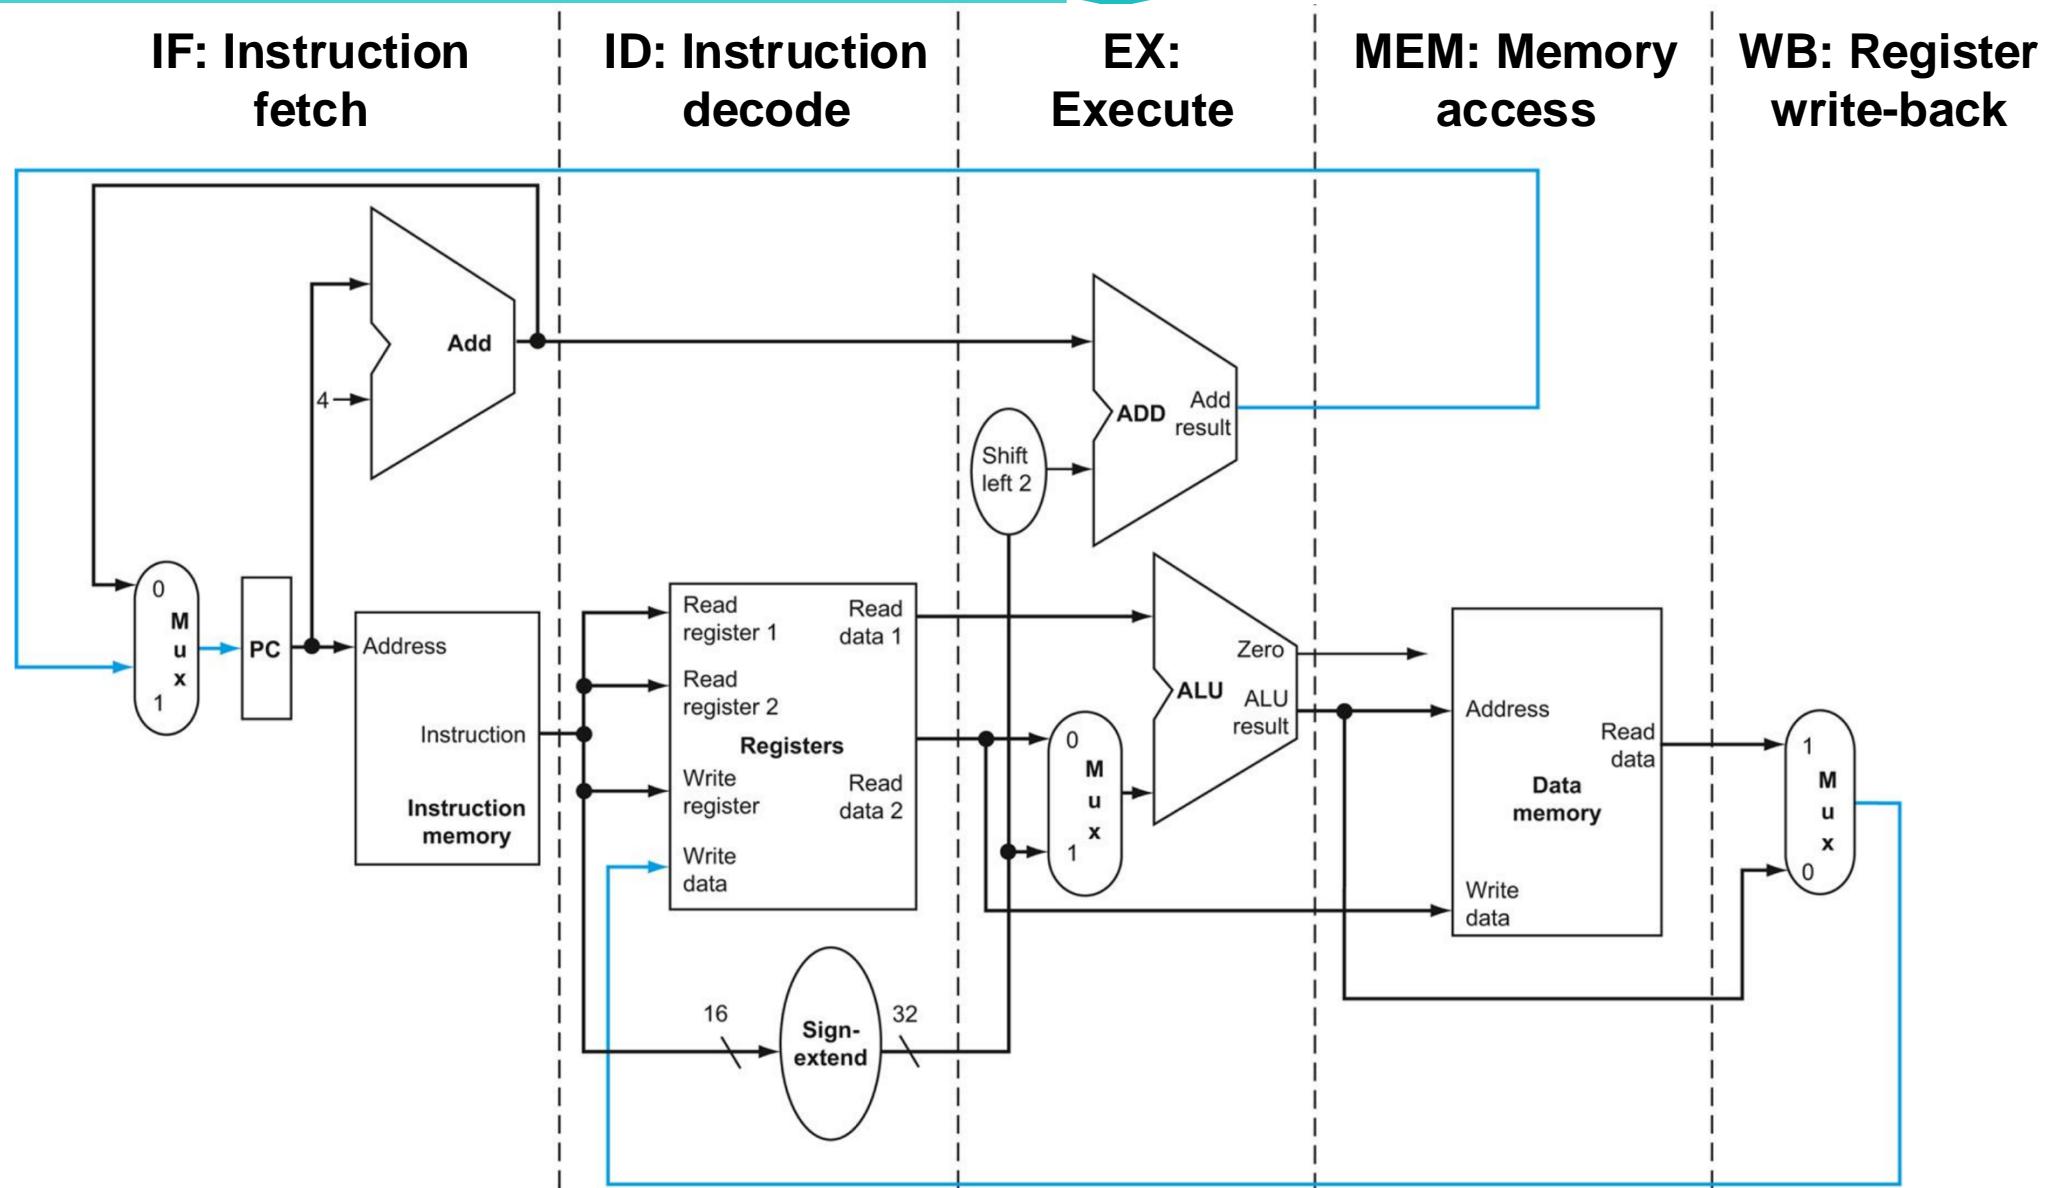
# Pipeline Summary



- Pipelining improves performance by increasing instruction throughput
  - Executes multiple instructions in parallel
  - Each instruction has the same latency
- The speedup depends on longest pipe stage and the number of pipe stages
  - Length of pipe stage  $\downarrow \rightarrow$  speedup  $\uparrow$
  - # of pipe stage  $\uparrow \rightarrow$  speedup  $\uparrow$
  - Instruction set design affects complexity of pipeline implementation

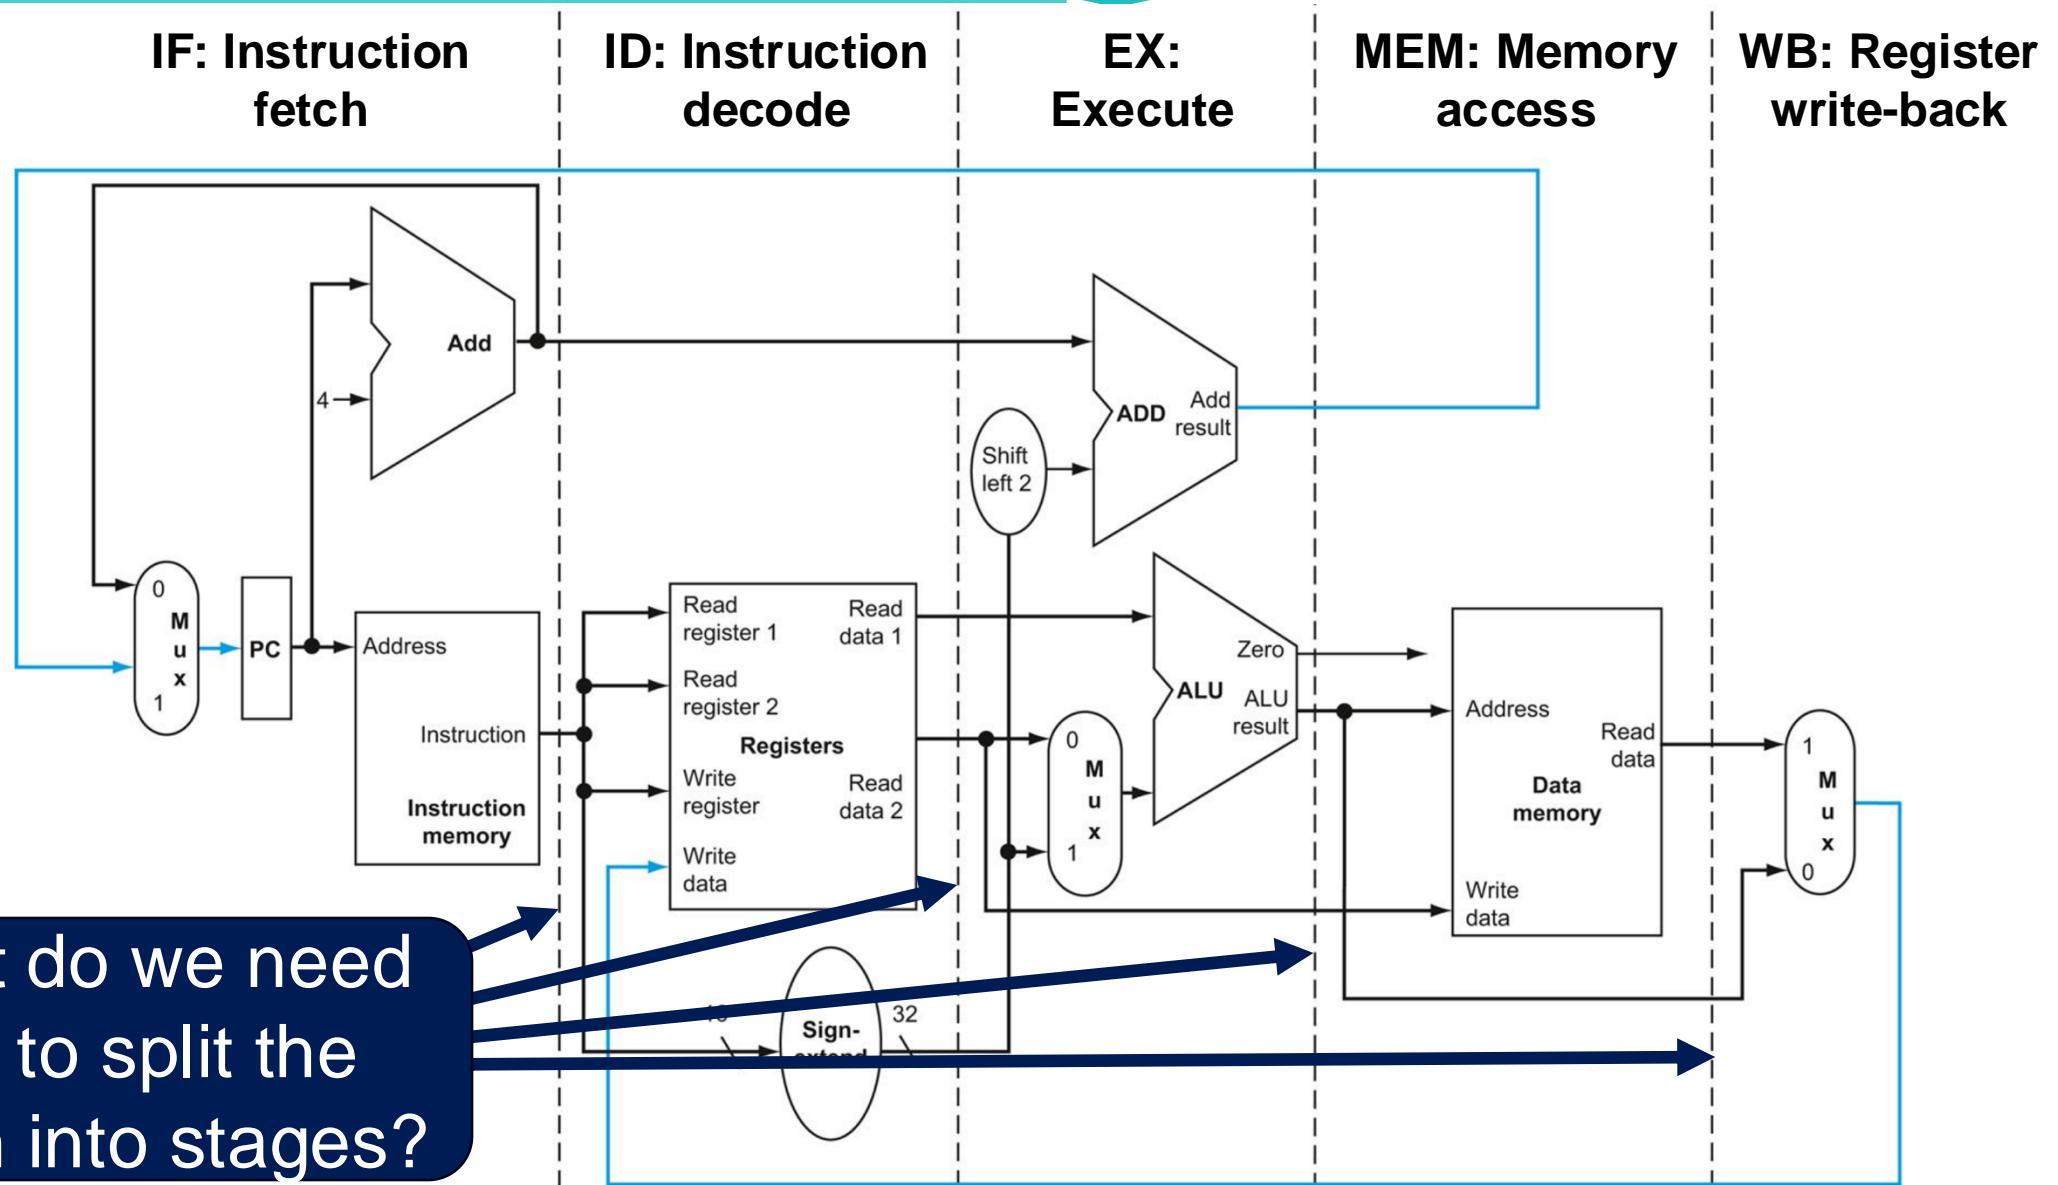
# Pipelined Datapath

# MIPS Pipelined Datapath



# MIPS Pipelined Datapath

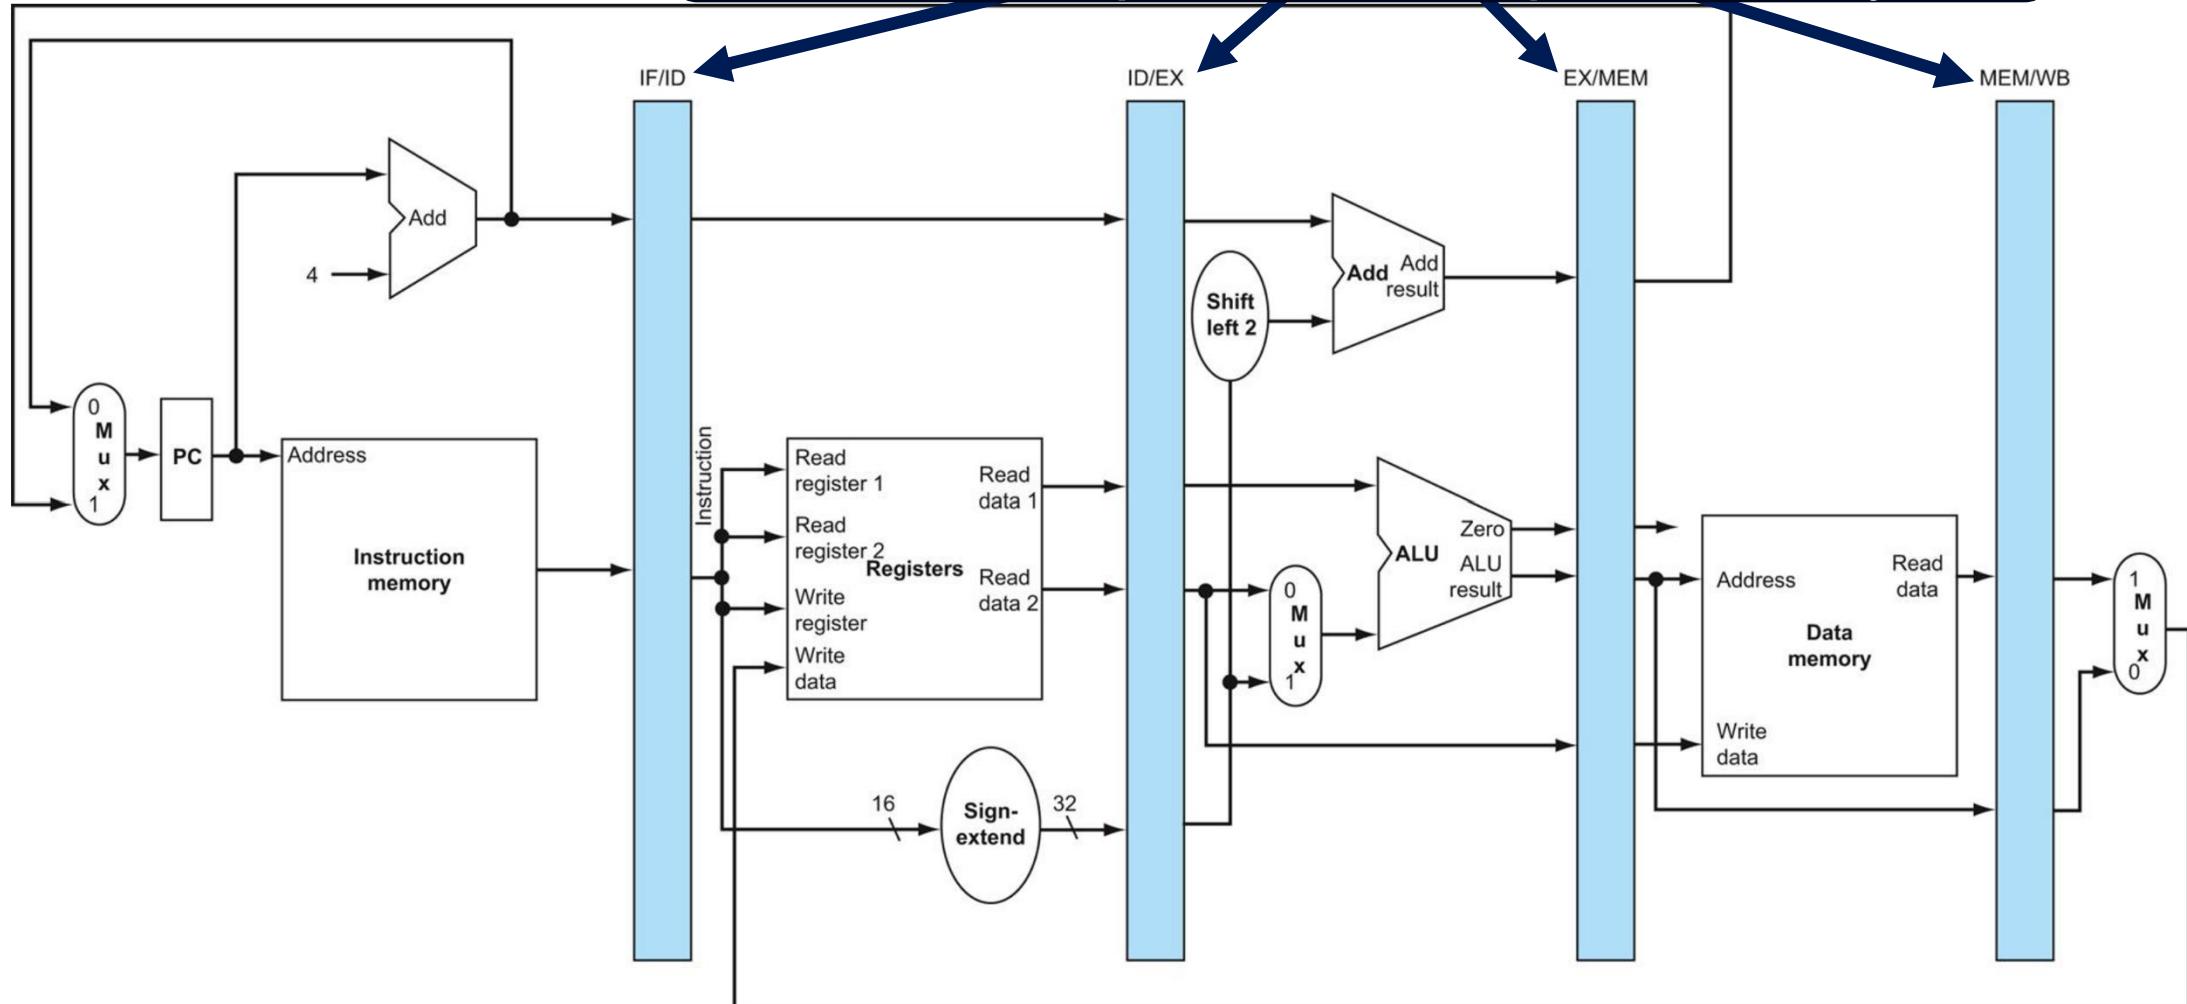
34



# The Pipelined Version of the Datapath

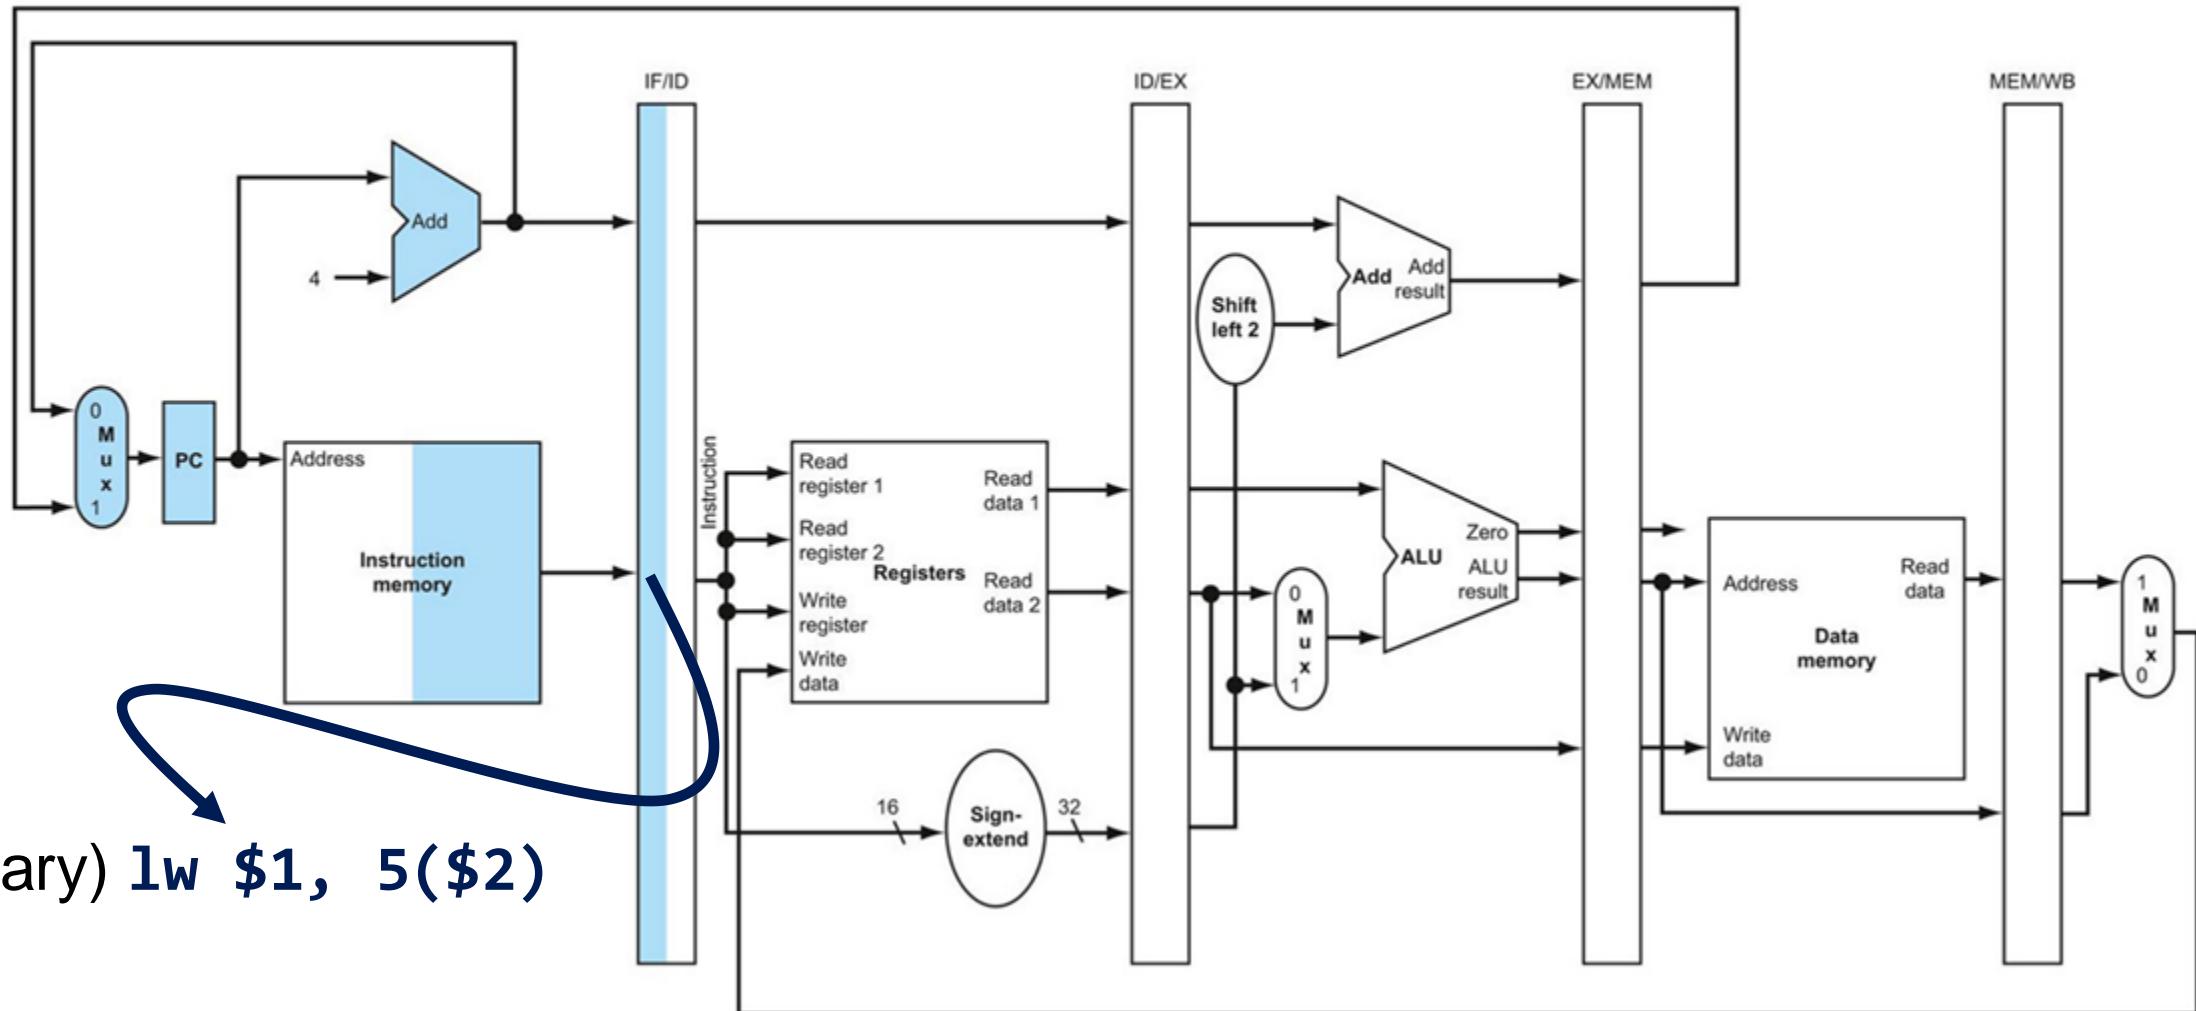
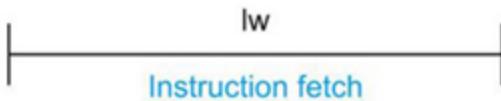
35

Registers between stages to hold information produced in previous cycle



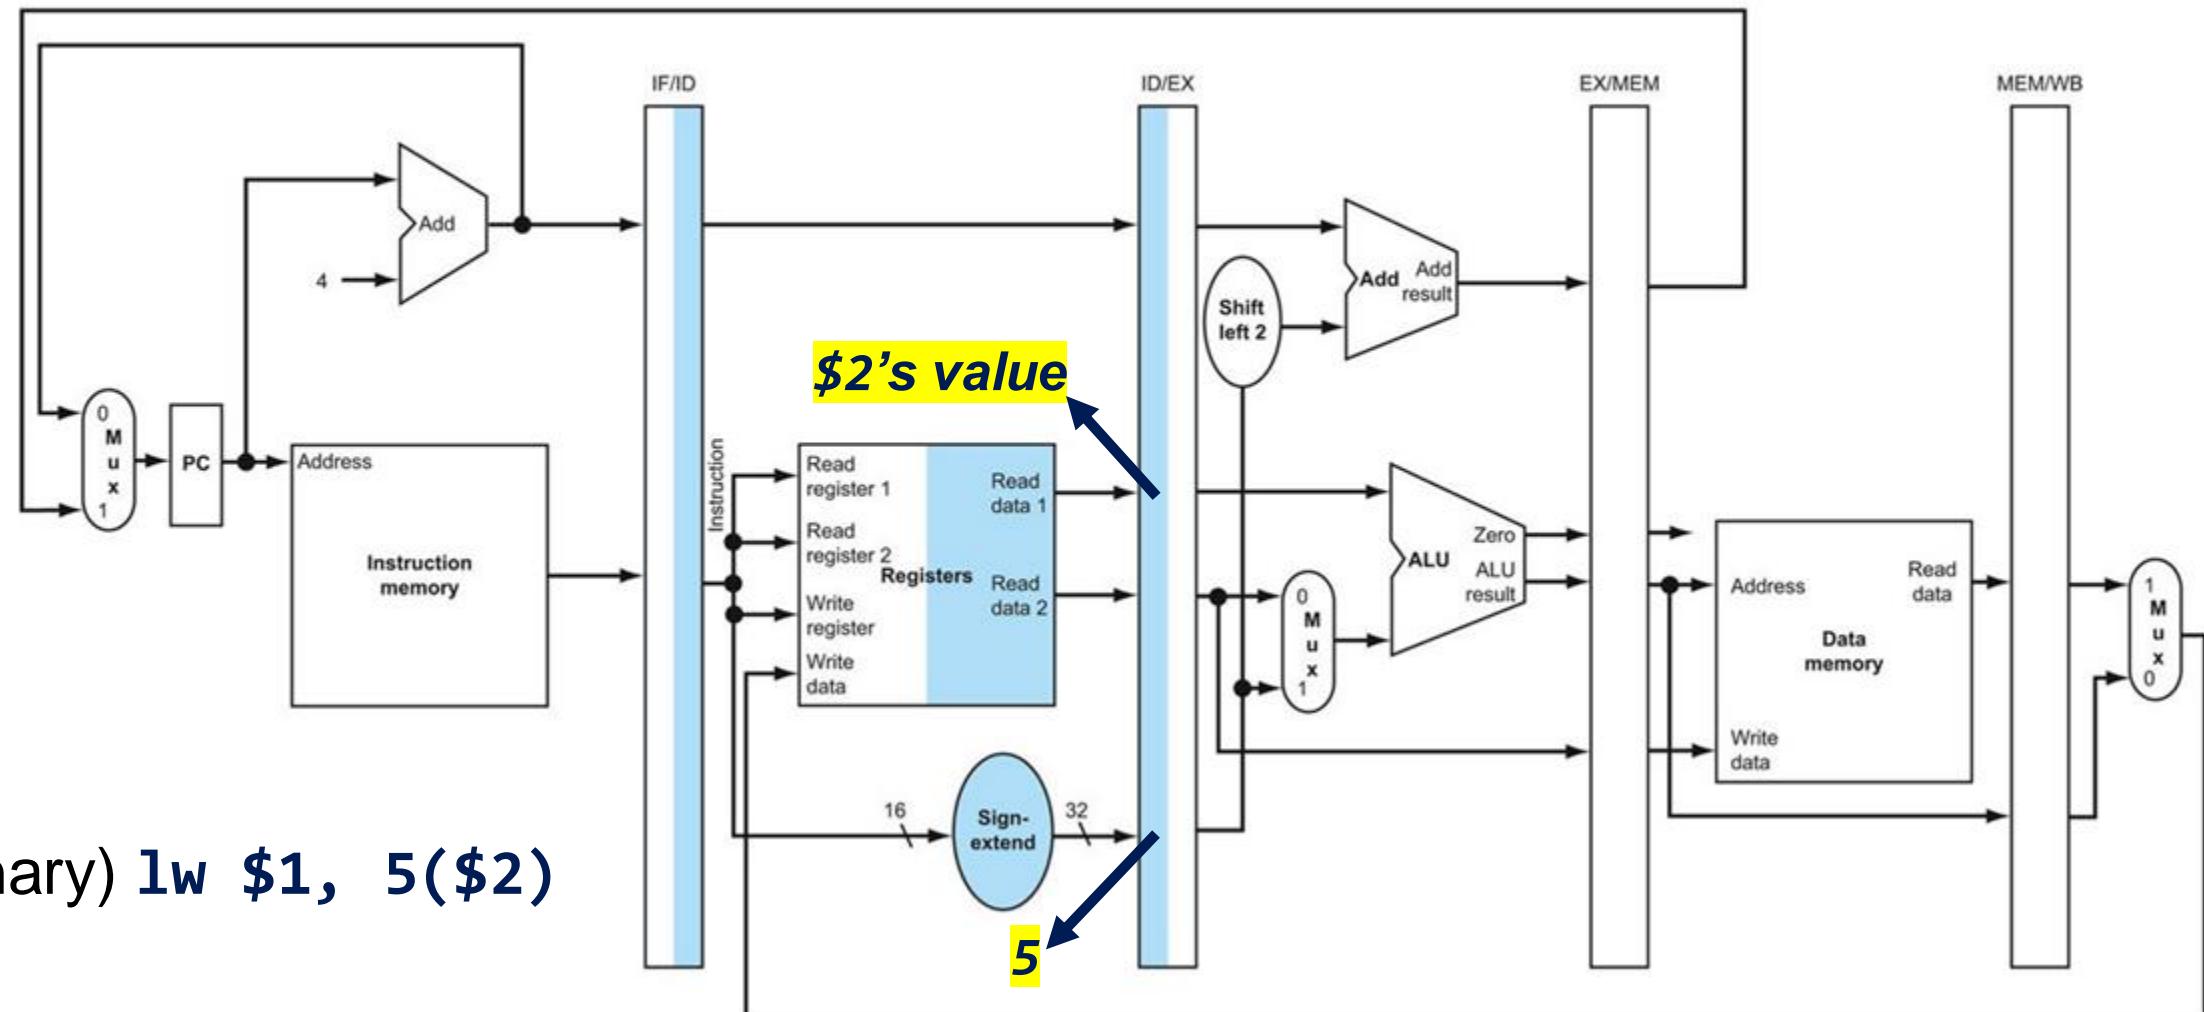
# The Pipelined Version of the Datapath

36



# The Pipelined Version of the Datapath

37

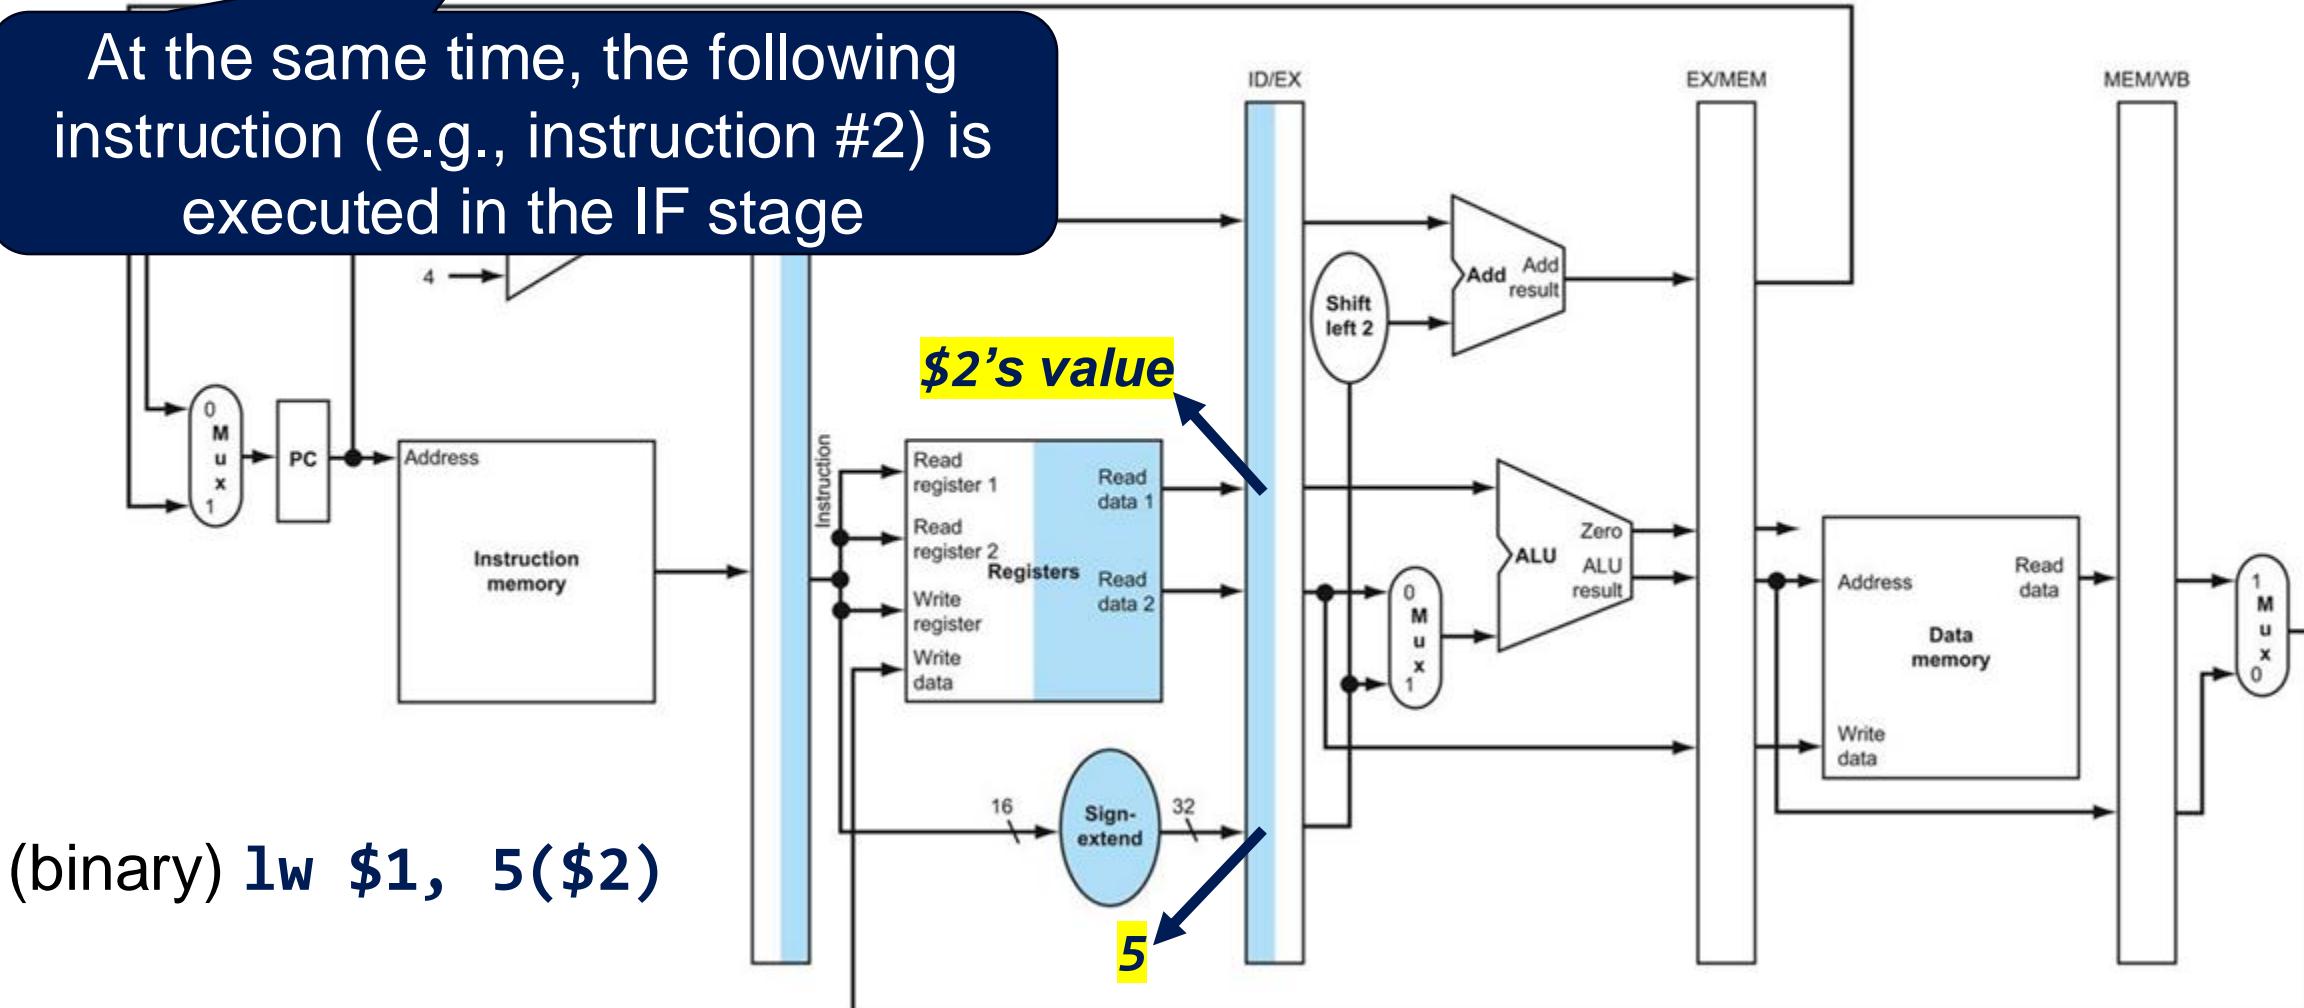


# The Pipelined Version of the Datapath

38

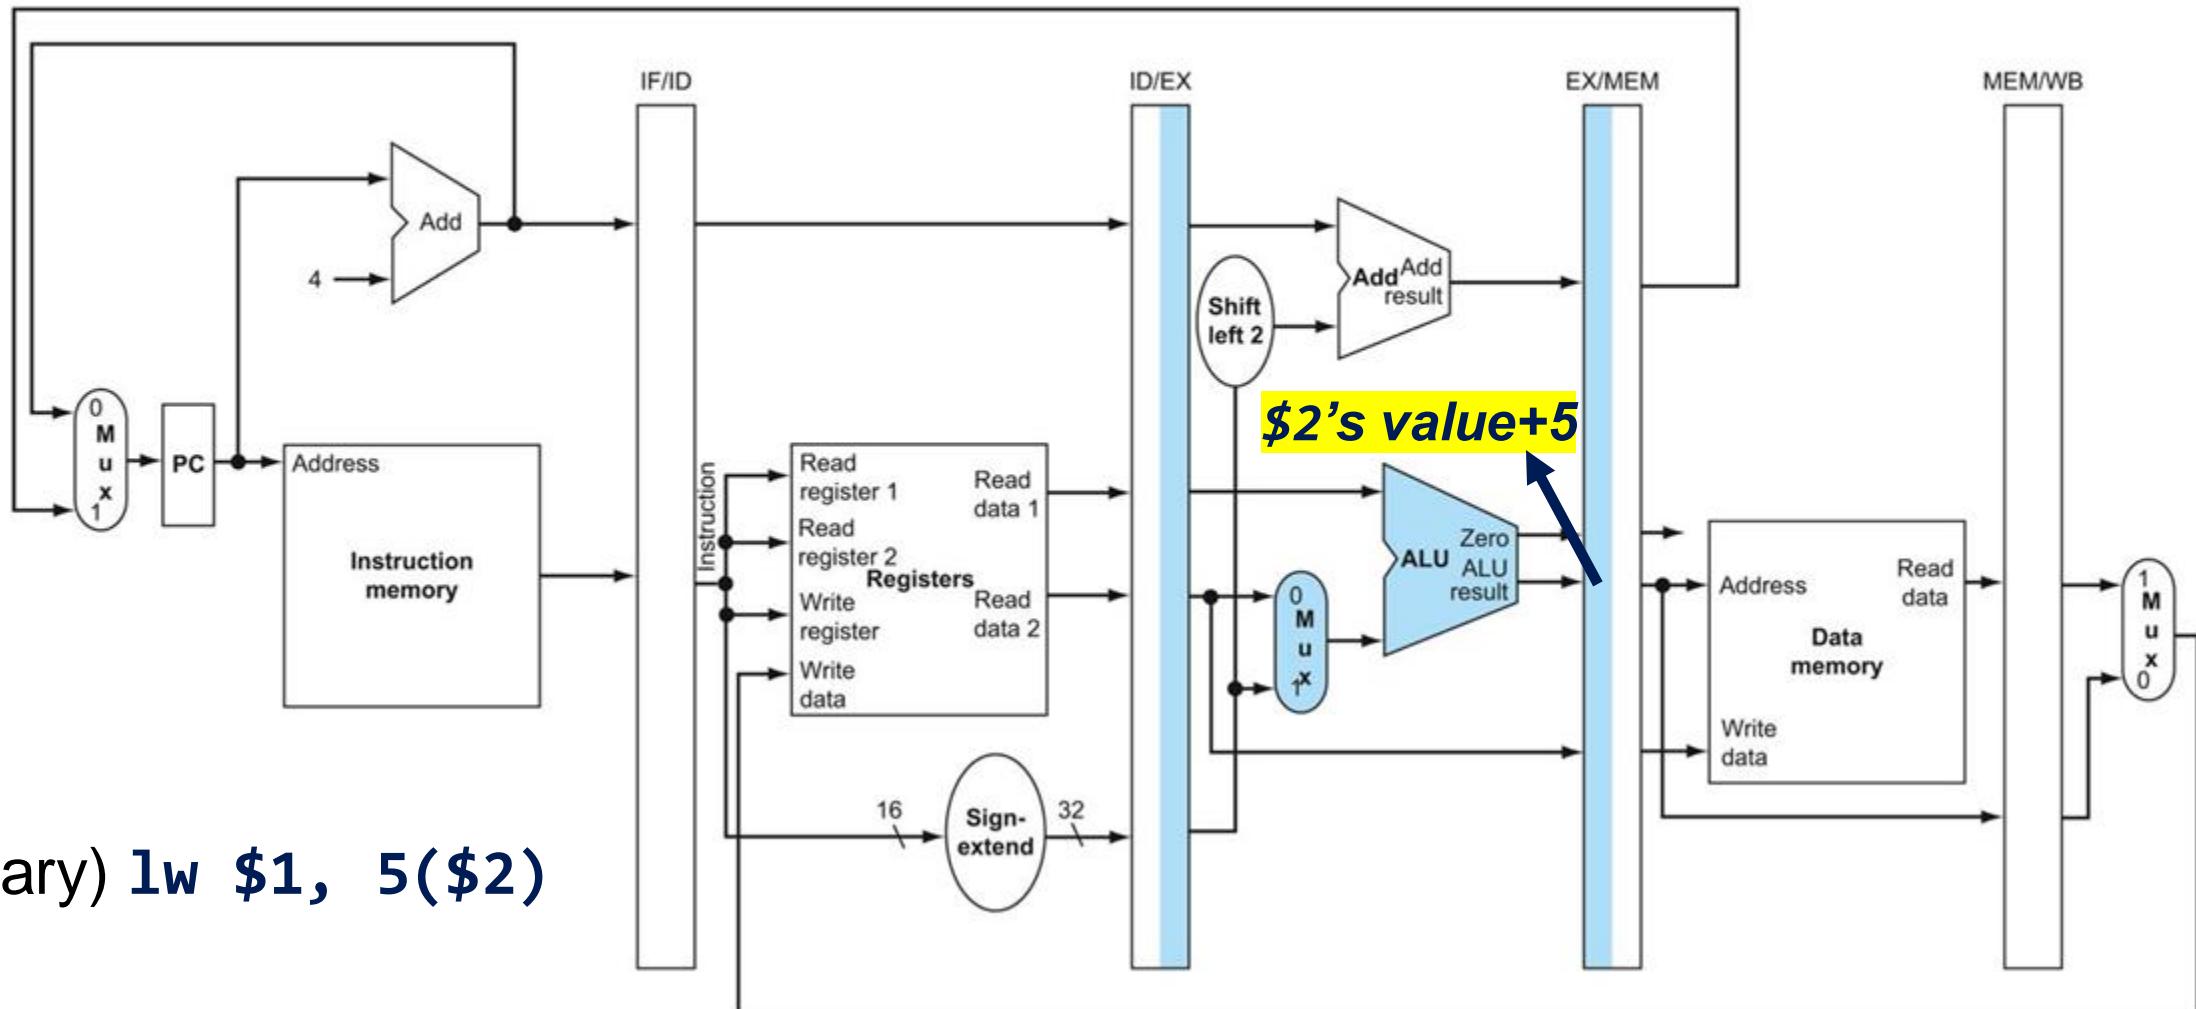


At the same time, the following instruction (e.g., instruction #2) is executed in the IF stage



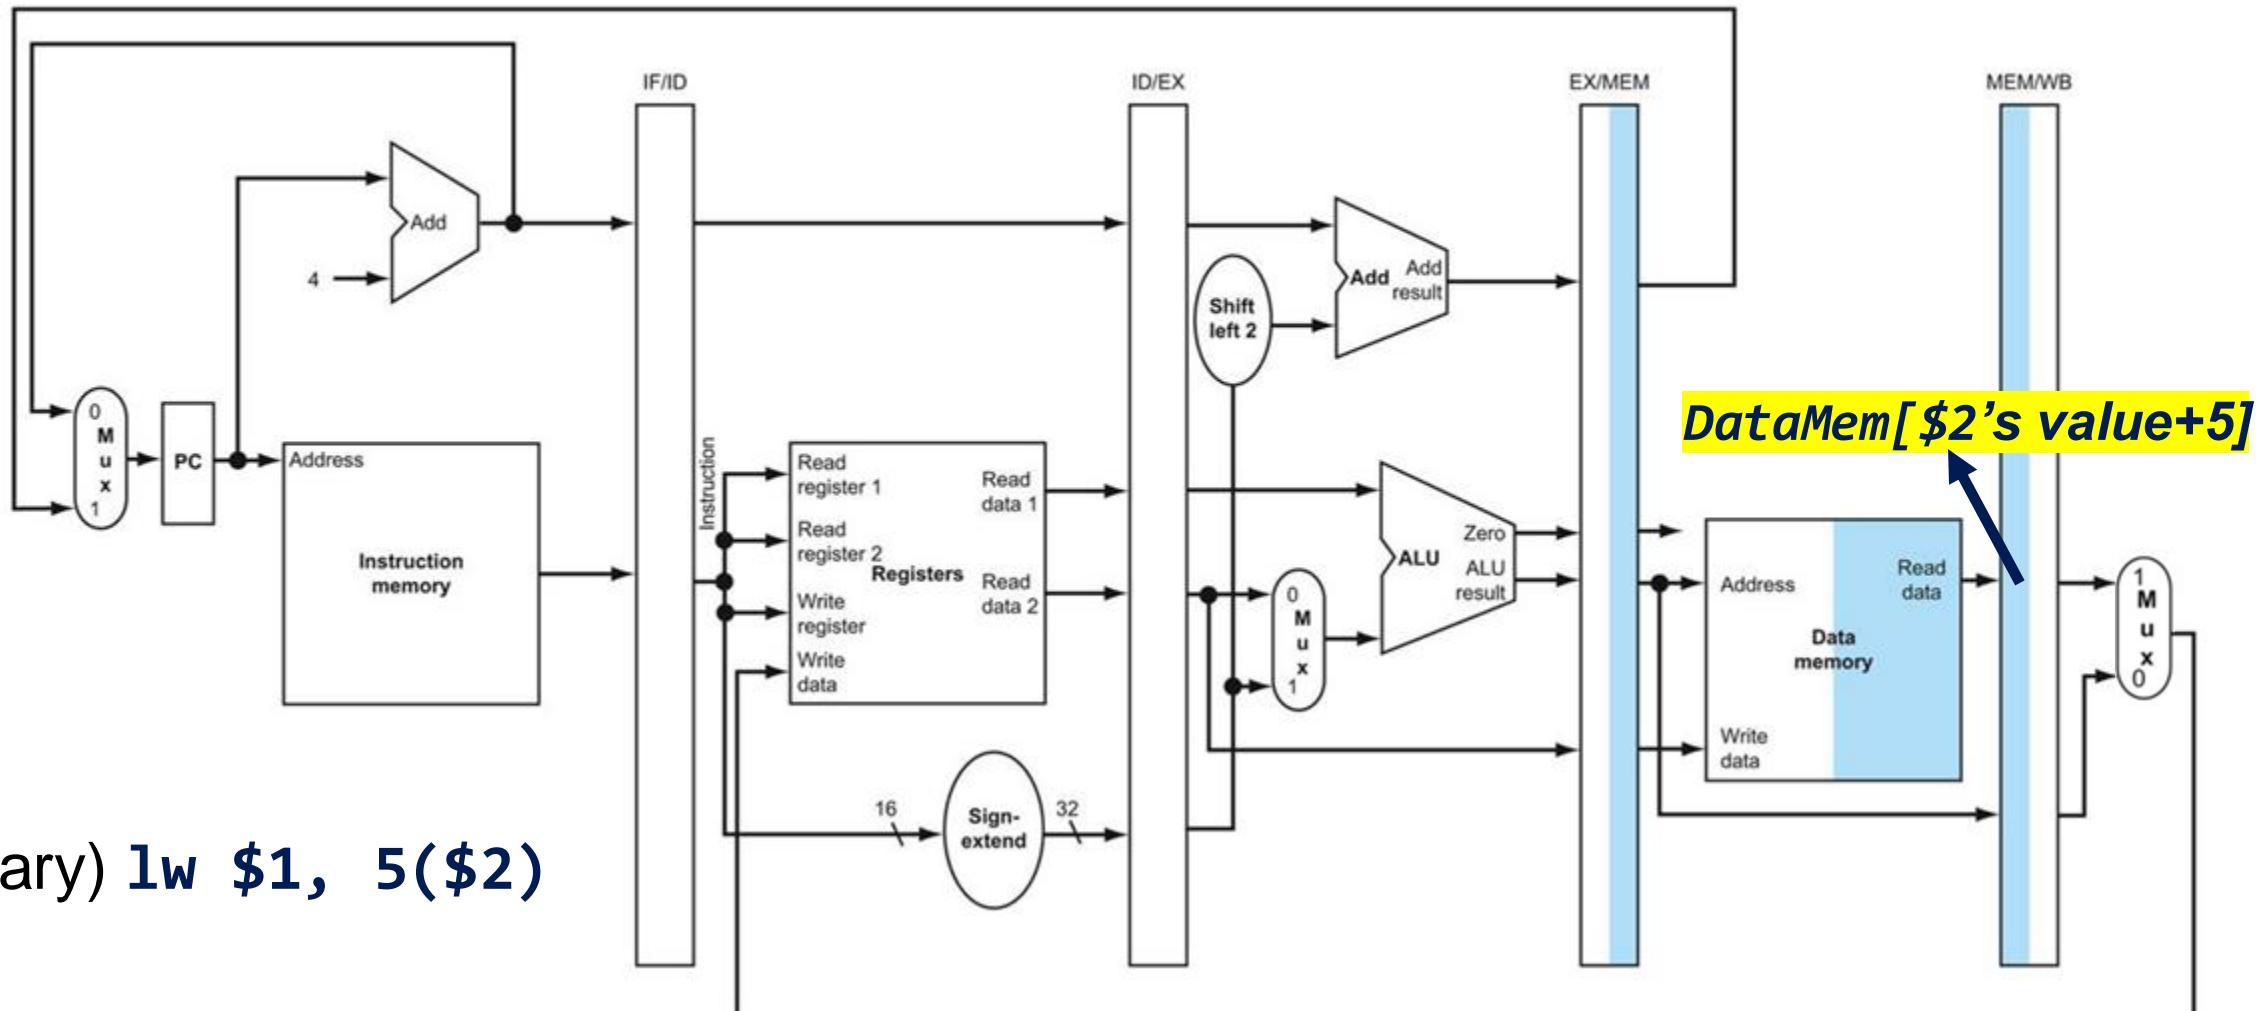
# The Pipelined Version of the Datapath

39



# The Pipelined Version of the Datapath

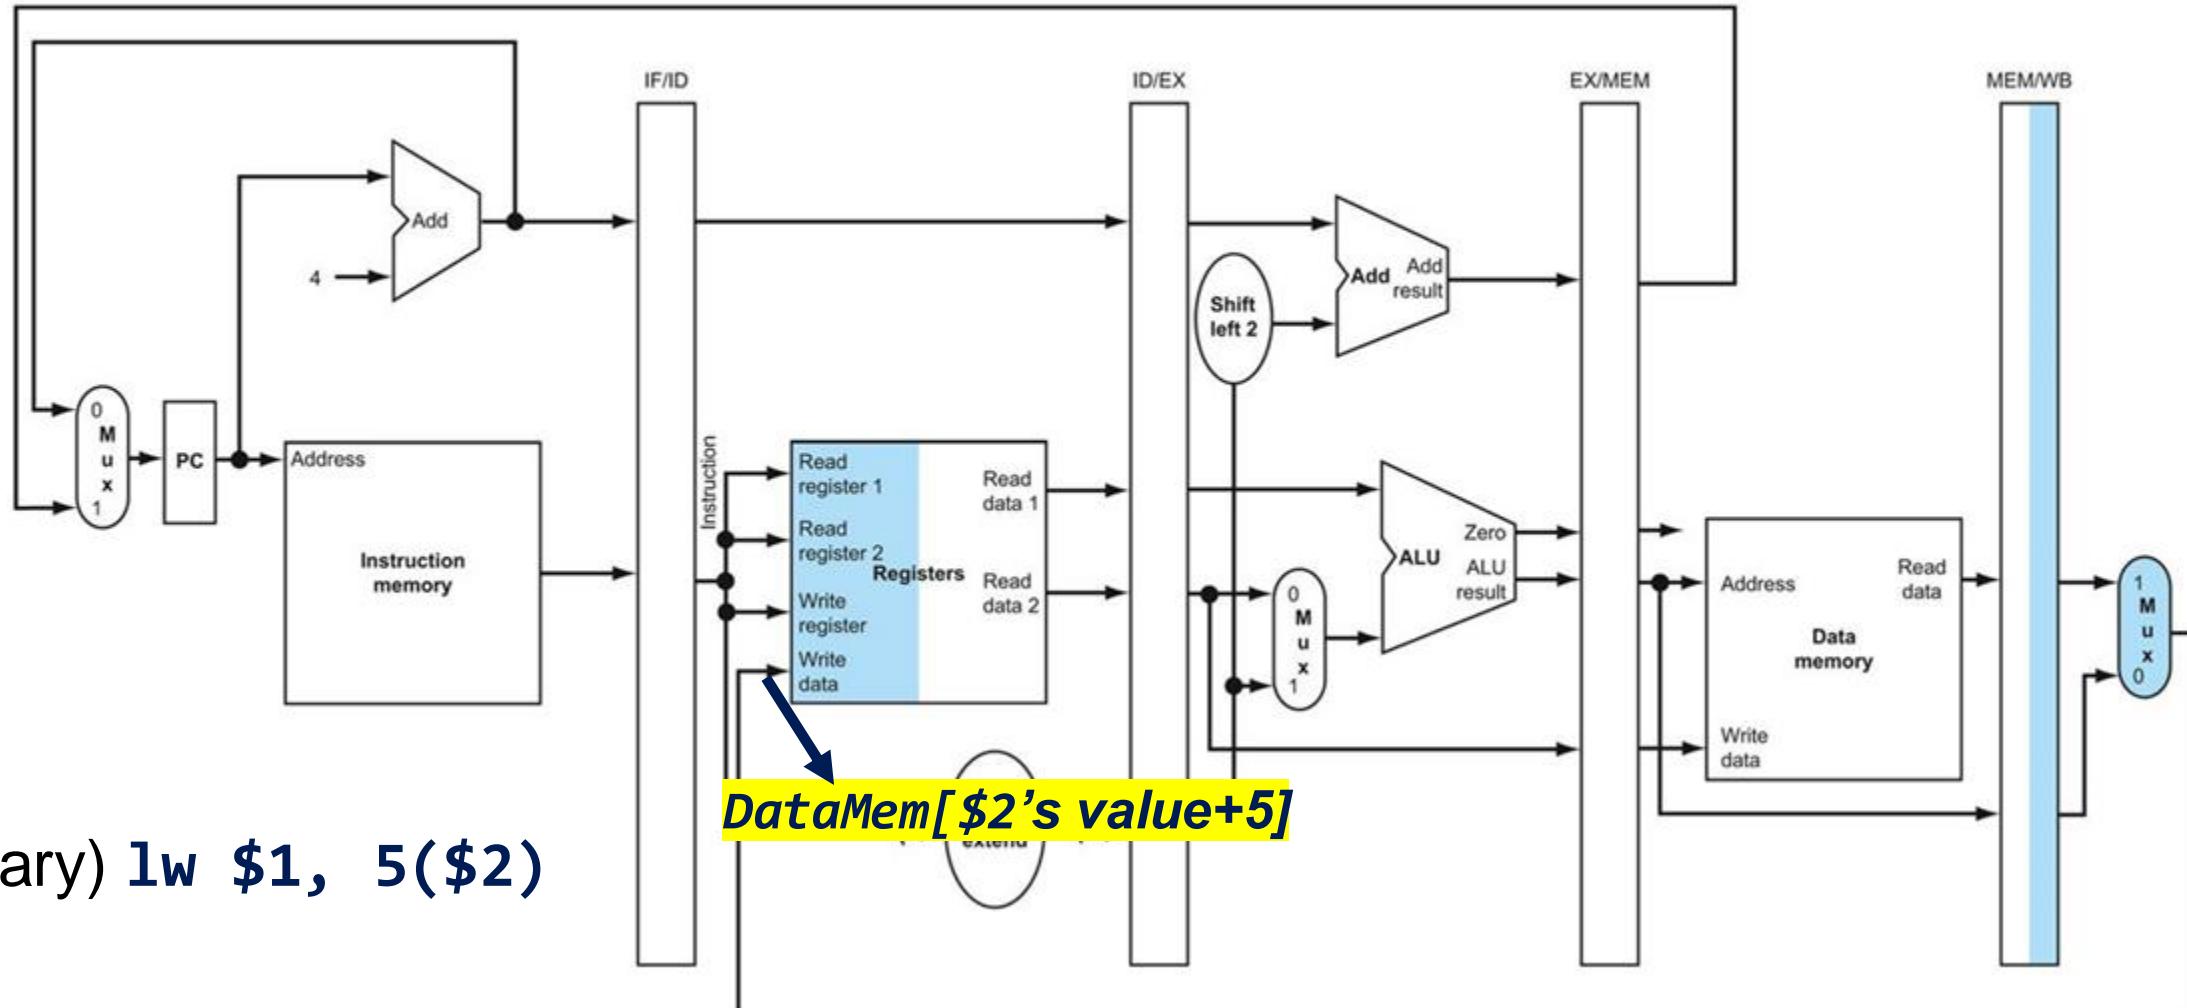
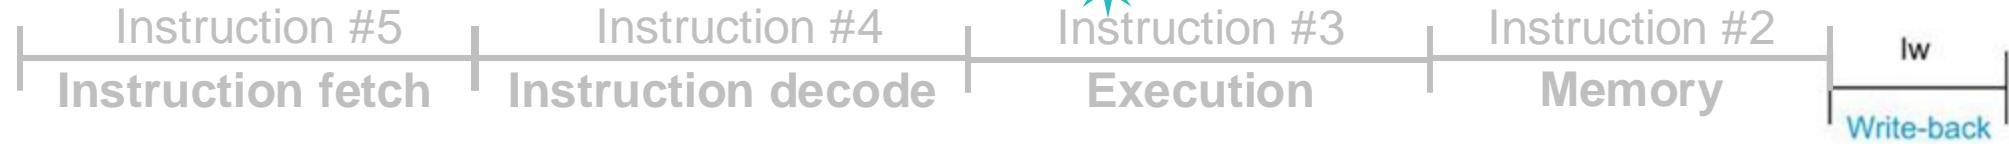
40



(binary) **Iw \$1, 5(\$2)**

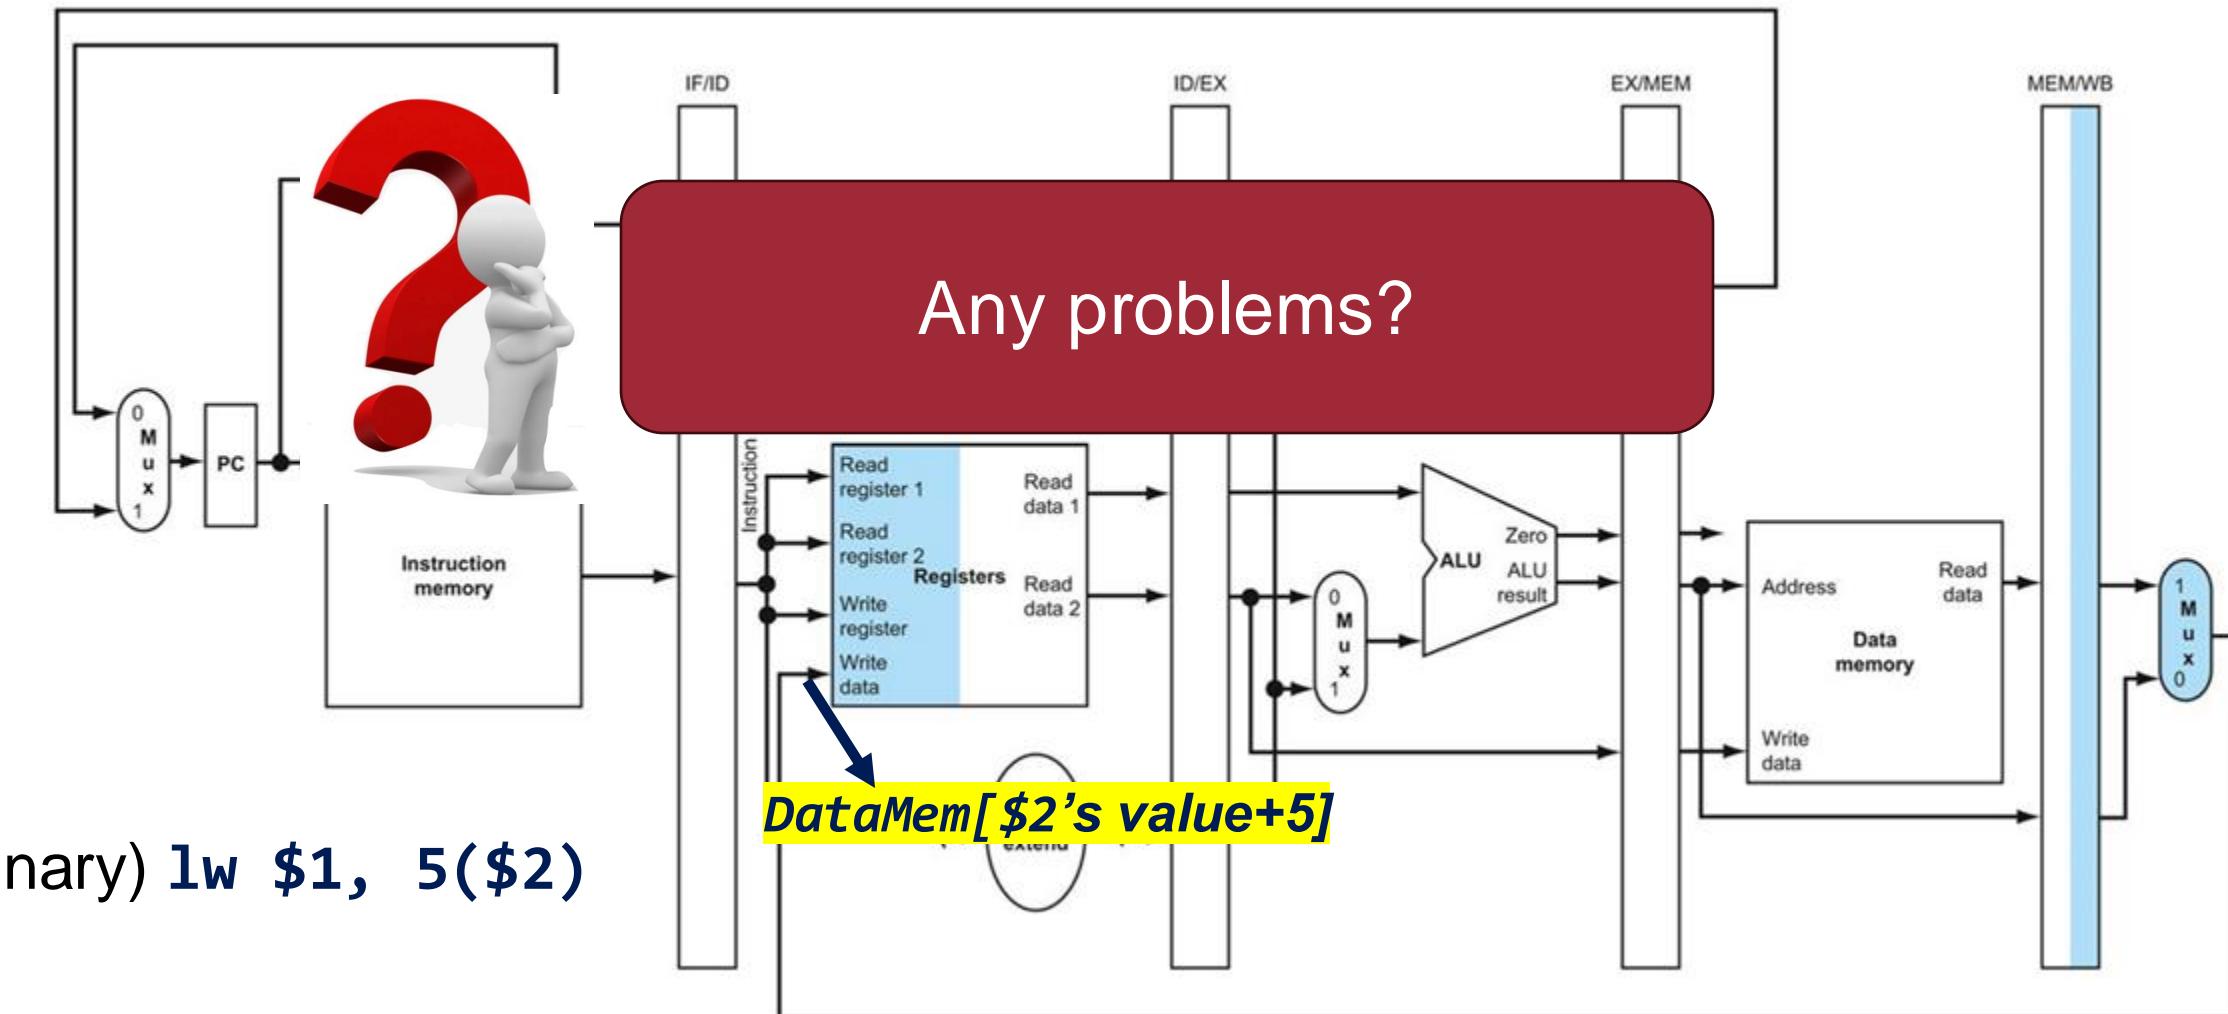
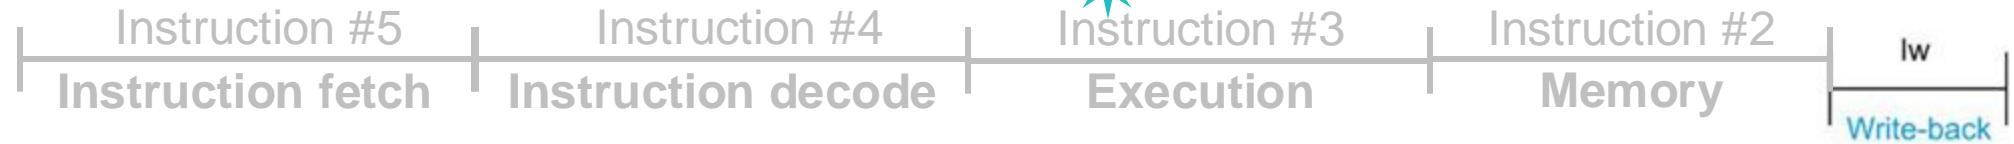
# The Pipelined Version of the Datapath

41



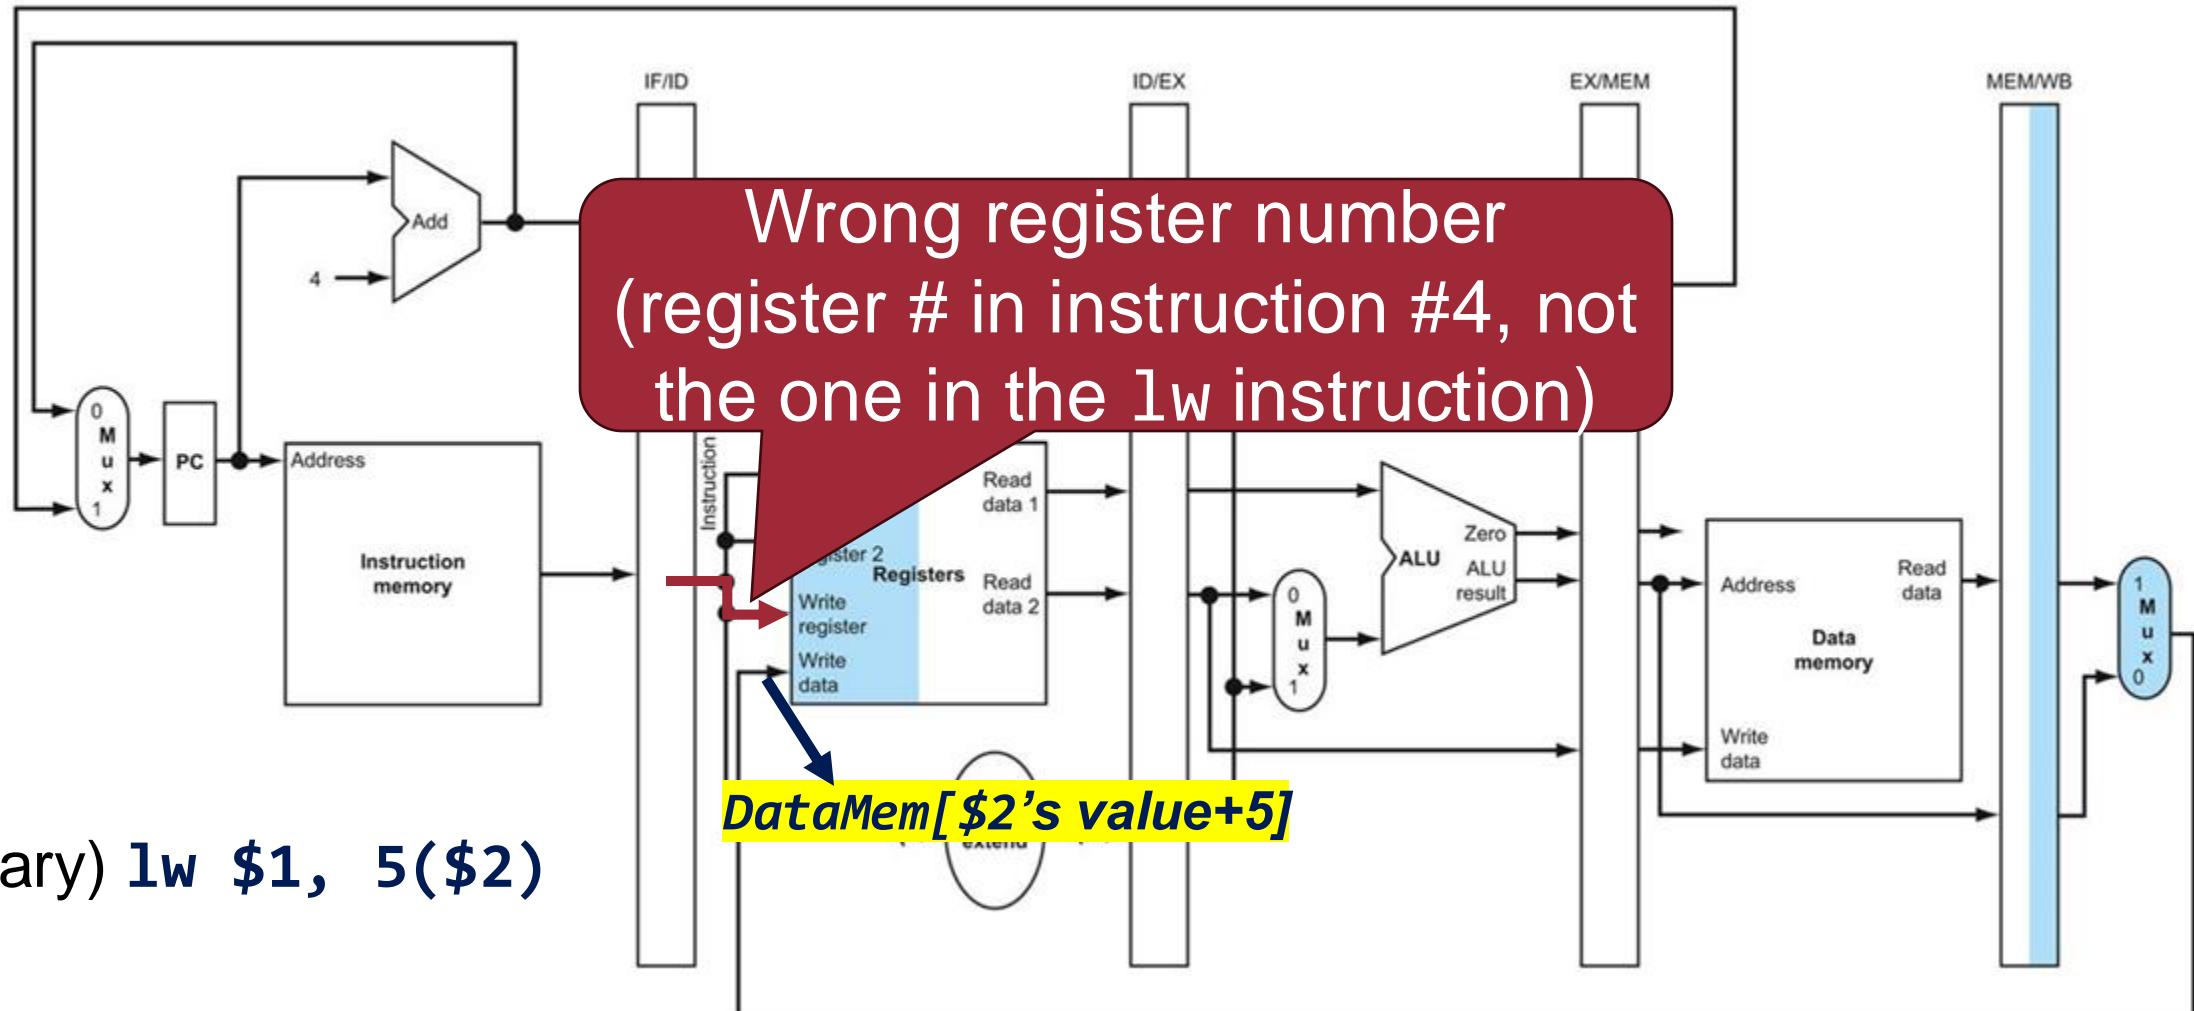
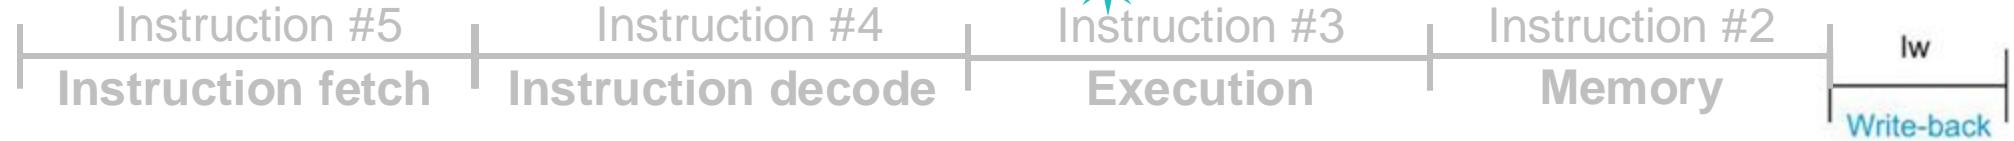
# The Pipelined Version of the Datapath

42



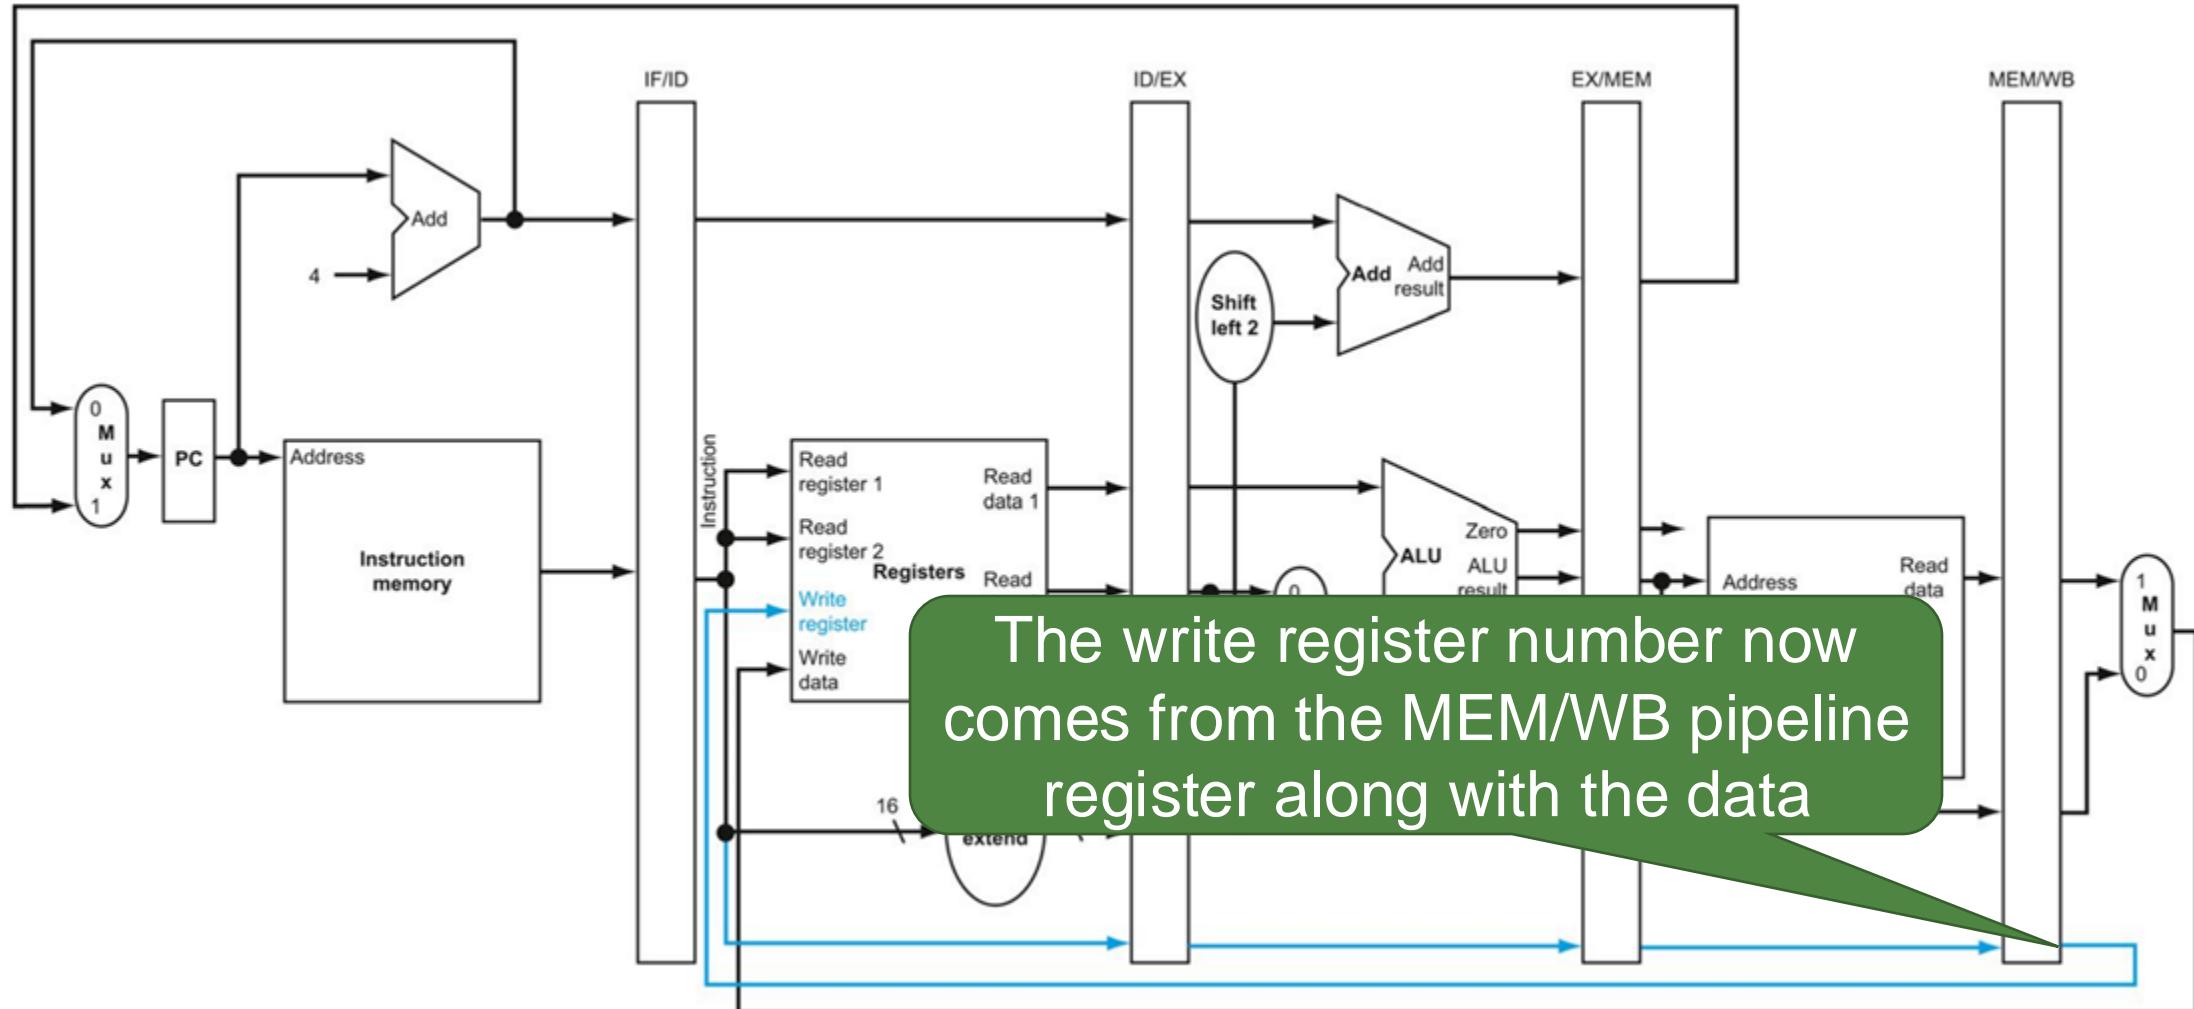
# The Pipelined Version of the Datapath

43



# Solution: Corrected Datapath for Load

44



# Pipelining Example

45

```
lw    $10, 20($1)
sub   $11, $2, $3
add   $12, $3, $4
lw    $13, 24($1)
add   $14, $5, $6
```

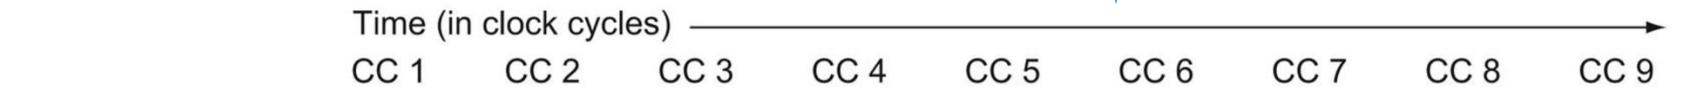


***Assumption: there is no hazard***

Later, we will cover about hazards!

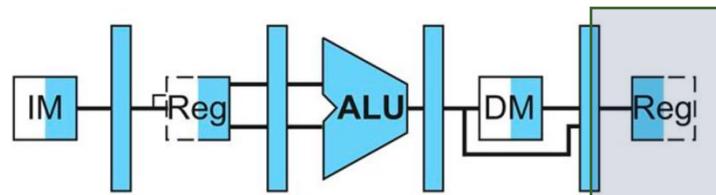
# Pipeline Diagram of Five Instructions

46

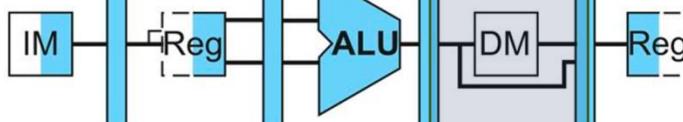


Program  
execution  
order  
(in instructions)

lw \$10, 20(\$1)



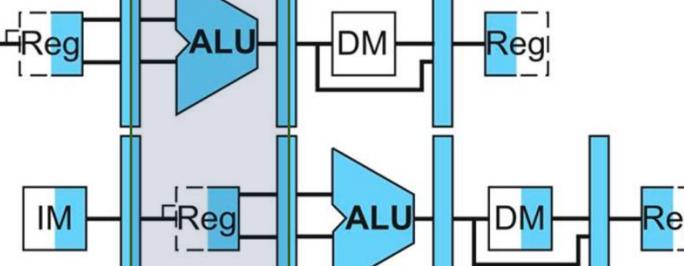
sub \$11, \$2, \$3



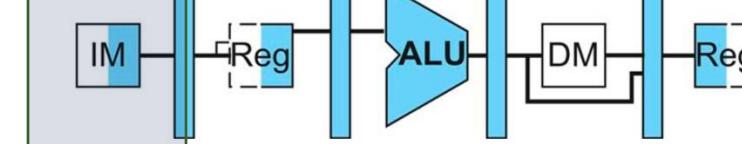
add \$12, \$3, \$4



lw \$13, 24(\$1)



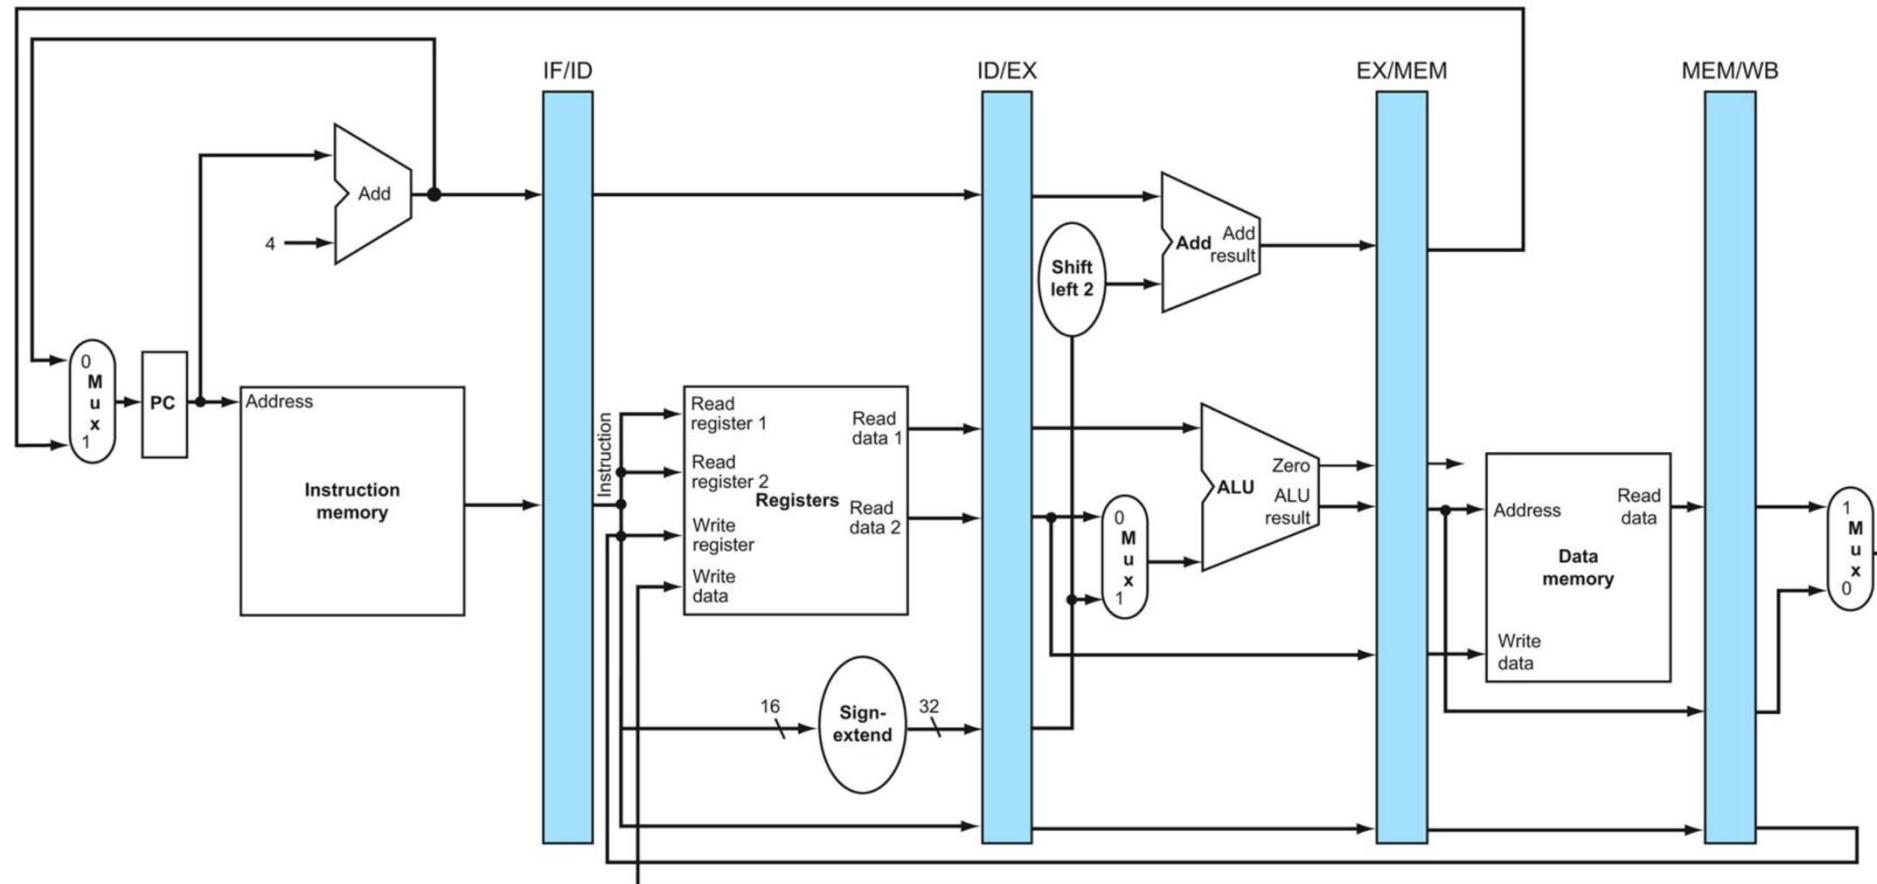
add \$14, \$5, \$6



At the cycle 5, we can fully  
utilize the hardware resources

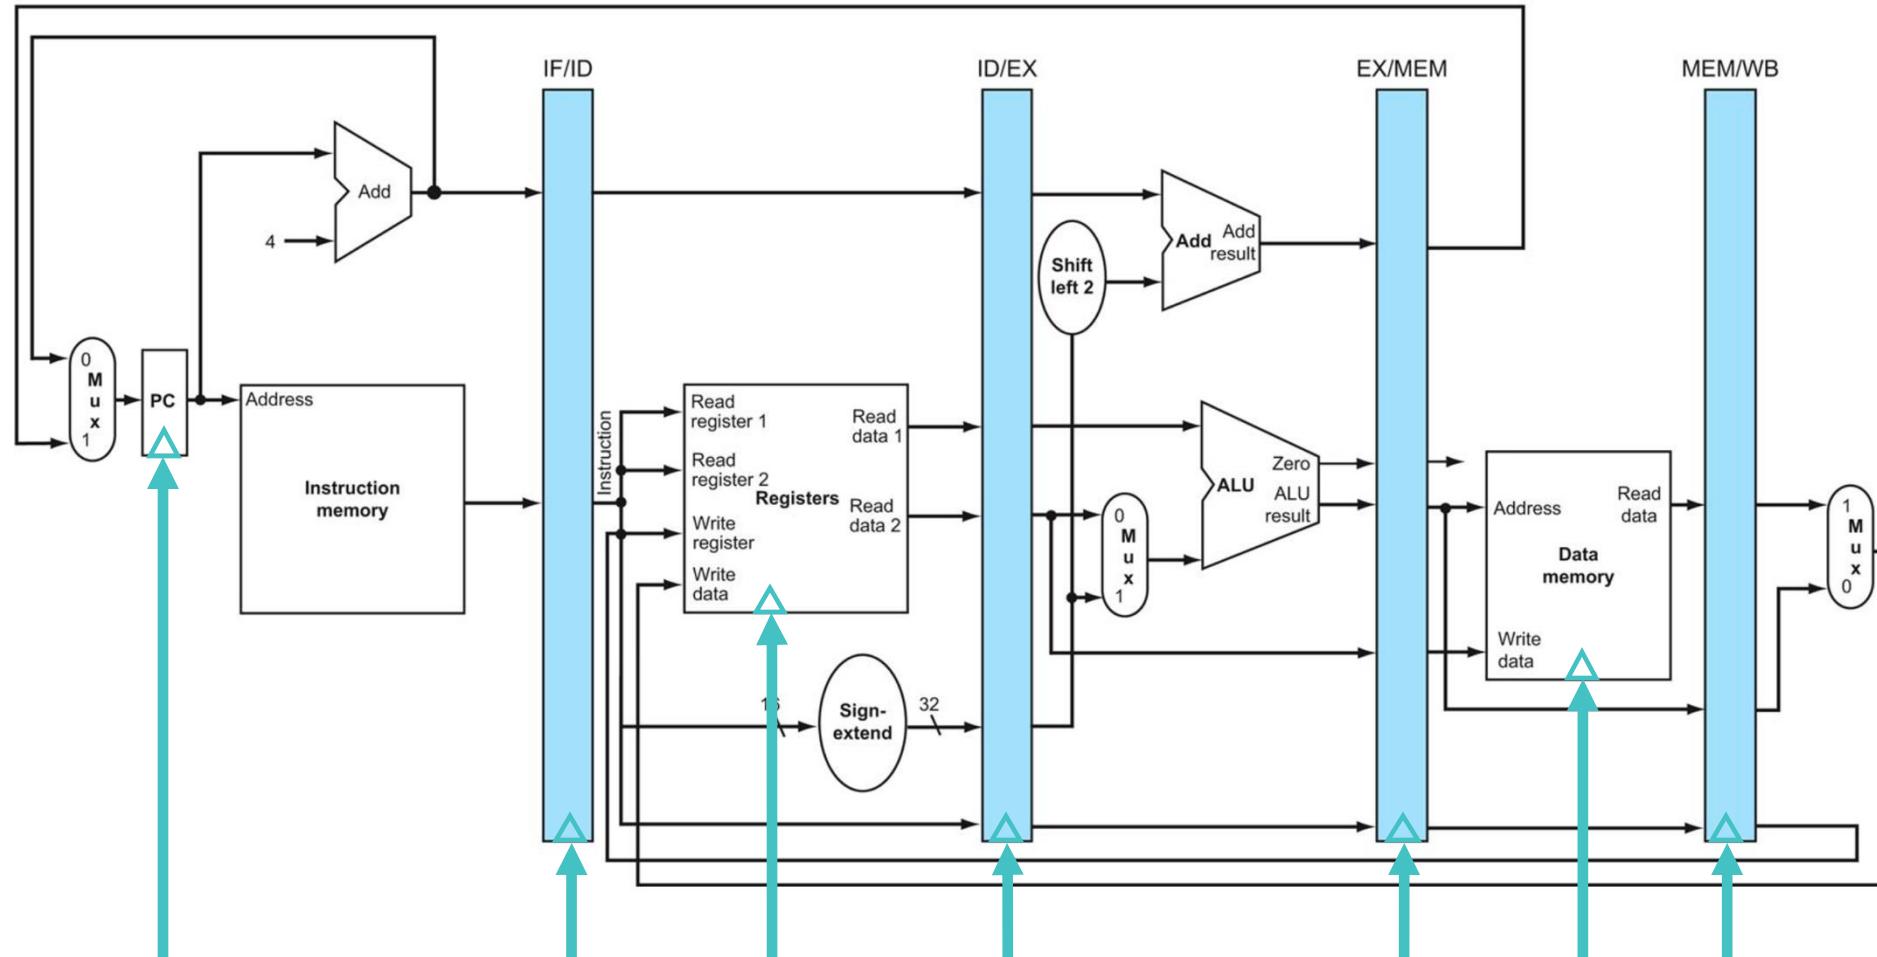
# Clock Cycle 5 of the Pipeline

47



# Clock Cycle 5 of the Pipeline

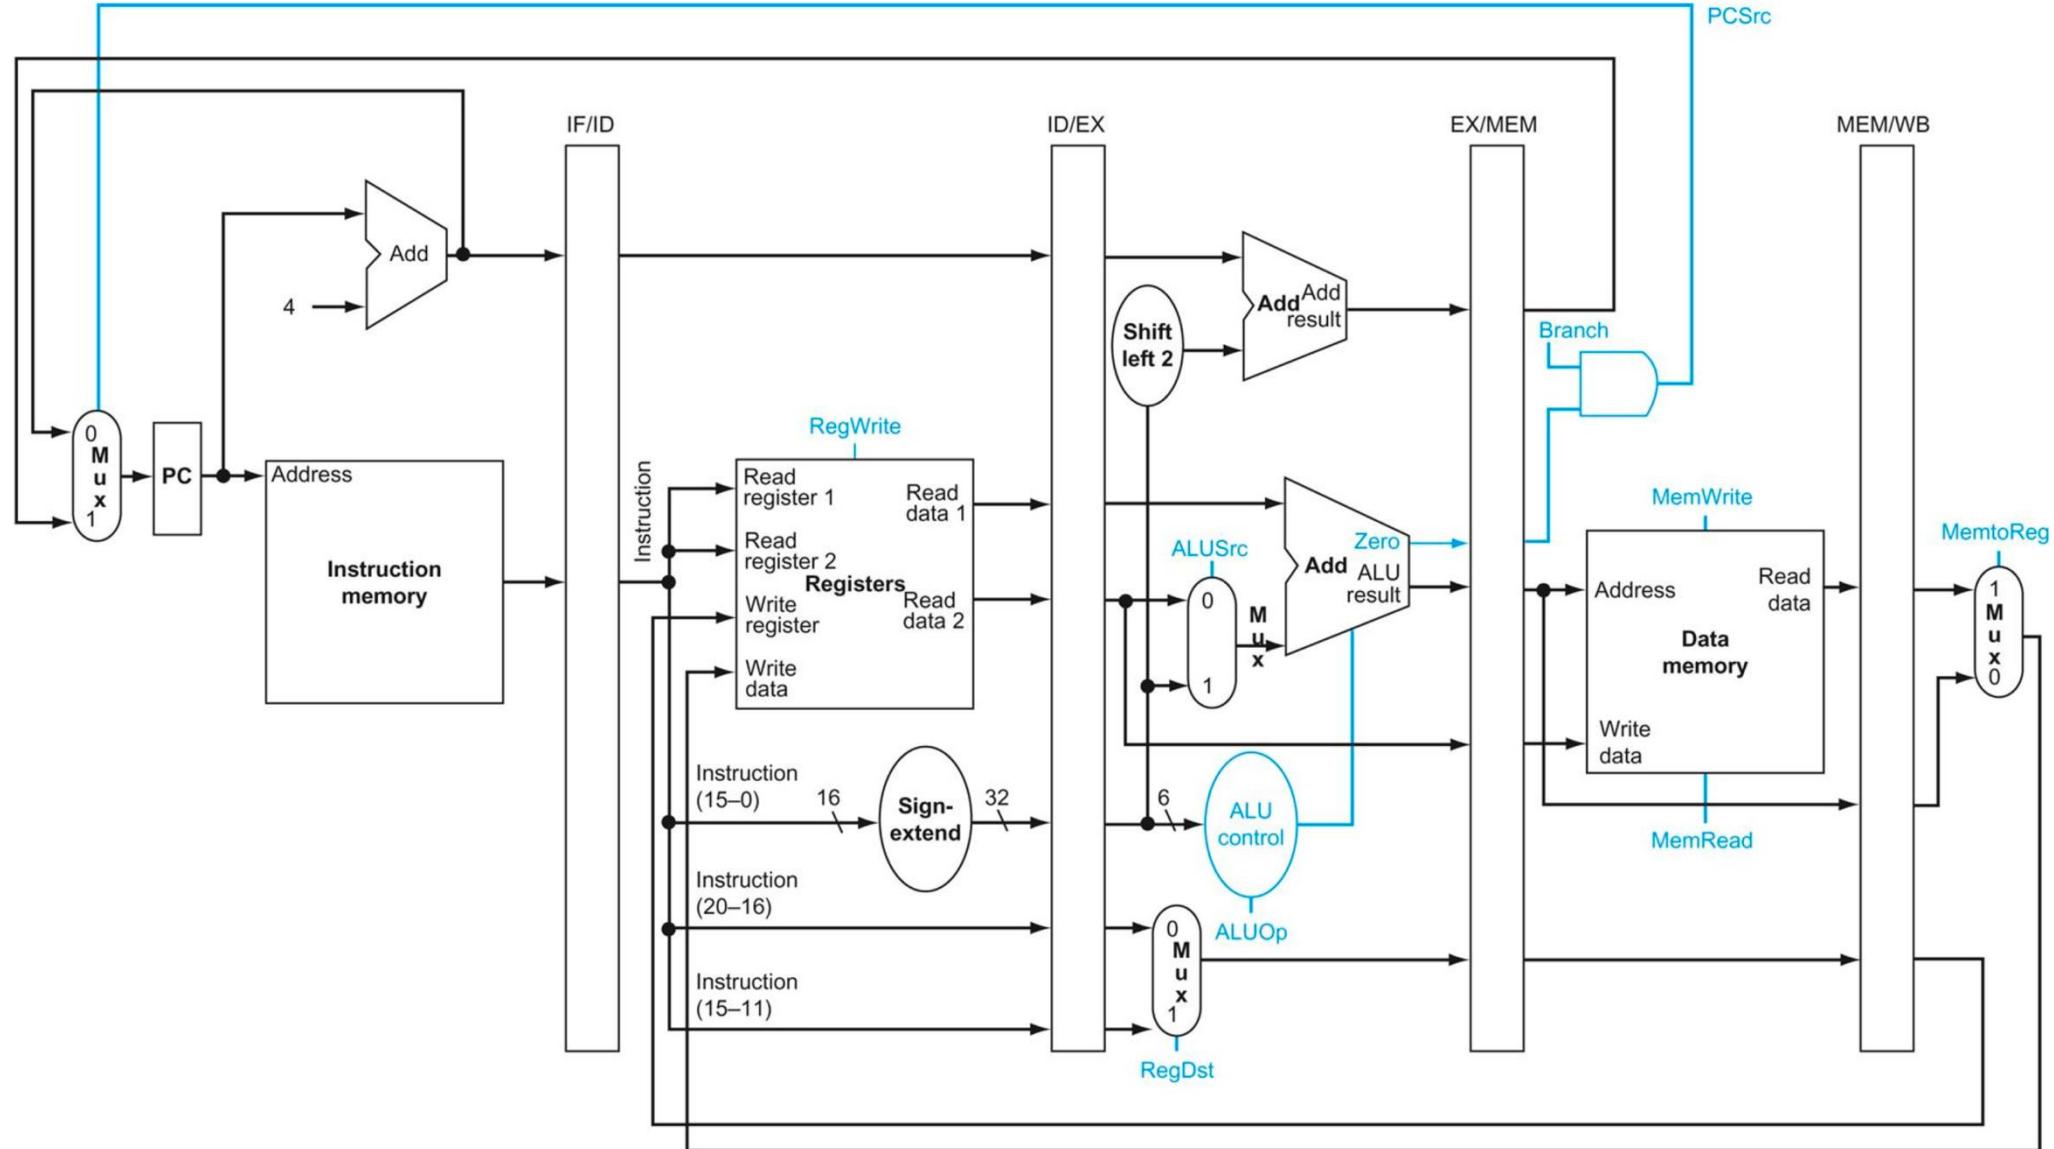
48



# Pipelined Control

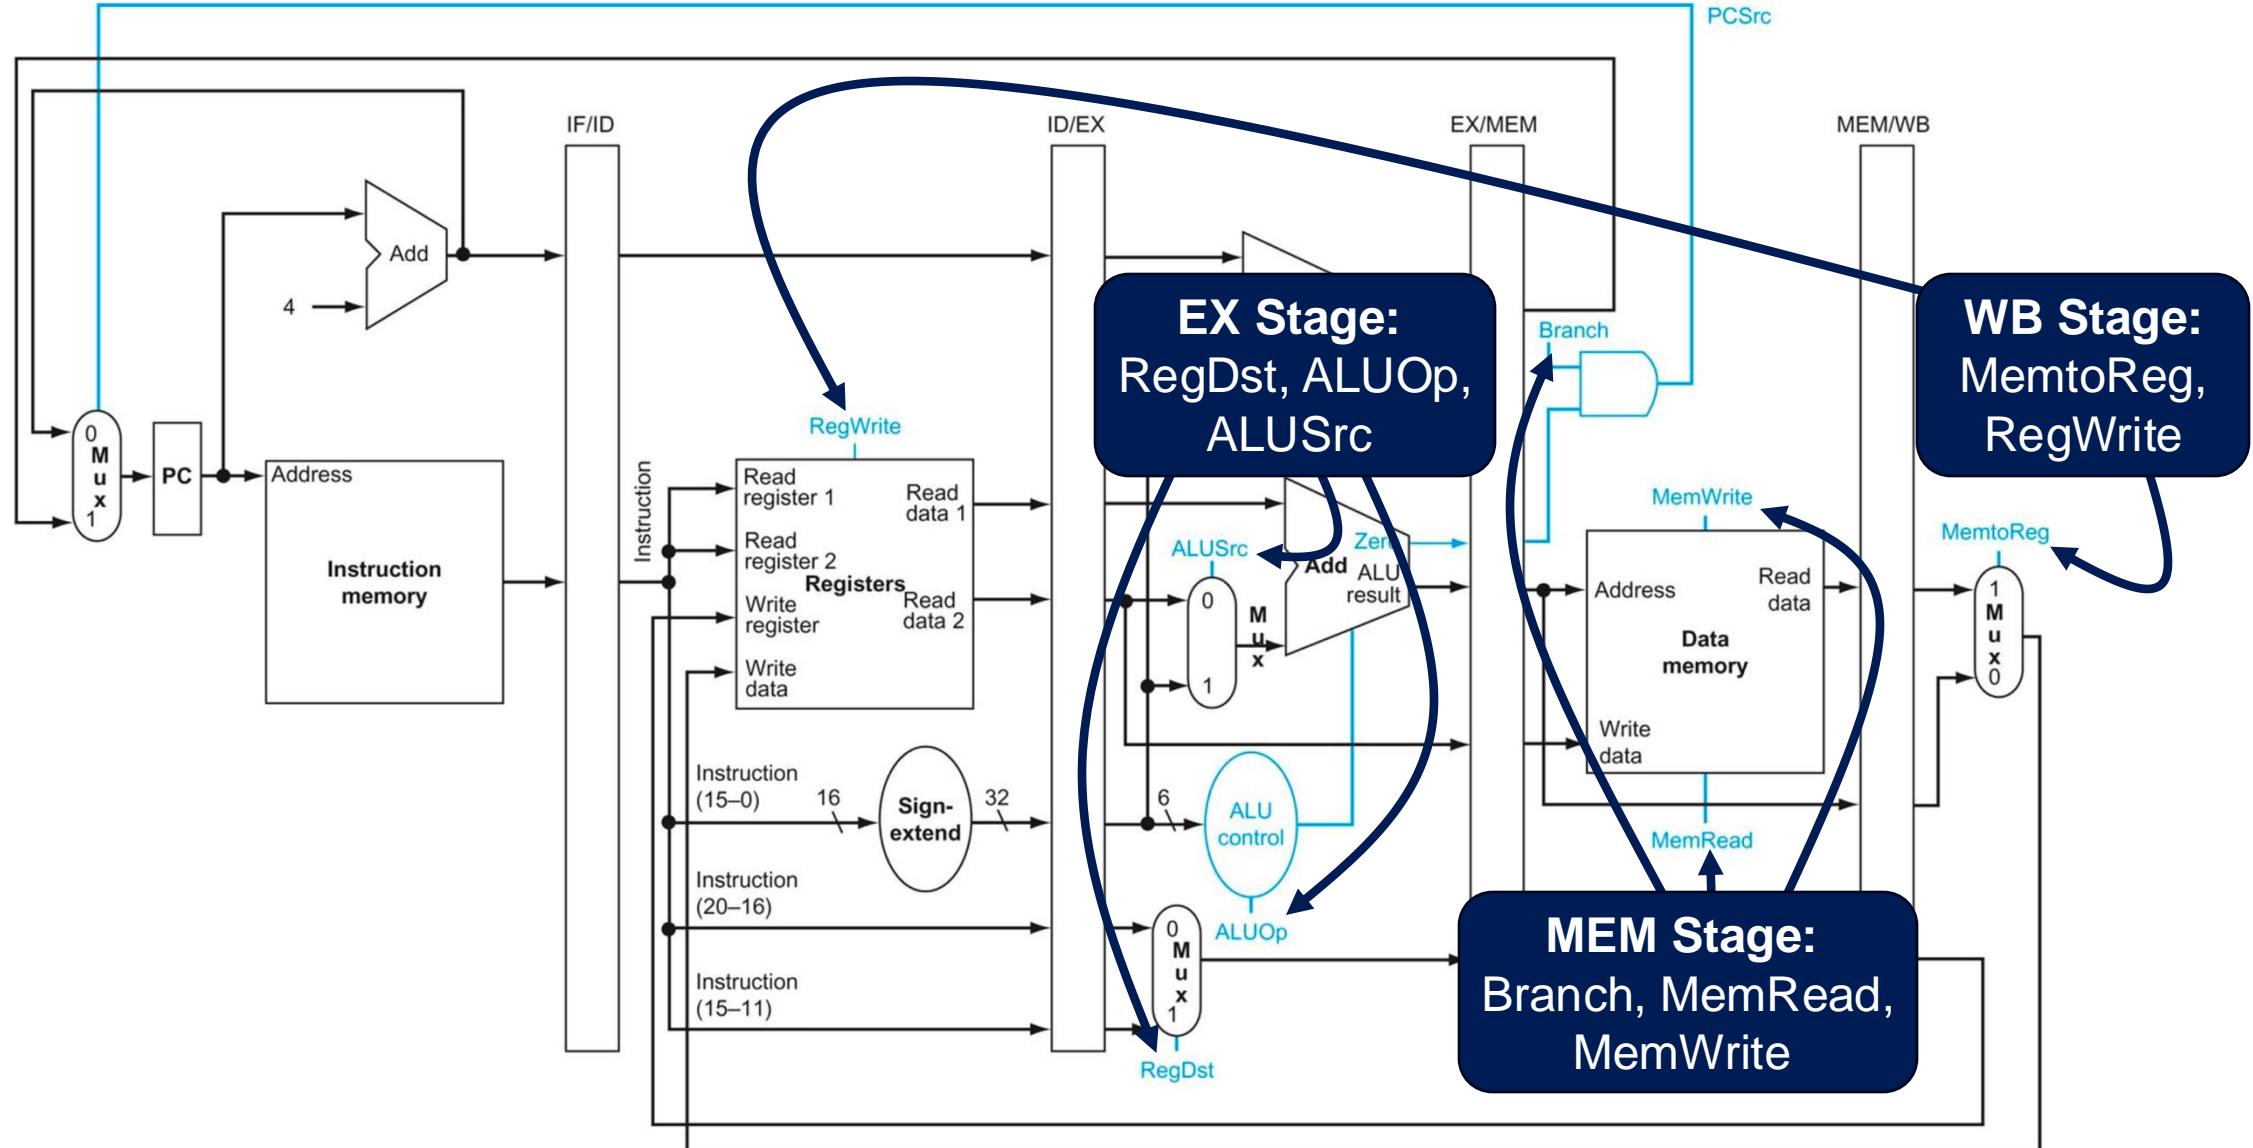
# Pipeline Control Signals

50



# Pipeline Control Signals

51



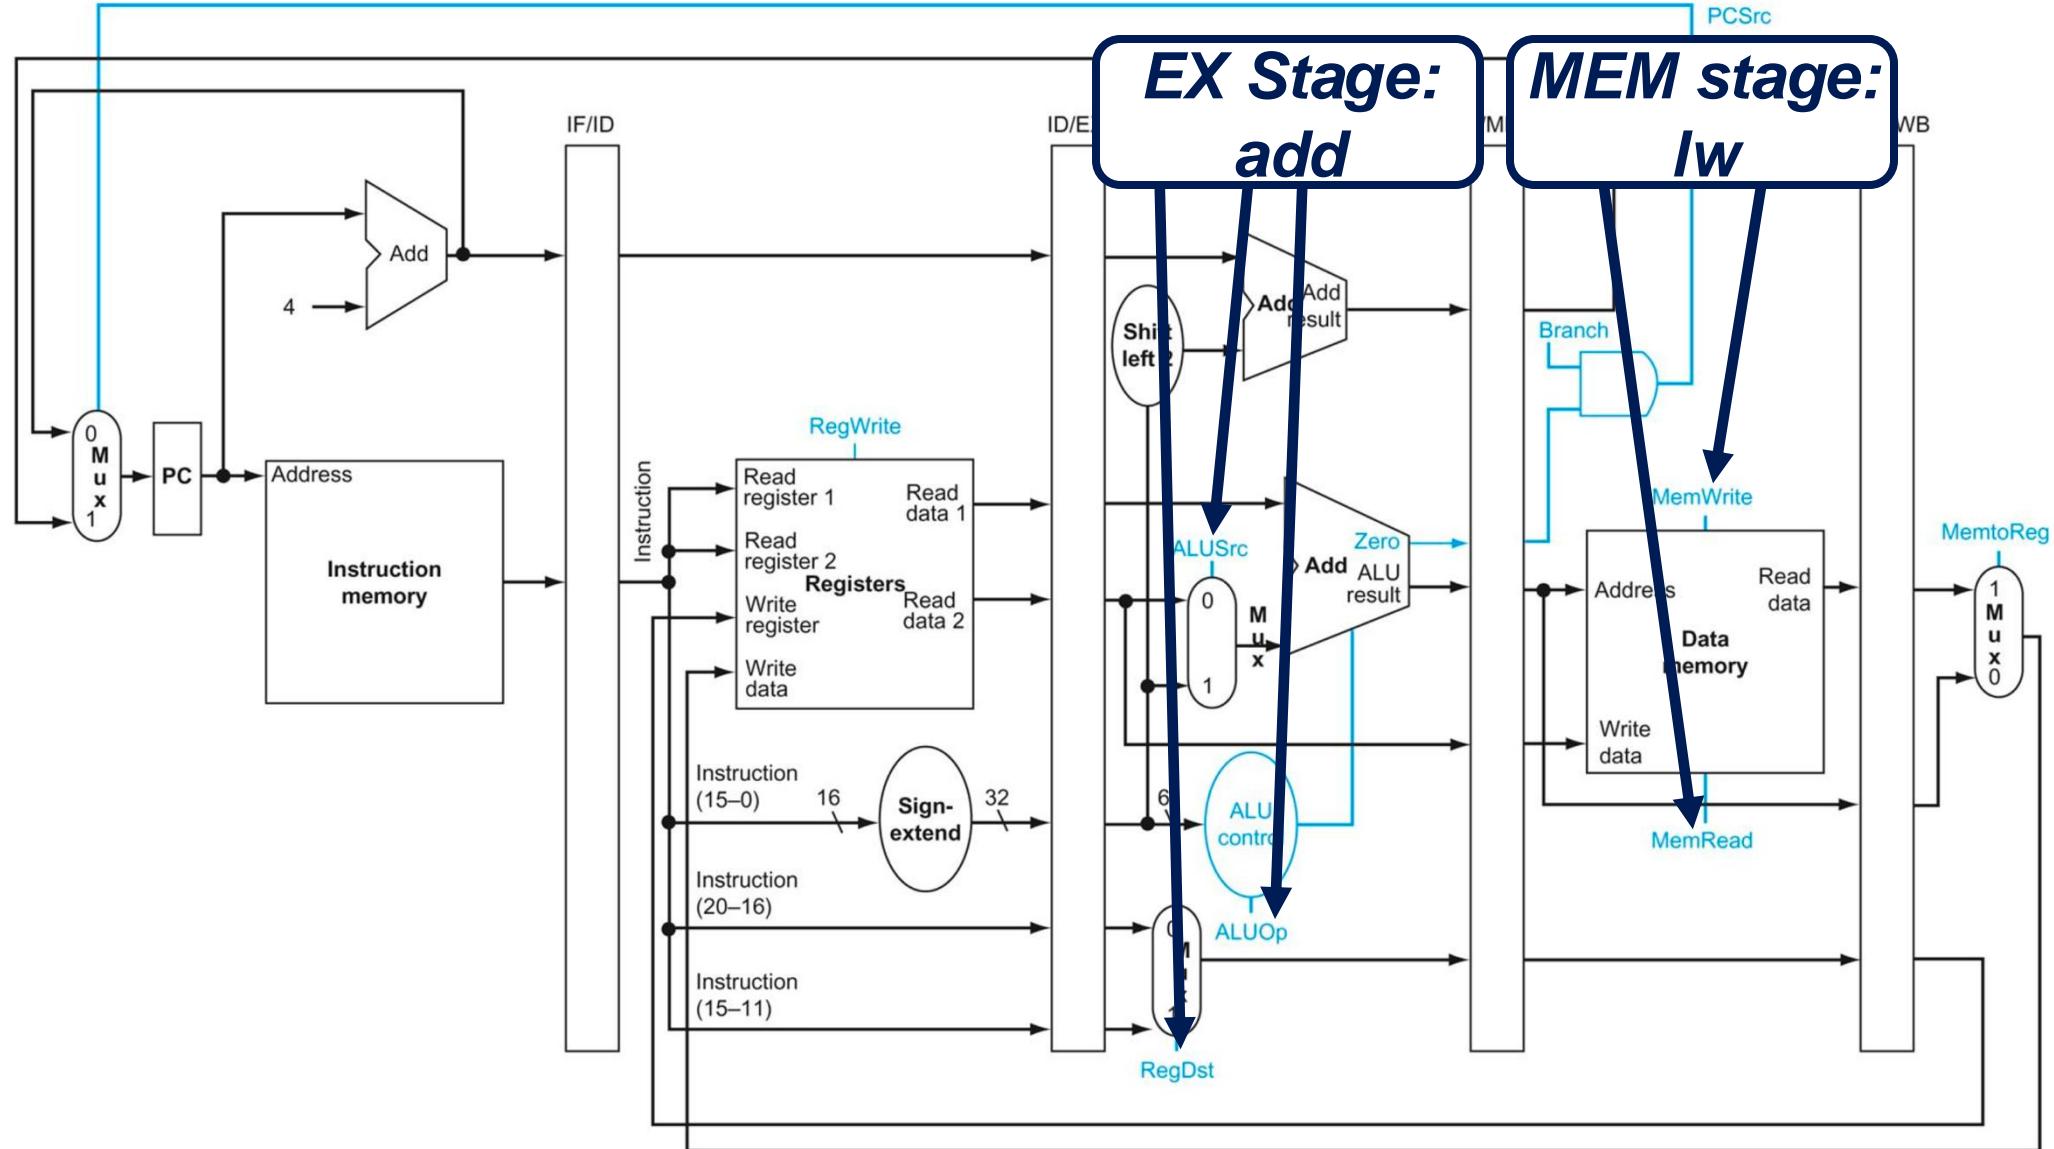
# Pipeline Control Signals

52

Signal name	Effect when deasserted (0)	Effect when asserted (1)
RegDst	The register destination number for the Write register comes from the rt field (bits 20:16).	The register destination number for the Write register comes from the rd field (bits 15:11).
RegWrite	None.	The register on the Write register input is written with the value on the Write data input.
ALUSrc	The second ALU operand comes from the second register file output (Read data 2).	The second ALU operand is the sign-extended, lower 16 bits of the instruction.
PCSrc	The PC is replaced by the output of the adder that computes the value of PC + 4.	The PC is replaced by the output of the adder that computes the branch target.
MemRead	None.	Data memory contents designated by the address input are put on the Read data output.
MemWrite	None.	Data memory contents designated by the address input are replaced by the value on the Write data input.
MemtoReg	The value fed to the register Write data input comes from the ALU.	The value fed to the register Write data input comes from the data memory.

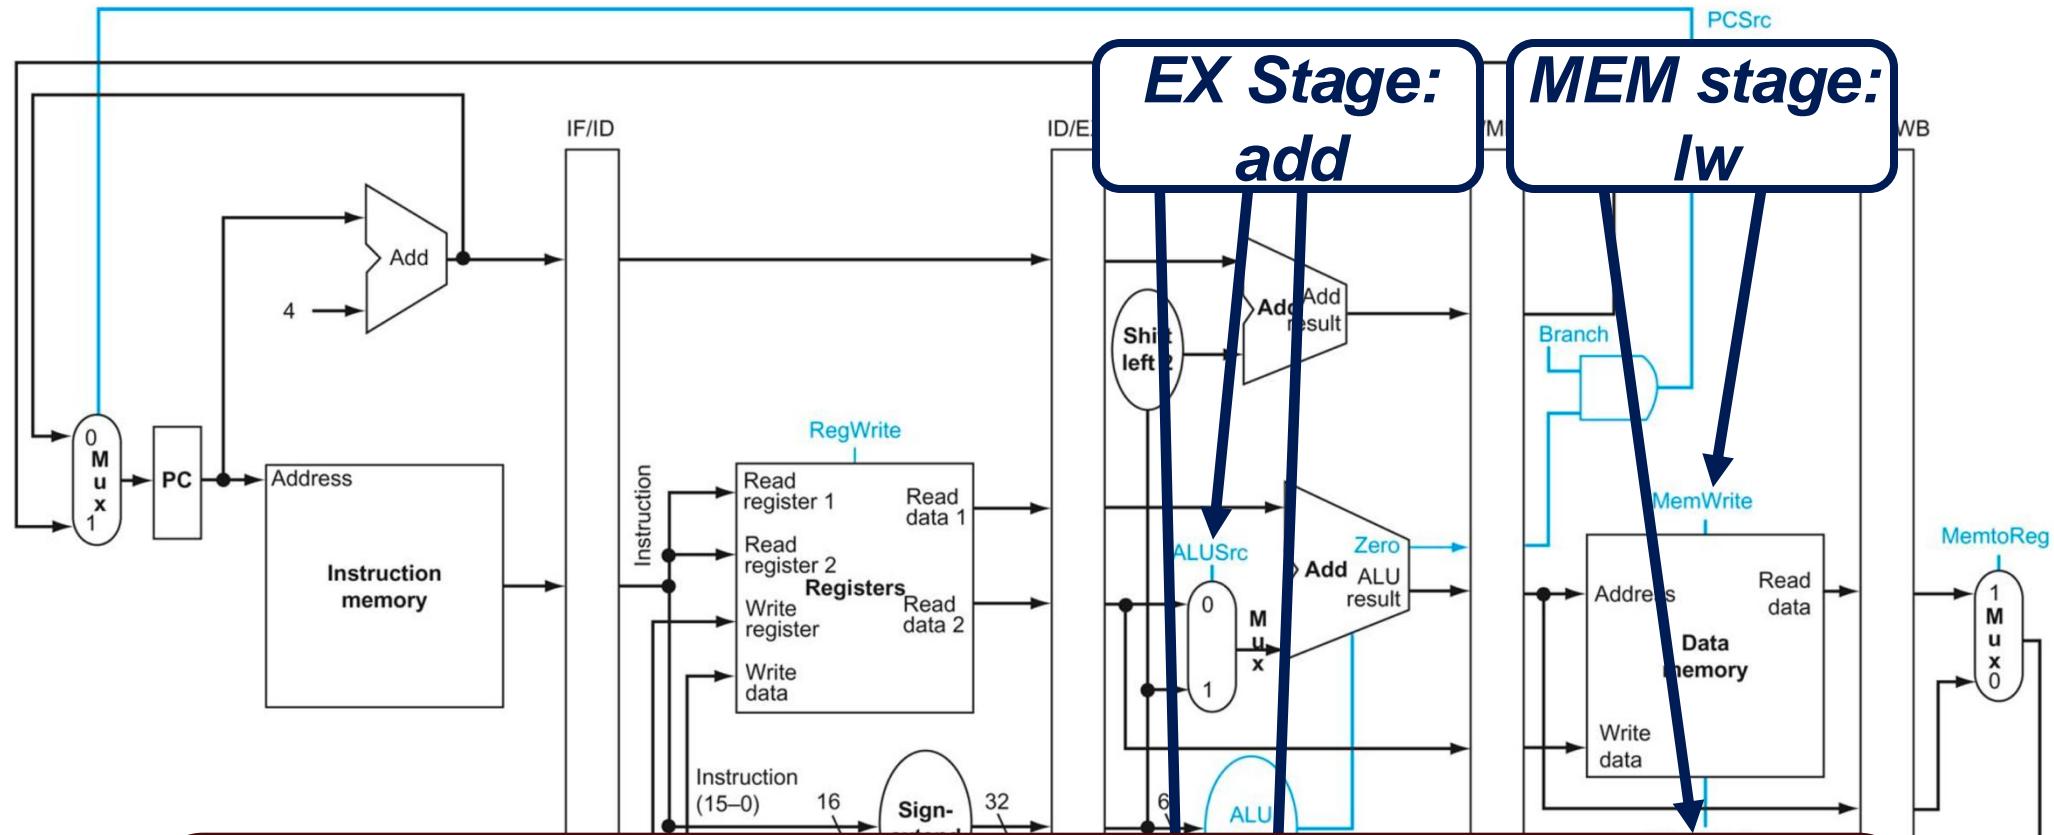
# Motivation: Pipelined Control

53



# Motivation: Pipelined Control

54



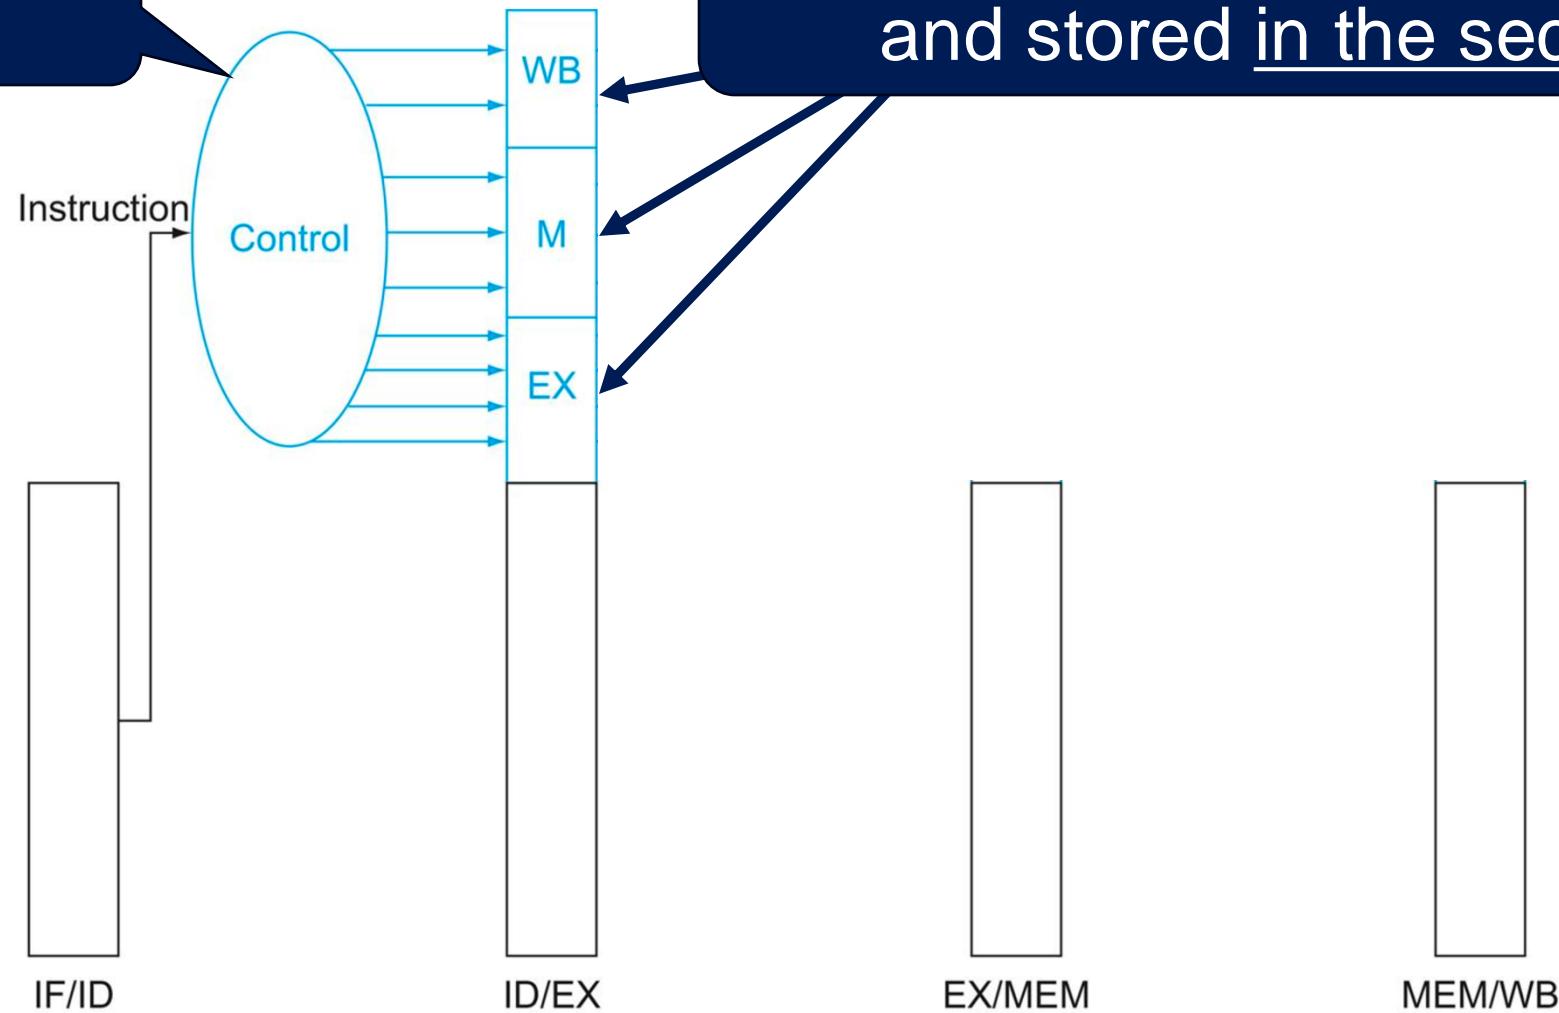
- Instructions are executed simultaneously at each stage
  - The required control signals differ for each stage
- So, How do we generate and deliver these signals?

# Idea: Pipelined Control

55

Combinational circuit

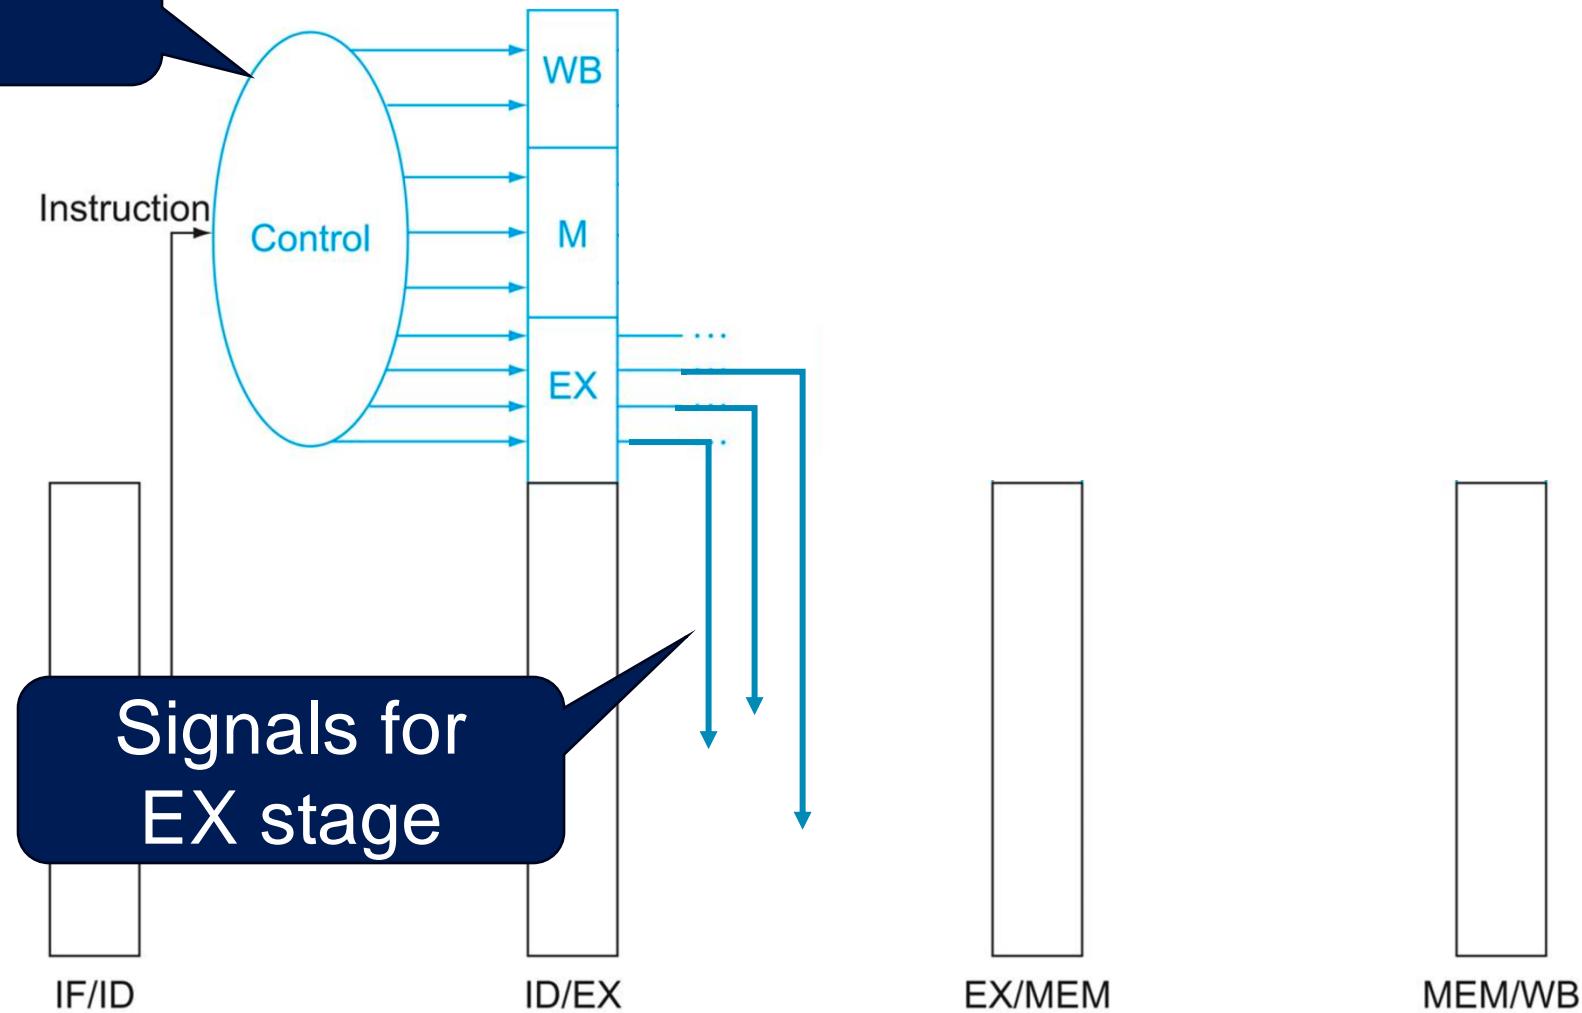
The signals for all stages are generated and stored in the second stage



# Idea: Pipelined Control

56

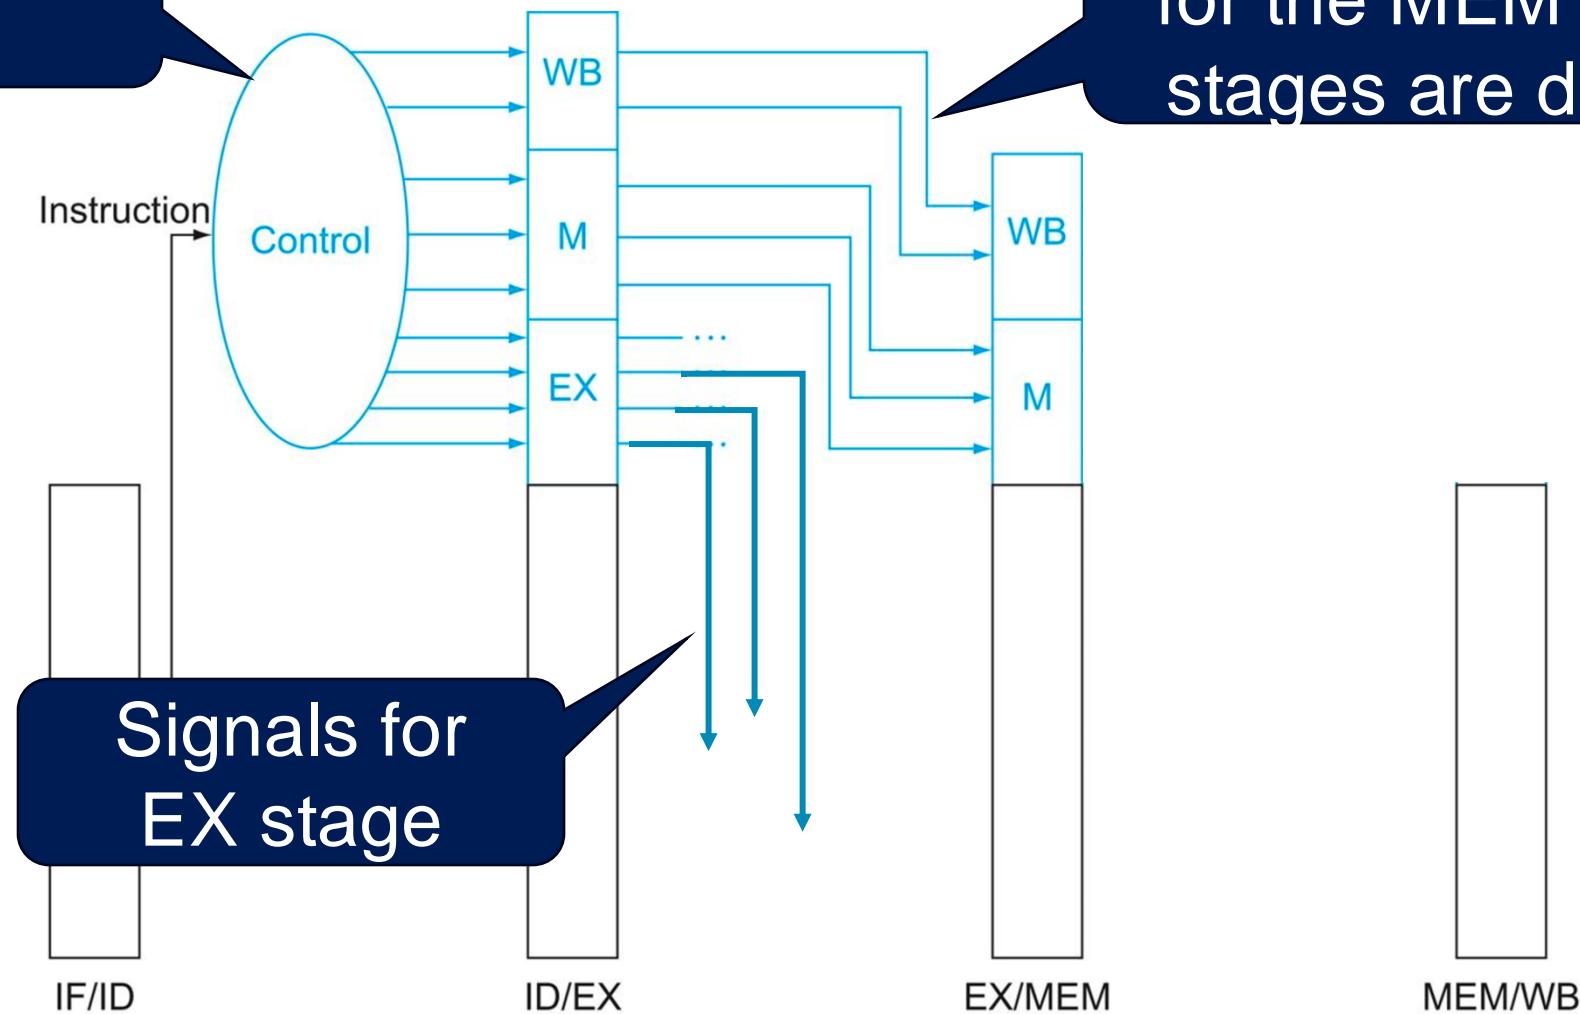
Combinational  
circuit



# Idea: Pipelined Control

57

Combinational circuit



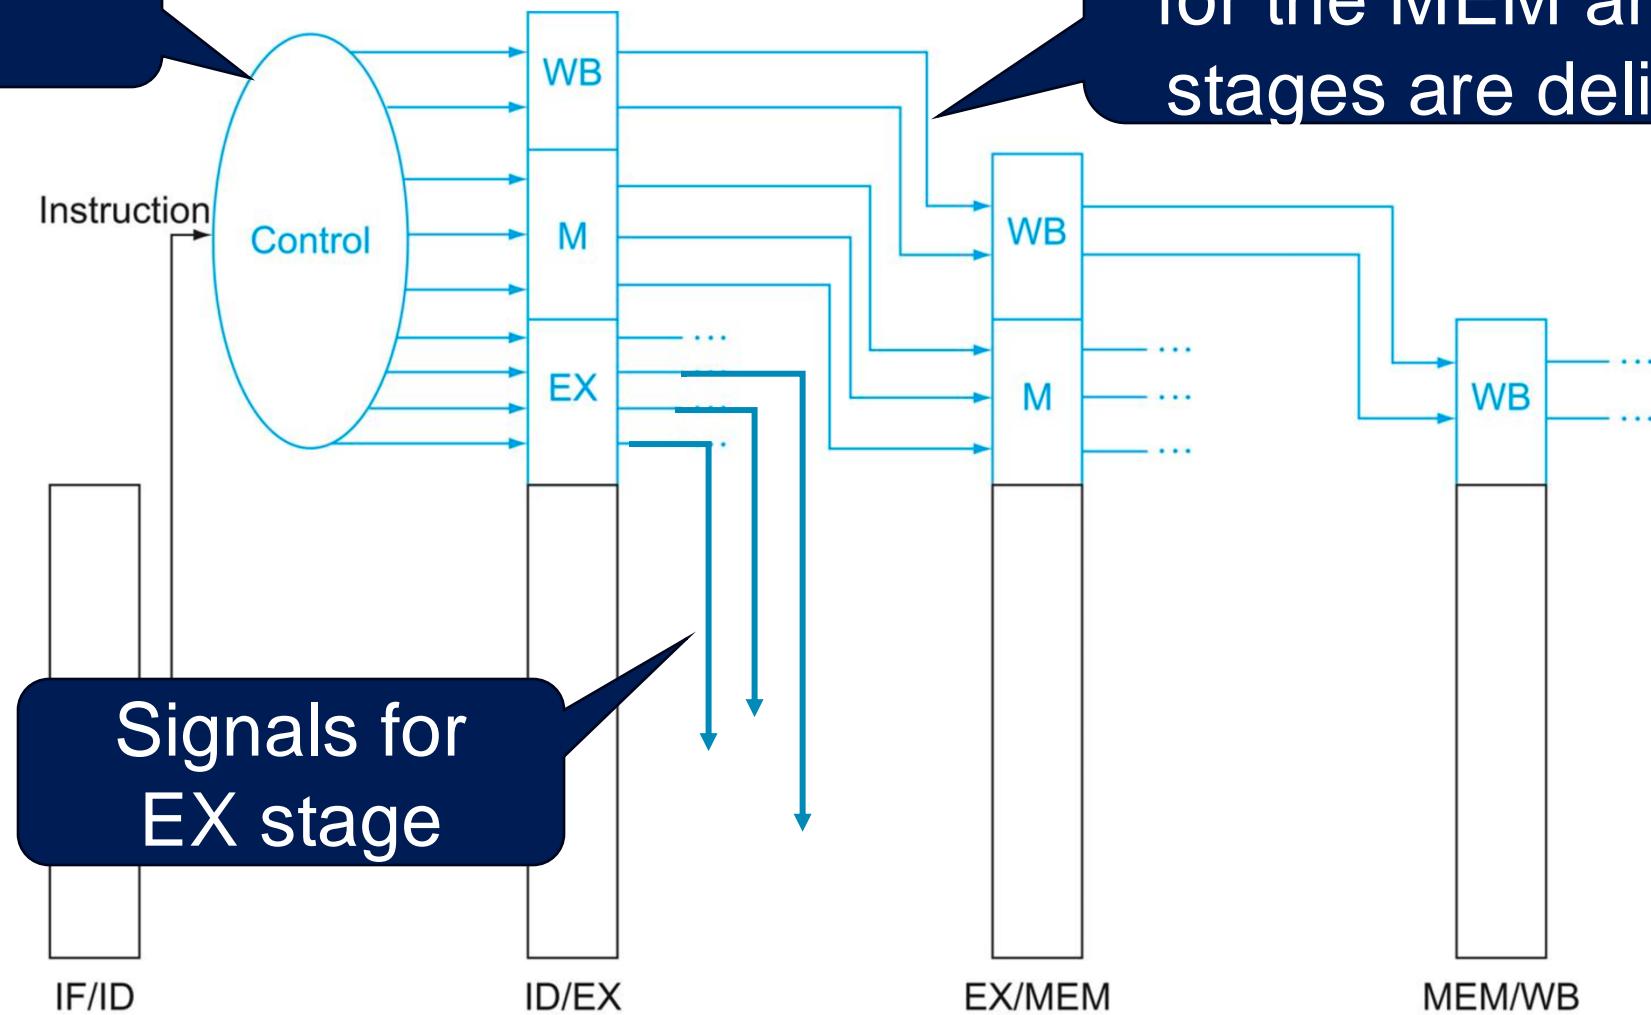
The signals required for the MEM and WB stages are delivered

Signals for EX stage

# Idea: Pipelined Control

58

Combinational circuit



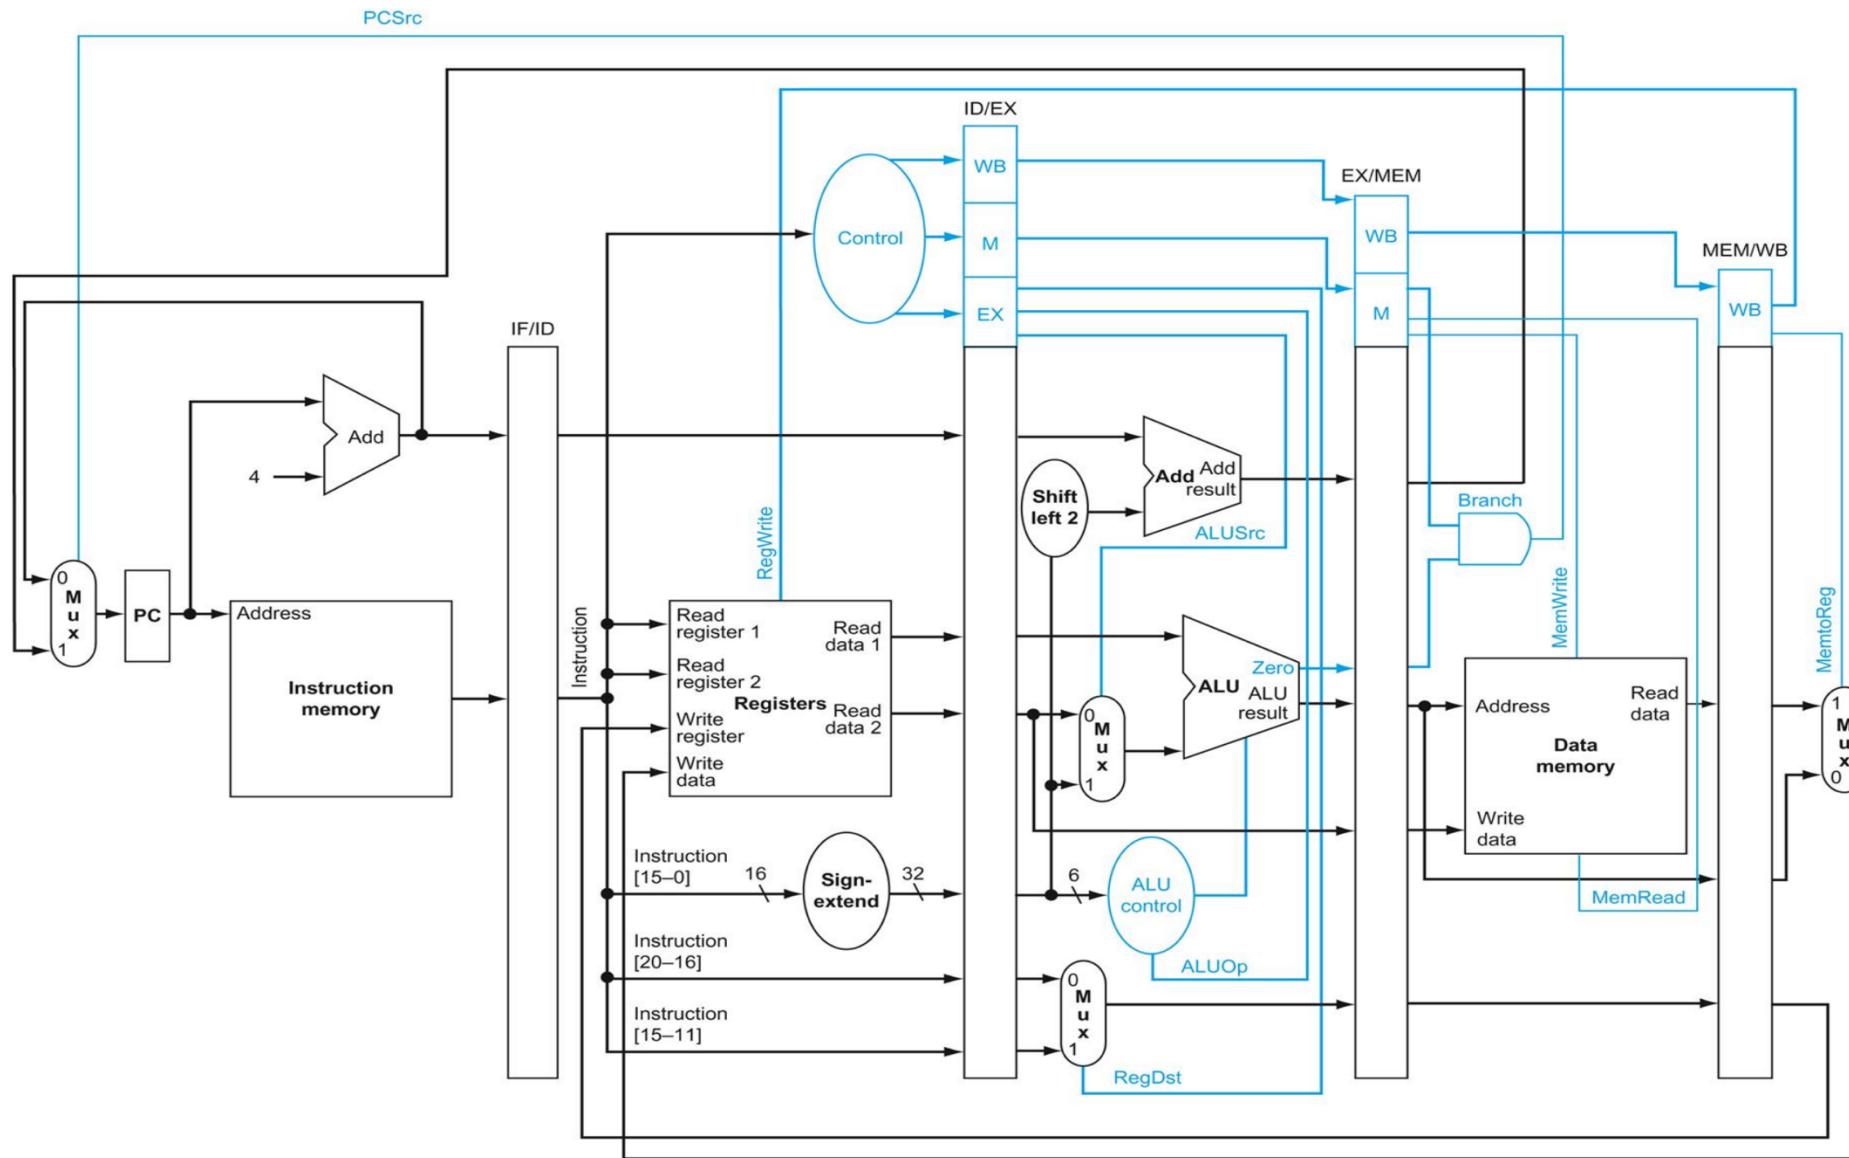
# Pipelined Control Signals

59

Instruction	Execution/address calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	RegDst	ALUOp1	ALUOp0	ALUSrc	Branch	Mem-Read	Mem-Write	Reg-Write	Memto-Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

# The Pipelined Datapath with the Control

60



# Pipelining Example

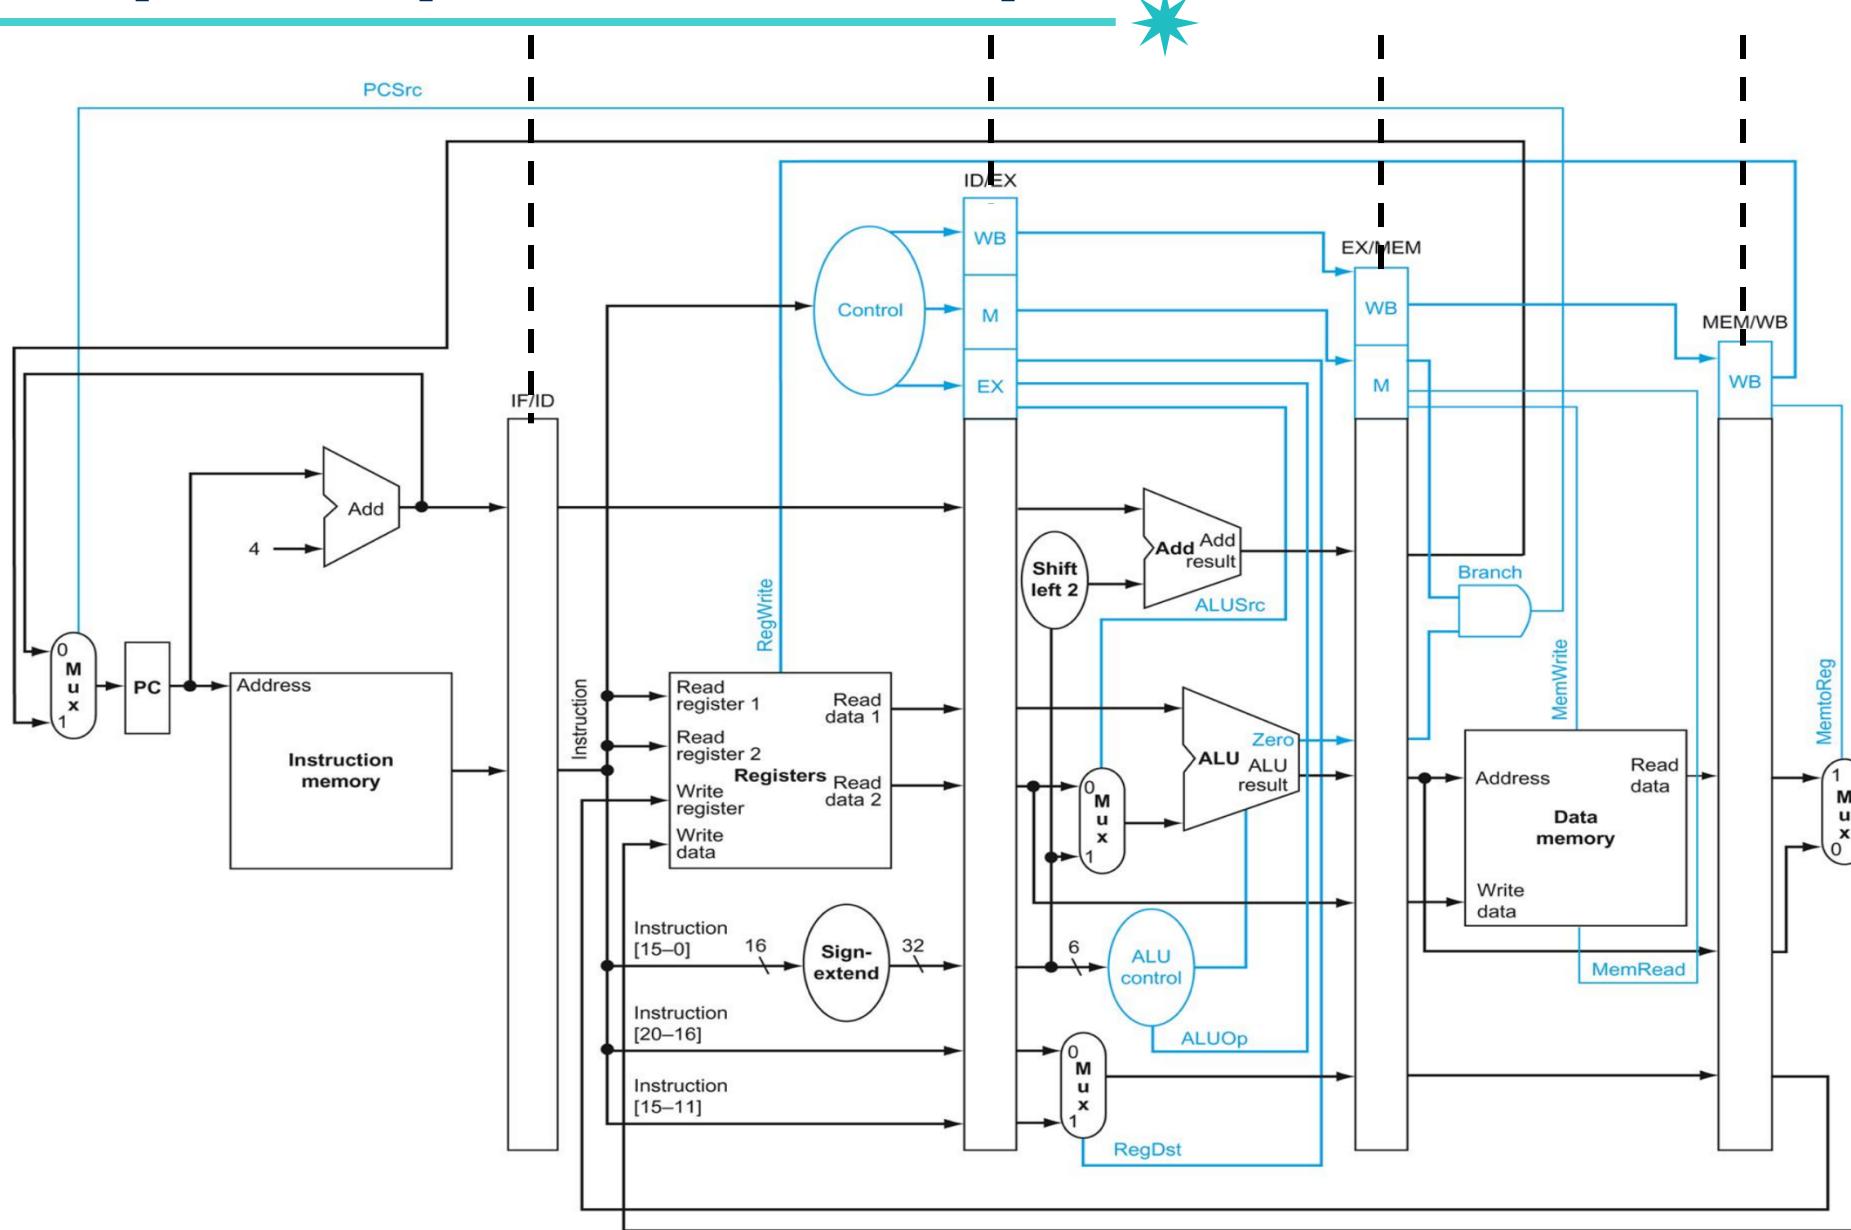


```
lw    $10, 20($1)
sub   $11, $2, $3
add   $12, $4, $5
or    $13, $6, $7
add   $14, $8, $9
```

***Assumption: there is no hazard***

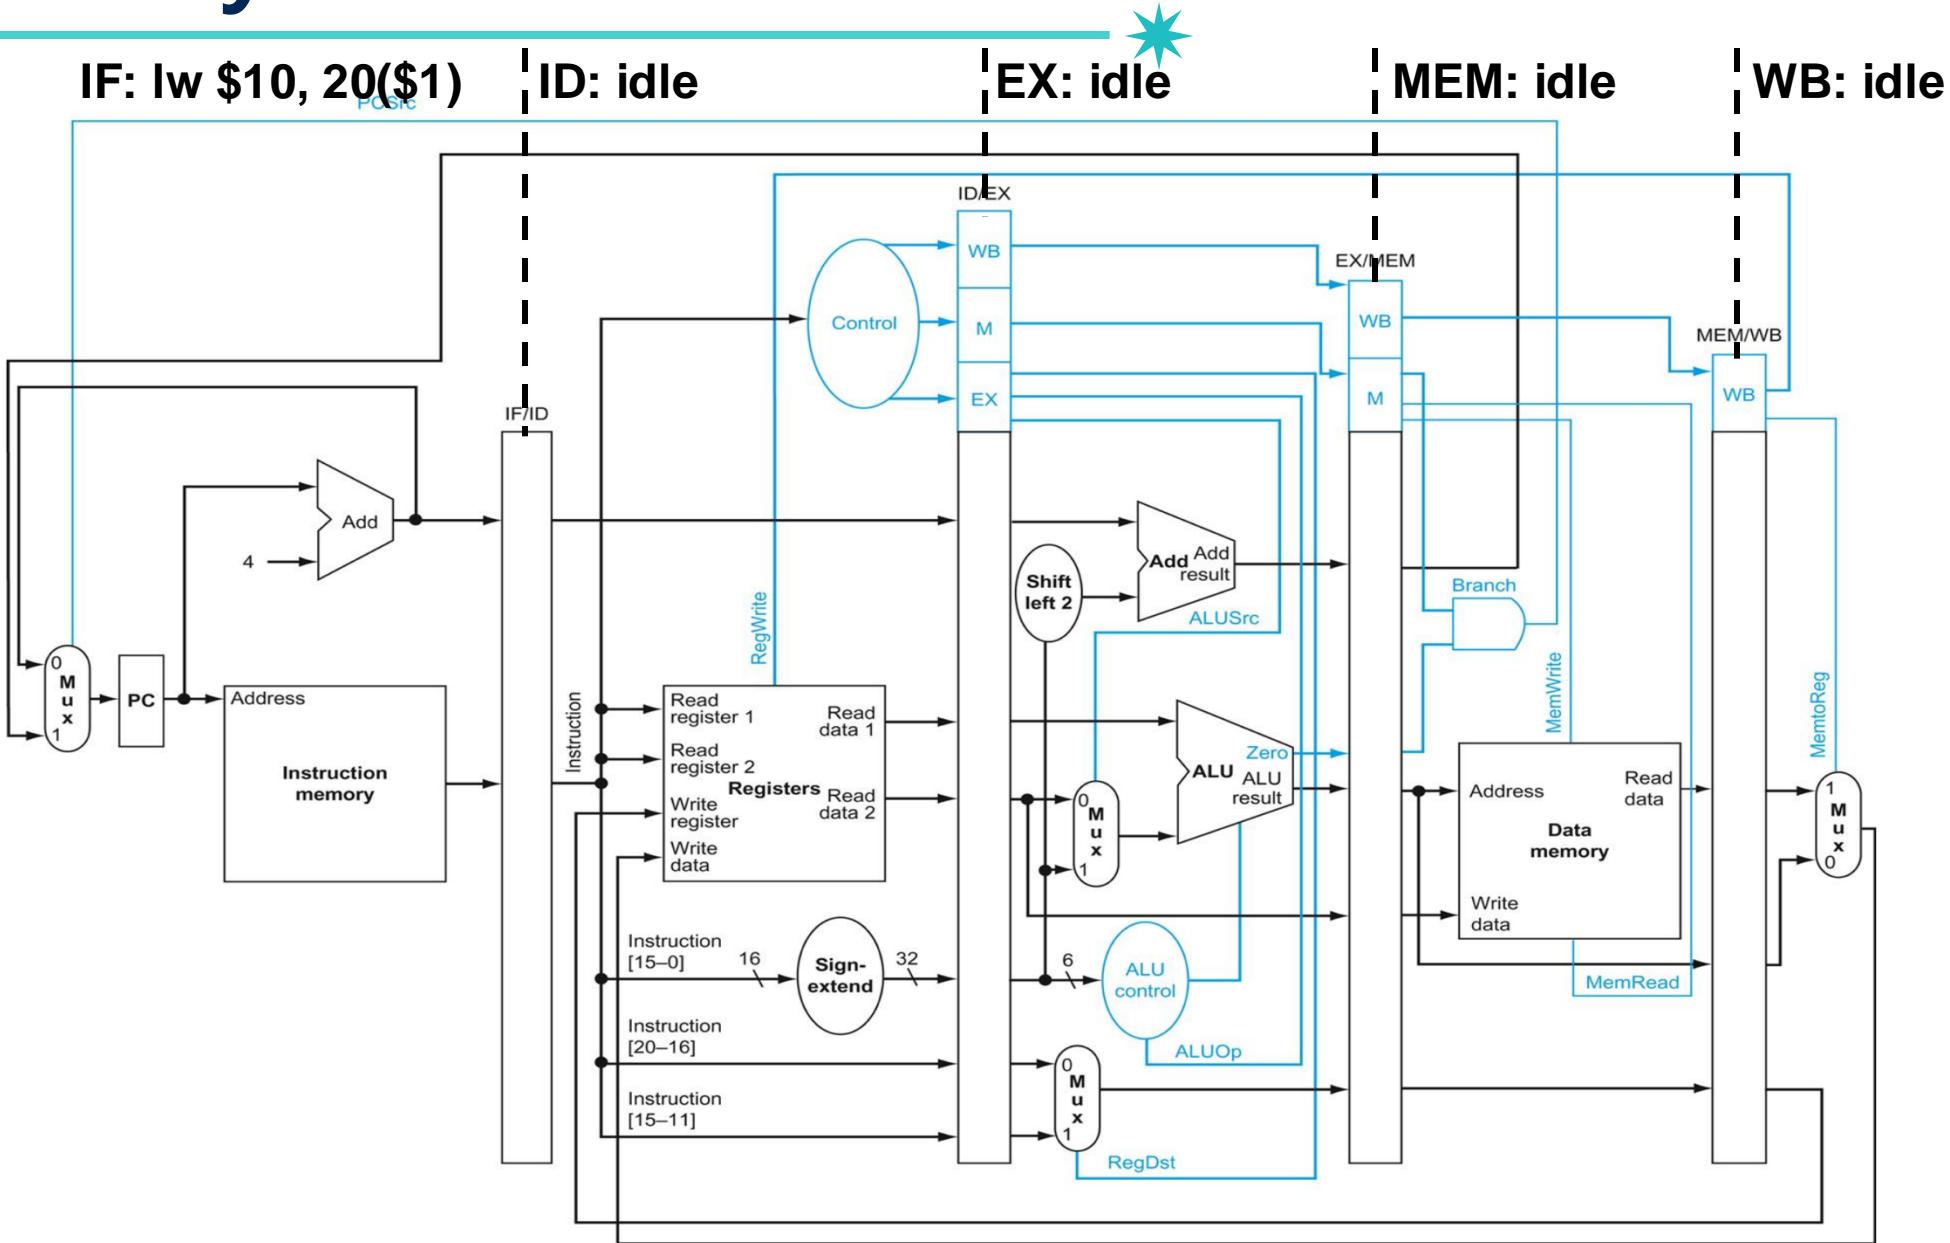
Later, we will cover about hazards!

# Example: Pipelined Datapath with the Control



# Clock Cycle #1

63



# Clock Cycle #2

64

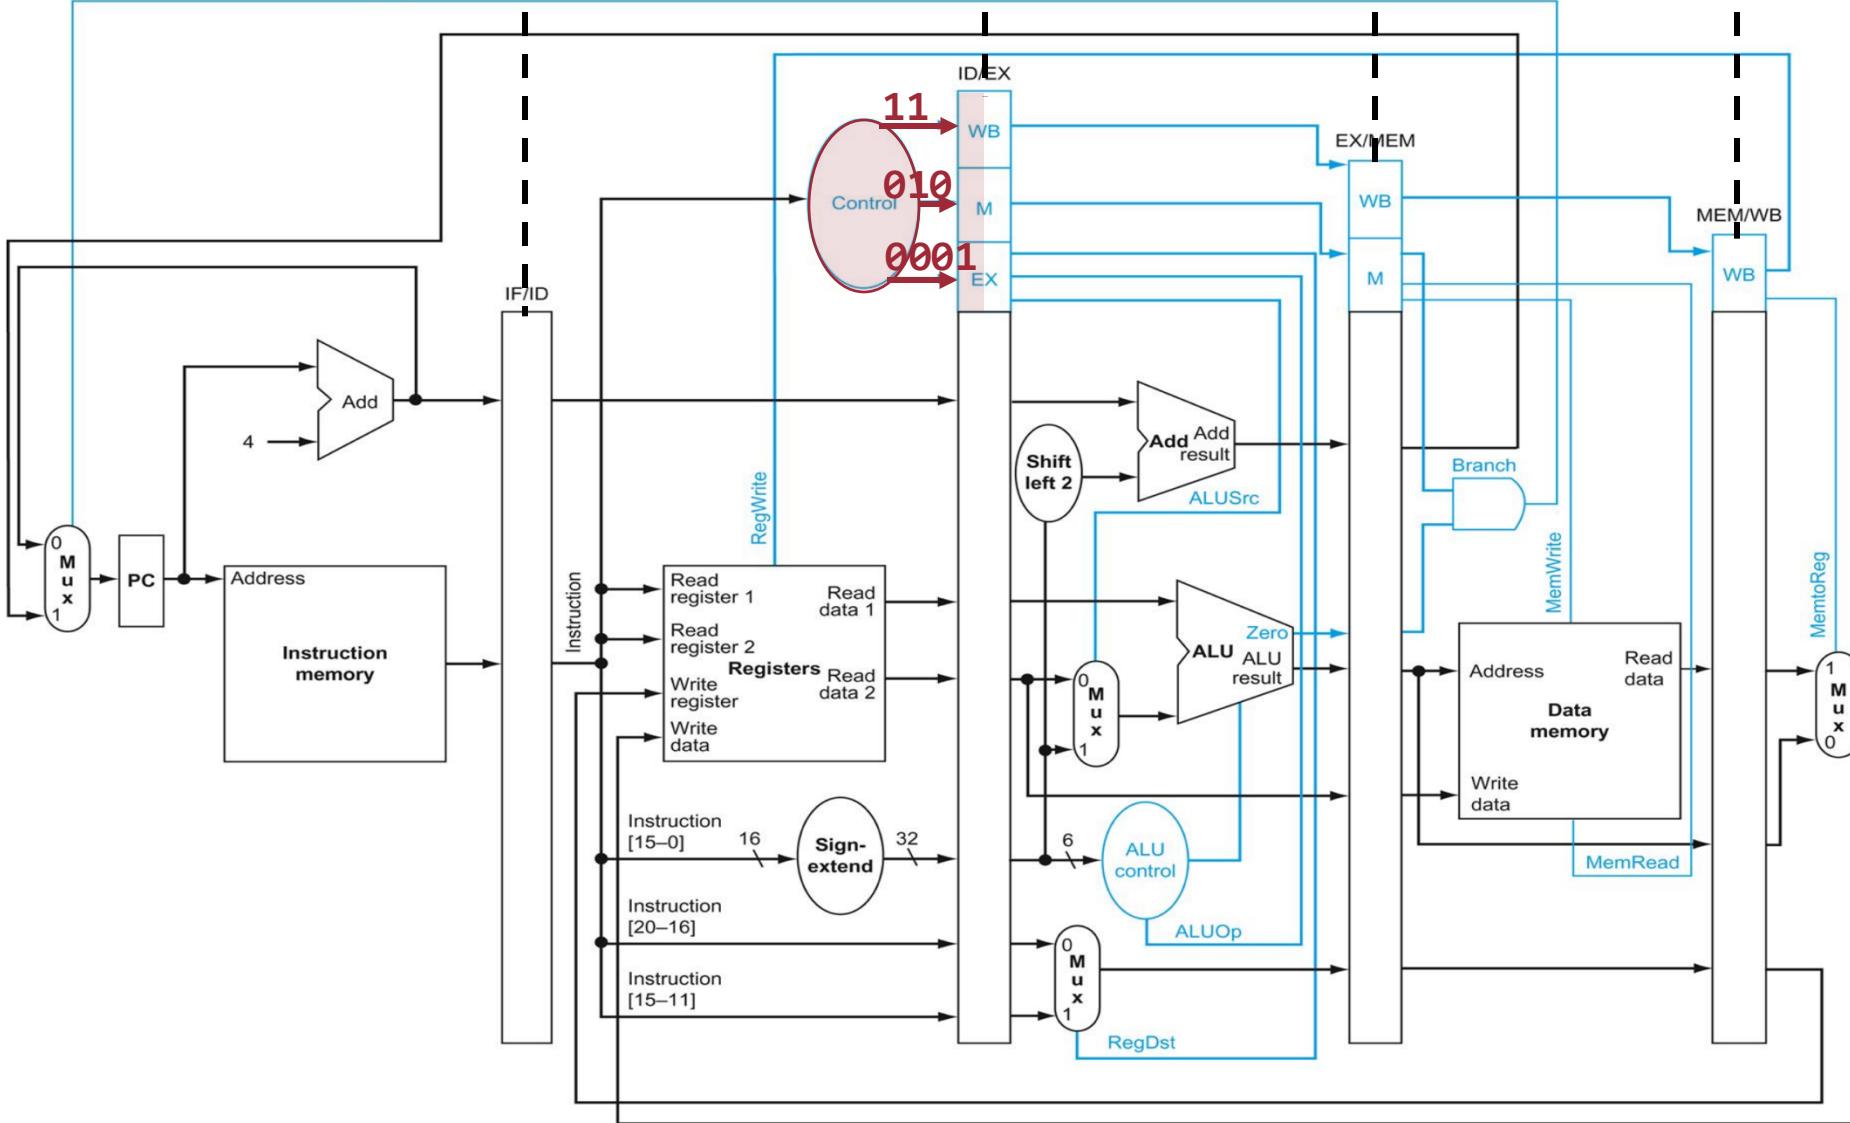
IF: sub \$11, \$2, \$3

ID: lw \$10, 20(\$1)

EX: idle

MEM: idle

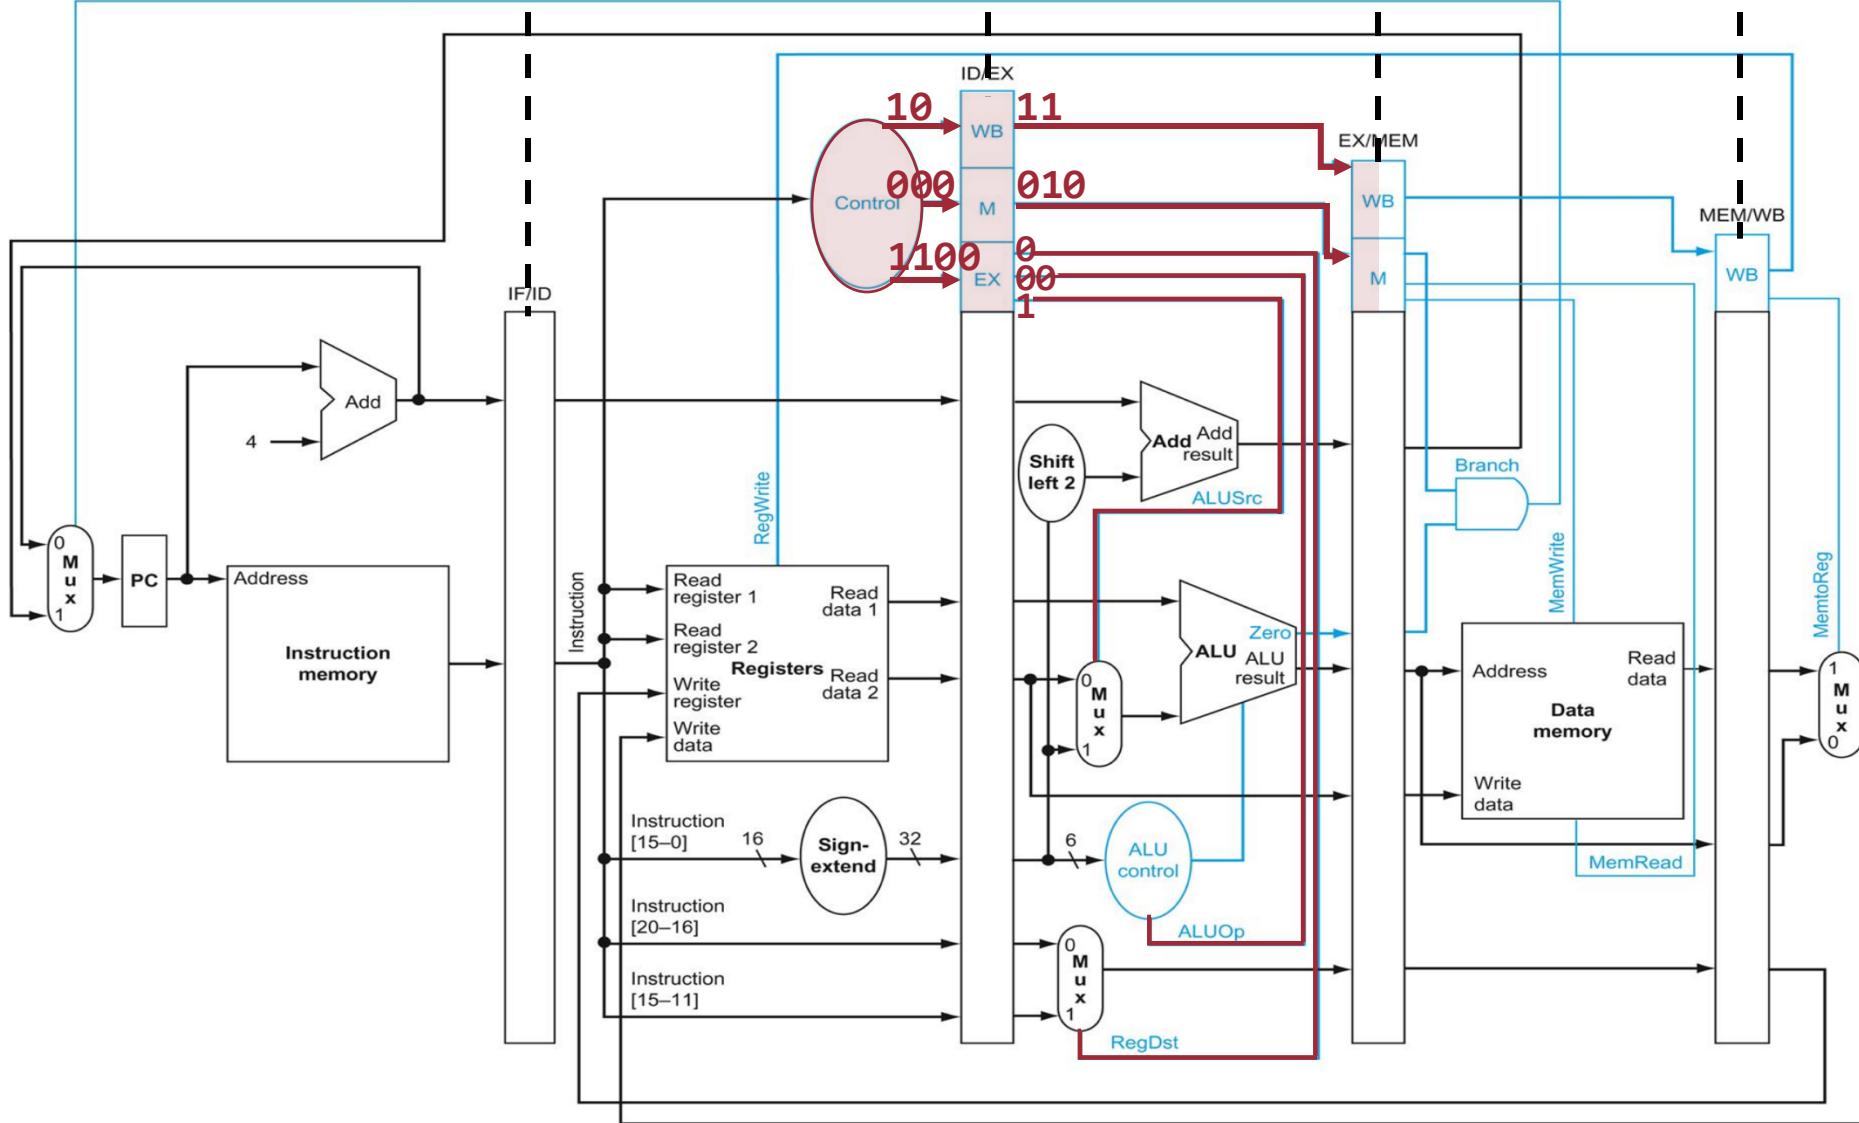
WB: idle



# Clock Cycle #3

65

IF: and \$12, \$4, \$5  
 PC<sub>src</sub>  
 ID: sub \$11, \$2, \$3  
 EX: lw \$10, 20(\$1)  
 MEM: idle  
 WB: idle



# Clock Cycle #4

66

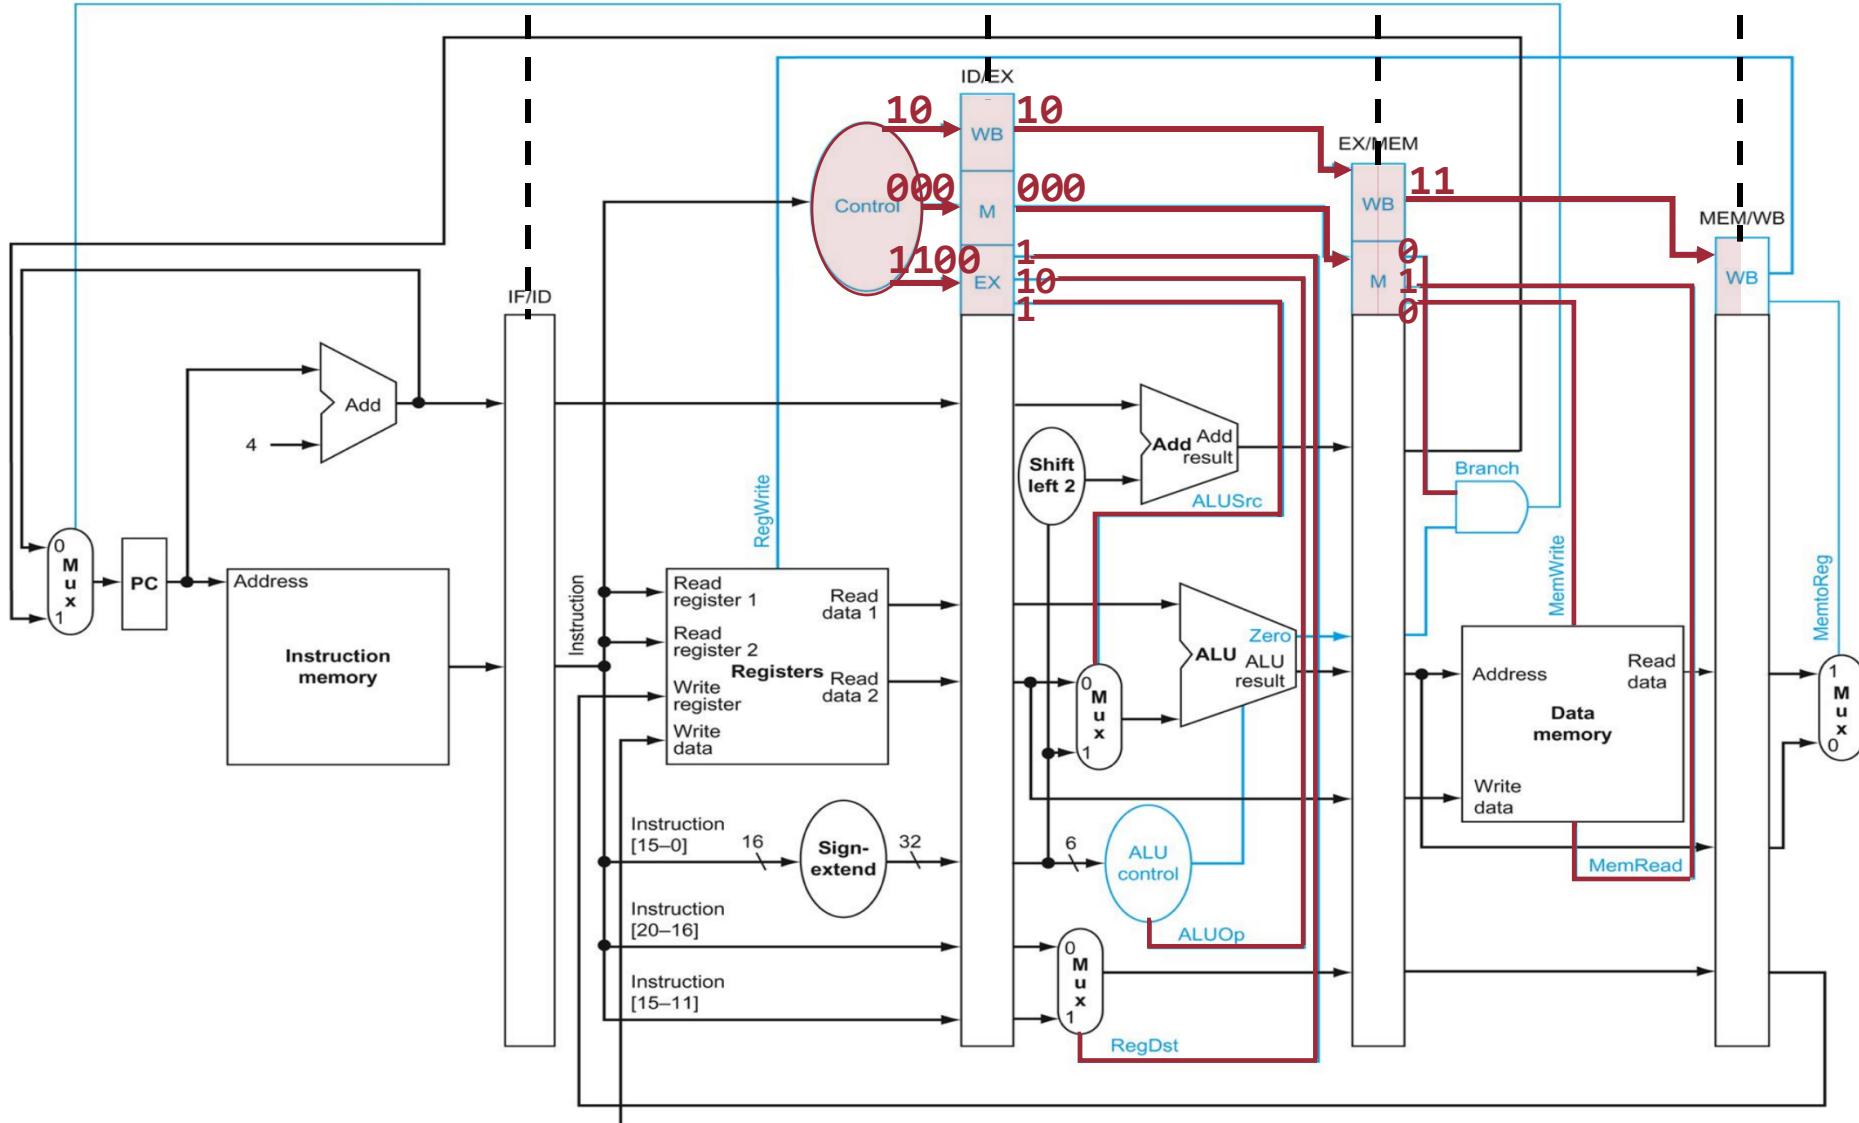
IF: or \$13, \$6, \$7  
PCSrc

ID: and \$12, \$4, \$5

EX: sub\$11, \$2, \$3

MEM: lw\$10,20(\$1)

WB: idle

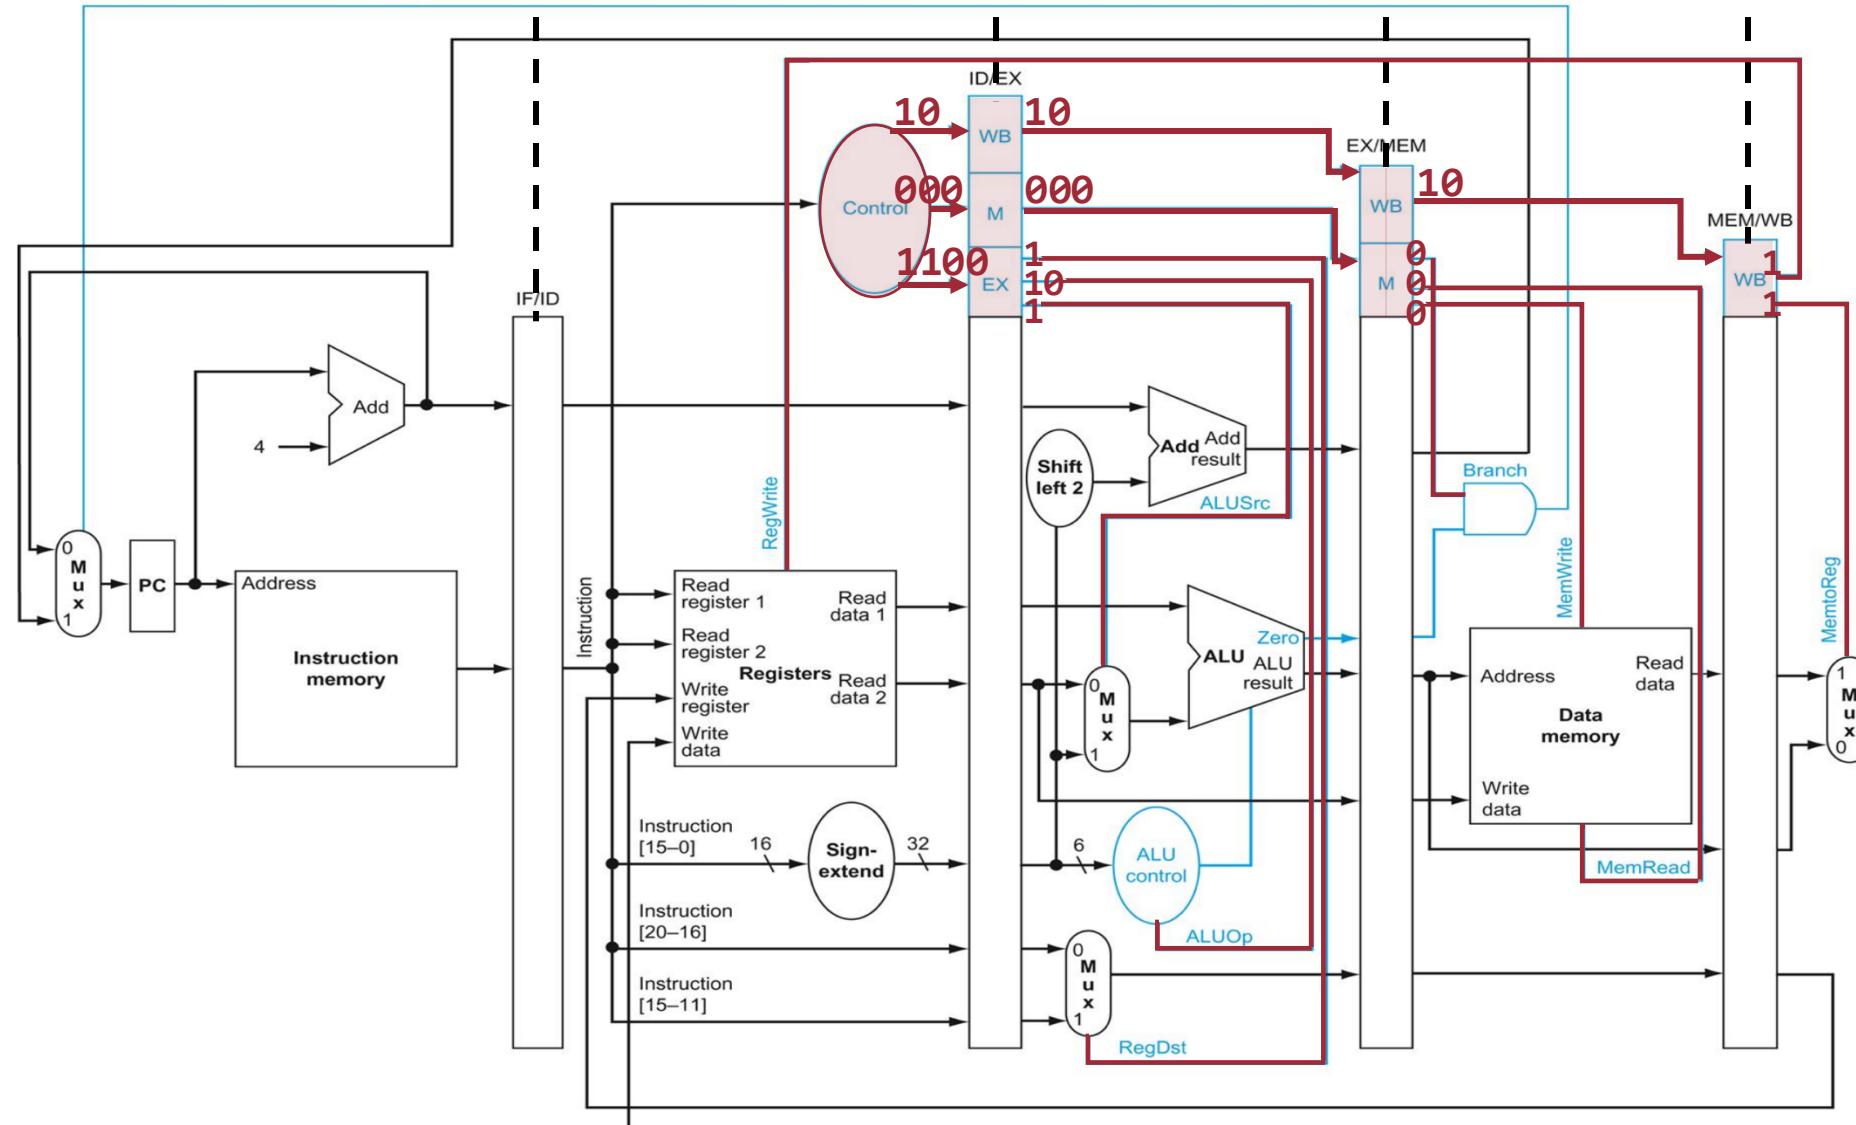


# Clock Cycle #5

IF: add \$14, \$8, \$9  
PCsrc

ID: or \$13, \$6, \$7

EX: and \$12, \$4, \$5  
MEM: sub \$11, \$2, \$3  
WB: lw \$10, 20(\$1)



# Question?