

Seongil Wi



HW2

2

Due date: 11/26, 11:59PM

Quiz 2



This is the last quiz of this semester

• Date: 11/26 (TUE.), Class time

- Scope:
 - Logic Design Basics
 - Processor (1) ~ Processor (6)
- Brint your own pen!
- T/F problems + 2~4 computation problems

Where are We?



Situations that prevent starting the next instruction in the next cycle

Hazard #1: Structural hazard

Conflict for use of a hardware resource

Solution:

- Stall
- Resource duplication

Hazard #2: Data hazard

An instruction cannot execute because data is not yet available

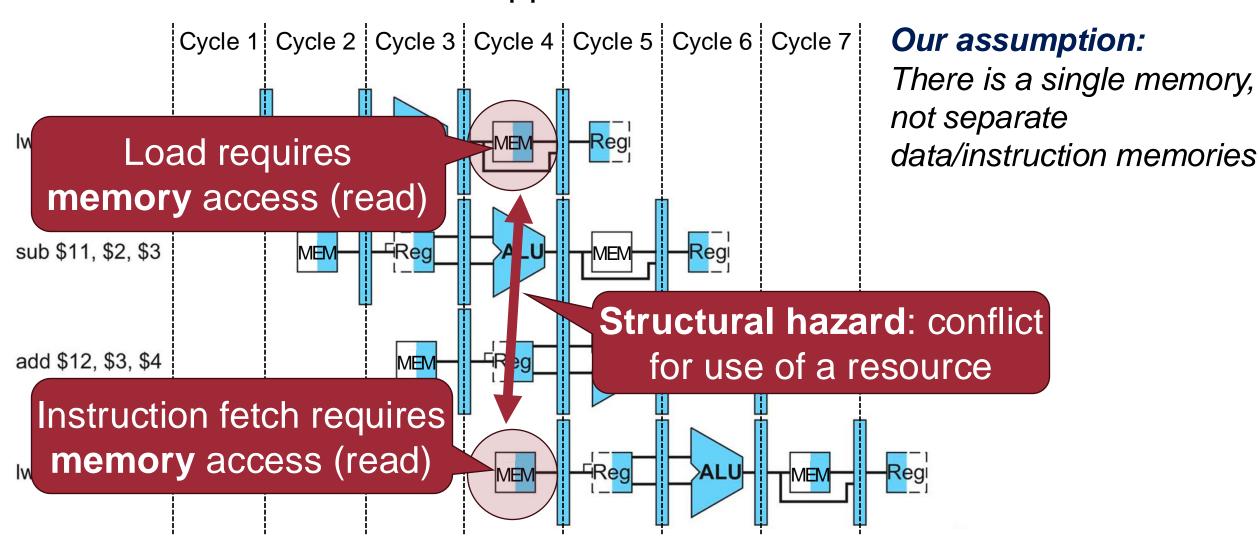
Solution:

- Stall
- Forwarding
- Compiler optimization

Hazard #3: Control hazard

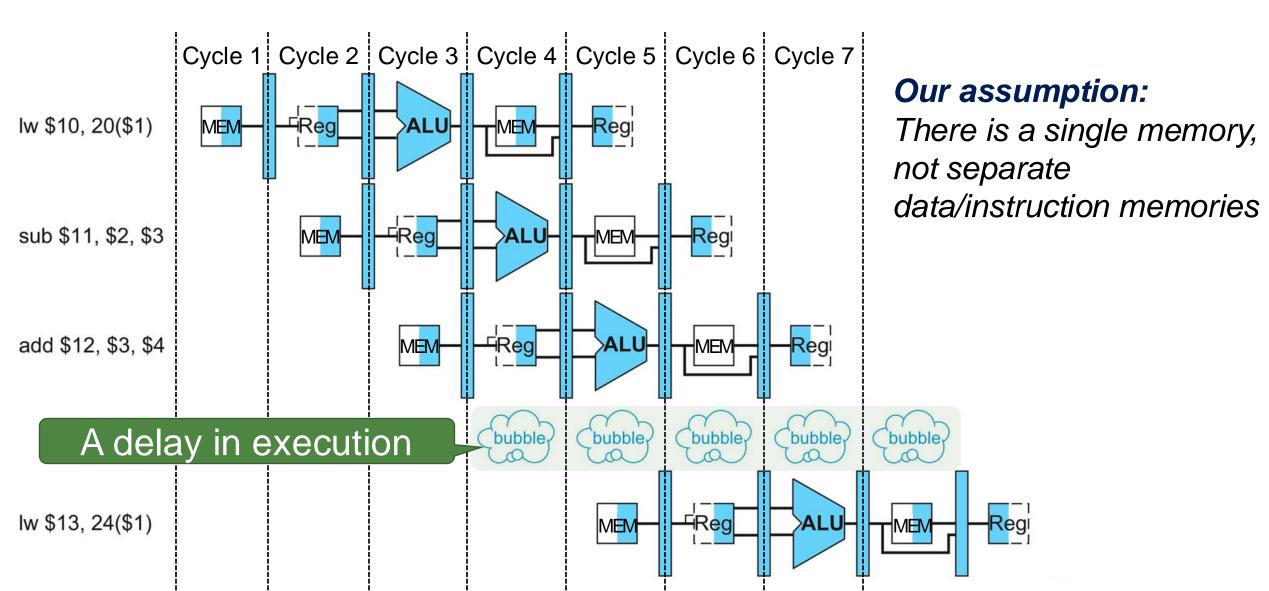
Recap: Structural Hazard

The **hardware** cannot support the combination of instructions



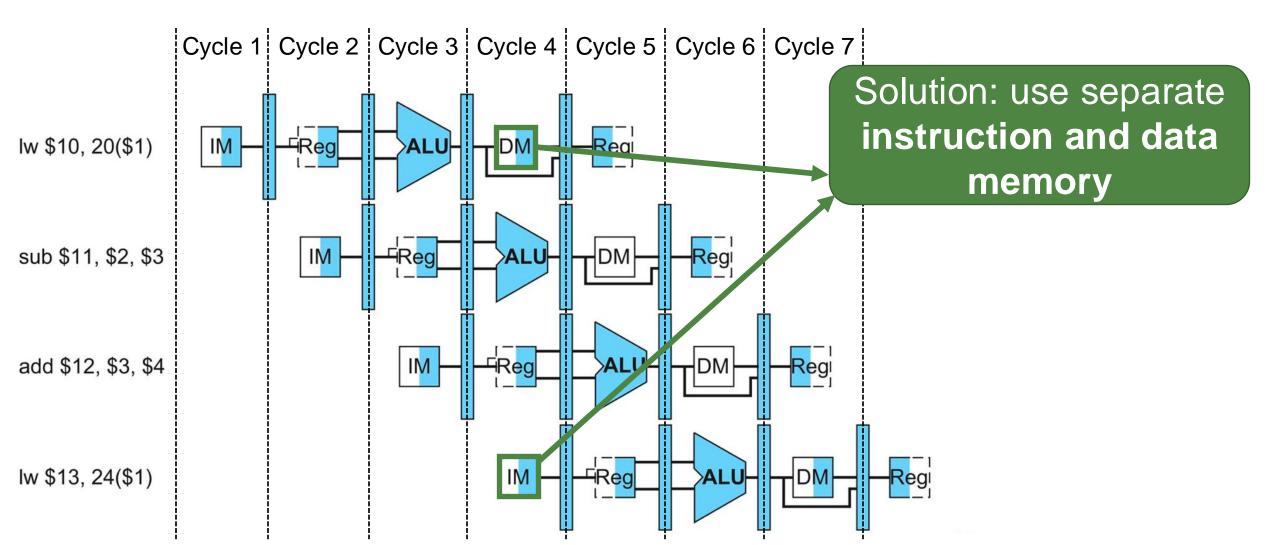
Recap: Pipeline Stall





Recap: Resource Duplication

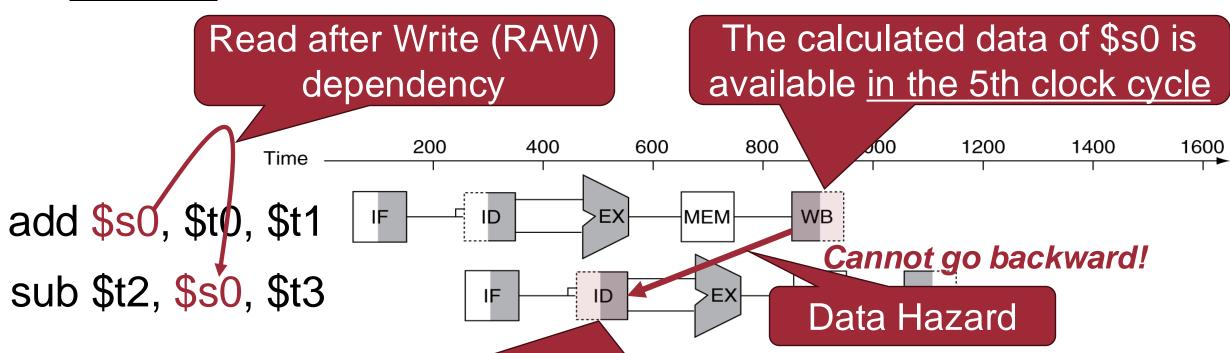




Recap: Data Hazard



A planned instruction cannot execute because data is not yet available

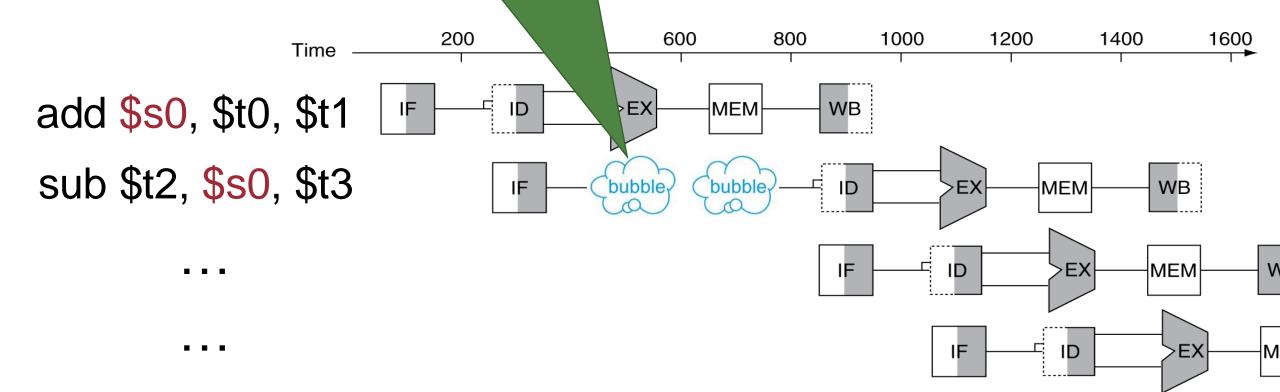


The data of \$s0 calculated in the previous instruction is needed in the 3rd clock cycle

Recap: Pipeline Stall



A delay in execution (2 clock cycle)

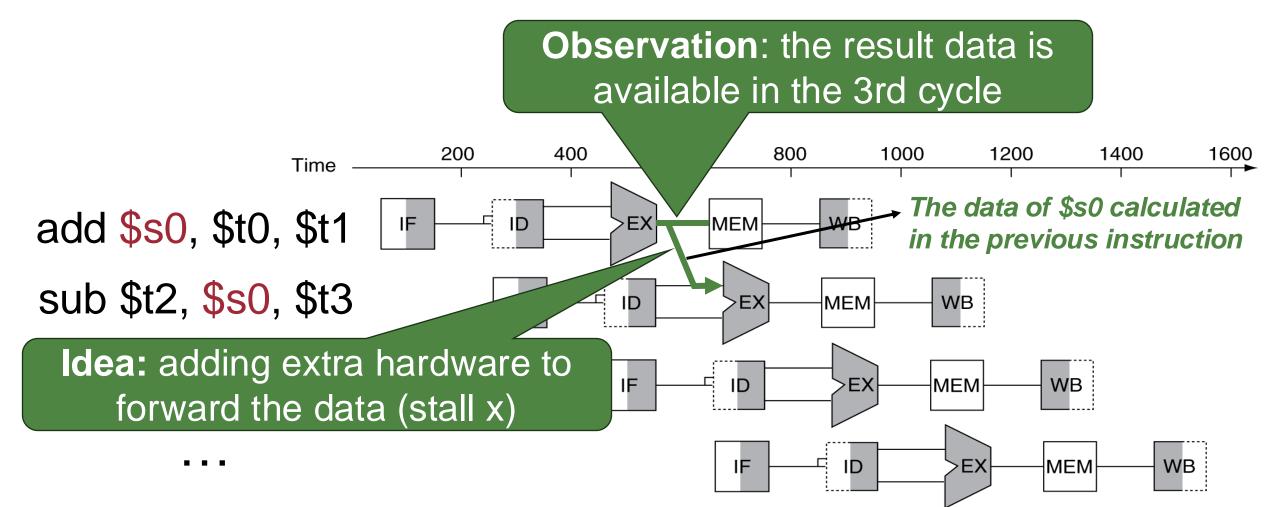


1

Recap: Forwarding

*

Use result when it is computed!



Recap: Compiler Optimization

- Reorder code to avoid use of load result in the next instruction
- C code for v[3] = v[0] + v[1]; v[4] = v[0] + v[2];

```
lw $t1, 0($t0)
lw $t2, 4($t0)
1 stall add $t3, $t1; $t2
sw $t3, 12($t0)
lw $t4 8($t0)
1 stall add $t5, $t1; $t4
sw $t5, 16($t0)
```

Idea: code reordering to avoid stalls (by compiler)

Compiler requires knowledge of the pipeline structure!

13 cycles





Situations that prevent starting the next instruction in the next cycle

Hazard #1: Structural hazard

Conflict for use of a hardware resource

Solution:

- Stall
- Resource duplication

Hazard #2: Data hazard

An instruction cannot execute because data is not yet available

Hazard #3: Control hazard

Solution:

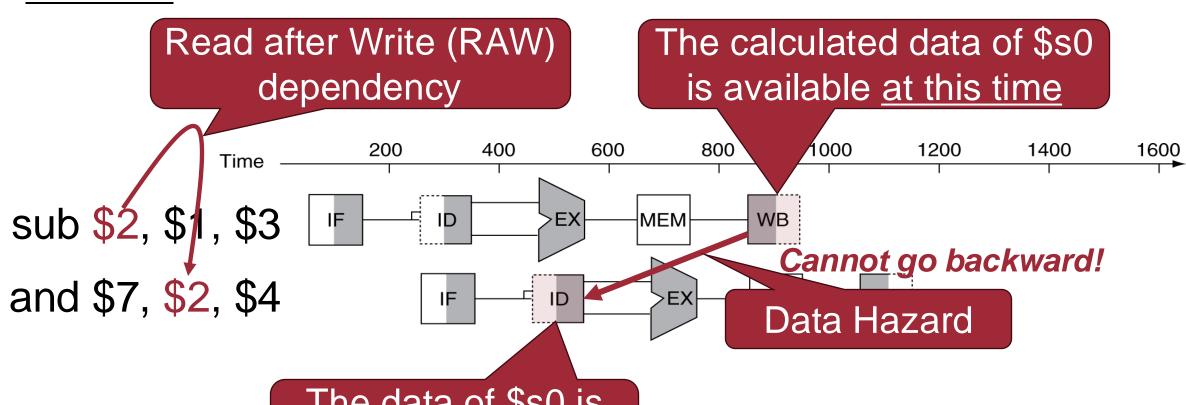
- Stall
- Forwarding
- Compiler optimization

Let's look at the hardware details

Data Hazard



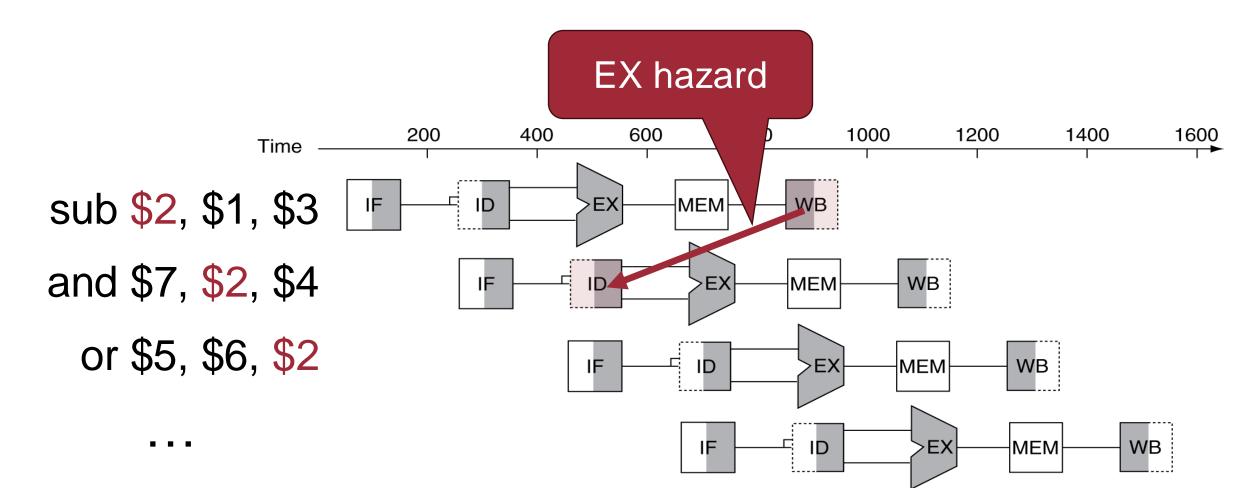
A planned instruction cannot execute because data is not yet available



The data of \$s0 is needed at this time

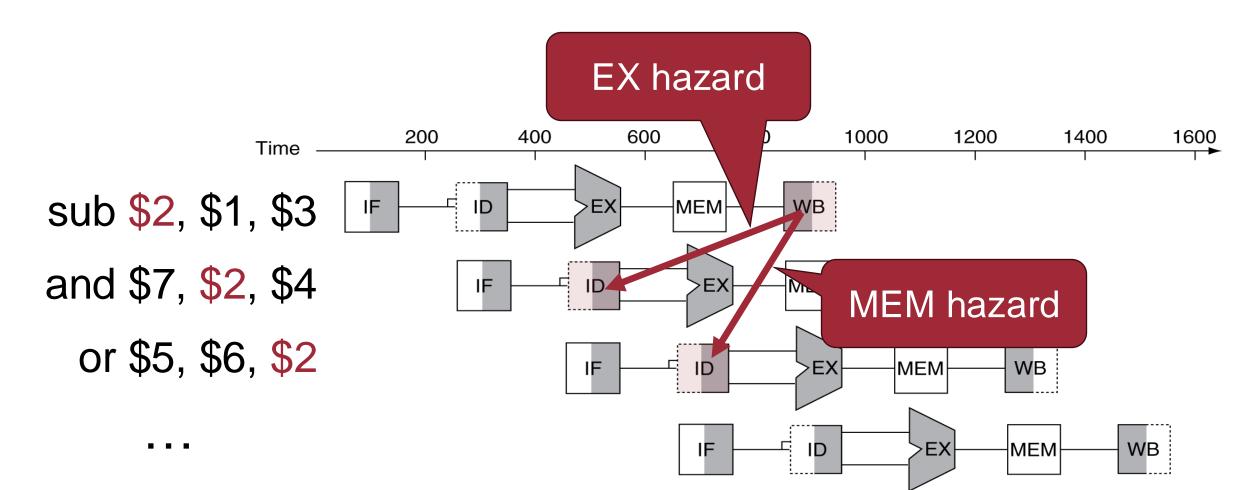
Data Hazard #1: EX Hazard

A planned instruction cannot execute because data is not yet available



Data Hazard #2: MEM Hazard

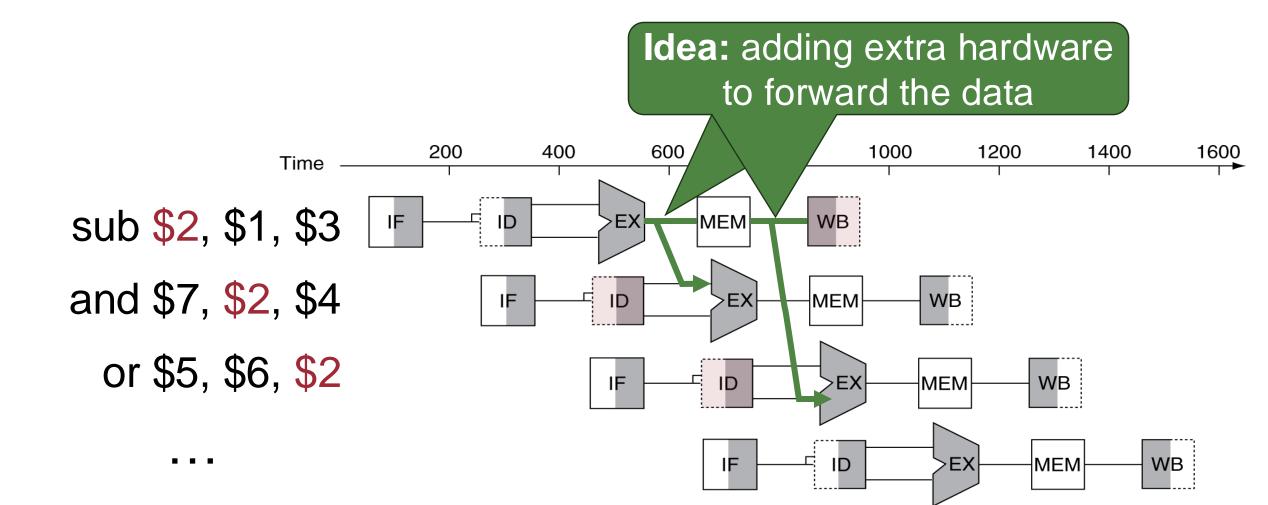
A planned instruction cannot execute because data is not yet available



Solution: Forwarding

- *

Use result when it is computed!

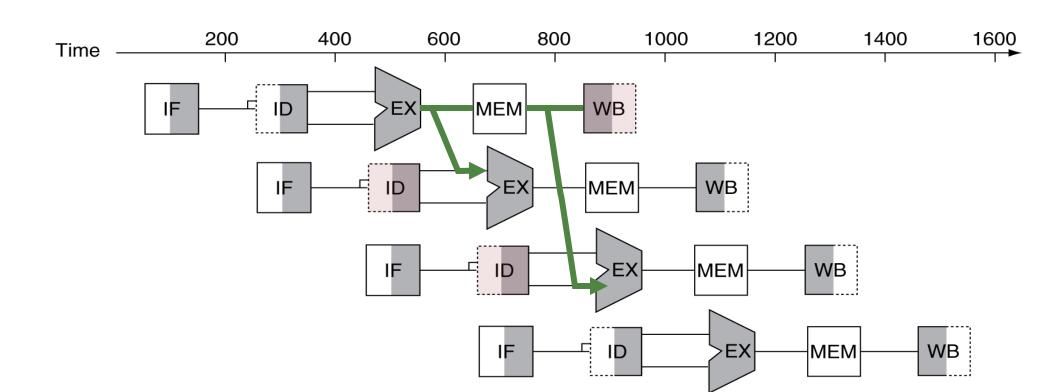


Today's Topic

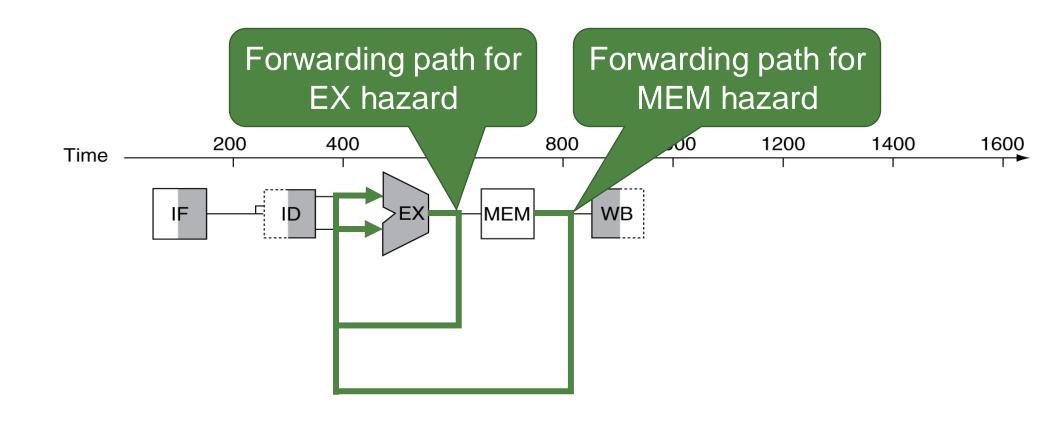


- *
- How should the hardware be modified to support forwarding?
- How should the hardware be modified to detect data hazards?
- How should the hardware be modified to support stalls?

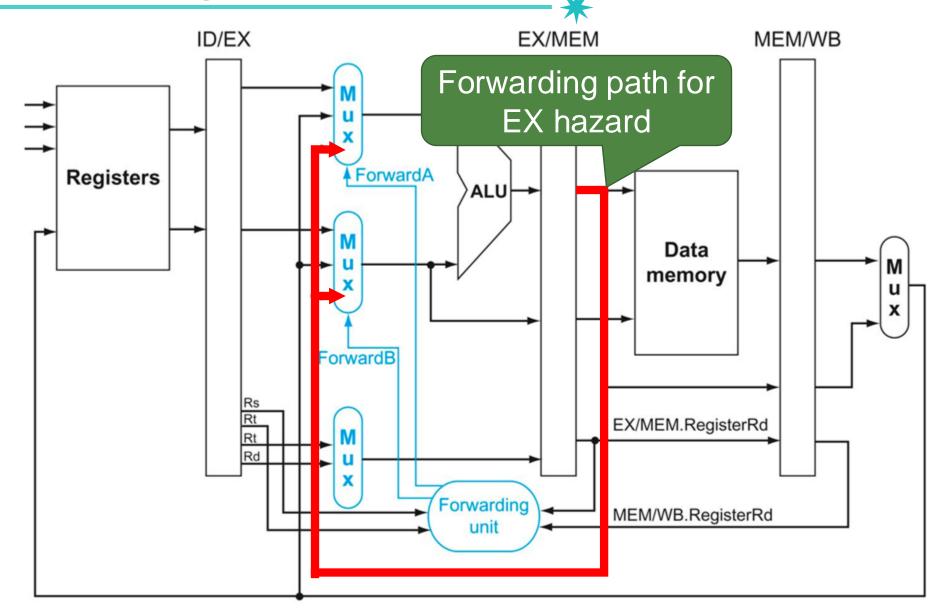
Forwarding Paths Overview



Forwarding Paths Overview

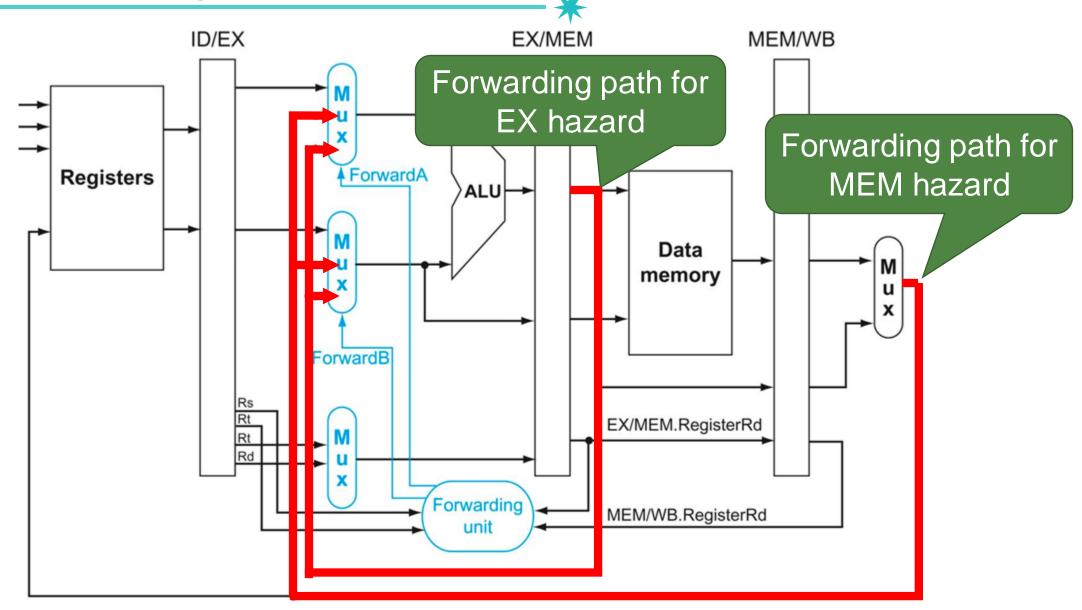


Forwarding Paths: Details



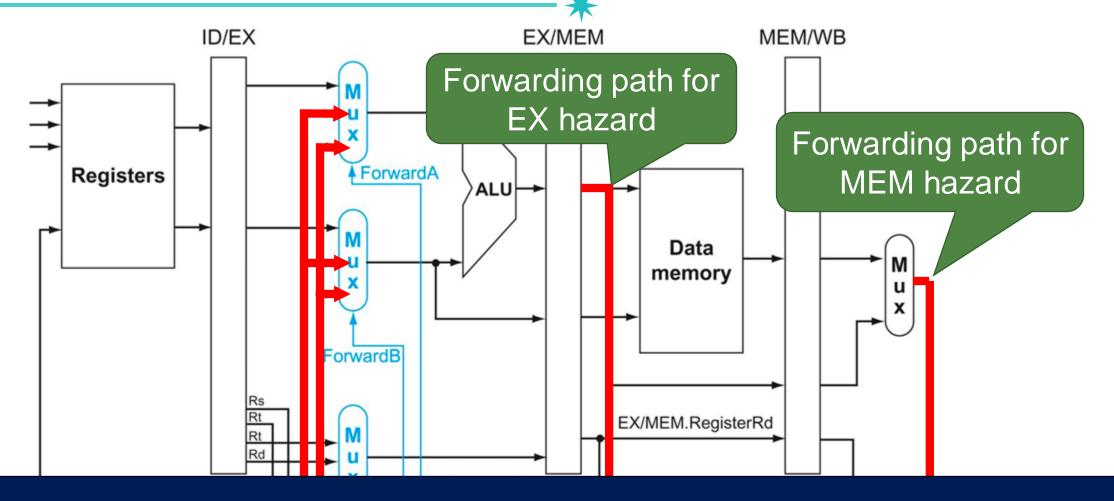
21

Forwarding Paths: Details

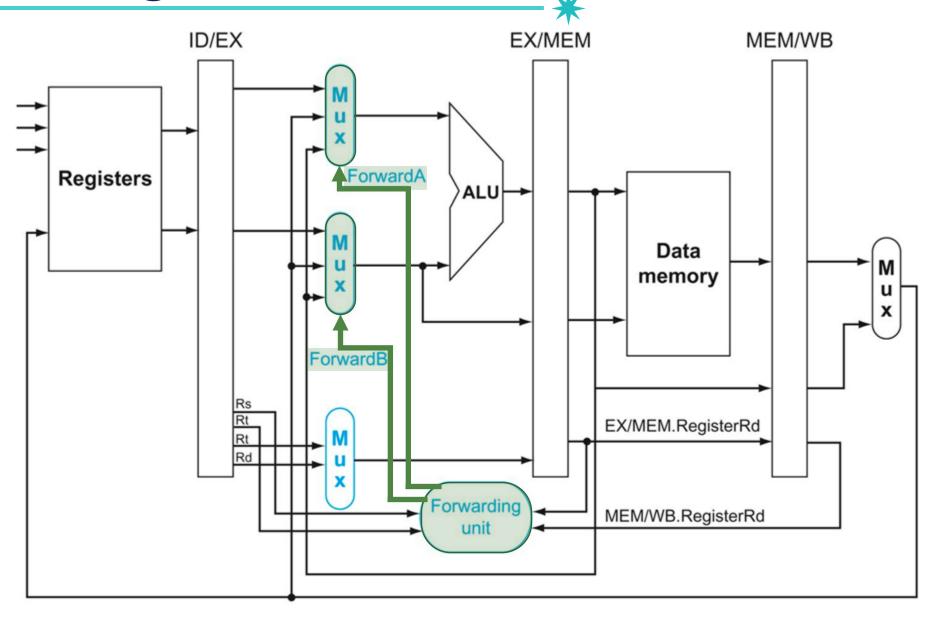


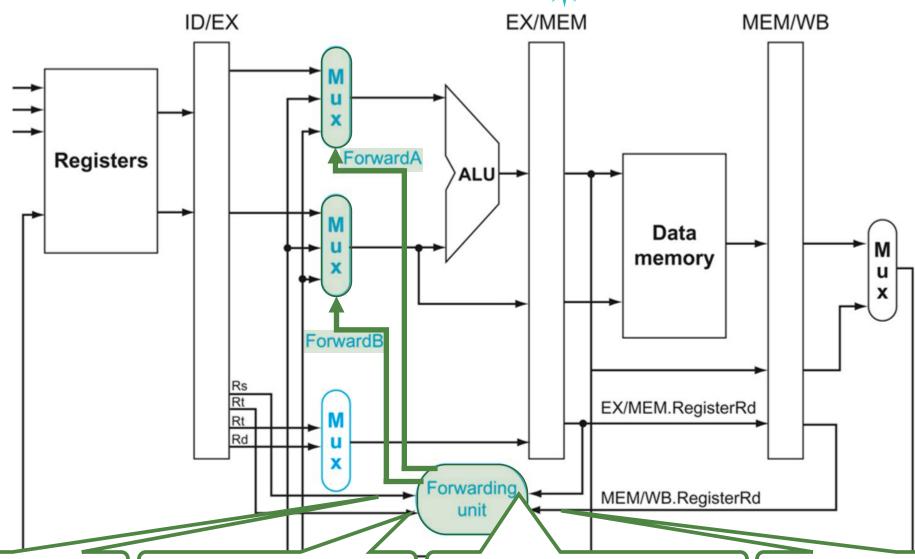






How to detecting the need to forward?



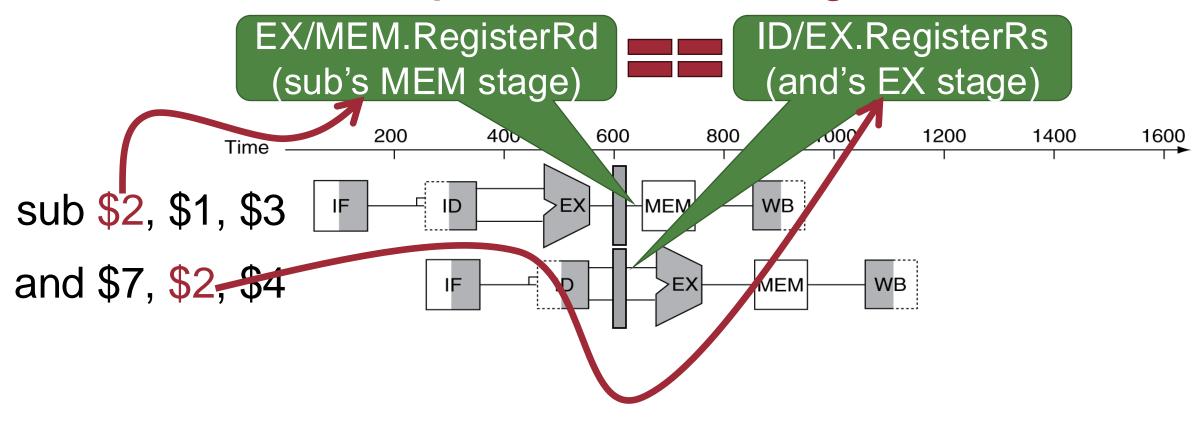


ID/EX.RegisterRs ↓

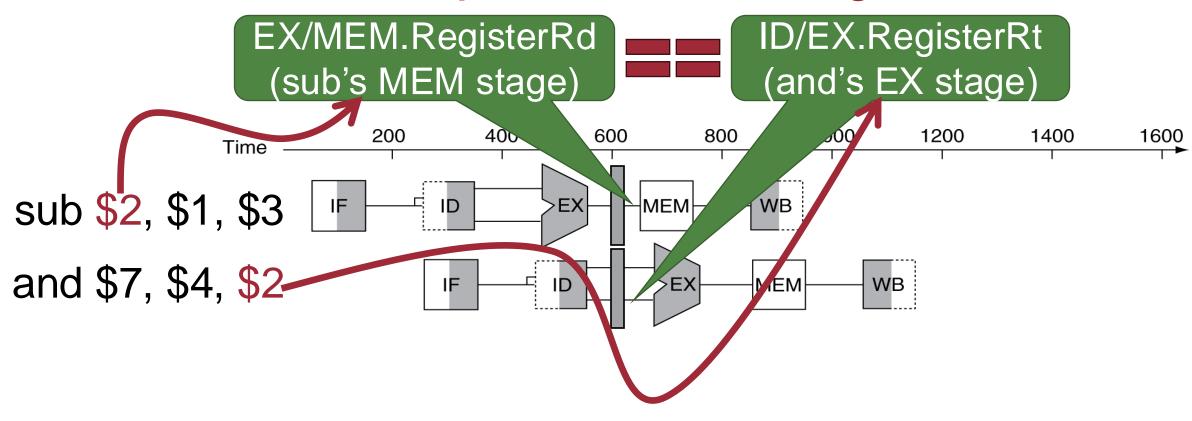
ID/EX.RegisterRt

EX/MEM.RegisterRd | MEM/WB.RegisterRd

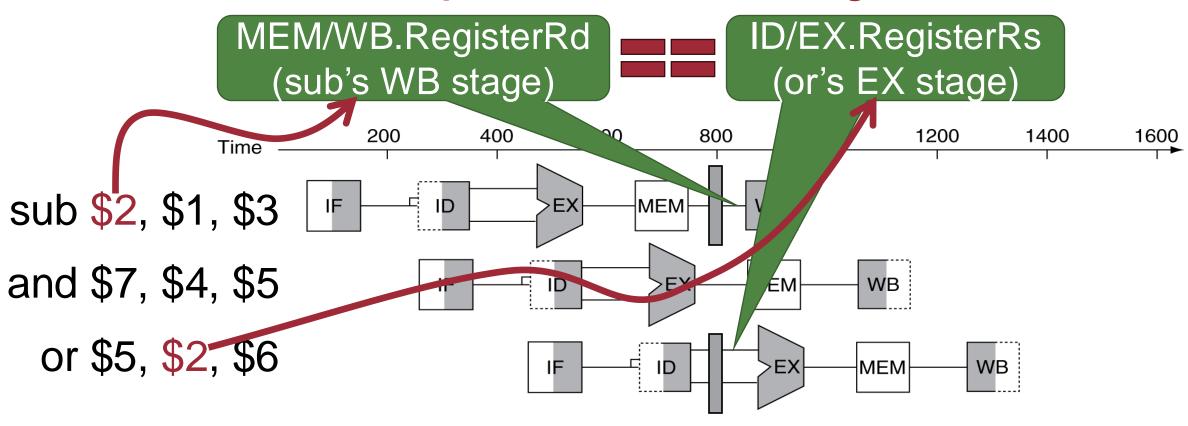
EX hazard detection! First operand ← EX/MEM.RegisterRd



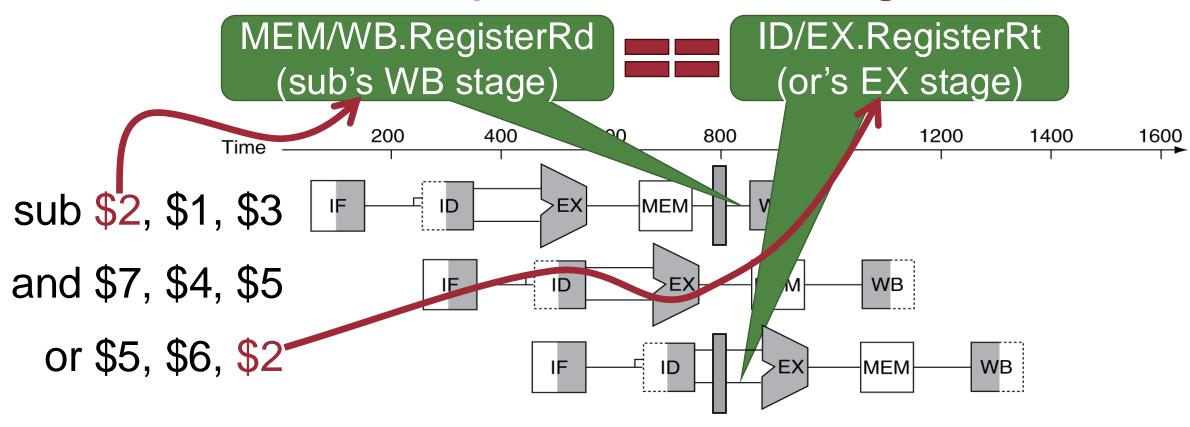
EX hazard detection! Second operand ← EX/MEM.RegisterRd



MEM hazard detection! First operand ← MEM/WB.RegisterRd



MEM hazard detection! Second operand ← MEM/WB.RegisterRd



Additional Conditions

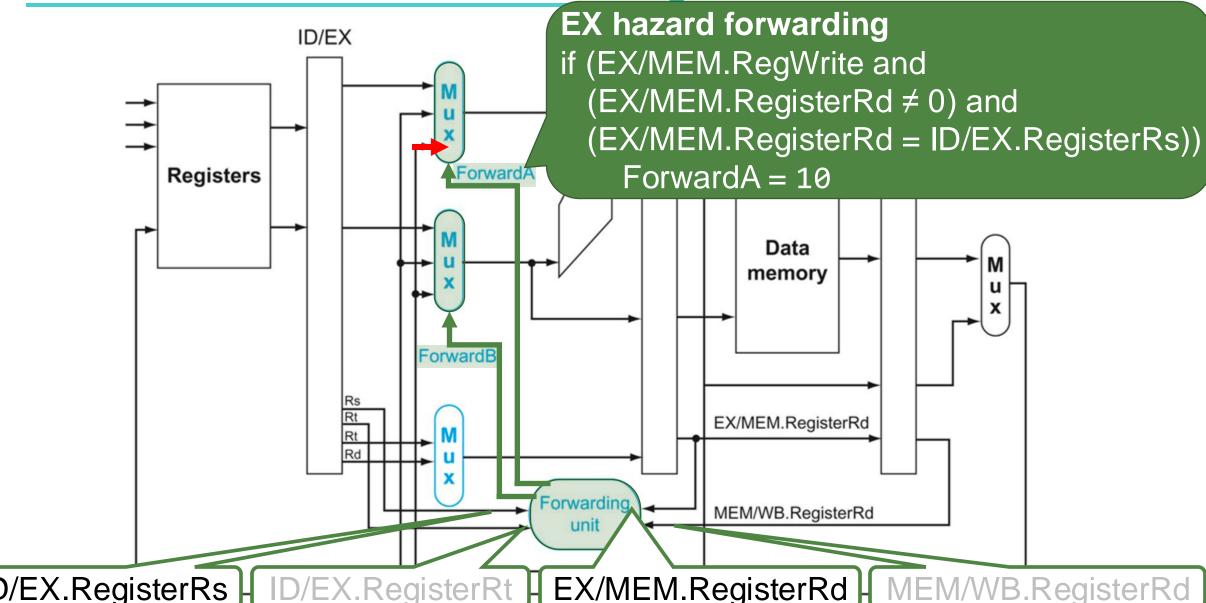




- But only if forwarding instruction will write to a register!
 - -EX/MEM.RegWrite = 1
 - -MEM/WB.RegWrite = 1
- And only if rd for that instruction is not \$zero
 - -EX/MEM.RegisterRd ≠ 0
 - -MEM/WB.RegisterRd ≠ 0

30

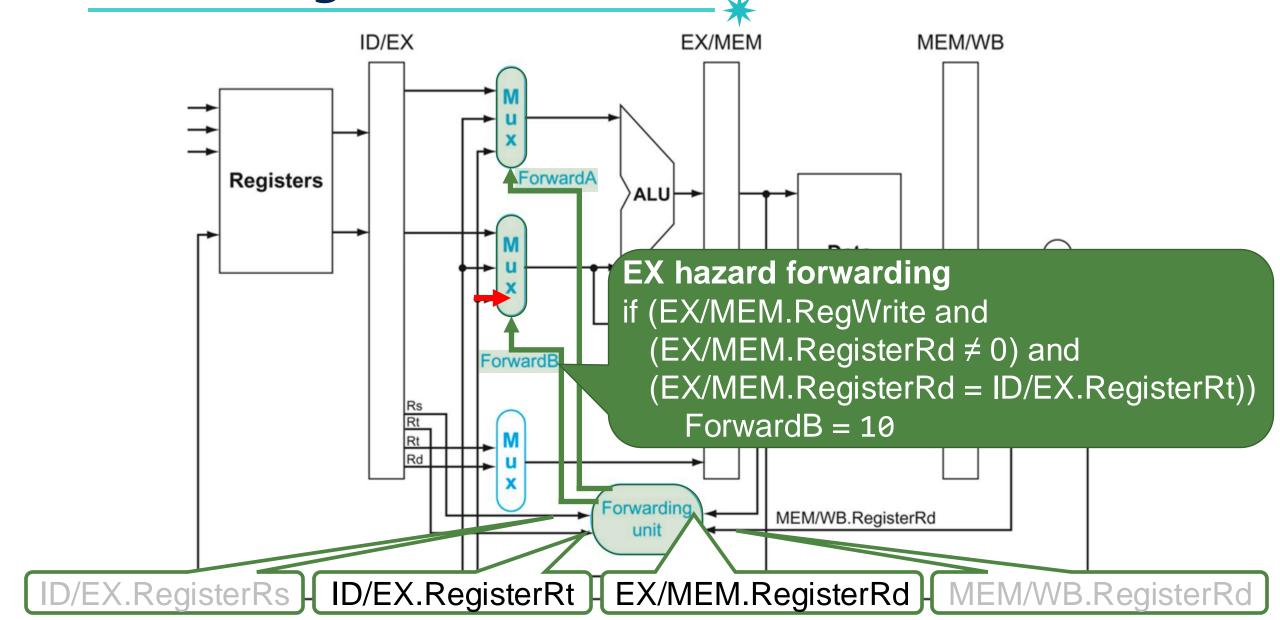
Detecting the Need to Forward: Details



ID/EX.RegisterRs

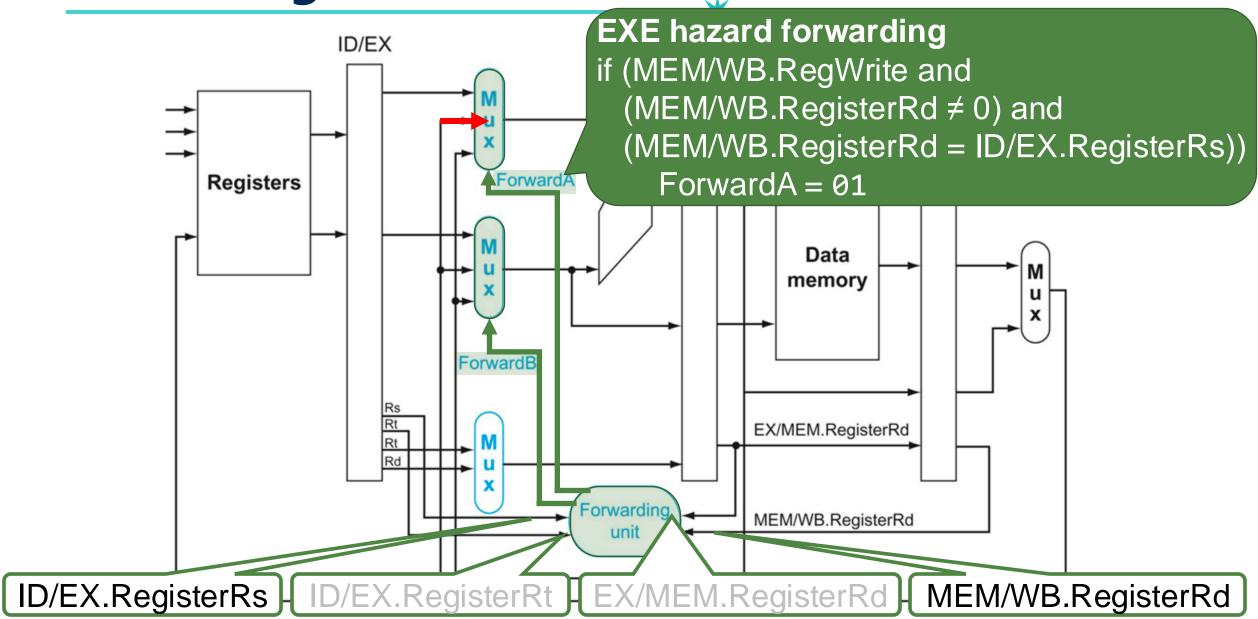
Detecting the Need to Forward: Details





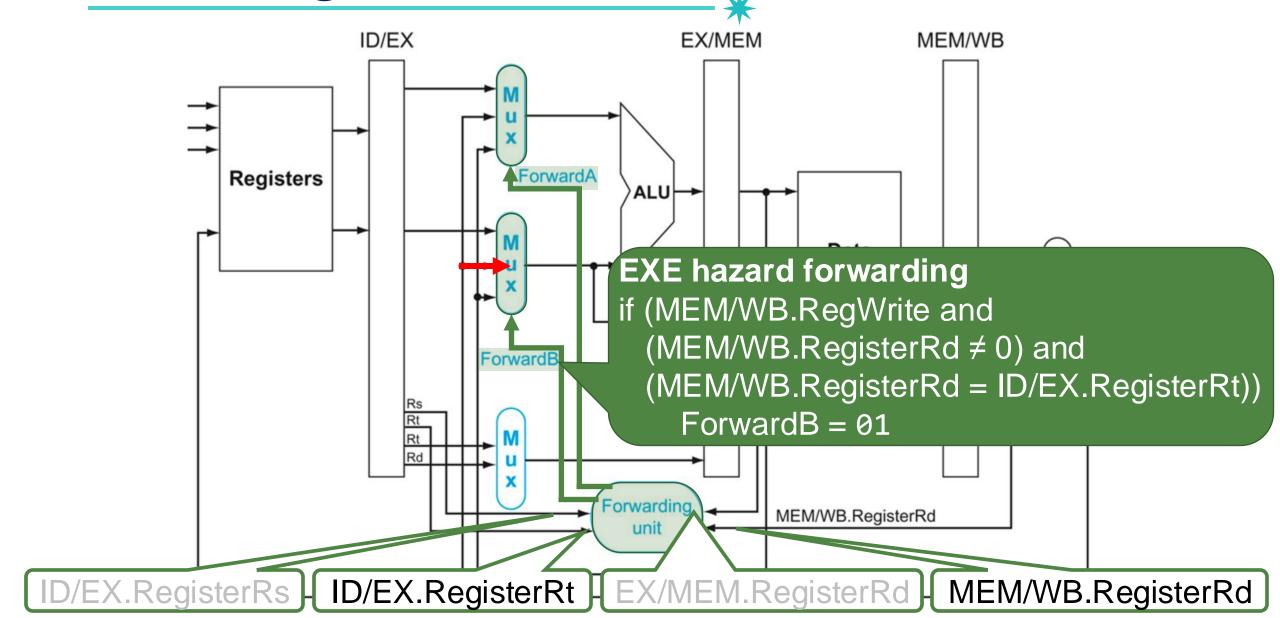






Detecting the Need to Forward: Details





Summary: Forwarding Conditions

- EX hazard: Forward <u>EX/MEM value</u> to EX stage
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRs)) ForwardA = 10
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRt)) ForwardB = 10
- MEM hazard: Forward <u>MEM/WB value</u> to EX stage
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRs)) ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRt)) ForwardB = 01

35

Double Data Hazard

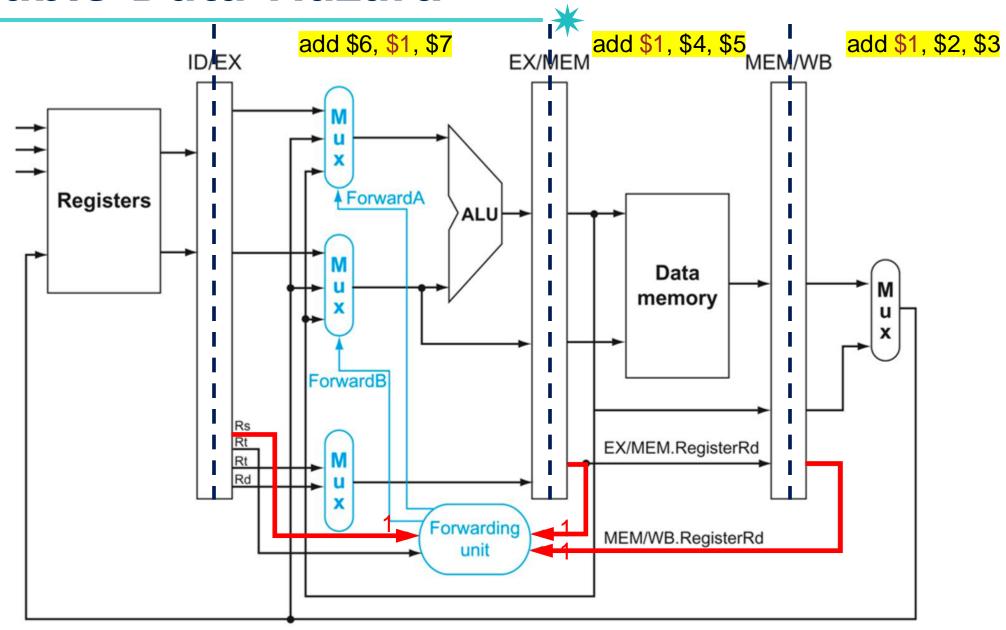


add \$1, \$2, \$3 add \$1, \$4, \$5 add \$6, \$1, \$7



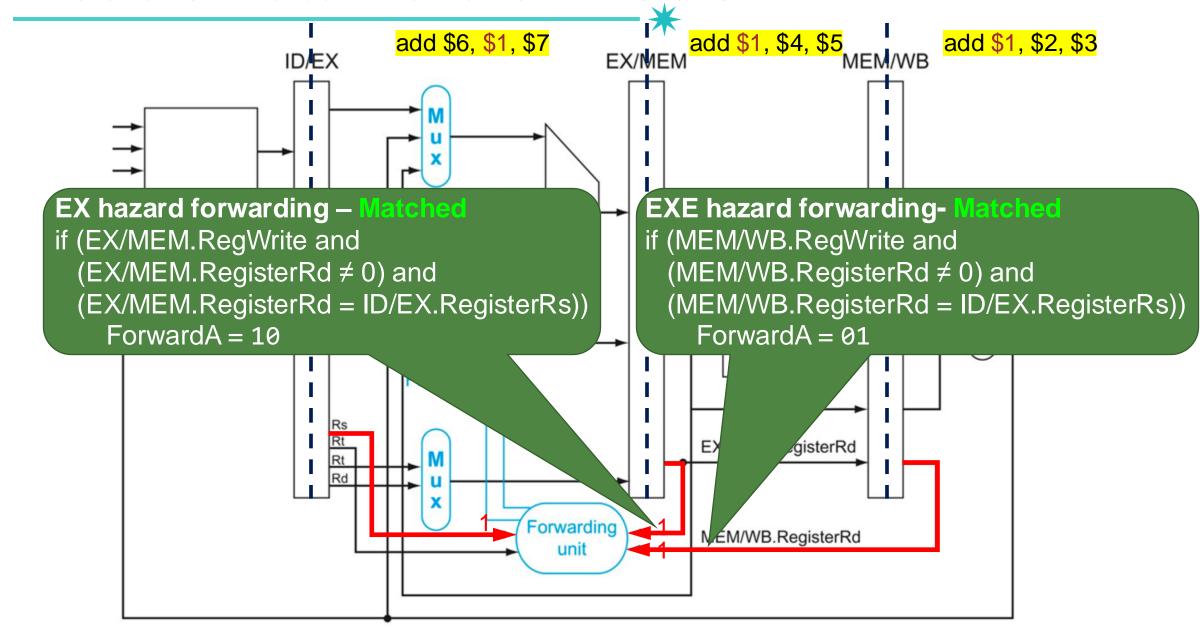
Any problem?

Double Data Hazard

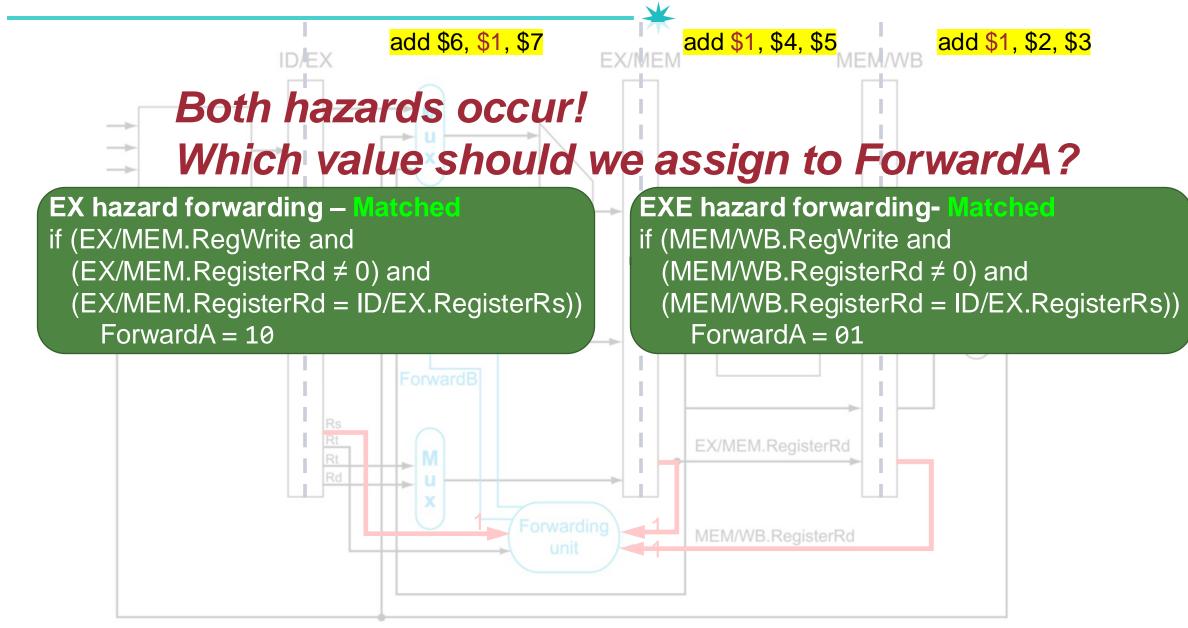


Double Data Hazard: Problem

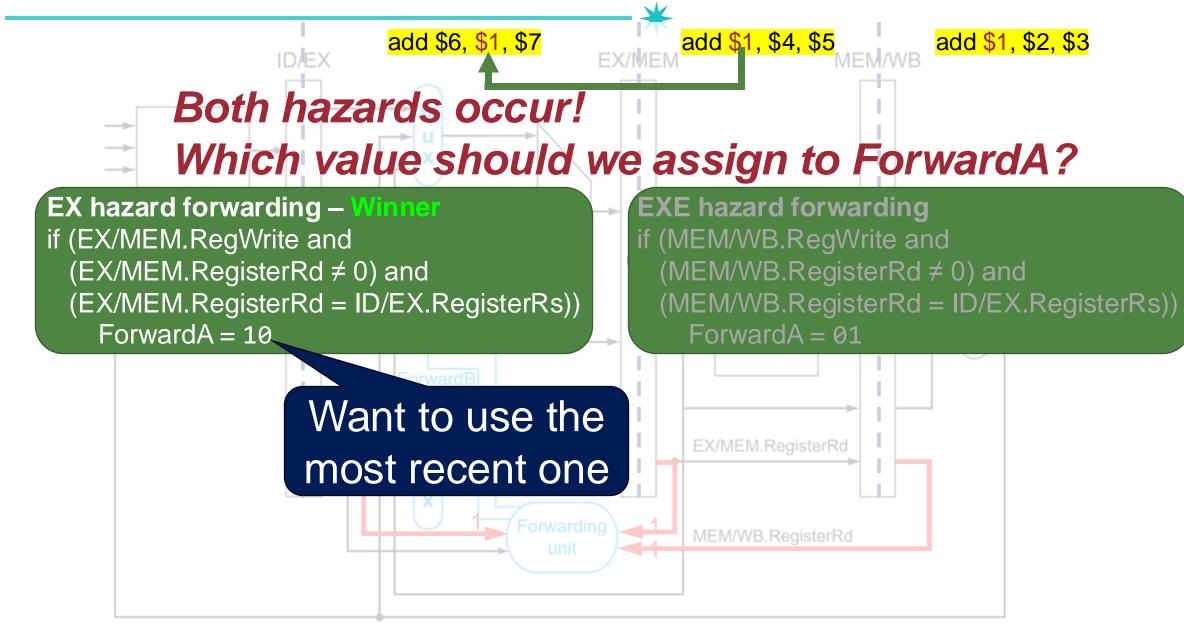




Double Data Hazard: Problem



Double Data Hazard: Problem



Revise the MEM Hazard Condition

- Only forward <u>if EX hazard condition isn't true</u>
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)

and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)

and (EX/MEM.RegisterRd = ID/EX.RegisterRs))

and (MEM/WB.RegisterRd = ID/EX.RegisterRs))

ForwardA = 01

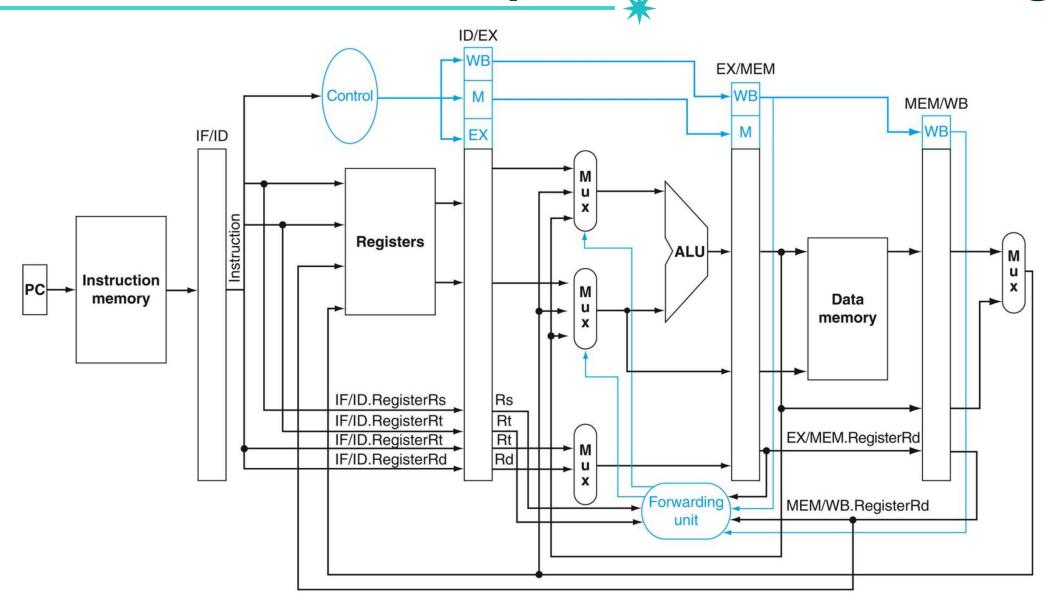
- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)

and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)

and (EX/MEM.RegisterRd = ID/EX.RegisterRt))

and (MEM/WB.RegisterRd = ID/EX.RegisterRt))

ForwardB = 01





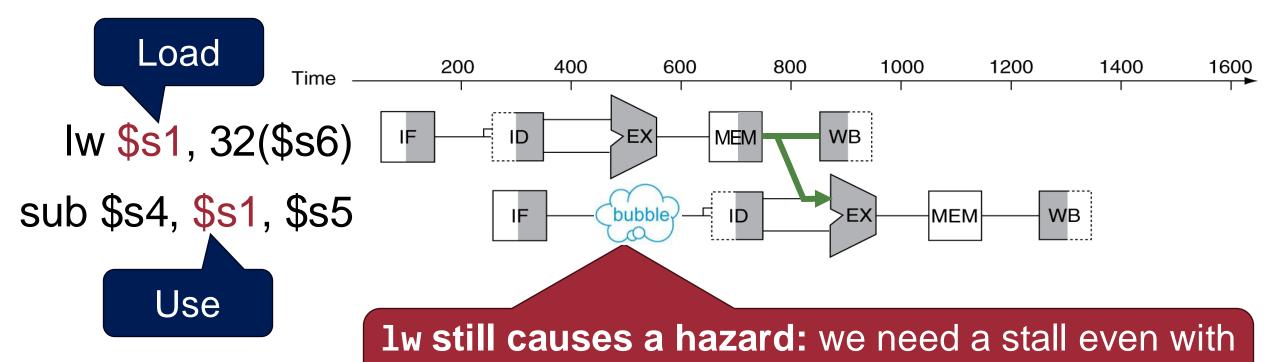
Recall: Three Types of Data Hazard

1. EX Hazard

2. MEM Hazard

3. Load-Use Hazard

Load-Use Data Hazard with Forwarding

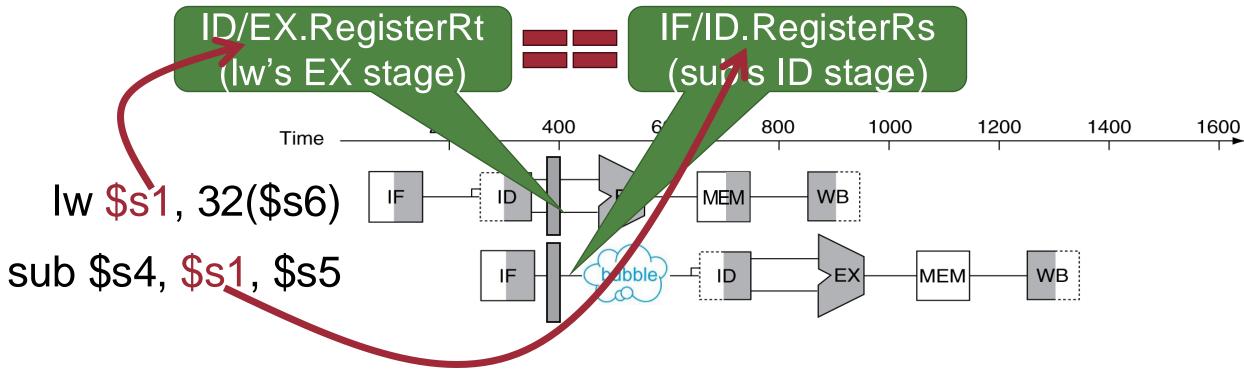


forwarding when a load tries to use the data

Load-Use Data Hazard Detection: Basic Idea

Load-Use hazard detection!

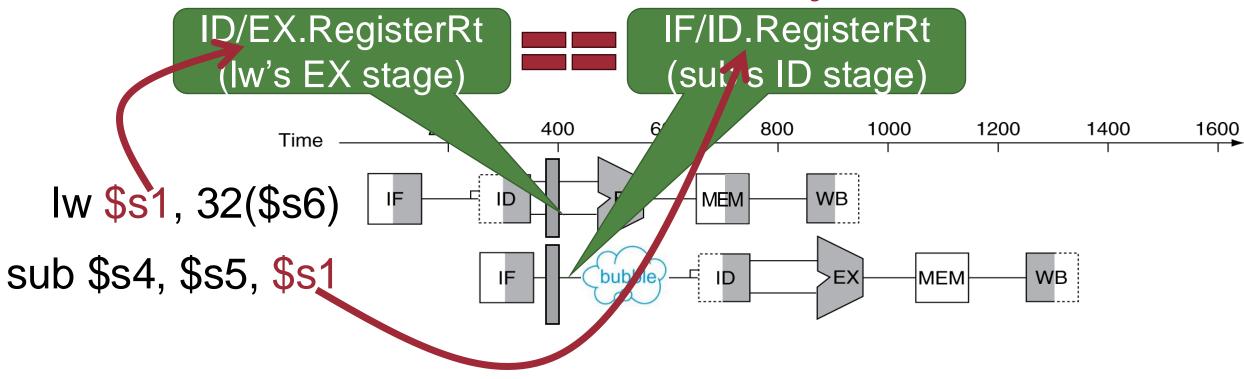
→ Insert bubble for one clock cycle



Load-Use Data Hazard Detection: Basic Idea

Load-Use hazard detection!

→ Insert bubble for one clock cycle



Additional Condition

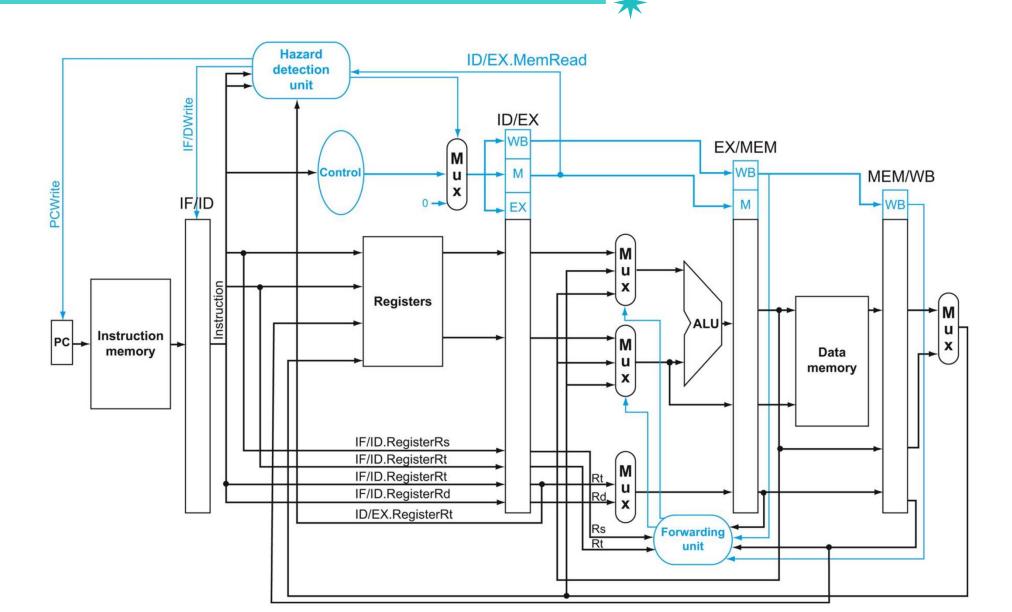


- But only if the instruction the ID/EX stage is a load instruction
 - -ID/EX.MemRead = 1

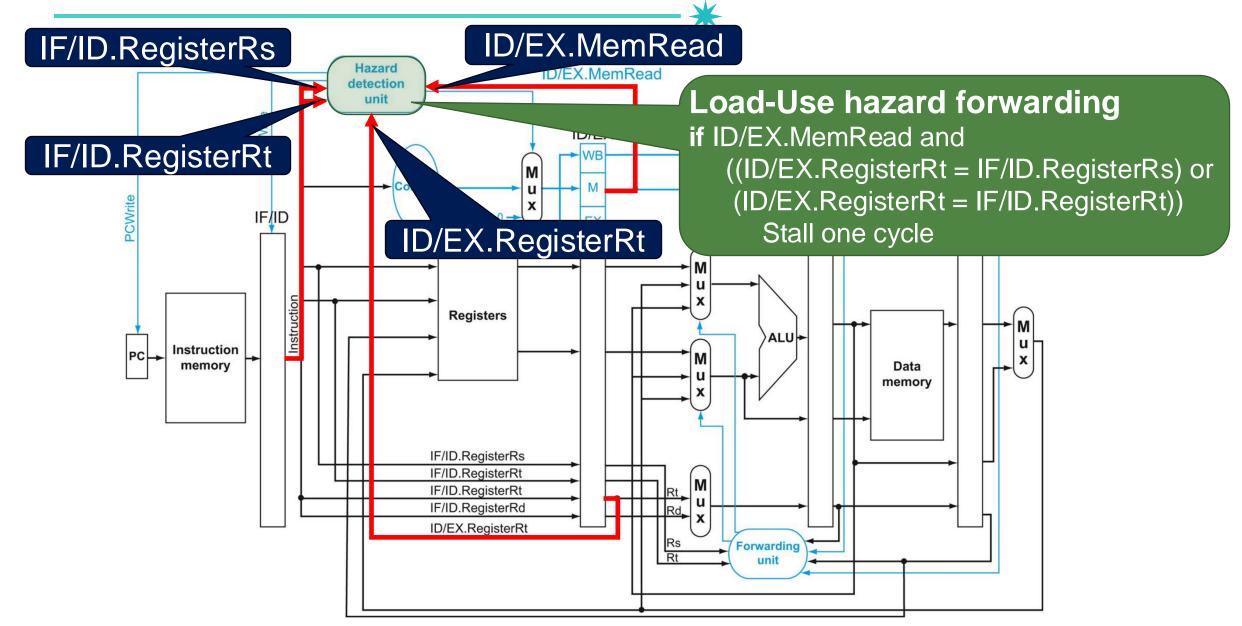
Load-Use Data Hazard Detection: Basic Idea

```
if ID/EX.MemRead and
    ((ID/EX.RegisterRt = IF/ID.RegisterRs) or
    (ID/EX.RegisterRt = IF/ID.RegisterRt))
Insert bubble for one clock cycle (Stall)
```

Load-Use Data Hazard Detection: Details 48



Load-Use Data Hazard Detection: Details 49



Load-Use Data Hazard Detection: Basic Idea

if ID/EX.MemRead and
 ((ID/EX.RegisterRt = IF/ID.RegisterRs) or
 (ID/EX.RegisterRt = IF/ID.RegisterRt))

Insert bubble for one clock cycle (Stall)

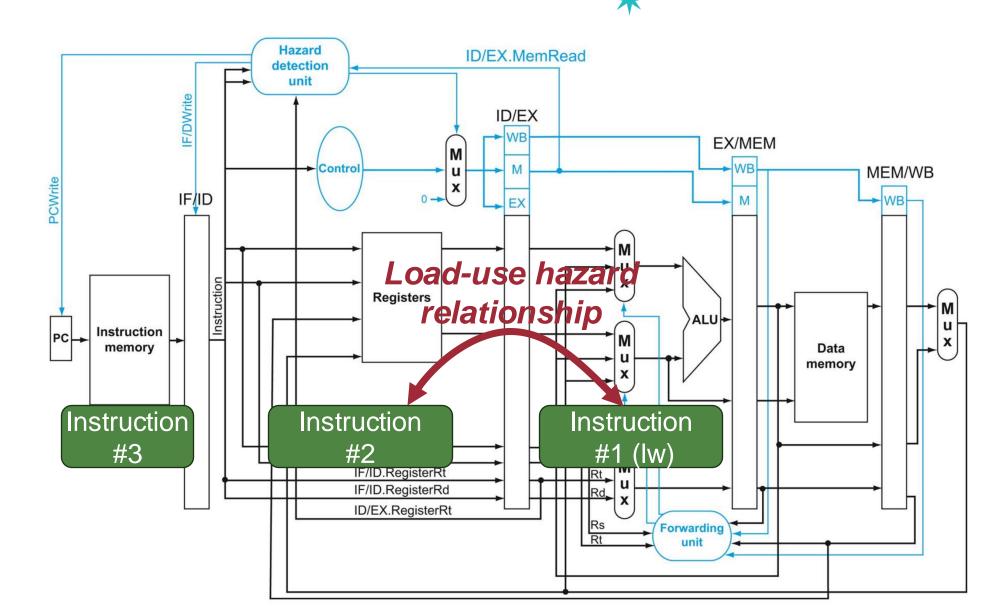
How to stall the pipeline?

How to Stall the Pipeline?

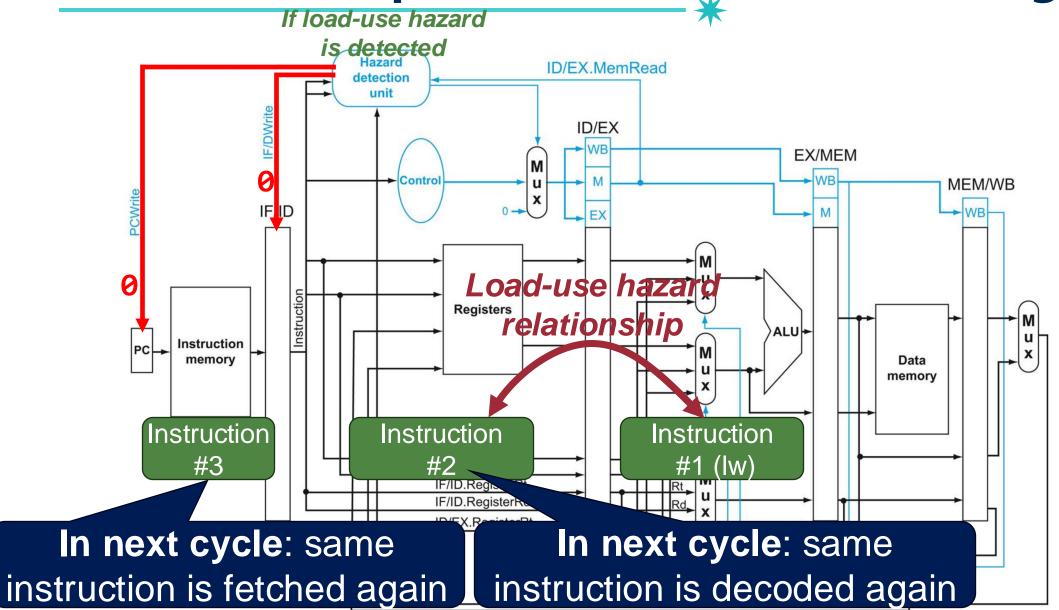
1) Prevent update of PC and IF/ID register

2) Force control values in ID/EX register to 0

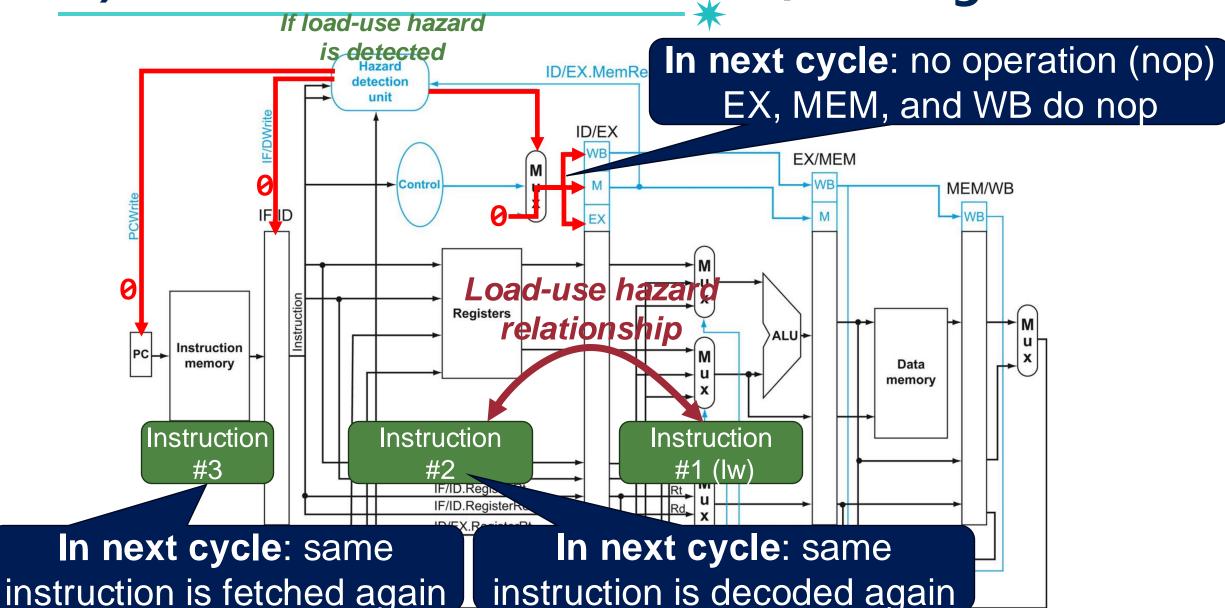
Stall Details: Current Cycle



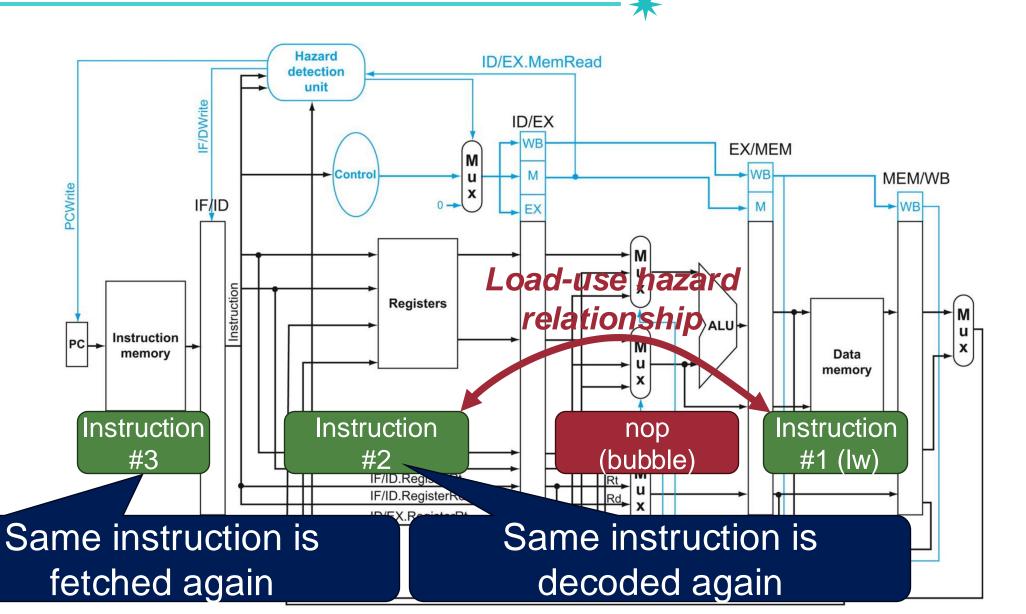
1) Prevent Update of PC and IF/ID Register



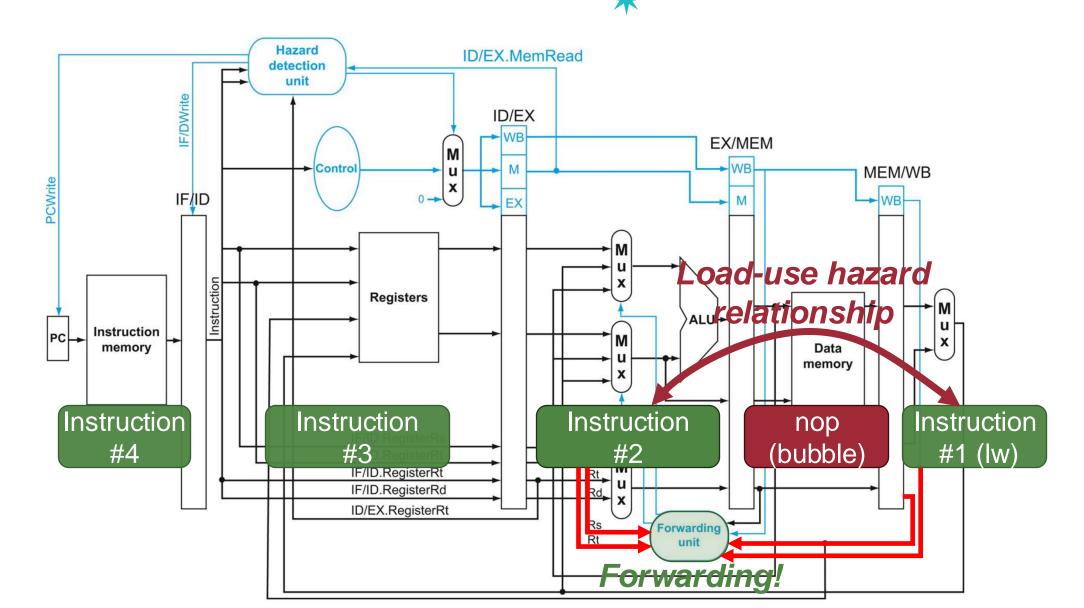
2) Force Control Values in ID/EX Register to 0



Stall Details: Next Cycle (Stall Inserted)



Stall Details: Next Next Cycle...



Summary: Stall the Pipeline

- 1) Prevent update of PC and IF/ID register
 - Same instruction is decoded again
 - Same instruction (with same PC) is fetched again
- 2) Force control values in ID/EX register to 0
 - -EX, MEM and WB do nop (no-operation)

Pipelining Hazards Summary

Situations that prevent starting the next instruction in the next cycle

Hazard #1: Structural hazard

Conflict for use of a hardware resource

Solution:

- Stall
- Resource duplication

Hazard #2: Data hazard

An instruction cannot execute because data is not yet available

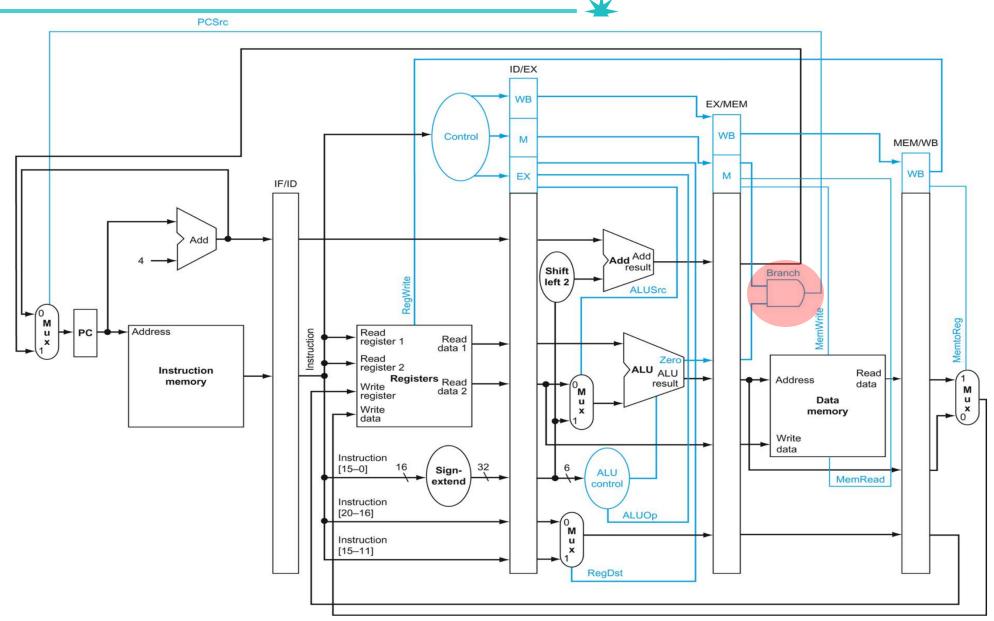
Solution:

- Stall
- Forwarding
- Compiler optimization

Hazard #3: Control hazard

Control Hazard

Recall: Branch Outcome Determined in MEM



Control Hazard

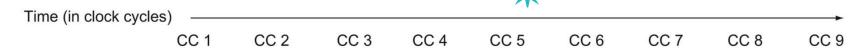


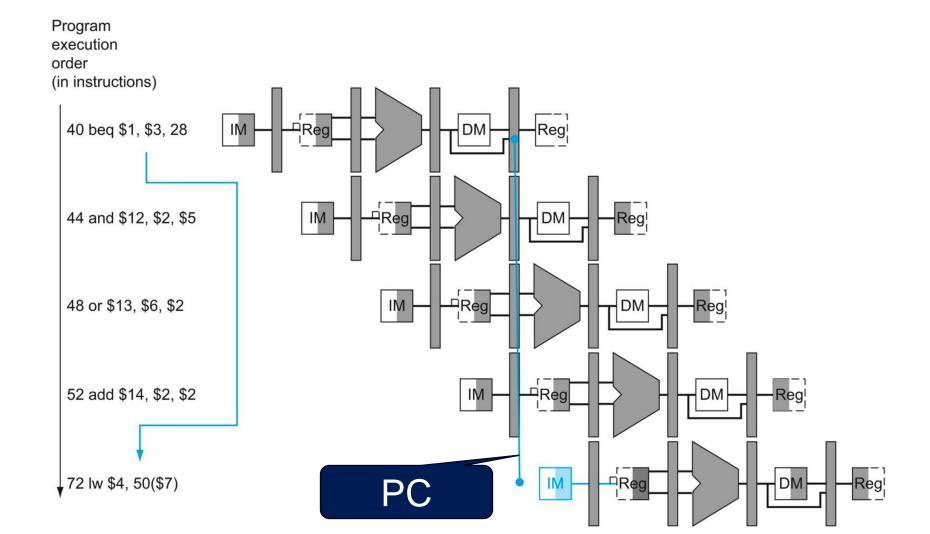


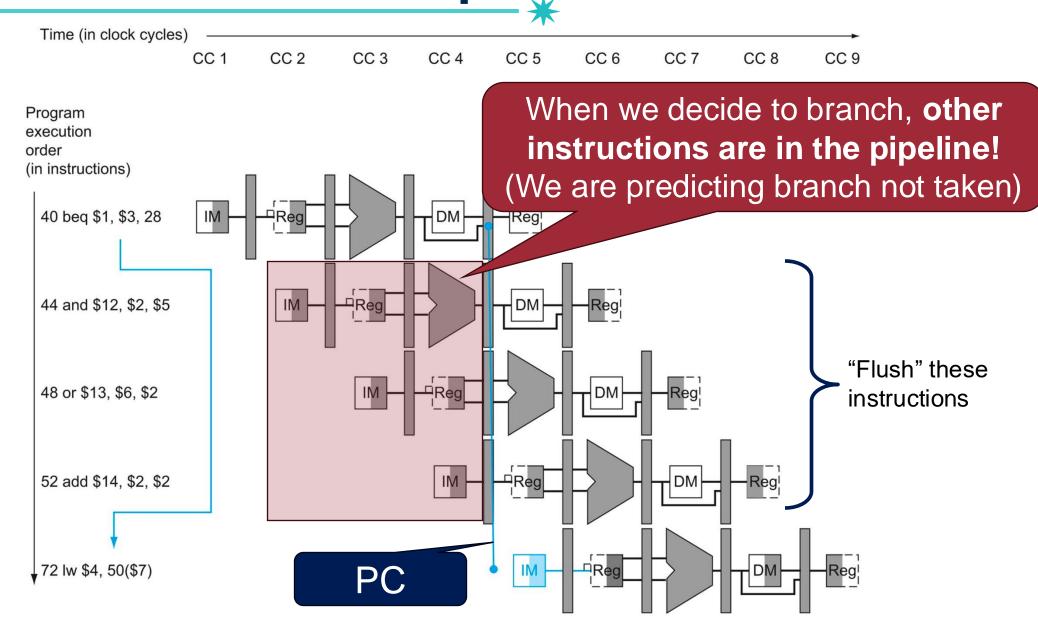
- Branch determines flow of control
 - -Fetching next instruction depends on branch outcome
 - -However, the branch decision is made in the MEM stage

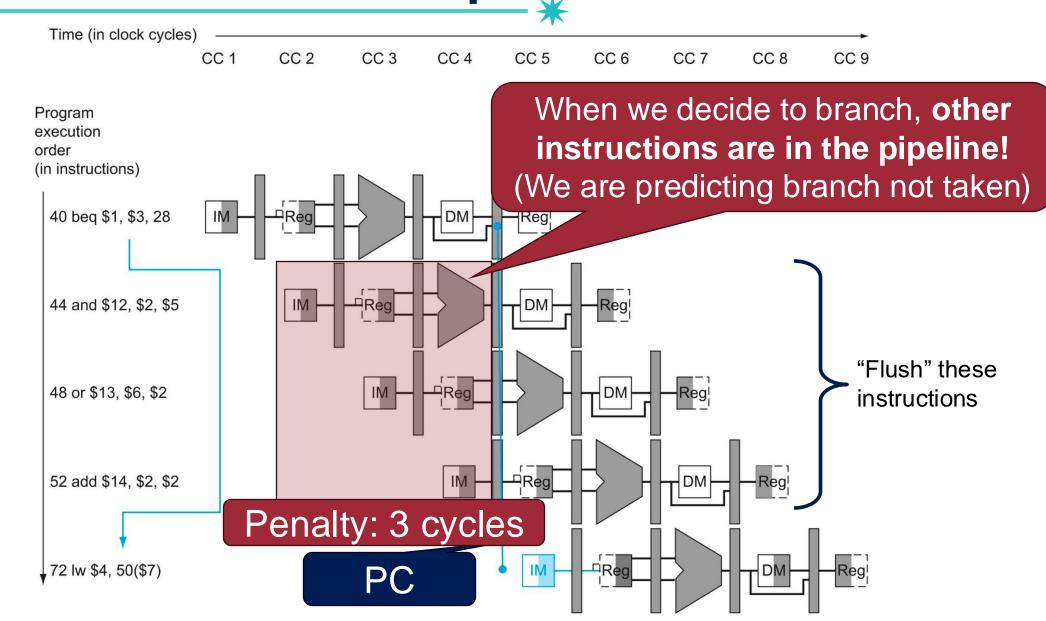


```
beq $1, $3, 28
and $12, $2, $5
or $13, $6, $2
                     Control flow
add $14, $2, $2
lw $4, 50($7)
```









Pipelining Hazards Summary

Situations that prevent starting the next instruction in the next cycle

Hazard #1: Structural hazard

Conflict for use of a hardware resource

Solution:

- Stall
- Resource duplication

Hazard #2: Data hazard

An instruction cannot execute because data is not yet available

Solution:

- Stall
- Forwarding
- Compiler optimization

Hazard #3: Control hazard

The next instruction is uncertain due to branching

Solution:

- Stall
- Optimized branch processing
- Branch prediction
- Delayed branch

Solution for Control Hazard: Stall

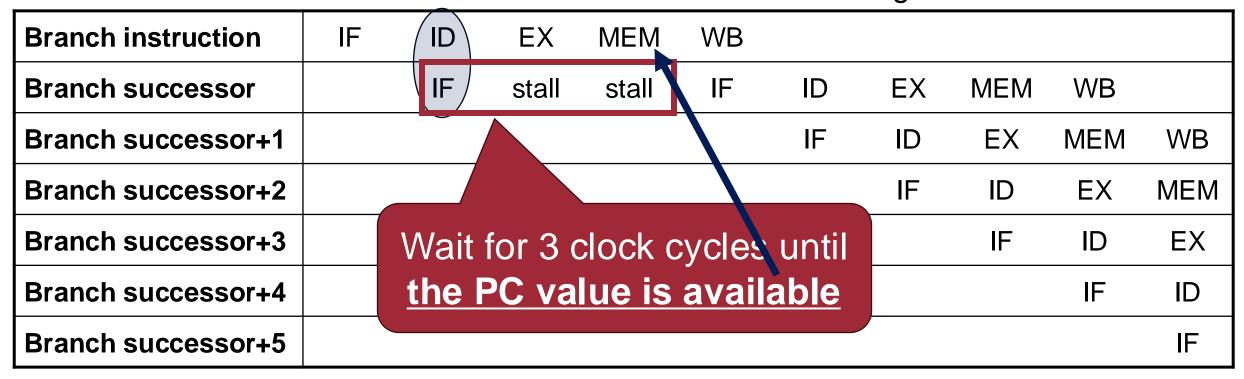


- Most naïve approach: stall on branch
 - Wait until branch outcome determined before fetching next instruction

Branch instruction	IF	ID	EX	MEM	WB					
Branch successor		IF	stall	stall	IF	ID	EX	MEM	WB	
Branch successor+1						IF	ID	EX	MEM	WB
Branch successor+2							IF	ID	EX	MEM
Branch successor+3								IF	ID	EX
Branch successor+4									IF	ID
Branch successor+5										IF

Solution for Control Hazard: Stall

- Most naïve approach: stall on branch
 - Wait until branch outcome determined before fetching next instruction



Work correctly, but is slow (3 cycle penalty regardless of whether a branch occurs)

Solution for Control Hazard: Stall

- For 5-stage pipeline, 3 cycle penalty for stall
- 15% branch frequency

What is the CPI?

$$1 + (0.15 \times 3) = 1.45$$

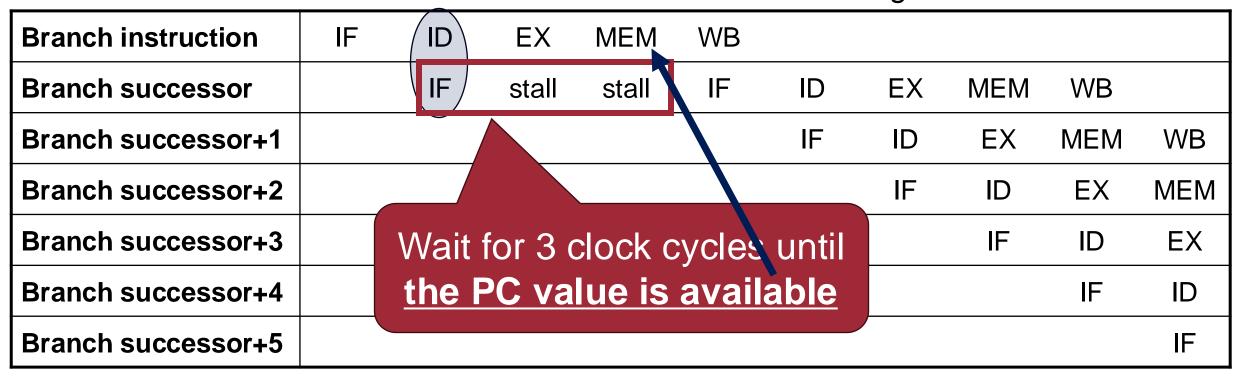
CPI in the ideal situation of pipelining

3 cycle penalty

Solution for Control Hazard: Stall

70

- Most naïve approach: stall on branch
 - Wait until branch outcome determined before fetching next instruction



Work correctly, but is slow (3 cycle penalty regardless of whether a branch occurs)

Solution for Control Hazard: Stall

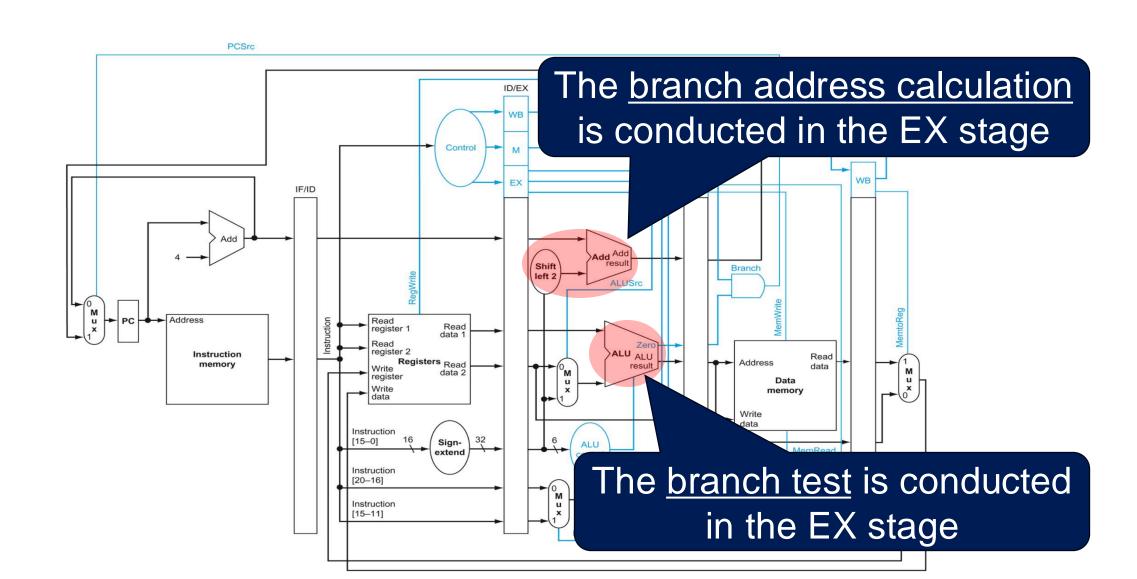




How can the penalty be reduced to less than 3 clock cycles?

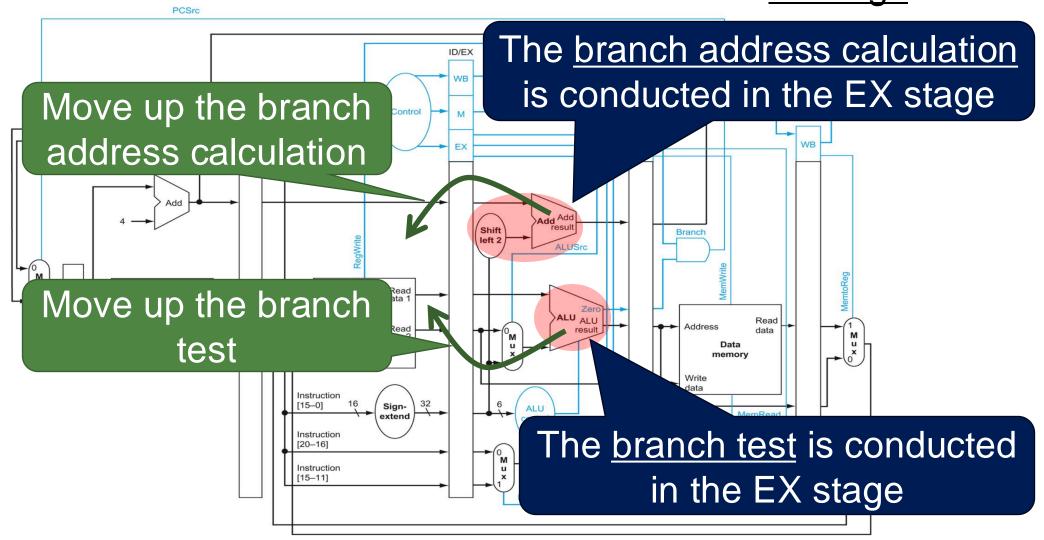
Optimized branch processing!
(H/W modification to achieve
3 cycle penalty → 1 cycle penalty)

Optimized Branch Processing: Observation



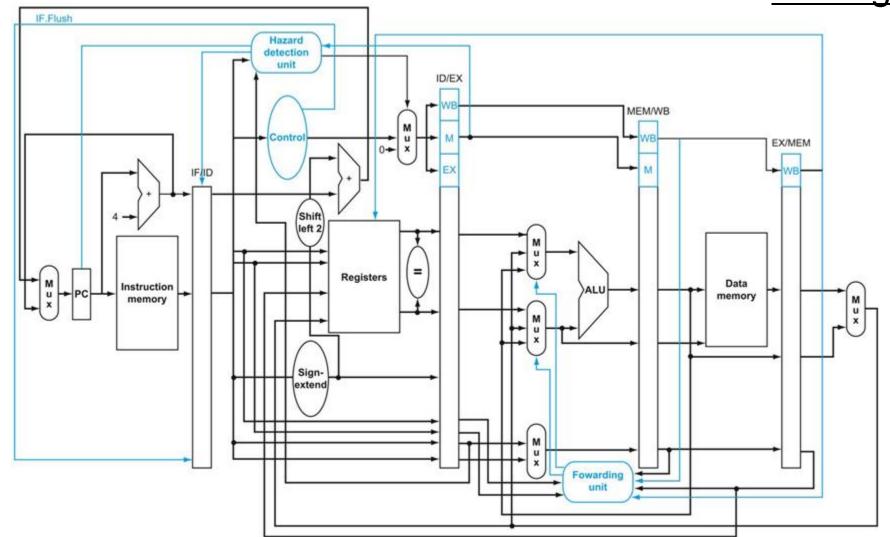
Optimized Branch Processing: Idea

Move hardware to determine branch outcome to the ID stage



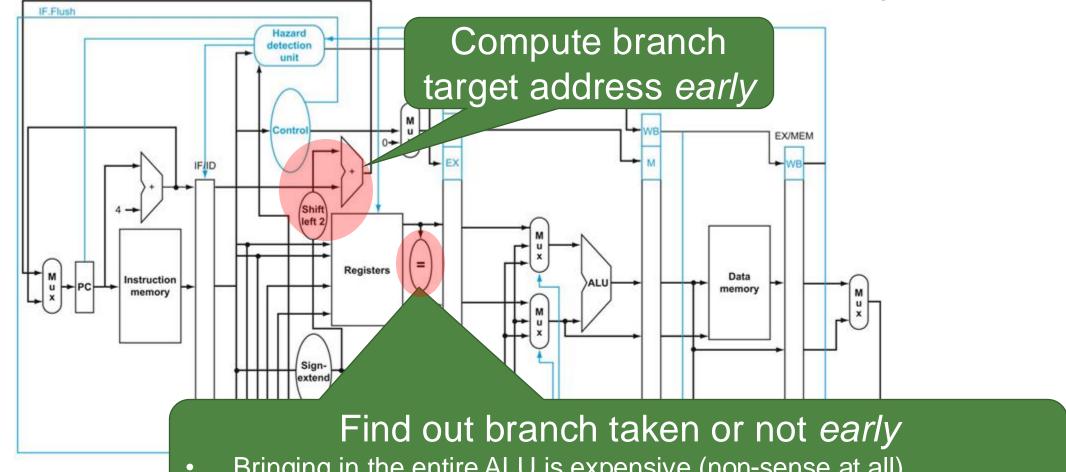
Solution for Control Hazard: Optimized * Branch Processing

Move hardware to determine branch outcome to the ID stage



Solution for Control Hazard: Optimized **Branch Processing**

Move hardware to determine branch outcome to the ID stage



- Bringing in the entire ALU is expensive (non-sense at all)
- Deciding whether to branch can be done with just an equality check

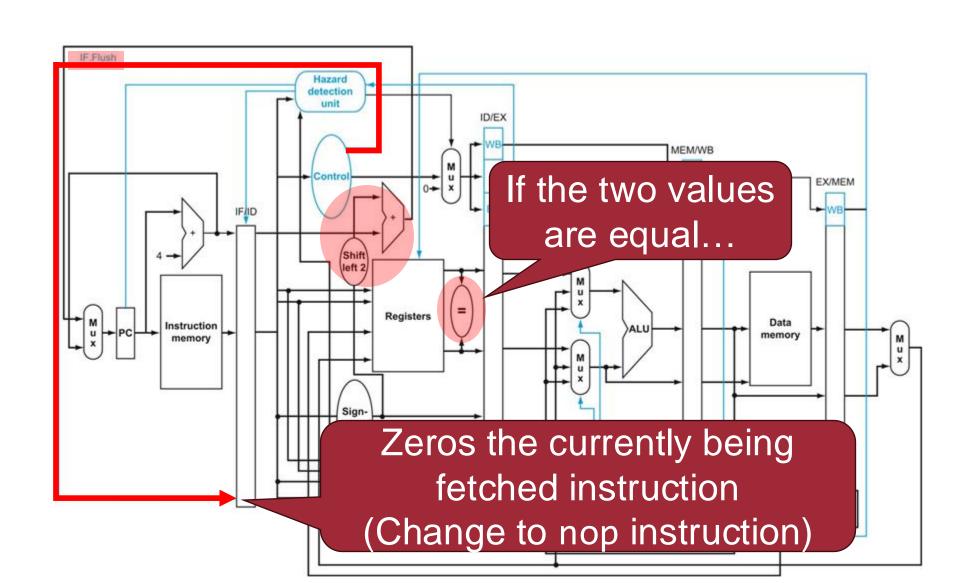
Solution for Control Hazard: Optimized * Branch Processing

Move hardware to determine branch outcome to the ID stage

- -Compute branch target address early
- -Find out branch taken or not early
 - Extra lightweight hardware for equality check
 - Equality can be tested by first XORing inputs and then ORing all the results

Reduce branch delay to one clock cycle (Flush one instruction if the branch is taken)

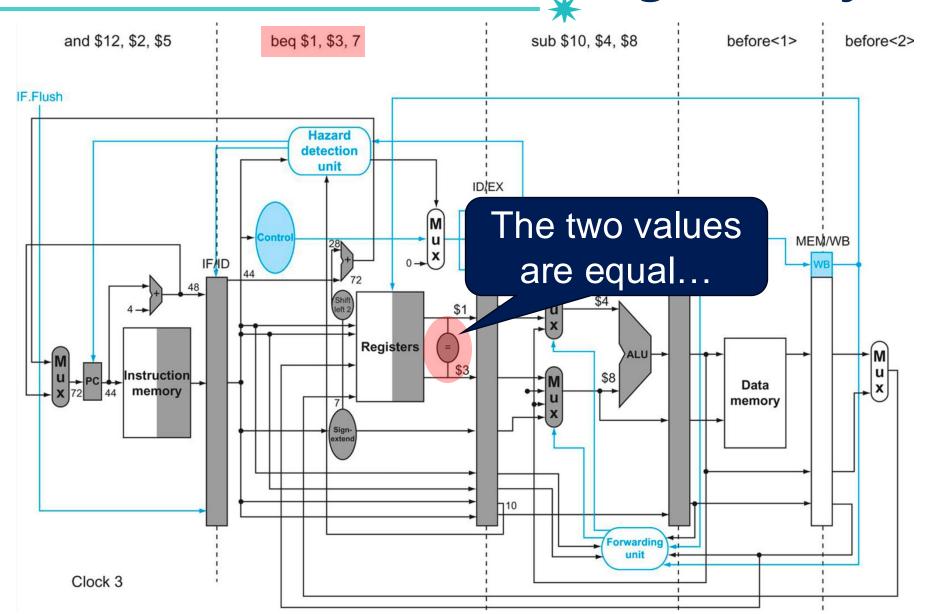
Optimized Branch Processing: Flushing Logic



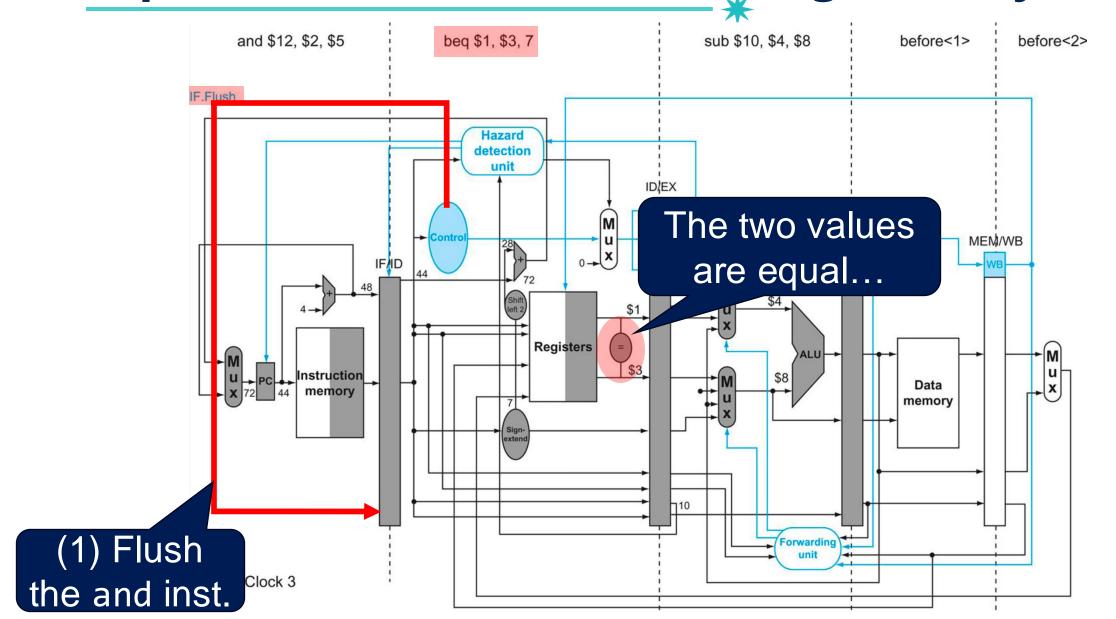
Optimized Branch Processing: Example

```
36: sub $10, $4, $8
40: beq $1, $3, 7
44: and $12, $2, $5
48: or $13, $2, $6
52: add $14, $4, $2
                         Control flow
56: slt $15, $6, $7
        $4, 50($7)
72: lw
```

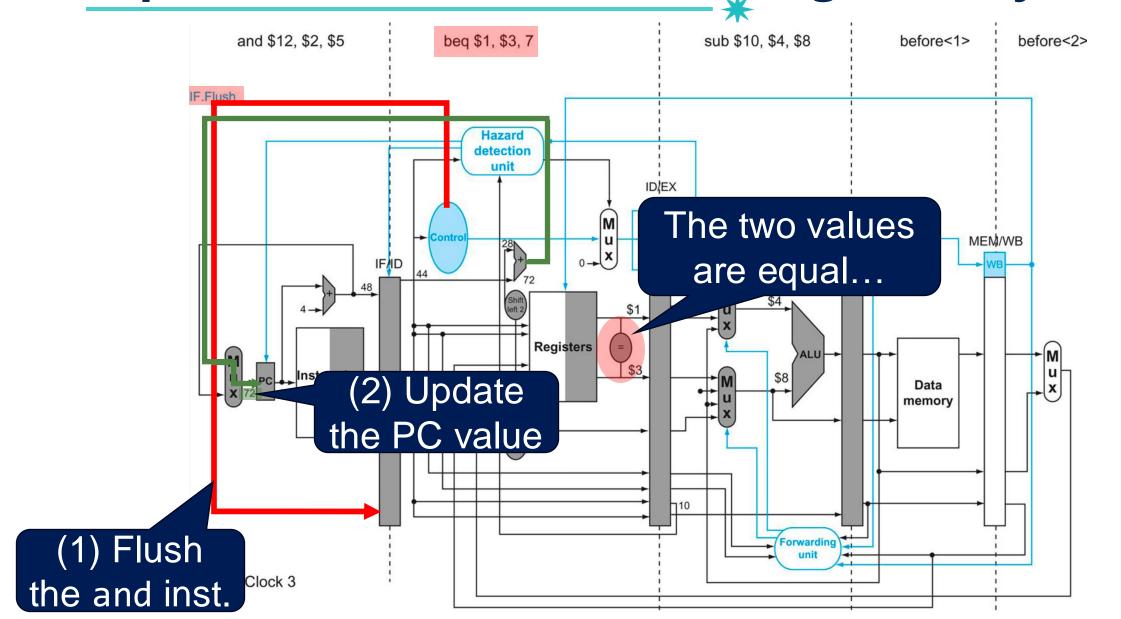
Optimized Branch Processing: 3rd Cycle



Optimized Branch Processing: 3rd Cycle

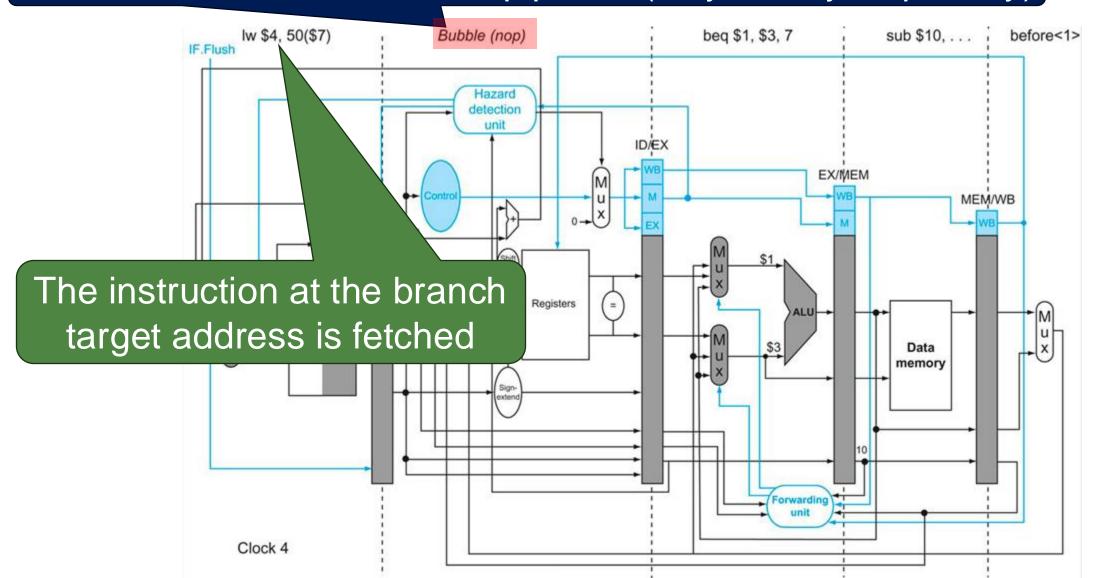


Optimized Branch Processing: 3rd Cycle



Optimized Branch Processing: 4th Cycle

and inst. was discarded from pipeline (only one cycle penalty)



Solution for Control Hazard: Optimized * Branch Processing



If there is a branch instruction, there is still a 1-cycle penalty.

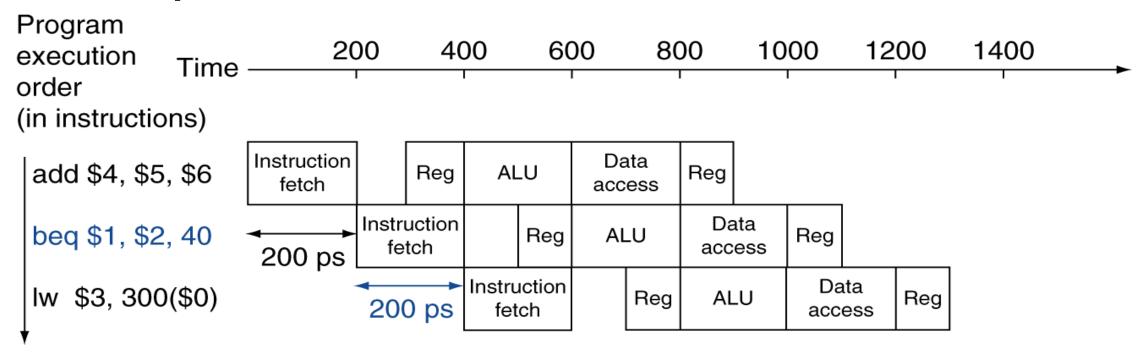
How can performance be further improved?

Branch prediction!

Solution for Control Hazard: Branch Prediction

Guess one direction then backup if wrong

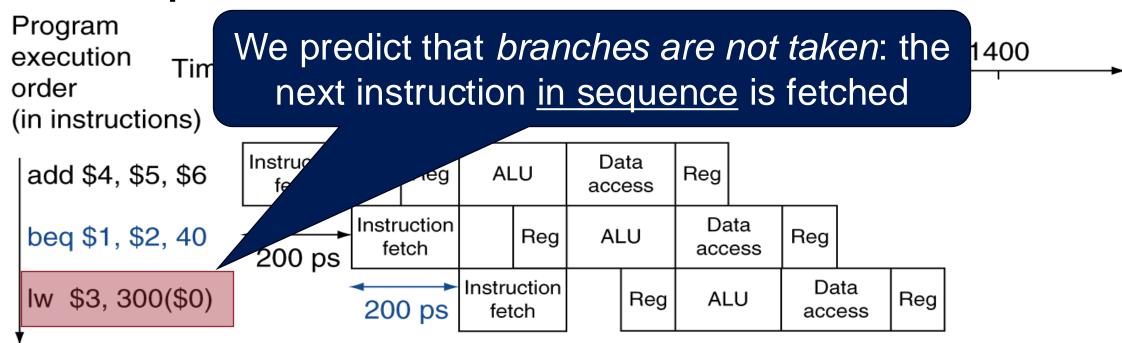
-Our prediction so far: a branch is not taken



Solution for Control Hazard: Branch Prediction

Guess one direction then backup if wrong

-Our prediction so far: a branch is not taken



Prediction No penalty

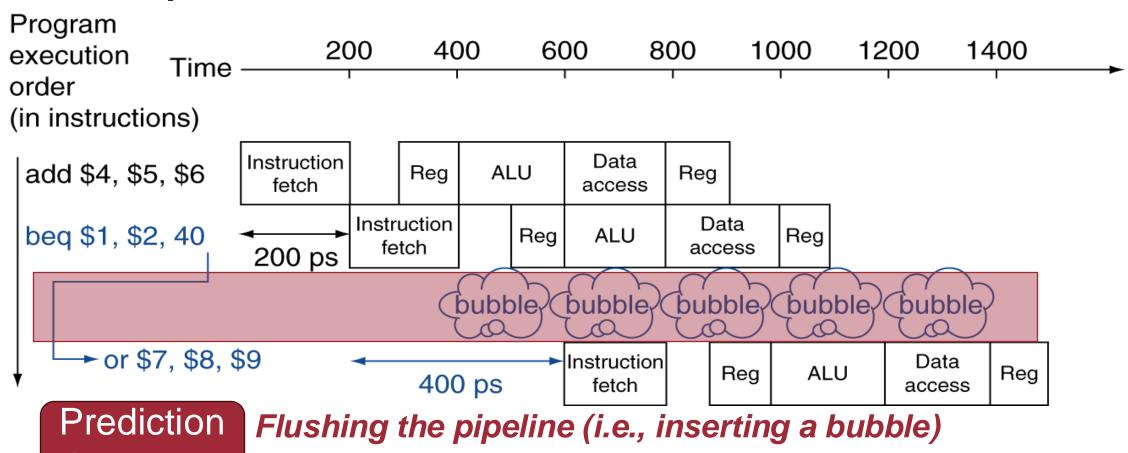
Solution for Control Hazard: Branch Prediction

Guess one direction then backup if wrong

incorrect

-Our prediction so far: a branch is not taken

when we are wrong



Types of Branch Prediction

Guess one direction then backup if wrong

- 1. Static branch prediction
 - A fixed decision based on certain *predefined rules*
 - Rule #1: Predicting Not Taken
 - Rule #2: Predicting *Taken*

Programs with few branches tend to have poor performance

Programs with many branches tend to have poor performance

A simple static prediction scheme will probably waste too much performance

Types of Branch Prediction

Guess one direction then backup if wrong

1. Static branch prediction

- A fixed decision based on certain <u>predefined rules</u>
- Rule #1: Predicting Not Taken
- Rule #2: Predicting Taken

2. Dynamic branch prediction

- Prediction of branches at runtime using runtime information

Dynamic Branch Prediction: Observation

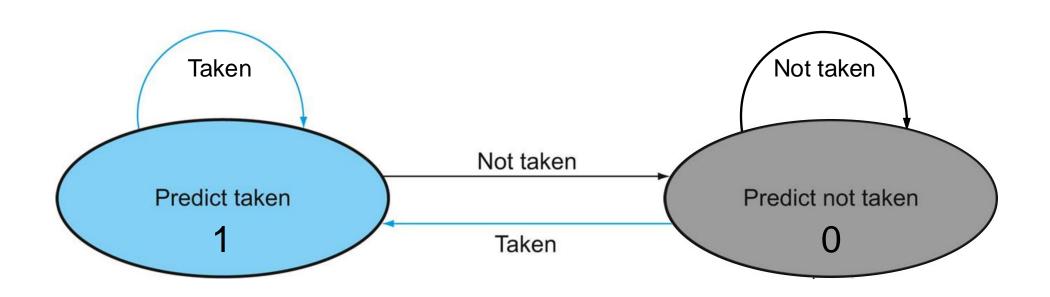
Branch history is important: if a branch was taken previously, it is likely that the branch will be taken again the next time

Dynamic Branch Prediction: Basic Idea



The memory contains one or more bits indicating whether the branch was recently taken or not

Dynamic Branch Prediction: 1-bit Prediction

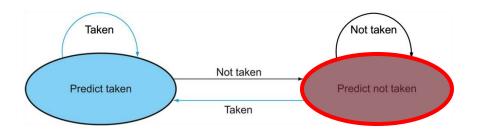


Dynamic Branch Prediction: Initial State

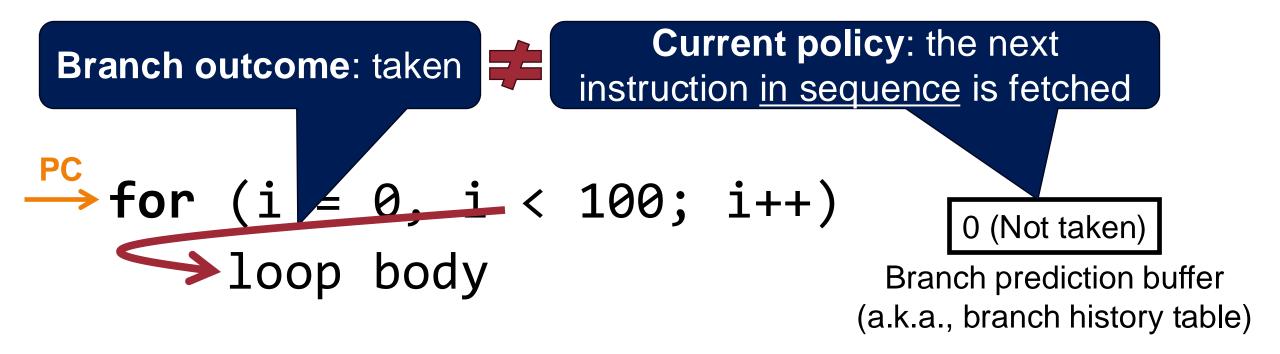
Current policy: the next instruction in sequence is fetched

0 (Not taken)

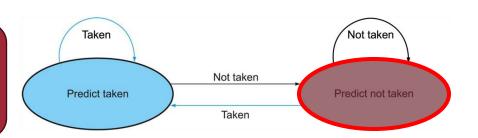
Branch prediction buffer (a.k.a., branch history table)



Dynamic Branch Prediction: 1st Iteration 98

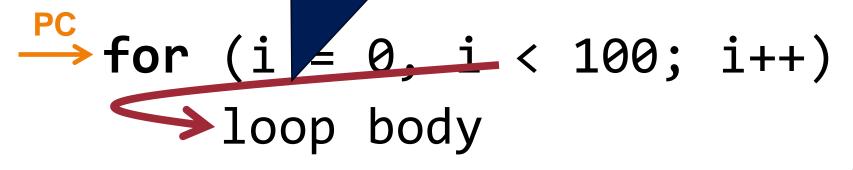


Mispredict cycle penalty)



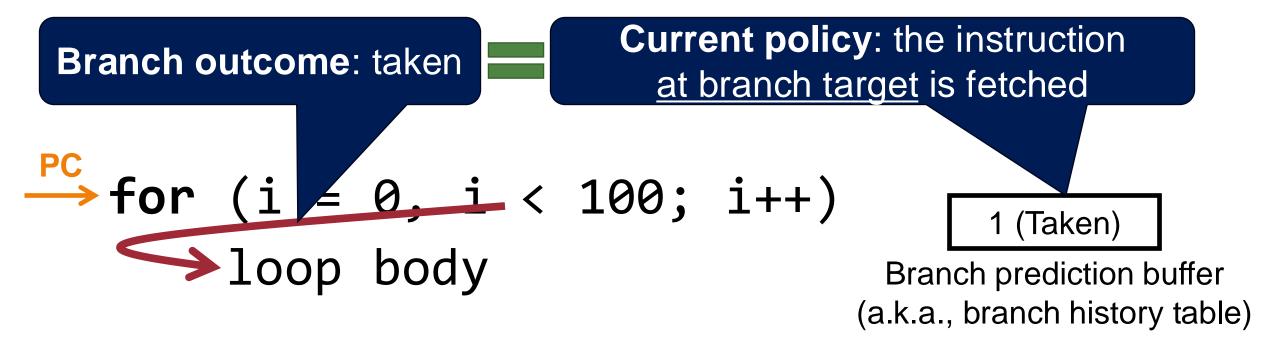
Dynamic Branch Prediction: Update Buffer

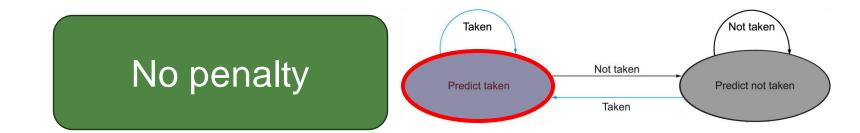
Branch outcome: taken



(Taken) Branch prediction buffer (a.k.a., branch history table) Not taken Taken Not taken Predict not taken

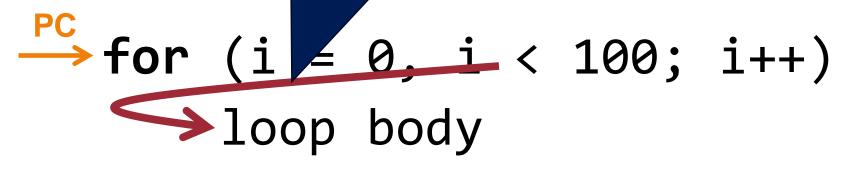
Dynamic Branch Prediction: 2nd Iteration 95





Dynamic Branch Prediction: Update Buffer

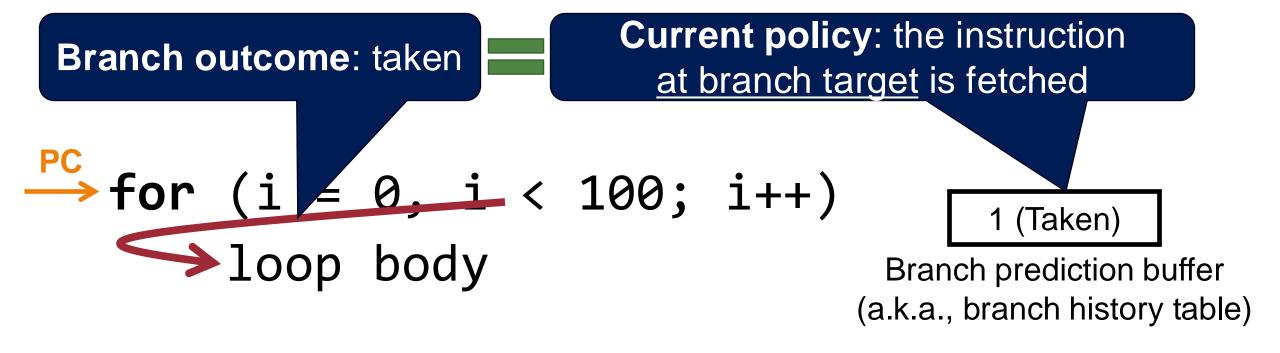
Branch outcome: taken

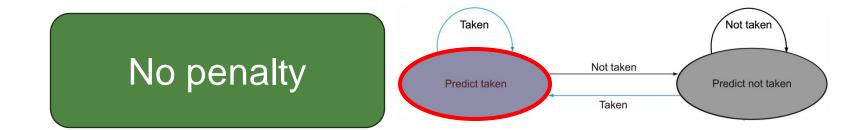


(Taken) Branch prediction buffer (a.k.a., branch history table) No change Not taken Not taken Predict not taken Taken

Dynamic Branch Prediction: 3rd Iteration 97





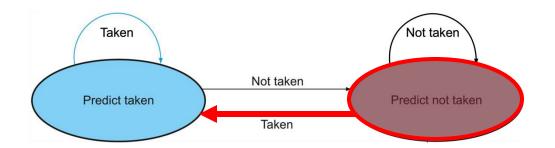


Dynamic Branch Prediction

- Prediction of branches <u>at runtime</u> using <u>runtime information</u>
- Use a branch prediction buffer (a.k.a., branch history table)
 - A bit says whether the branch was recently taken or not
 - If an instruction is predicted to be taken, fetching begins from the branch target
- Ideally, each branch instruction would have its own branch prediction buffer
 - However, due to the performance issue, the buffer table size is limited
 - Therefore, multiple branch instructions can share the same buffer
 - Indexed using the lower bits of the branch instruction's address

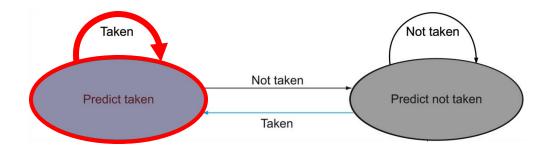


Initial state

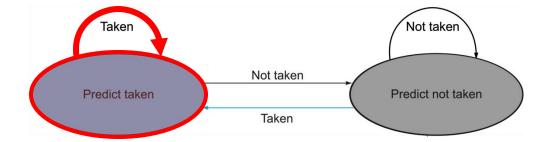


for	(i	= 0	, i	<	16	90;	; i-	++)/	
	for	(j	=	0,	j	<	3;	j++	-)
		100	ор	bo	dy				

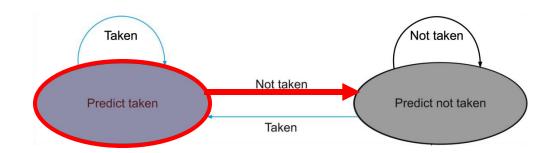
Time	Buffer State	Outcome	Result?
1	Not taken	Taken	Wrong
2	Taken	Taken	Correct



Time Buffer State		Outcome	Result?	
1	Not taken	Taken	Wrong	
2	Taken	Taken	Correct	
3	Taken	Taken	Correct	

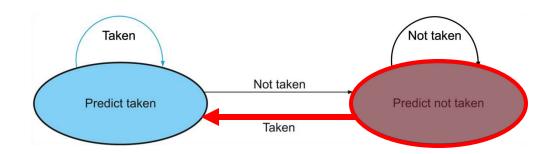


Time	Buffer State	Outcome	Result?
1	Not taken	Taken	Wrong
2	Taken	Taken	Correct
3	Taken	Taken	Correct
4	Taken	Not taken	Wrong



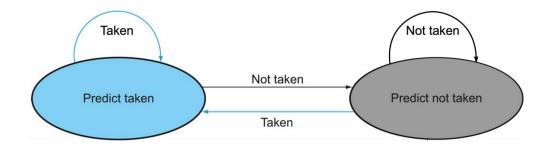
Mispredict as taken on last iteration of inner loop

	Time	Buffer State	Outcome	Result?
	1	Not taken	Taken	Wrong
	2	Taken	Taken	Correct
	3 Taken		Taken	Correct
4		Taken	Not taken	Wrong
	5	Not taken	Taken	Wrong
				·



Then mispredict as not taken on first iteration of inner loop next time around

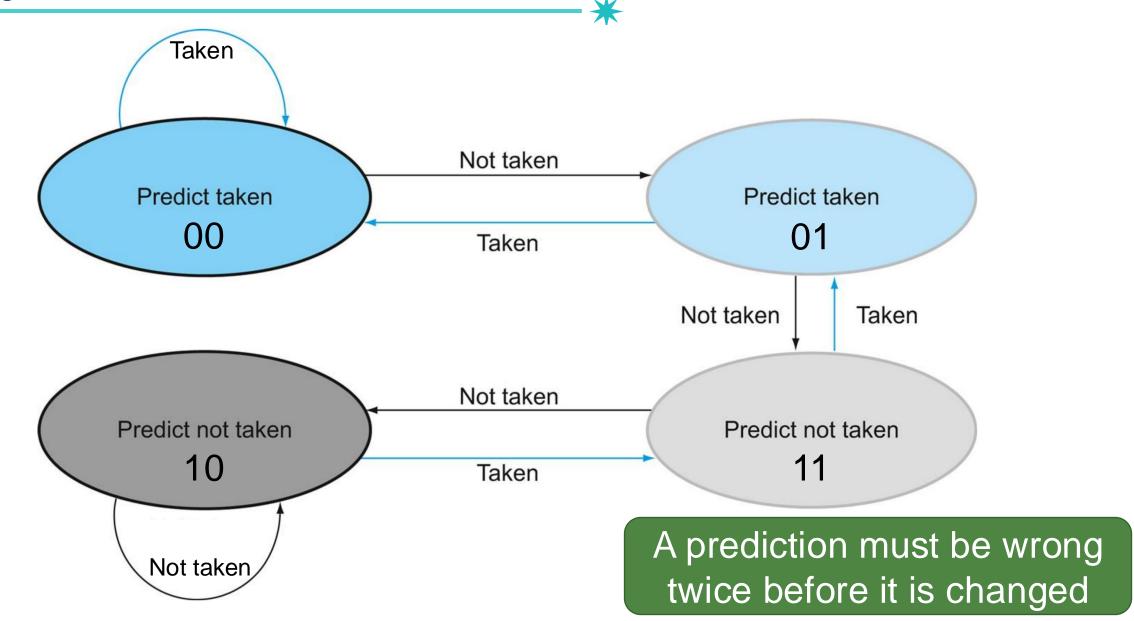
Problem: Inner loop branches mispredicted twice!



	Time	Buffer State	Outcome	Result?
		Not taken	Taken	Wrong
	2	Taken	Taken	Correct
	3	Taken	Taken	Correct
	4	Taken	Not taken	Wrong
	5	Not taken	Taken	Wrong
	6	Taken	Taken	Correct
•	7	Taken	Taken	Correct
	8	Taken	Not taken	Wrong
	9	Not taken	Taken	Wrong
	10	Taken	Taken	Correct
	11	Taken	Taken	Correct
	12	Taken	Not taken	Wrong







Dynamic Branch Prediction: 2-bit Prediction

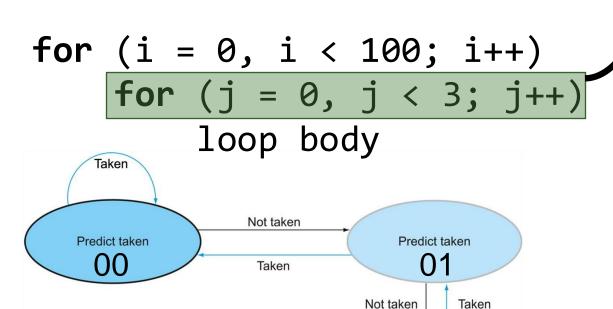
10

One misprediction each loop execution

Predict not taken

Initial state

Time	Buffer St	é	Outcome	Result?
1	Not taken (10)	Taken	Wrong



Not taken

Taken

Predict not taken

Not taken

Dynamic Branch Prediction: 2-bit Prediction

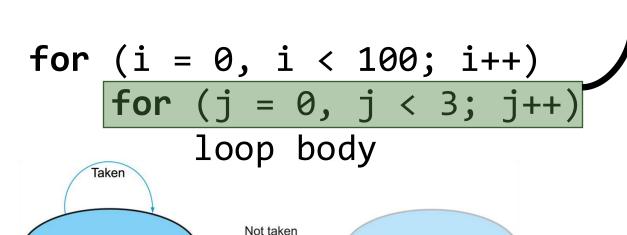
One misprediction each loop execution

Predict taken

Predict not taken

Taken

Not taken



Taken

Not taken

Taken

Predict taken

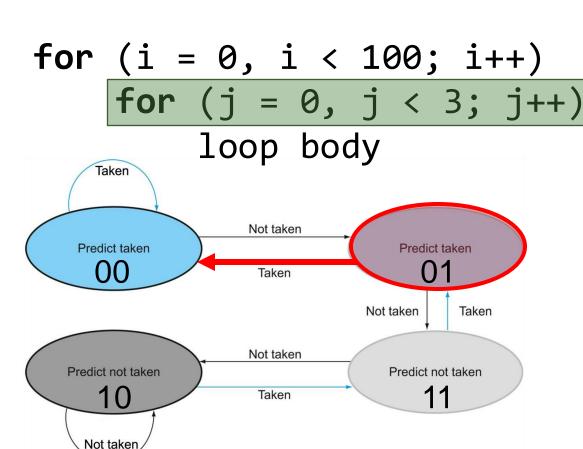
00

Predict not taken

Not taken

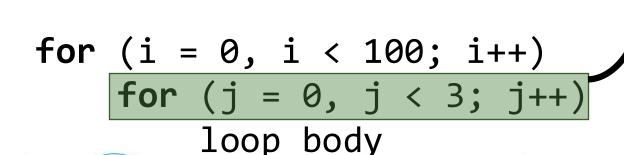
Time	Buffer State	Outcome	Result?
1	Not taken (10)	Taken	Wrong
2	Not taken (11)	Taken	Wrong

Dynamic Branch Prediction: 2-bit Prediction



Time	Buffer State	Outcome	Result?
1	Not taken (10)	Taken	Wrong
2	Not taken (11)	Taken	Wrong
3	Taken (01)	Taken	Correct

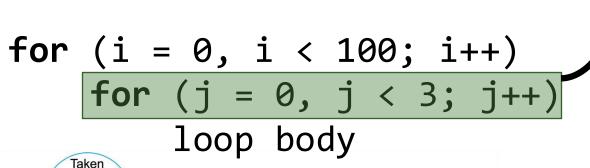
Dynamic Branch Prediction: 2-bit Prediction



*	Time	Buffer State	Outcome	Result?
	1	Not taken (10)	Taken	Wrong
	2	Not taken (11)	Taken	Wrong
	3	Taken (01)	Taken	Correct
	4	Taken (00)	Not taken	Wrong

Taken	•	
	Not taken	
Predict taken		Predict taken
00	Taken	01
		Not taken Taken
	Not taken	
Predict not taken	, tot tallon	Predict not taken
10	Taken	11
Not taken		

Dynamic Branch Prediction: 2-bit Prediction



>	Time	Buffer State	Outcome	Result?
	1	Not taken (10)	Taken	Wrong
	2	Not taken (11)	Taken	Wrong
	3	Taken (01)	Taken	Correct
	4	Taken (00)	Not taken	Wrong
	5	Taken (01)	Taken	Correct

Predict taken	Not taken Taken	Predict t	aken
		Not taken	Taken
Predict not taken 10 Not taken	Not taken Taken	Predict no 1	

Dynamic Branch Prediction: 2-bit Prediction

One mis

Taken		
	Not taken	
Predict taken	Taken	Predict taken 01
00	Takon	1
		Not taken Taken
	Not taken	
Predict not taken		Predict not taken
10	Taken	11
Not taken		

Time	Buffer State	Outcome	Result?
1	Not taken (10)	Taken	Wrong
2	Not taken (11)	Taken	Wrong
3	Taken (01)	Taken	Correct
4	Taken (00)	Not taken	Wrong
5	Taken (01)	Taken	Correct
0	Taken (00)	Taken	Correct
7	Taken (00)	Taken	Correct
8	Taken (00)	Not taken	Wrong
9	Taken (01)	Taken	Correct
10	Taken (00)	Taken	Correct
11	Taken (00)	Taken	Correct
12	Taken (00)	Not taken	Wrong

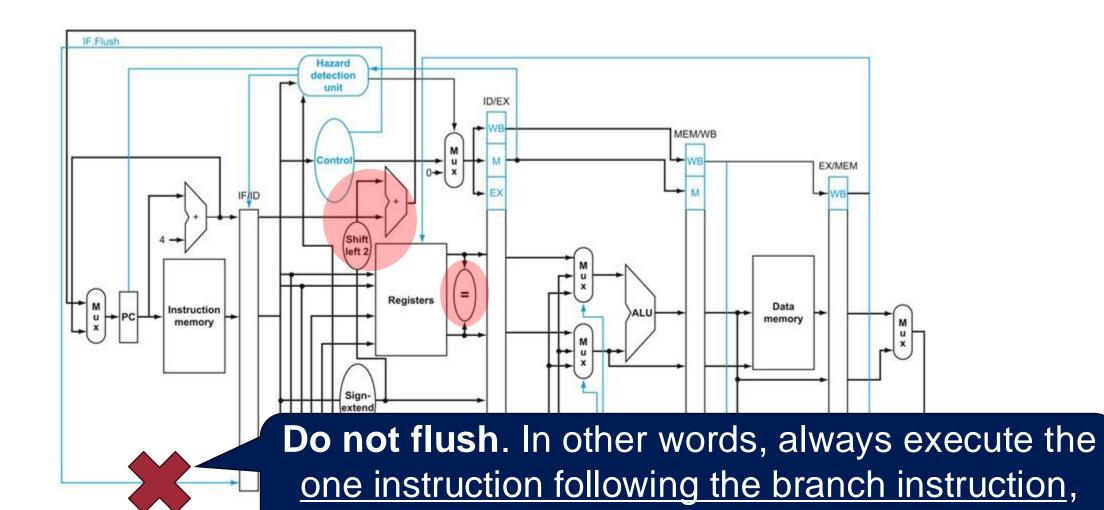
Solution for Control Hazard: Branch Prediction



Prediction is effective, but we don't have a hardware expert. In this situation, is there any way to resolve the control hazard in software manner?

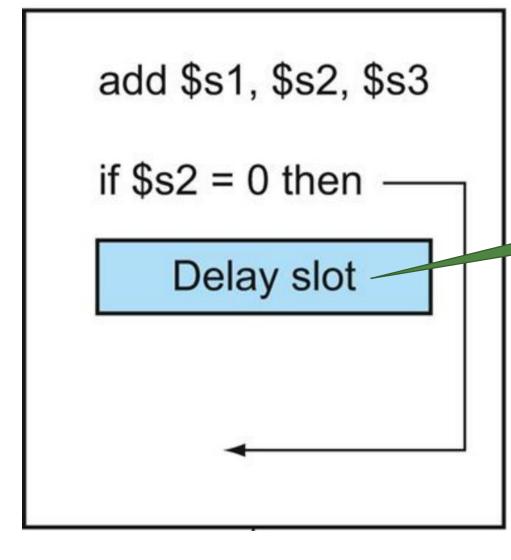
Delayed Branch! (Compiler-driven optimization)

Our Assumption on HW: Do not Flush



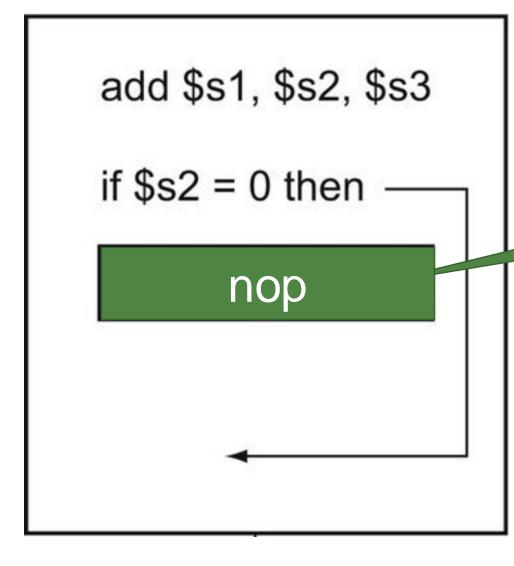
regardless of whether the branch is taken or not

Solution for Control Hazard: Delayed Branch



Compiler insert the **branch delay slot** with instruction that is not affected by the branch

Solution for Control Hazard: Delayed Branch

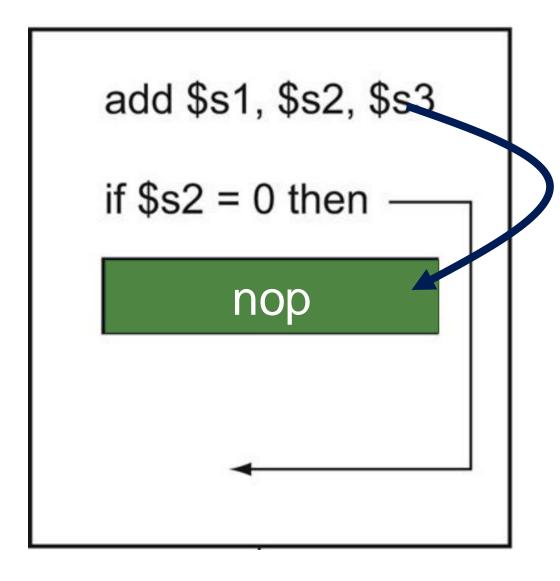


Compiler insert the **branch delay slot** with instruction that is not affected by the branch

The simplest method is to insert a nop instruction (1 cycle penalty)

Solution for Control Hazard: Delayed Branck

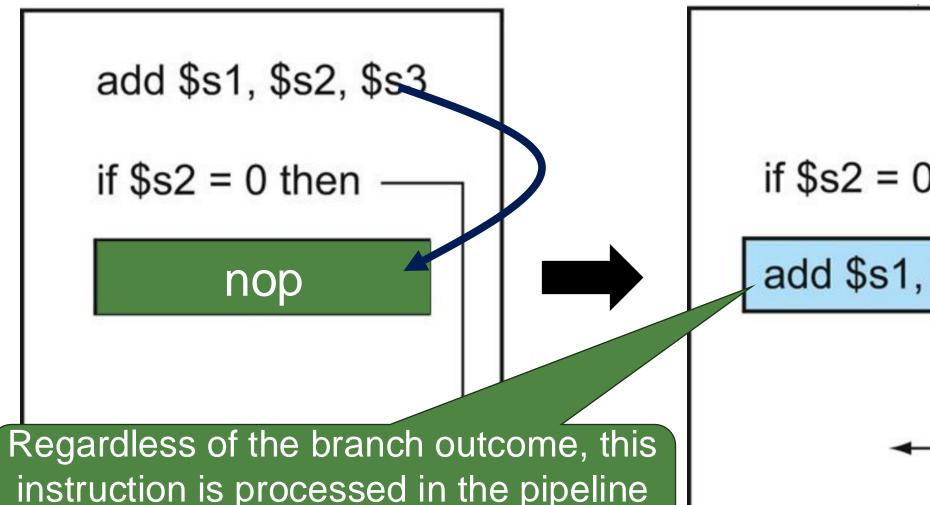
(with advanced compiler)



Code reordering (move this instruction, which is not affected by the branch)

Solution for Control Hazard: Delayed Branck

(with advanced compiler)

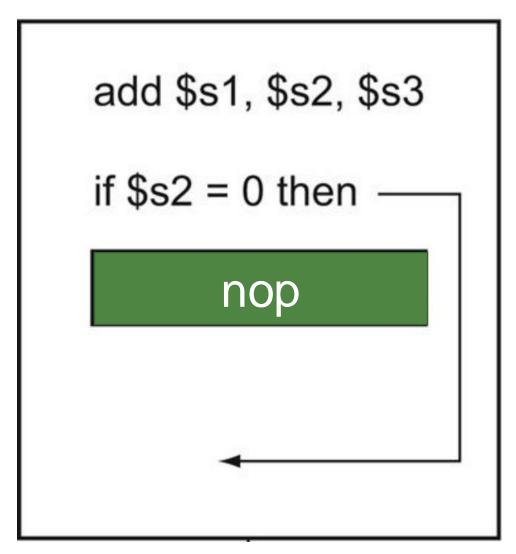


(no penalty)

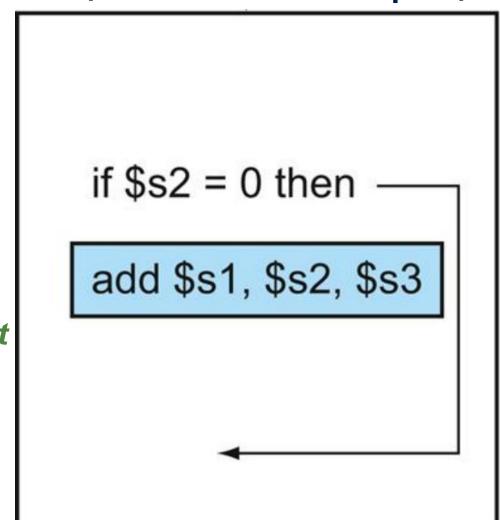
if \$s2 = 0 then add \$s1, \$s2, \$s3

Solution for Control Hazard: Delayed Branck

(with advanced compiler)



Equivalent meaning







Situations that prevent starting the next instruction in the next cycle

Hazard #1: Structural hazard

Conflict for use of a hardware resource

Solution:

- Stall
- Resource duplication

Hazard #2: Data hazard

An instruction cannot execute because data is not yet available

Solution:

- Stall
- Forwarding
- Compiler optimization

Hazard #3: Control hazard

The next instruction is uncertain due to branching

Solution:

- Stall
- Optimized branch processing
- Branch prediction
- Delayed branch

Question?