


CSE261: Computer Architecture

20. Parallel Architectures

Seongil Wi

Uniprocessor Performance



Constrained by power
(and long memory latency)

How we address the power wall?


⇒ Parallel architectures

Parallel Processors



- Goal: connecting multiple computers to get higher performance
+ availability, power efficiency...

Uniprocessor to Multiprocessors



Uniprocessor: one processor (core) per chip

Uniprocessor to Multiprocessors



Control
(brain)

Datapath

Processor

Uniprocessor to Multiprocessors



Control
(brain)

Datapath

Processor #1

Control
(brain)

Datapath

Processor #3

Control
(brain)

Datapath

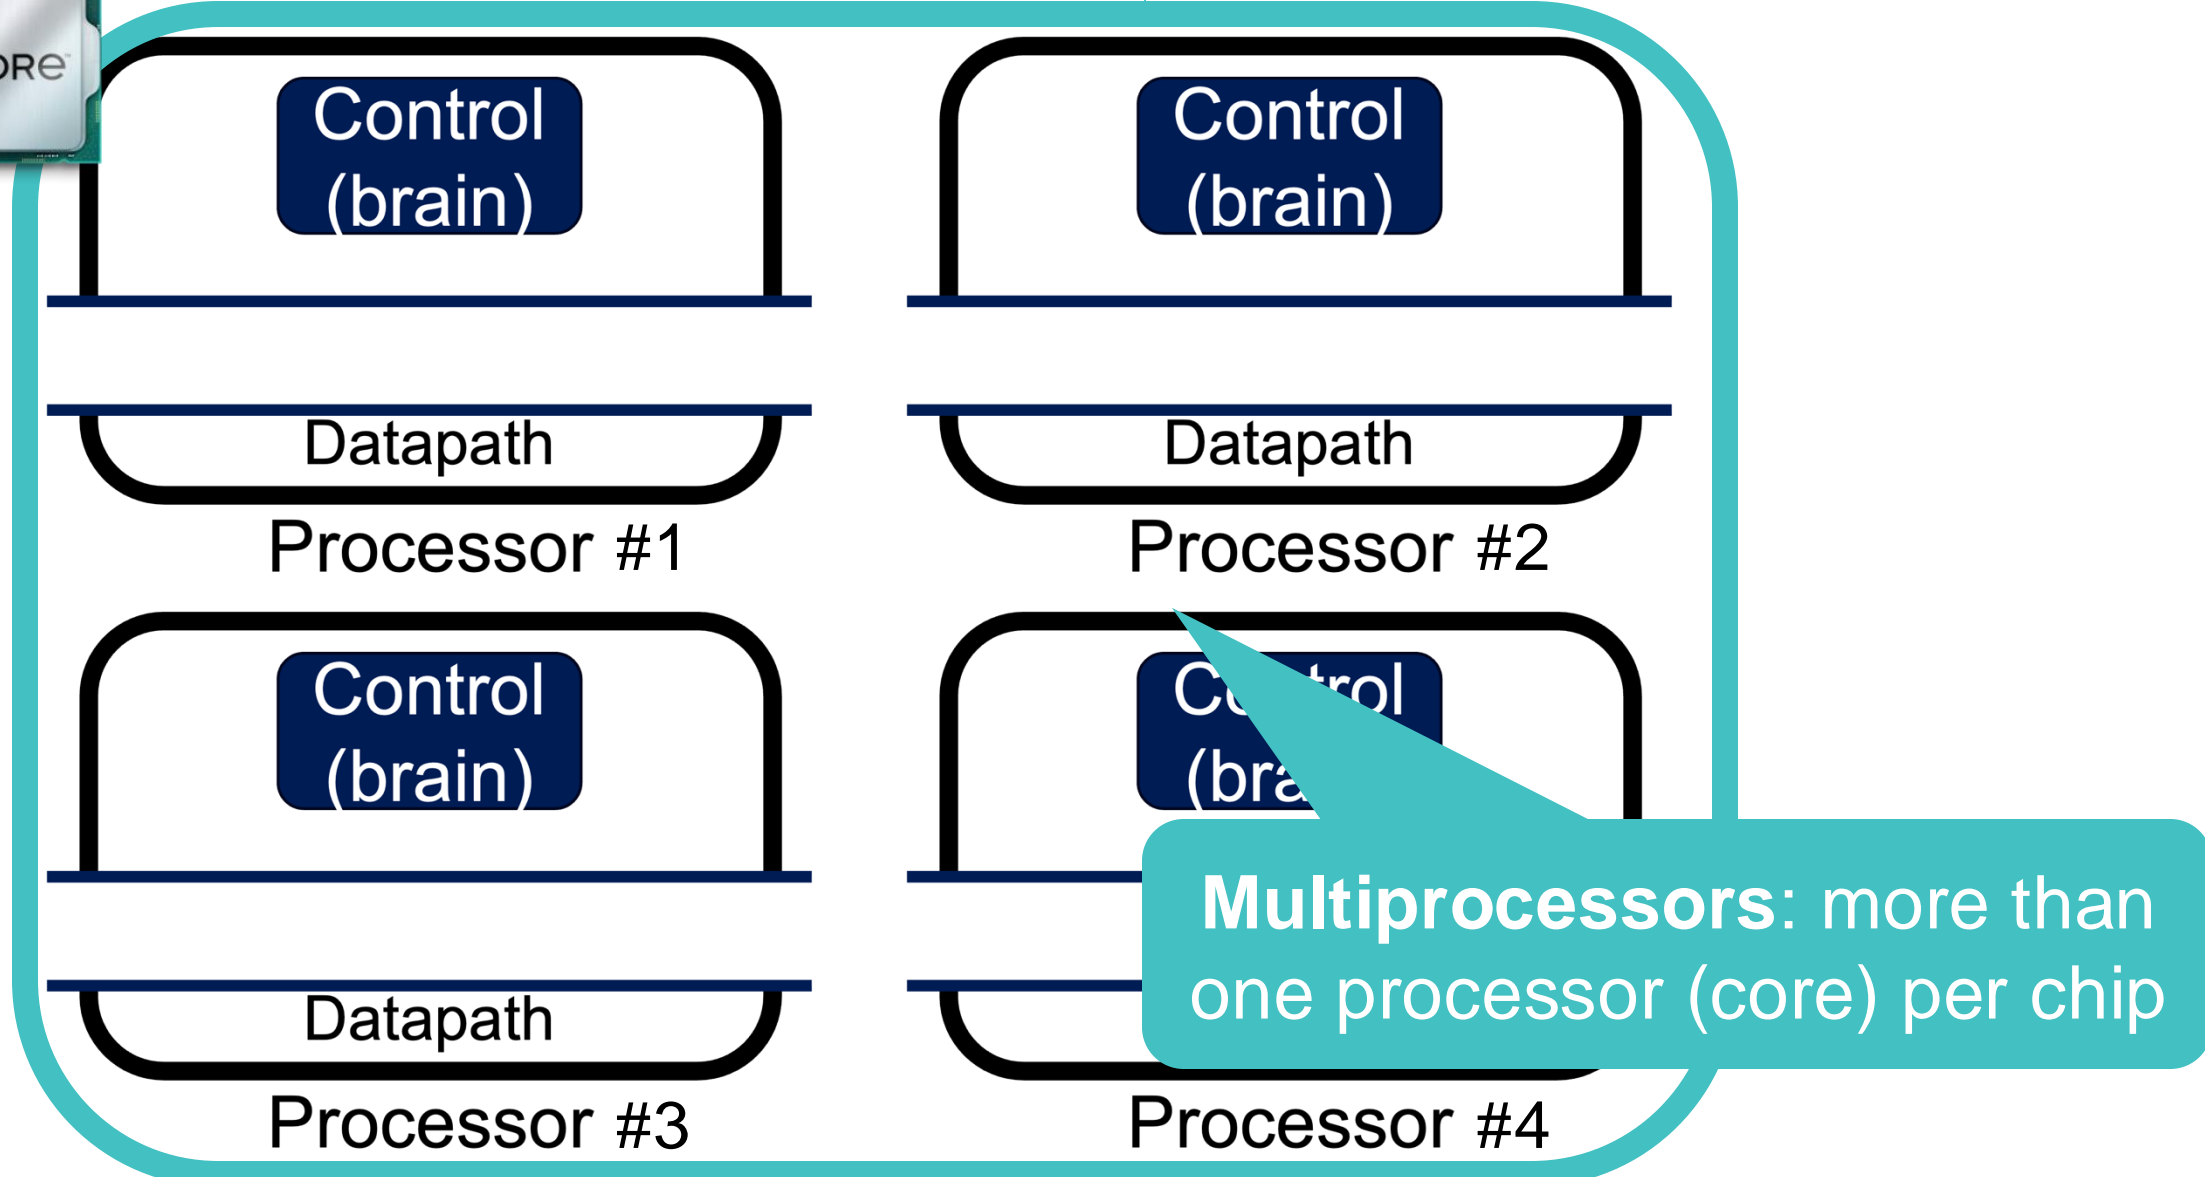
Processor #2

Control
(brain)

Datapath

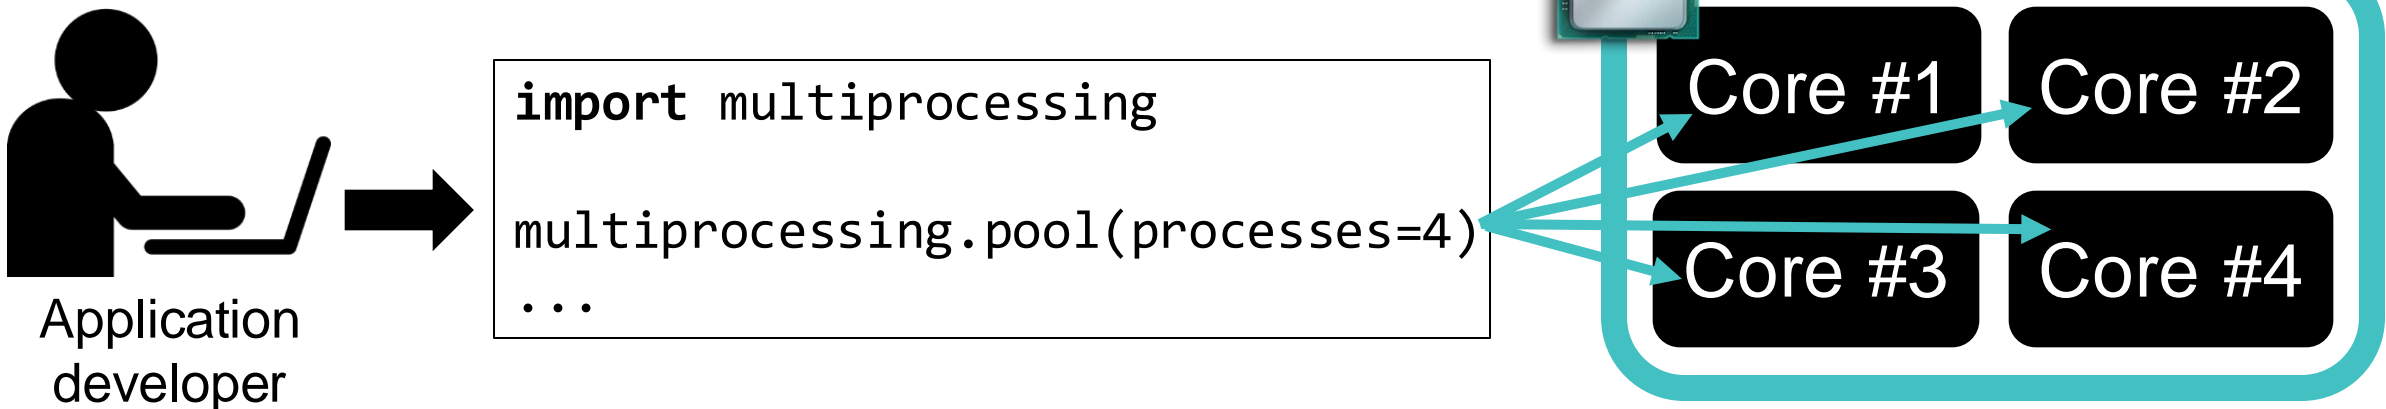
Processor #4

Uniprocessor to Multiprocessors



Multiprocessors

- Multicore microprocessors
 - More than one processor (core) per chip
- Requires explicitly **parallel programming**
 - Programming for performance
 - Load balancing
 - Optimizing communication and synchronization



The Challenge with Parallelism

- The challenge is not the hardware, but the **software**
- Why is it difficult to write parallel processing programs (that are fast)?
 - **Partitioning**: the task must be broken into eight equal-sized pieces
 - **Coordination**: synchronize between tasks
 - **Communication overhead**: spend time to communicate each other
 - **Amdahl's law**: sequential part can limit speed up!



Application
developer

```
import multiprocessing
multiprocessing.pool(processes=4)
...
```



The challenge is
getting worse with
more processors

Pitfall: Amdahl's Law (Amdahl's Argument)




"The overall performance improvement gained by optimizing a single part of a system is limited by the fraction of time that the improved part is actually used"

The **speedup** of parallel computations can be estimated:

$$\text{Speedup} = \frac{1}{(1-P) + \frac{P}{n}}$$

P: parallel
portion

n: # of
processor cores



Gene Myron Amdahl
(1922-2015)

Pitfall: Amdahl's Law (Amdahl's Argument)

"The overall performance improvement gained by optimizing a single part of a system is limited by the fraction of time that the improved part is actually used"

The **speedup** of parallel computations can be estimated:

$$\text{Speedup} = \frac{1}{(1-P) + \frac{P}{n}}$$

1-P: time spent executing the sequential portion


P: parallel portion

P/n: time spent executing the parallel portion

n: # of processor cores

Pitfall: Amdahl's Law

14



Reducing the sequential portion

Instruction and Data Streams



- According to the Flynn's taxonomy of computers, ...
 - Mike Flynn, "Very High-Speed Computing Systems", Proc. Of IEEE, 1966

Very High-Speed Computing Systems

MICHAEL J. FLYNN, MEMBER, IEEE

Abstract—Very high-speed computers may be classified as follows:

- 1) Single Instruction Stream—Single Data Stream (SISD)
- 2) Single Instruction Stream—Multiple Data Stream (SIMD)
- 3) Multiple Instruction Stream—Single Data Stream (MISD)
- 4) Multiple Instruction Stream—Multiple Data Stream (MIMD).

"Stream," as used here, refers to the sequence of data or instructions as seen by the machine during the execution of a program.

The constituents of a system: storage, execution, and instruction handling (branching) are discussed with regard to recent developments and/or systems limitations. The constituents are discussed in terms of concurrent SISD

systems (CDC 6600 series and, in particular, IBM Model 90 series), since multiple stream organizations usually do not require any more elaborate components.

Representative organizations are selected from each class and the arrangement of the constituents is shown.

INTRODUCTION

MANY SIGNIFICANT scientific problems require the use of prodigious amounts of computing time. In order to handle these problems adequately, the large-scale scientific computer has been developed. This computer addresses itself to a class of problems characterized by having a high ratio of computing requirement to input/output requirements (essentially, the fast situation

Manuscript received June 30, 1966; revised August 16, 1966. This work was performed under the auspices of the U. S. Atomic Energy Commission.

The author is with Northwestern University, Evanston, Ill., and

Instruction and Data Streams



- According to the Flynn's taxonomy of computers, ...

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345

Instruction and Data Streams



- According to the Flynn's taxonomy of computers, ...

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345

Instruction and Data Streams

- According to the Flynn's taxonomy of computers, ...

		Data Stream	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345

Single Instruction
Single Data stream

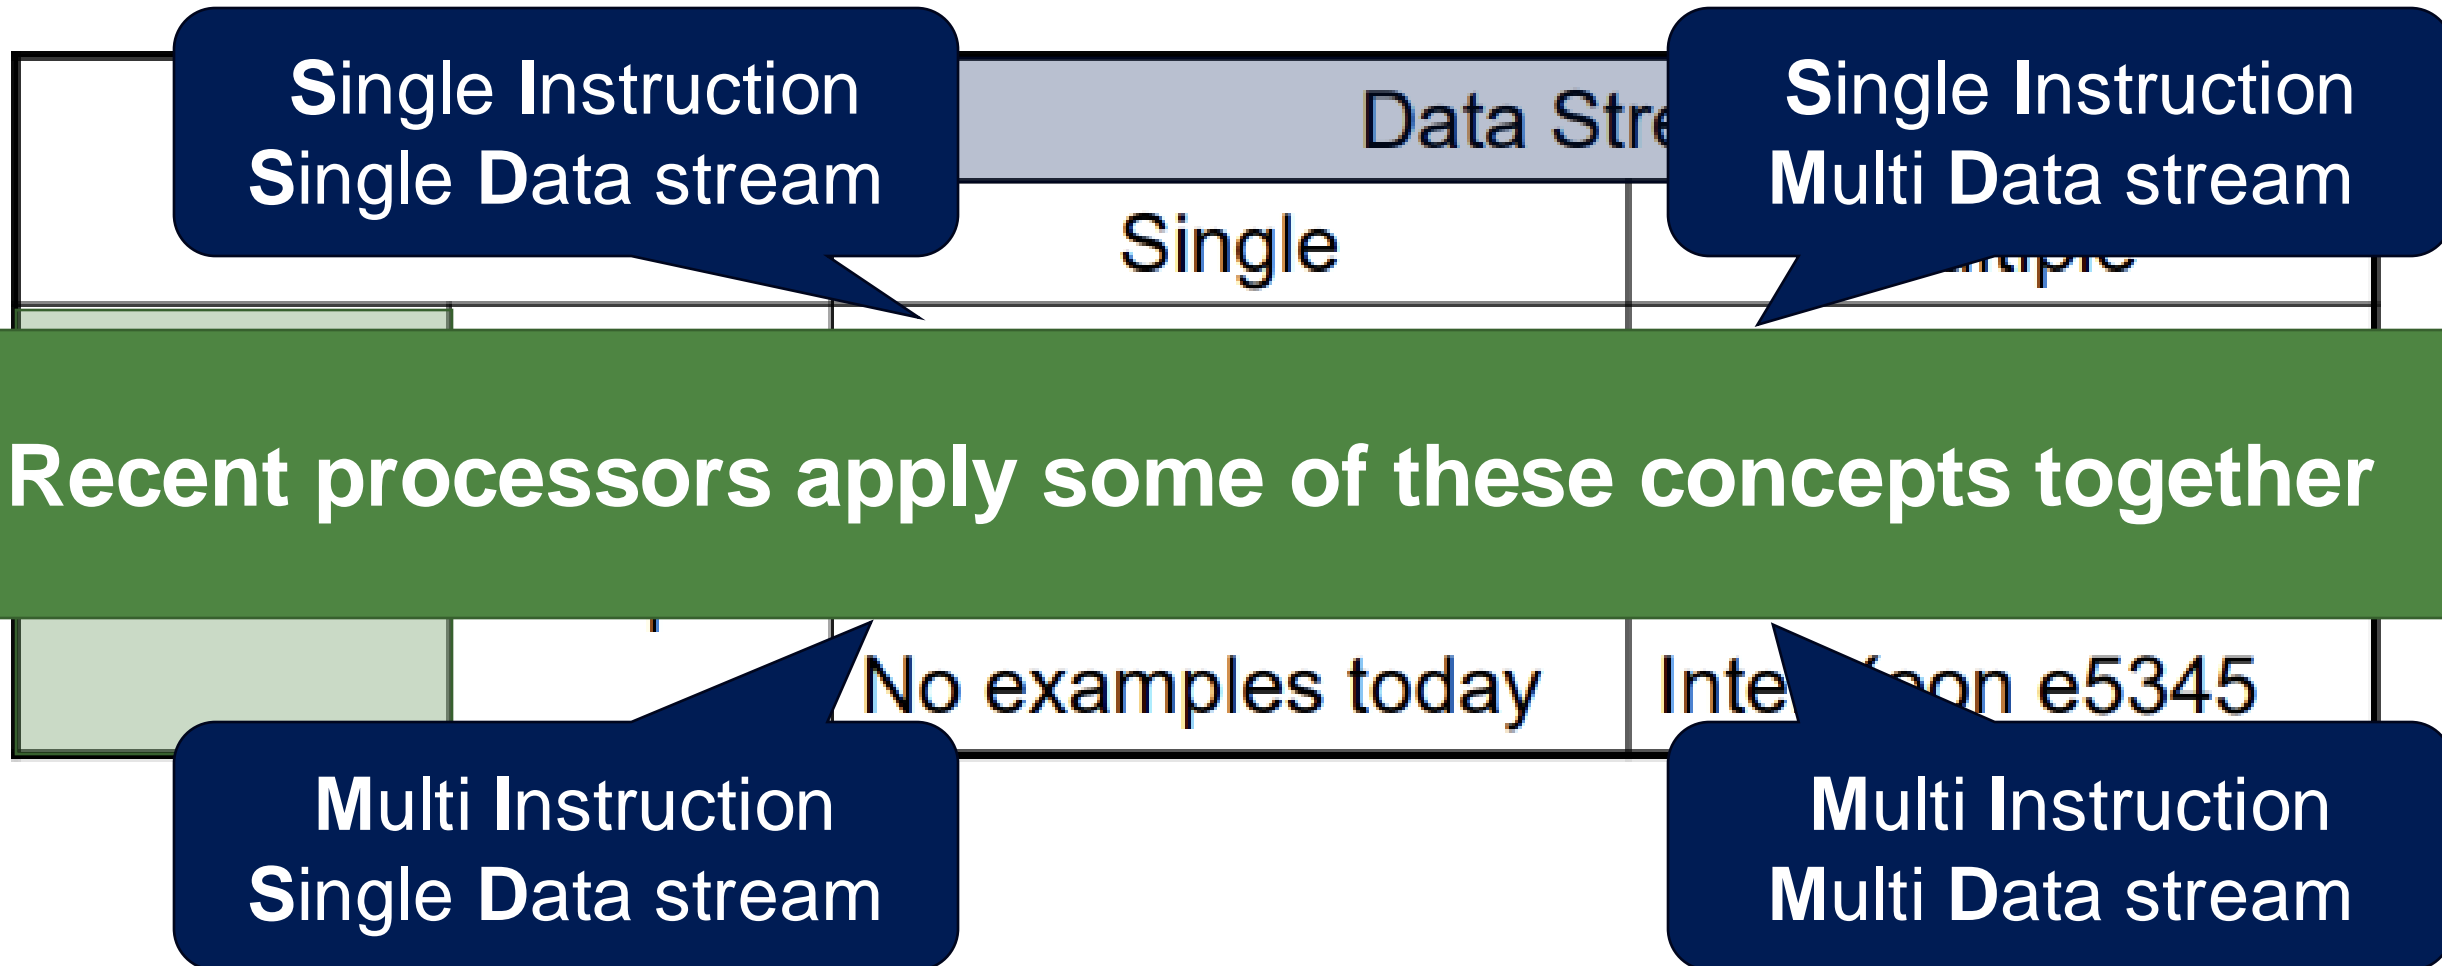
Single Instruction
Multi Data stream

Multi Instruction
Single Data stream

Multi Instruction
Multi Data stream

Instruction and Data Streams

- According to the Flynn's taxonomy of computers, ...



Instruction and Data Streams

- According to the Flynn's taxonomy of computers, ...

		Data Stream	
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345

Single Instruction
Single Data stream


Single Instruction
Multi Data stream

Multi Instruction
Single Data stream

Multi Instruction
Multi Data stream

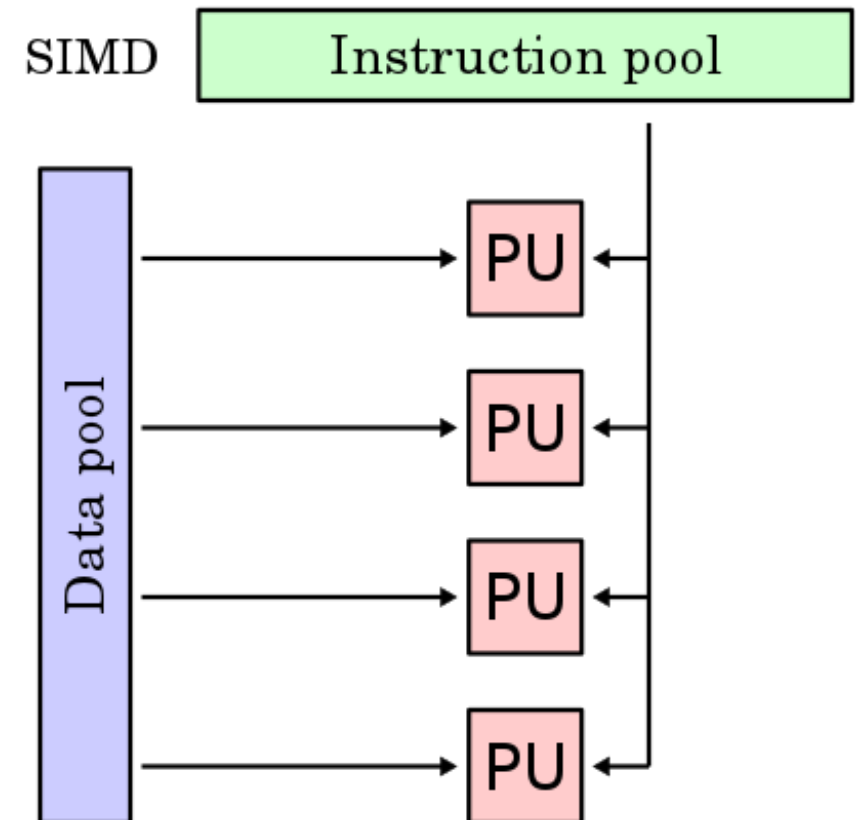
Single Instruction Single Data stream (SISD)

- Uniprocessor
- Disadvantages
 - Limited speed due to being a single core
 - Pipelining can be implemented, but only one instruction will be executed at a time



Single Instruction Multi Data stream (SIMD)

- A single instruction is executed on multiple different pieces of data

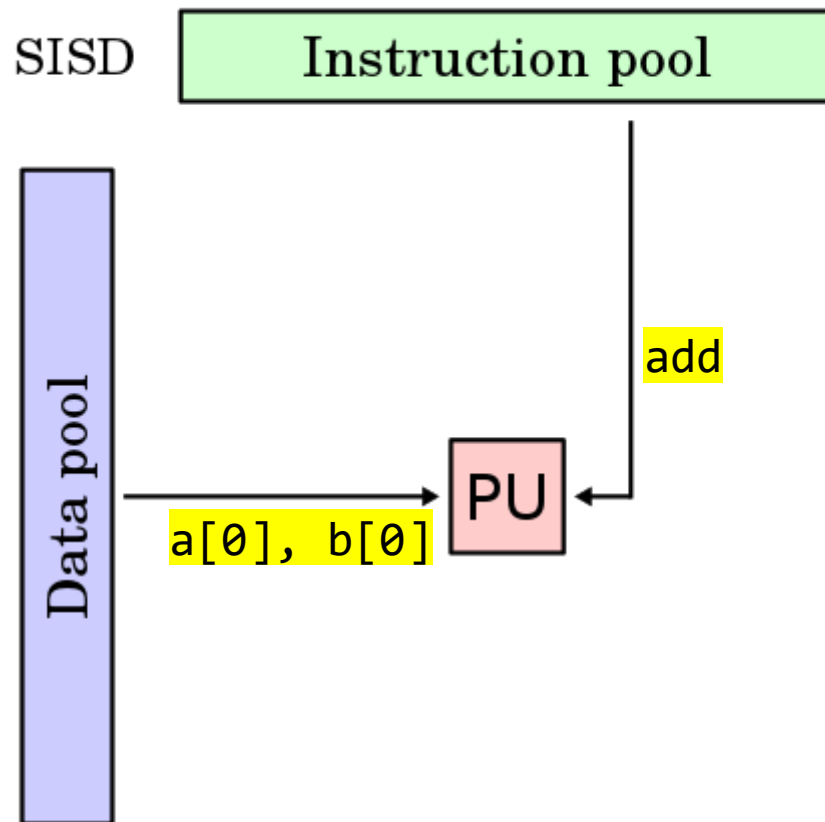


SIMD: Motivating Example

```
for (int i = 0; i < 4; i++) {  
    c[i] = a[i] + b[i];  
}
```

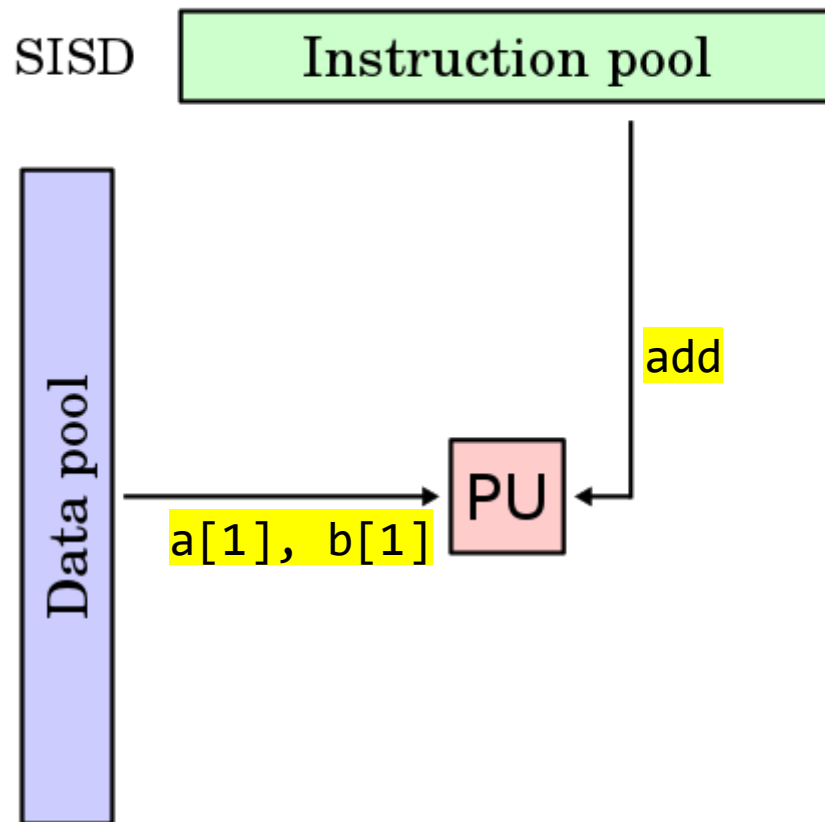
SIMD: Motivating Example

```
for (int i = 0; i < 4; i++) {  
    c[i] = a[i] + b[i];  
}
```



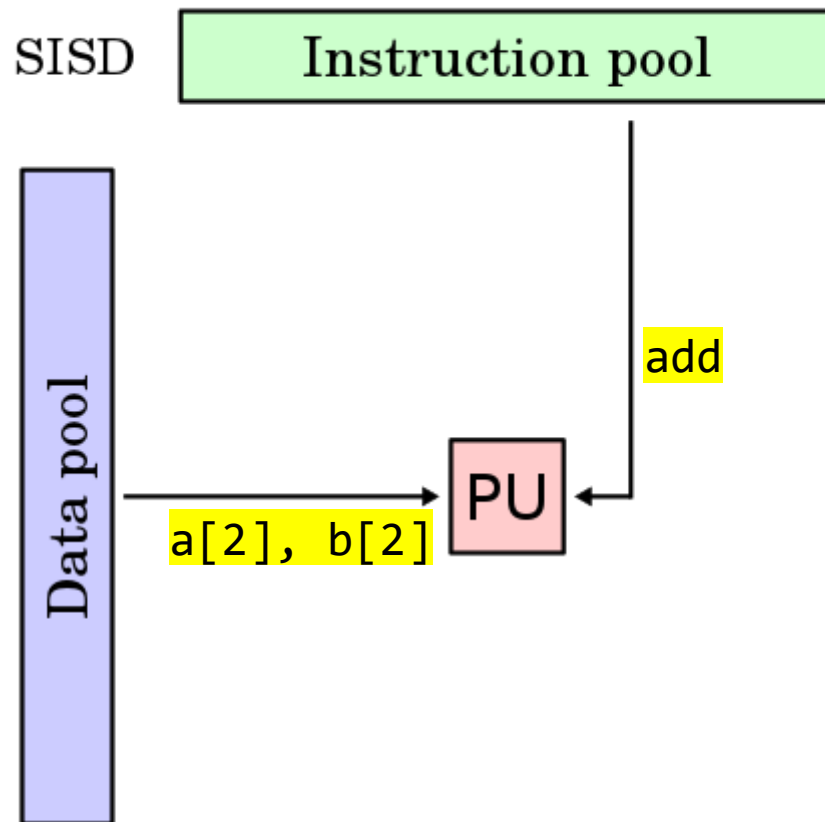
SIMD: Motivating Example

```
for (int i = 0; i < 4; i++) {  
    c[i] = a[i] + b[i];  
}
```



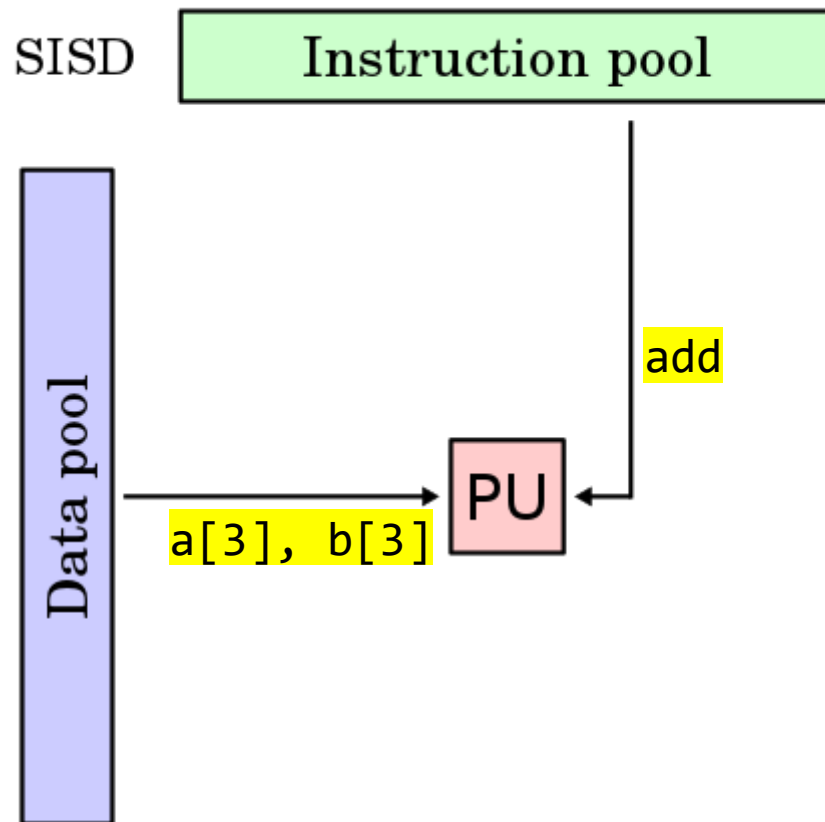
SIMD: Motivating Example

```
for (int i = 0; i < 4; i++) {  
    c[i] = a[i] + b[i];  
}
```



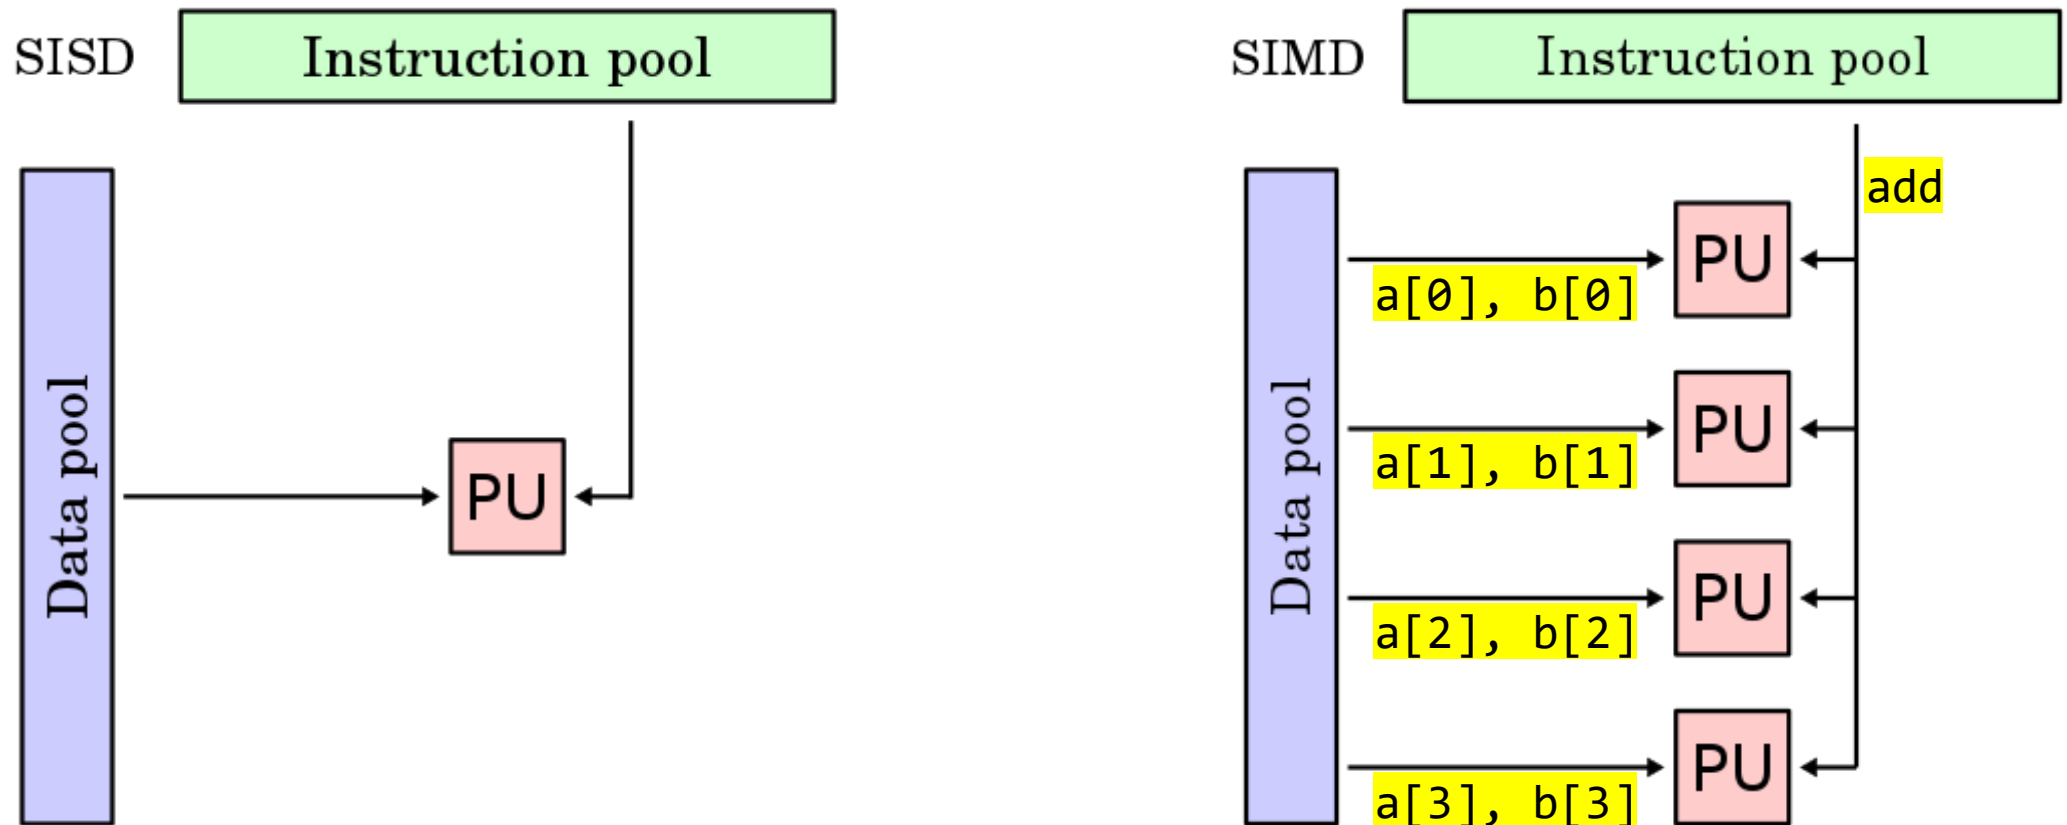
SIMD: Motivating Example

```
for (int i = 0; i < 4; i++) {  
    c[i] = a[i] + b[i];  
}
```



SIMD: Motivating Example

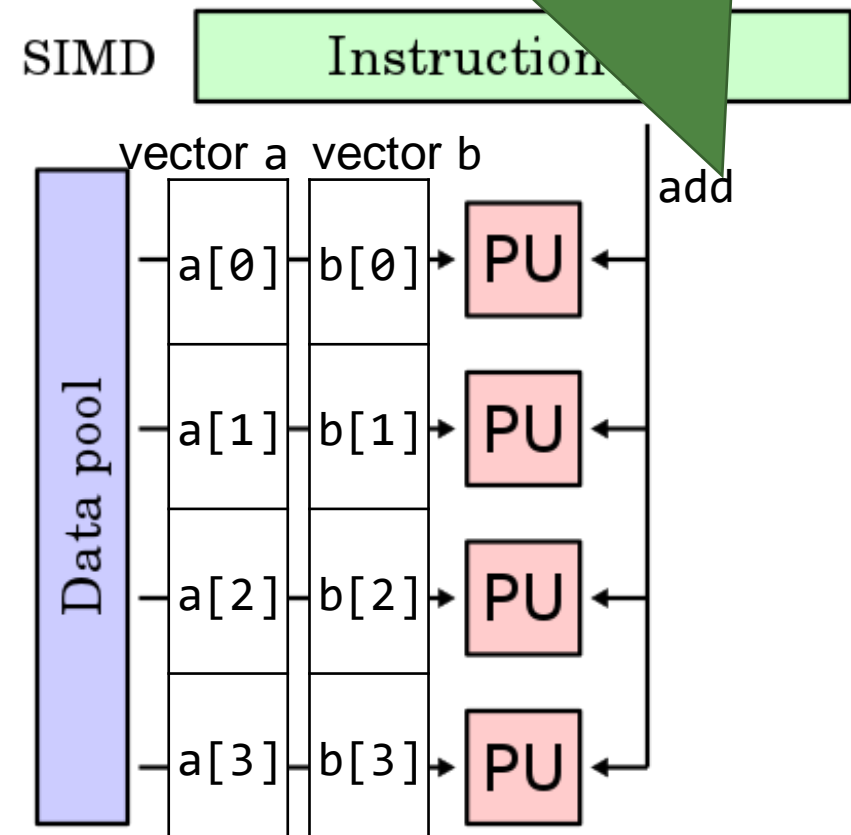

```
for (int i = 0; i < 4; i++) {  
    c[i] = a[i] + b[i];  
}
```



SIMD: Motivating Example

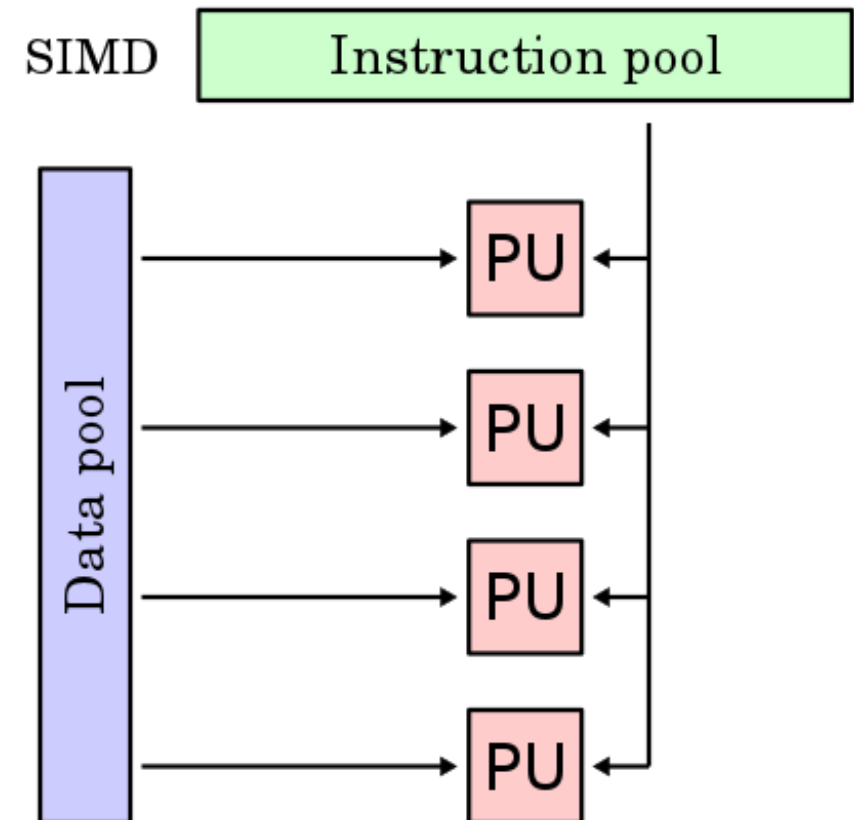
```
for (int i = 0; i < N; i++)  
    c[i] = a[i] + b[i];  
}
```

Single control unit dispatches the same (single) instruction



Single Instruction Multi Data stream (SIMD)

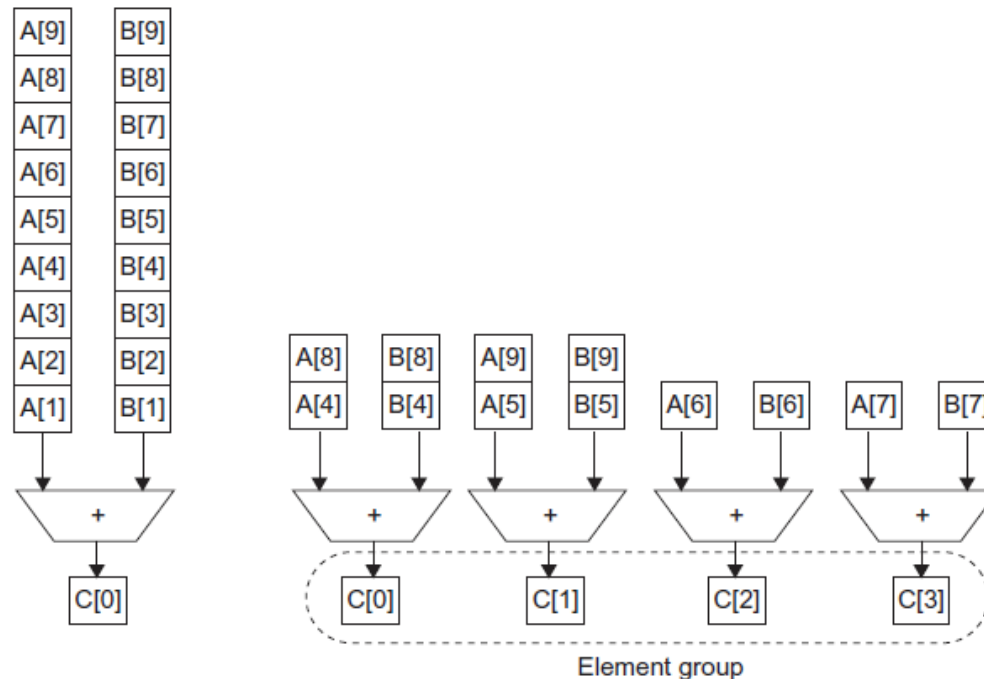
- A single instruction is executed on multiple different pieces of data
- Excels for computations with regular structure (e.g., linear algebra)
- Works best for data-parallel applications
- **Primary application:** vector processing
 - e.g., numpy in Python



Vector Processing



- Highly pipelined function units
- Stream data from/to vector registers to units
 - Data collected from memory into registers
 - Operate on them sequentially using pipelined execution units
 - Results stored from registers to memory




Example: 128-bit SIMD Vectors



- Data types: anything that fits into 16 bytes, e.g.,



4x words



2x double words





16x bytes

- Data bytes must be **contiguous in memory** and **aligned**
- Instructions operate in parallel on data in this 16-byte register
 - add, subtract, multiply, etc.
- Additional instructions needed for
 - Masking data
 - Moving data from one part of a register to another

Computing with SIMD Vector Units

- SISD processing
 - One operation produces **one result**
- SIMD vector units
 - One operation produces **multiple results**



Vector Extension to MIPS (VMIPS)



- 32 vector registers: `v0` to `v31` (64 64-bit elements)
- Vector instructions
 - `lv, sv`: load/store vector
 - `addv.d`: add vectors of double
 - `addvs.d`: add scalar to each element of vector of double

Example: DAXPY ($Y = a * X + Y$)



- Address of X in \$s0 and Y in \$s1

SISD

```
l.d    $f0,a($sp)      ;load scalar a
addiu  $t0,$s0,#512     ;upper bound of what to load
loop:  l.d    $f2,0($s0) ;load x(i)
      mul.d  $f2,$f2,$f0 ;a × x(i)
      l.d    $f4,0($s1) ;load y(i)
      add.d  $f4,$f4,$f2 ;a × x(i) + y(i)
      s.d    $f4,0($s1) ;store into y(i)
      addiu  $s0,$s0,#8  ;increment index to x
      addiu  $s1,$s1,#8  ;increment index to y
      subu   $t0,r4,$s0  ;compute bound
      bne    $t0,$zero,loop ;check if done
```

Example: DAXPY ($Y = a * X + Y$)

1. Load for one element ($X[n]$)

2. Compute $a * X[n]$

3. Load for one element ($Y[n]$)

4. Compute $a * X[n] + Y[n]$

5. Store for one element

```

loop:  l.d    $f0,a($sp)      ;load scalar a
      addiu $t0,$s0,#512    ;upper bound of what to load
      l.d    $f2,0($s0)     ;load x(i)
      mul.d  $f2,$f2,$f0    ;a × x(i)
      l.d    $f4,0($s1)     ;load y(i)
      add.d  $f4,$f4,$f2    ;a × x(i) + y(i)
      s.d    $f4,0($s1)     ;store into y(i)
      addiu  $s0,$s0,#8     ;increment index to x
      addiu  $s1,$s1,#8     ;increment index to y
      subu   $t0,r4,$s0     ;compute bound
      bne    $t0,$zero,loop ;check if done
  
```

Example: DAXPY ($Y = a * X + Y$)

- Address of X in \$s0 and Y in \$s1

SISD

```

l.d    $f0,a($sp)      ;load scalar a
addiu  $t0,$s0,#512    ;upper bound of what to load
loop:  l.d    $f2,0($s0) ;load x(i)
      mul.d   $f2,$f2,$f0 ;a × x(i)
      l.d    $f4,0($s1) ;load y(i)
      add.d   $f4,$f4,$f2 ;a × x(i) + y(i)
      s.d    $f4,0($s1) ;store into y(i)
      addiu  $s0,$s0,#8  ;increment index to x
      addiu  $s1,$s1,#8  ;increment index to y
      subu   $t0,$t0,$s0 ;compute bound
      bne    $t0,$zero,loop ;check if done

```

SIMD

```

l.d    $f0,a($sp)      ;load scalar a
lv     $v1,0($s0)      ;load vector x
mulvs.d $v2,$v1,$f0    ;vector-scalar multiply
lv     $v3,0($s1)      ;load vector y
addv.d $v4,$v2,$v3     ;add y to product
sv     $v4,0($s1)      ;store the result

```

Example: DAXPY ($Y = a * X + Y$)

- Address of X in \$s0 and Y in \$s1

SISD

1. Load for one vector X

2. Compute $a * X$

3. Load for one vector Y

4. Compute $a * X + Y$

5. Store for one vector

```

1.d    $f0,a($sp)      ;load scalar a
addiu   $t0,$s0,#512    ;upper bound of what to load
        $f2,0($s0)      ;load x(i)
        $f2,$f2,$f0     ;a * x(i)
        $f4,0($s1)      ;load y(i)
        $f4,$f4,$f2     ;a * x(i) + y(i)
        $f4,0($s1)      ;store into y(i)
addiu   $s0,$s0,#8      ;increment index to x
addiu   $s1,$s1,#8      ;increment index to y
subu    $t0,$t0,$s0     ;compute bound
bne     $t0,$zero,loop  ;check if done
  
```

SIMD

```

1.d    $f0,a($sp)      ;load scalar a
lv      $v1,0($s0)      ;load vector x
mulvs.d $v2,$v1,$f0     ;vector-scalar multiply
lv      $v3,0($s1)      ;load vector y
addv.d  $v4,$v2,$v3     ;add y to product
sv      $v4,0($s1)      ;store the result
  
```

Example: DAXPY ($Y = a * X + Y$)

- Address of X in \$s0 and Y in \$s1

Almost 600
instructions

Vs.

6 instructions

SISD

```

l.d    $f0,a($sp)      ;load scalar a
addiu  $t0,$s0,#512    ;upper bound of what to load
loop:  l.d    $f2,0($s0) ;load x(i)
      mul.d   $f2,$f2,$f0 ;a × x(i)
      l.d    $f4,0($s1) ;load y(i)
      add.d   $f4,$f4,$f2 ;a × x(i) + y(i)
      s.d    $f4,0($s1) ;store into y(i)
      addiu  $s0,$s0,#8  ;increment index to x
      addiu  $s1,$s1,#8  ;increment index to y
      subu   $t0,r4,$s0  ;compute bound
      bne    $t0,$zero,loop ;check if done
  
```

SIMD

```

l.d    $f0,a($sp)      ;load scalar a
lv     $v1,0($s0)      ;load vector x
mulvs.d $v2,$v1,$f0    ;vector-scalar multiply
lv     $v3,0($s1)      ;load vector y
addv.d $v4,$v2,$v3     ;add y to product
sv     $v4,0($s1)      ;store the result
  
```

Example: DAXPY ($Y = a * X + Y$)

- Address of X in \$s0 and Y in \$s1

SISD

```

l.d    $f0,a($sp)      ;load scalar a
addiu  $t0,$s0,#512    ;upper bound of what to load
loop:  l.d    $f2,0($s0) ;load x(i)
      mul.d   $f2,$f2,$f0 ;a * x(i)
      l.d    $f4,0($s1) ;load y(i)
      add.d   $f4,$f4,$f2 ;a * x(i) + y(i)
      s.d    $f4,0($s1) ;store into y(i)
      addiu  $s0,$s0,#8  ;increment index to x
      addiu  $s1,$s1,#8  ;increment index to y
      subu   $t0,r4,$s0  ;compute bound
      bne    $t0,$zero,loop ;check if done
  
```

Stalls per vector element

SIMD

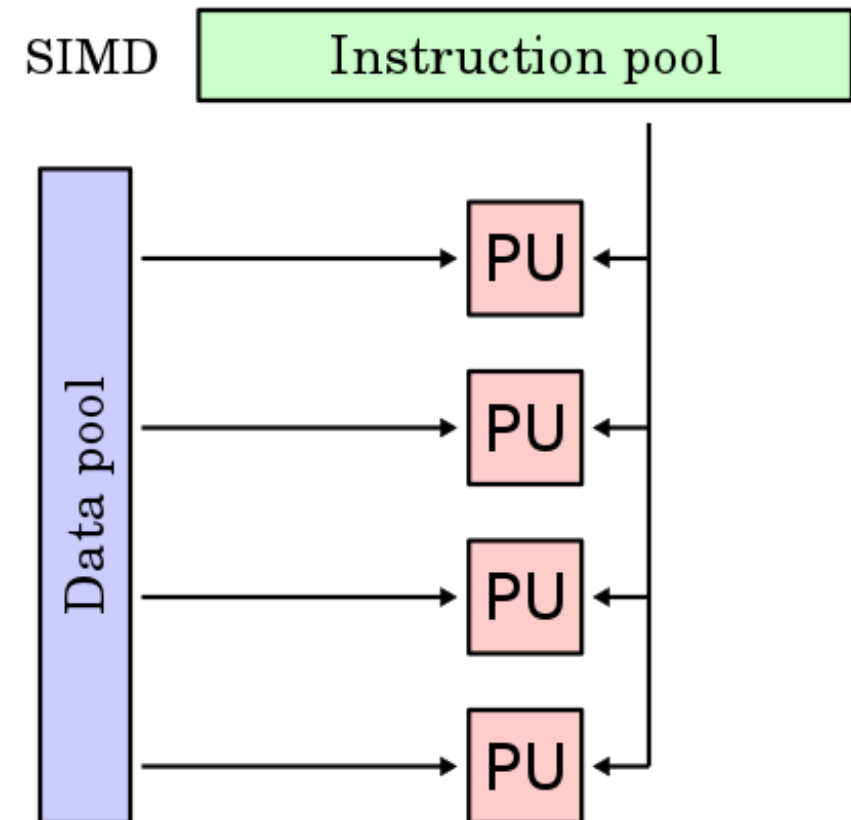
```

l.d    $f0,a($sp)      ;load scalar a
lv     $v1,0($s0)      ;load vector x
mulvs.d $v2,$v1,$f0    ;vector-scalar multiply
lv     $v3,0($s1)      ;load vector y
addv.d $v4,$v2,$v3     ;add y to product
sv     $v4,0($s1)      ;store the result
  
```

Stalls only once per vector operation




Single Instruction Multi Data stream (SIMD) ⁴¹

- A single instruction is executed on multiple different pieces of data
- Excels for computations with regular structure (e.g., linear algebra)
- Works best for data-parallel applications
- **Primary application:** vector processing
 - e.g., numpy in Python
- Modern GPUs, containing vector processors, are commonly SIMD systems
- **Disadvantages**
 - Limited to specific applications



SIMD Example

42

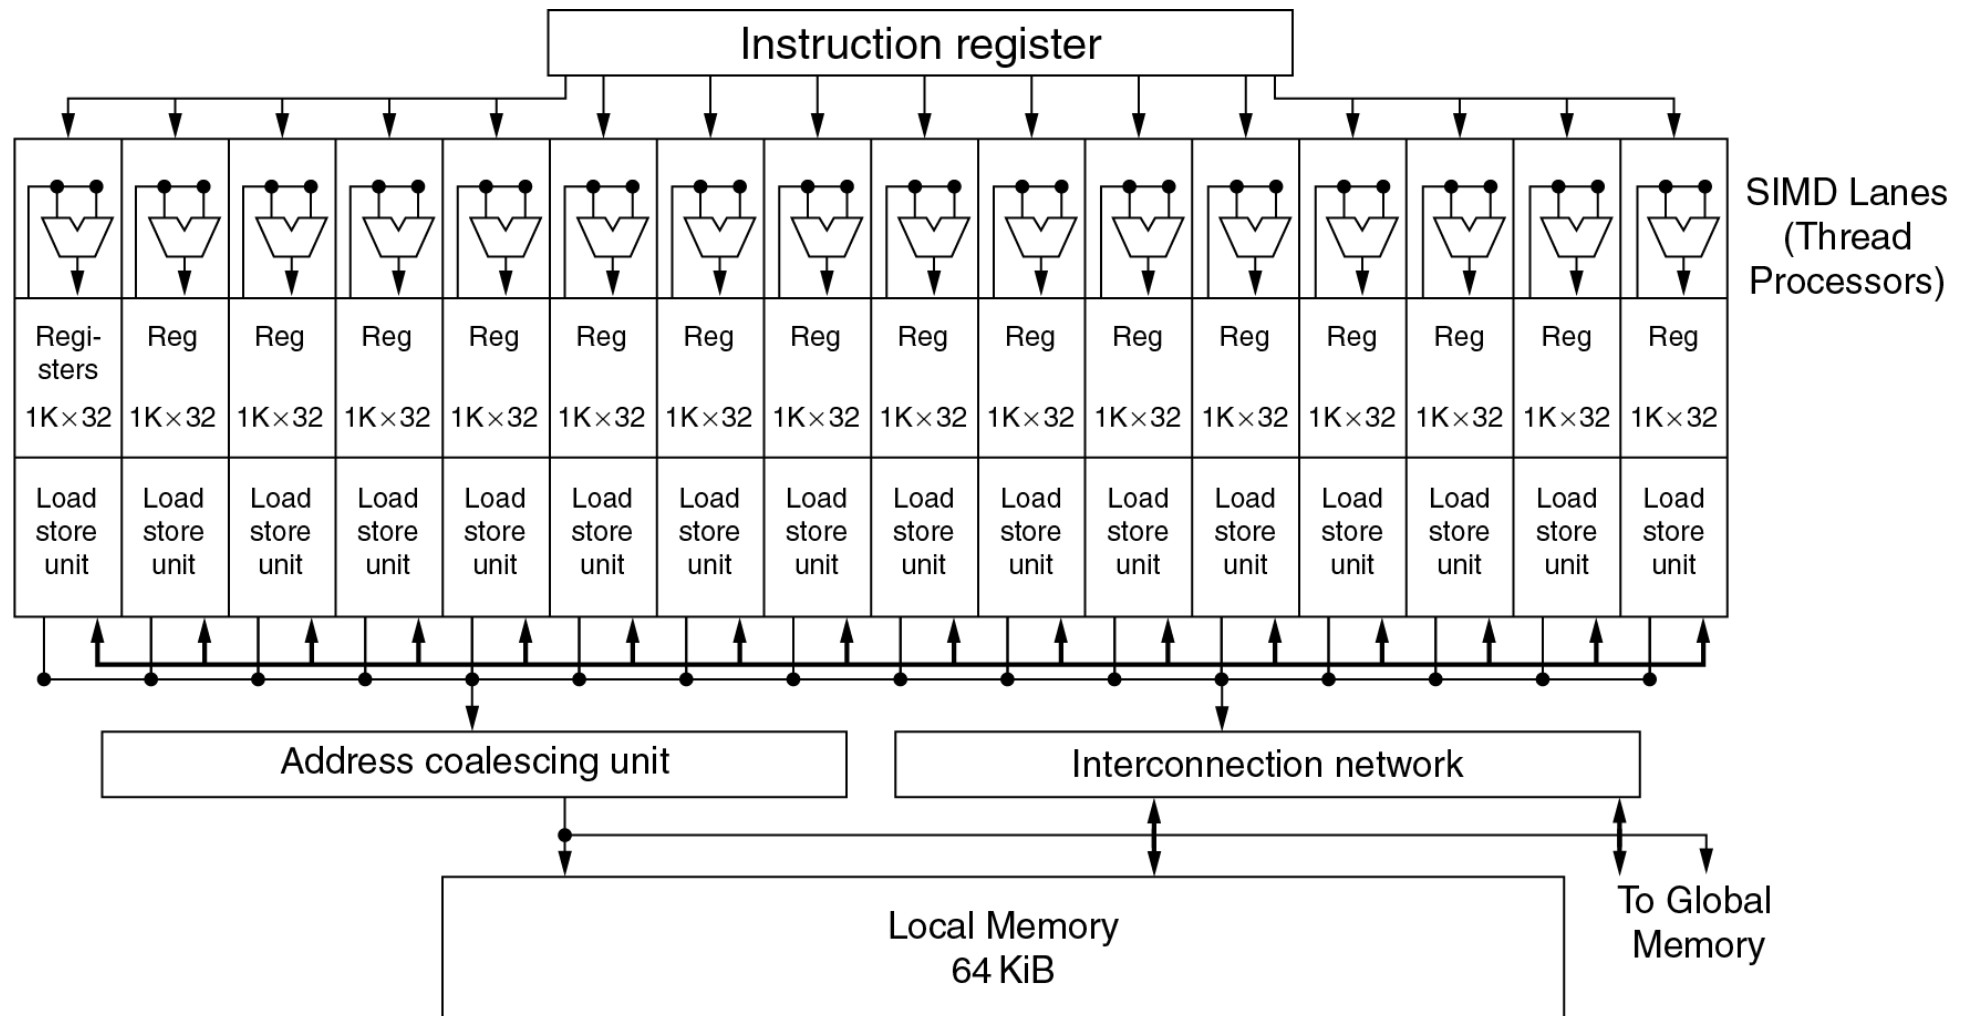


$X+2, Y-1$

Example: Nvidia Tesla




- Multiple (multithreaded) SIMD processors, each as shown:



Multi Instruction Multiple Data stream (MIMD) ⁴⁴

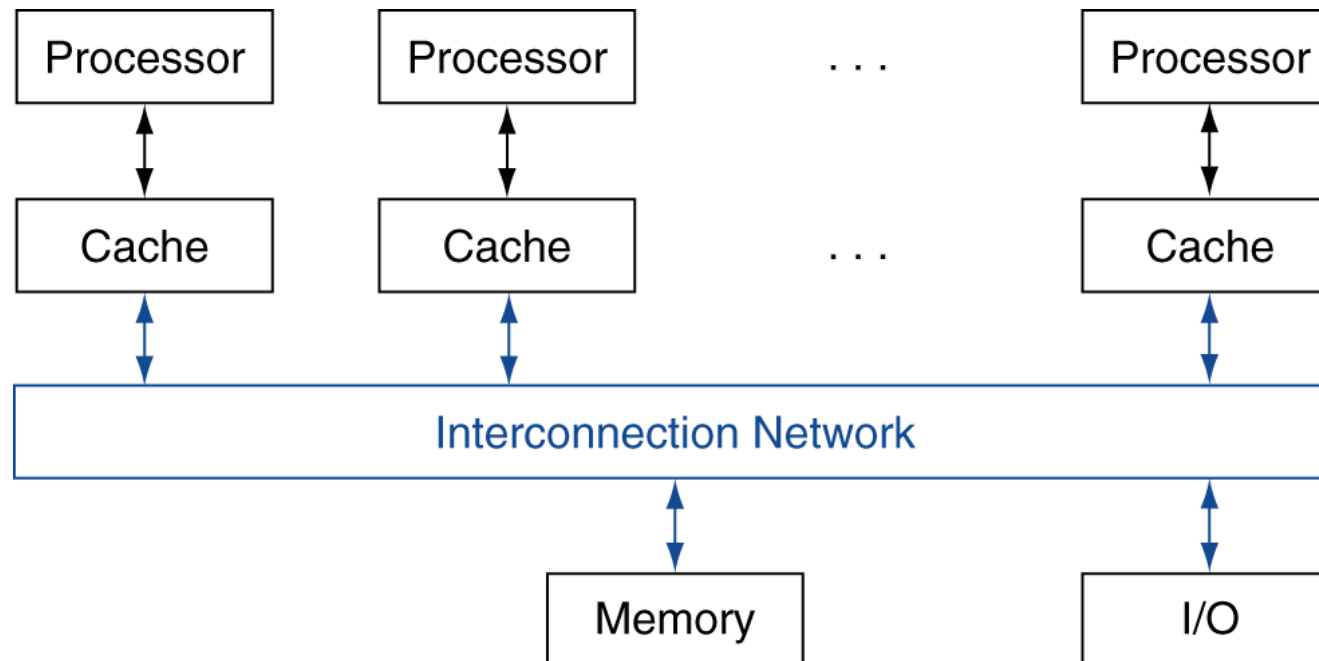
- Multiple instructions operate on different pieces of data, either *independently* or *as part of shared memory space*
 - Multiprocessor
 - Multithreaded processor
- Uses:
 - Most modern computers
 - Cluster



Shared Memory

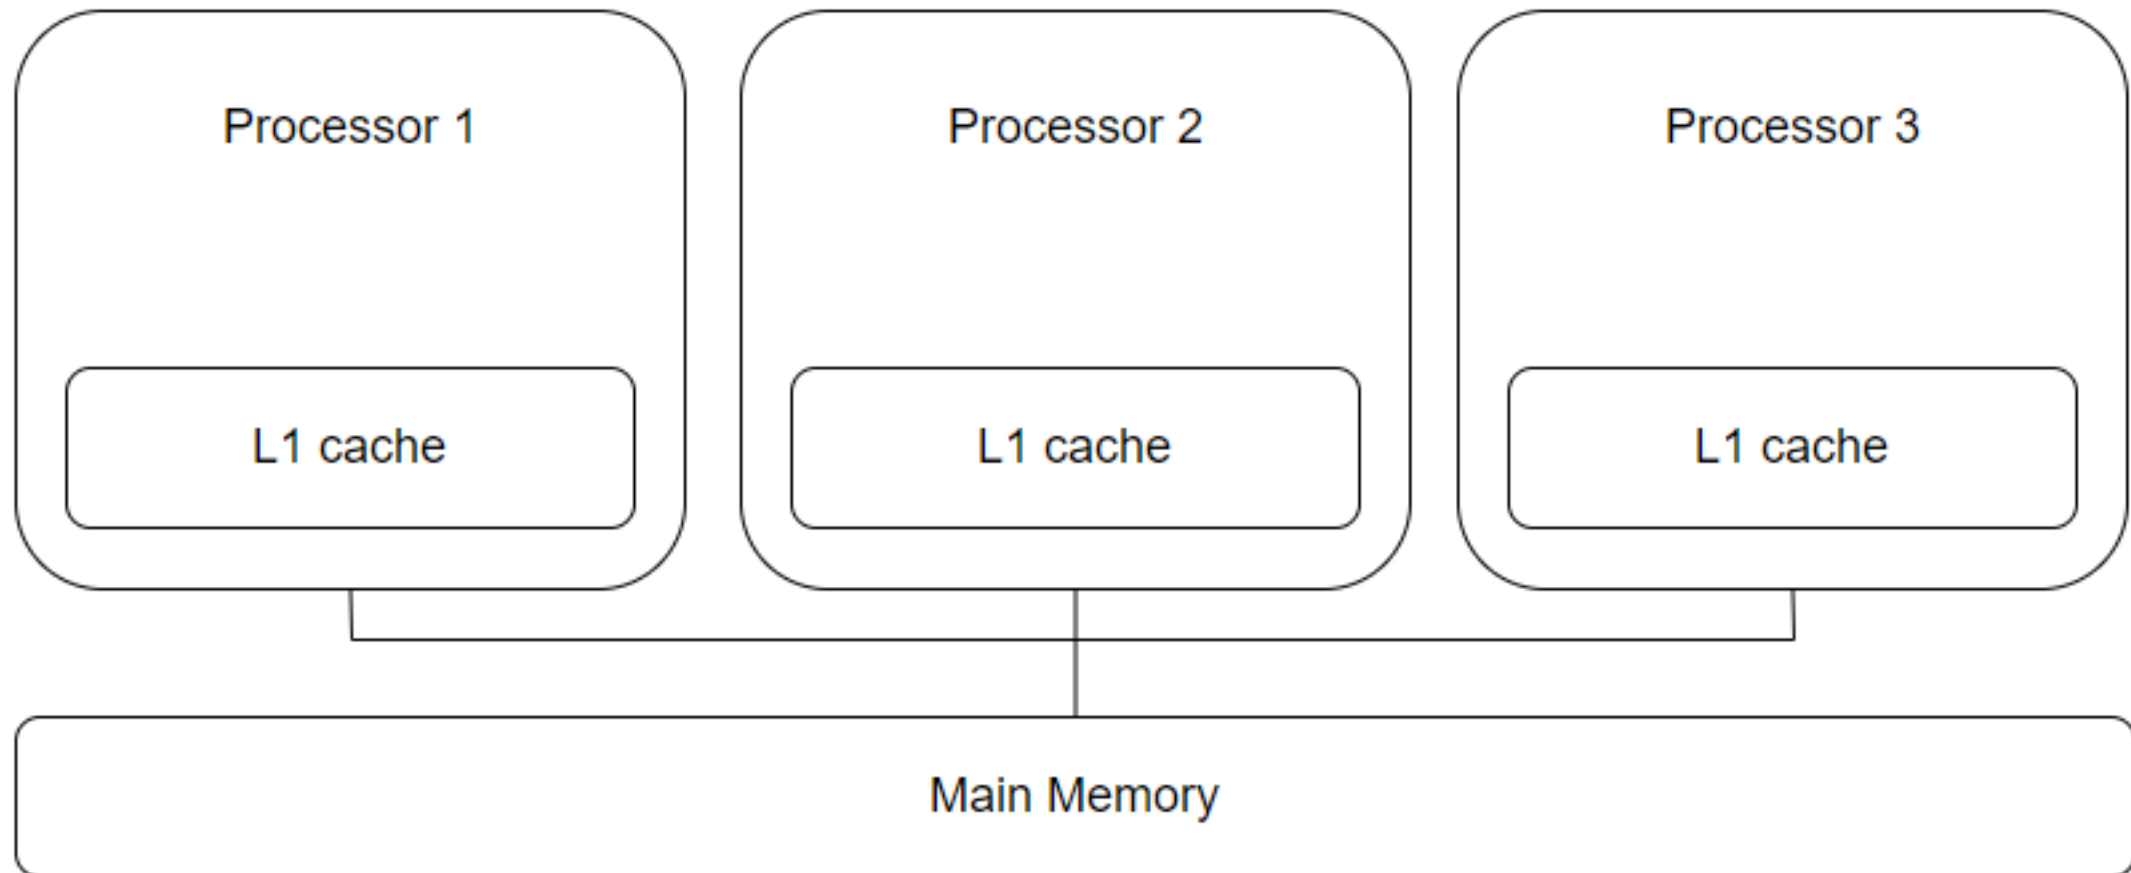


- **Shared Memory Multiprocessor (SMP)**
 - Hardware provides **single physical address space** for all processors
 - Synchronize ***shared variables*** using locks
 - Memory access time
 - UMA (uniform) vs. NUMA (nonuniform)



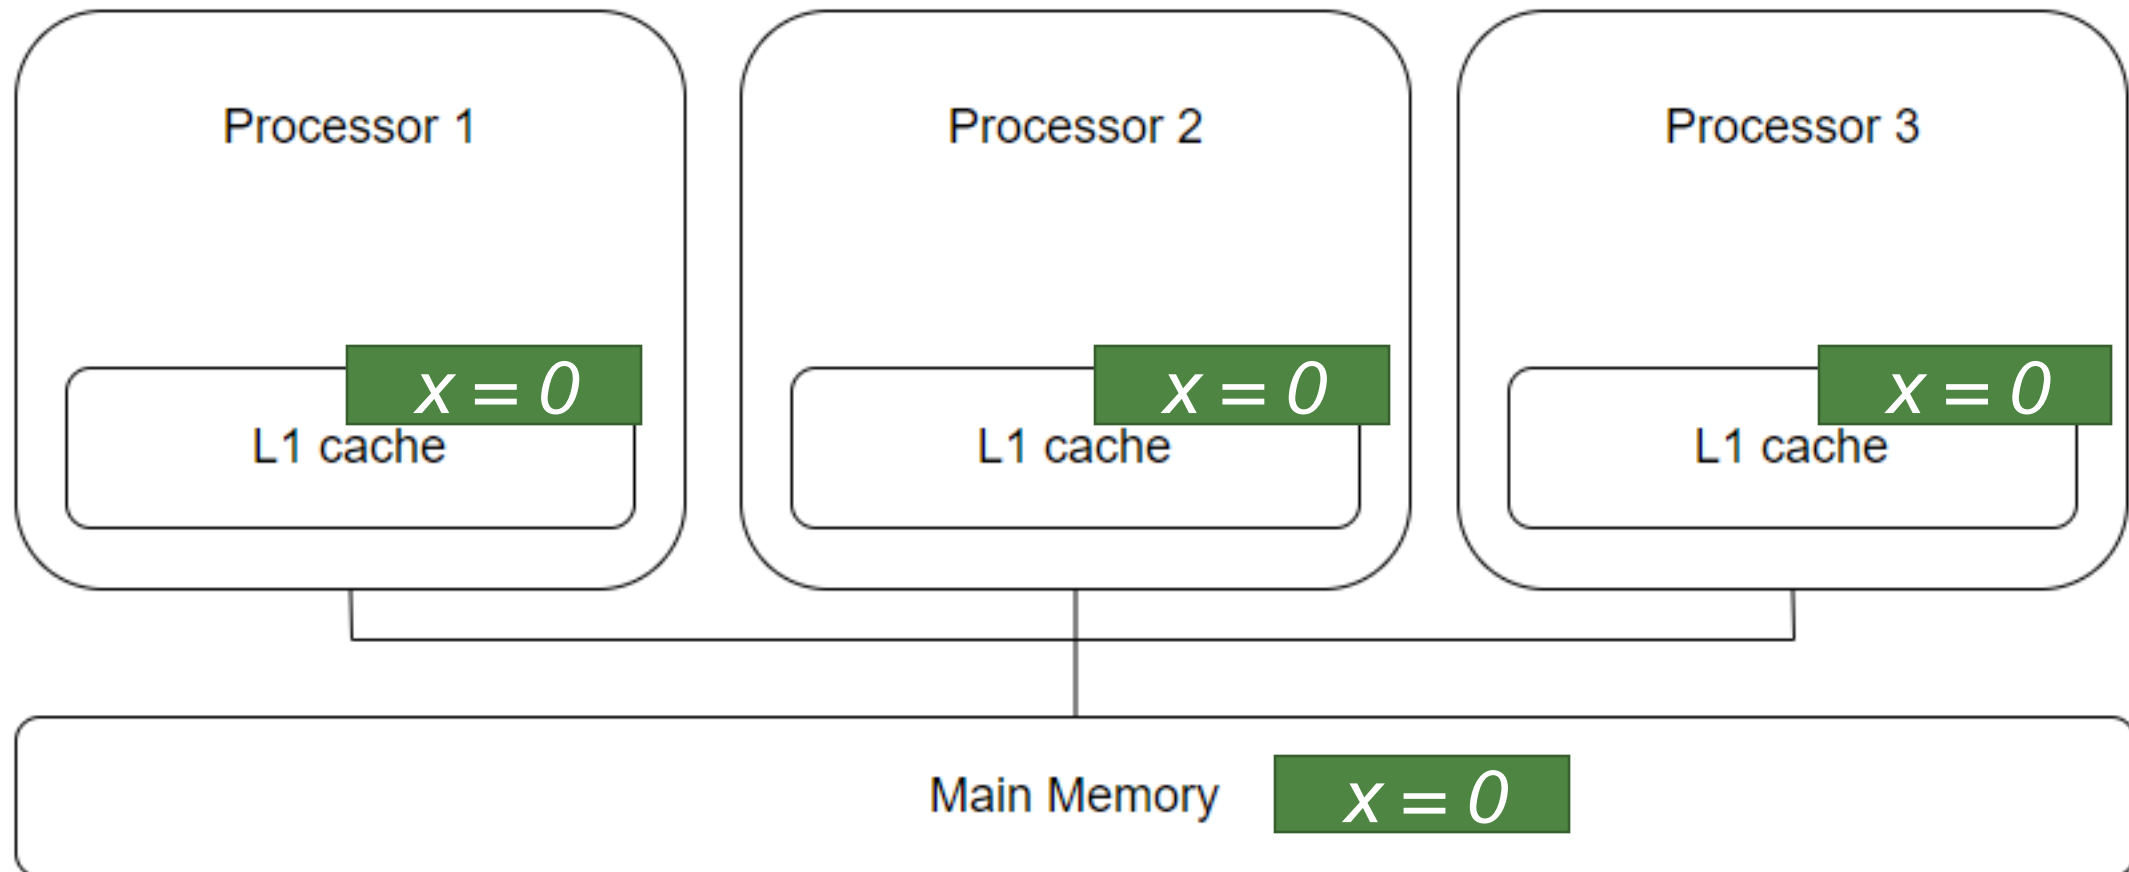
Cache Coherence Problems

- Multiple copies of same data may become inconsistent



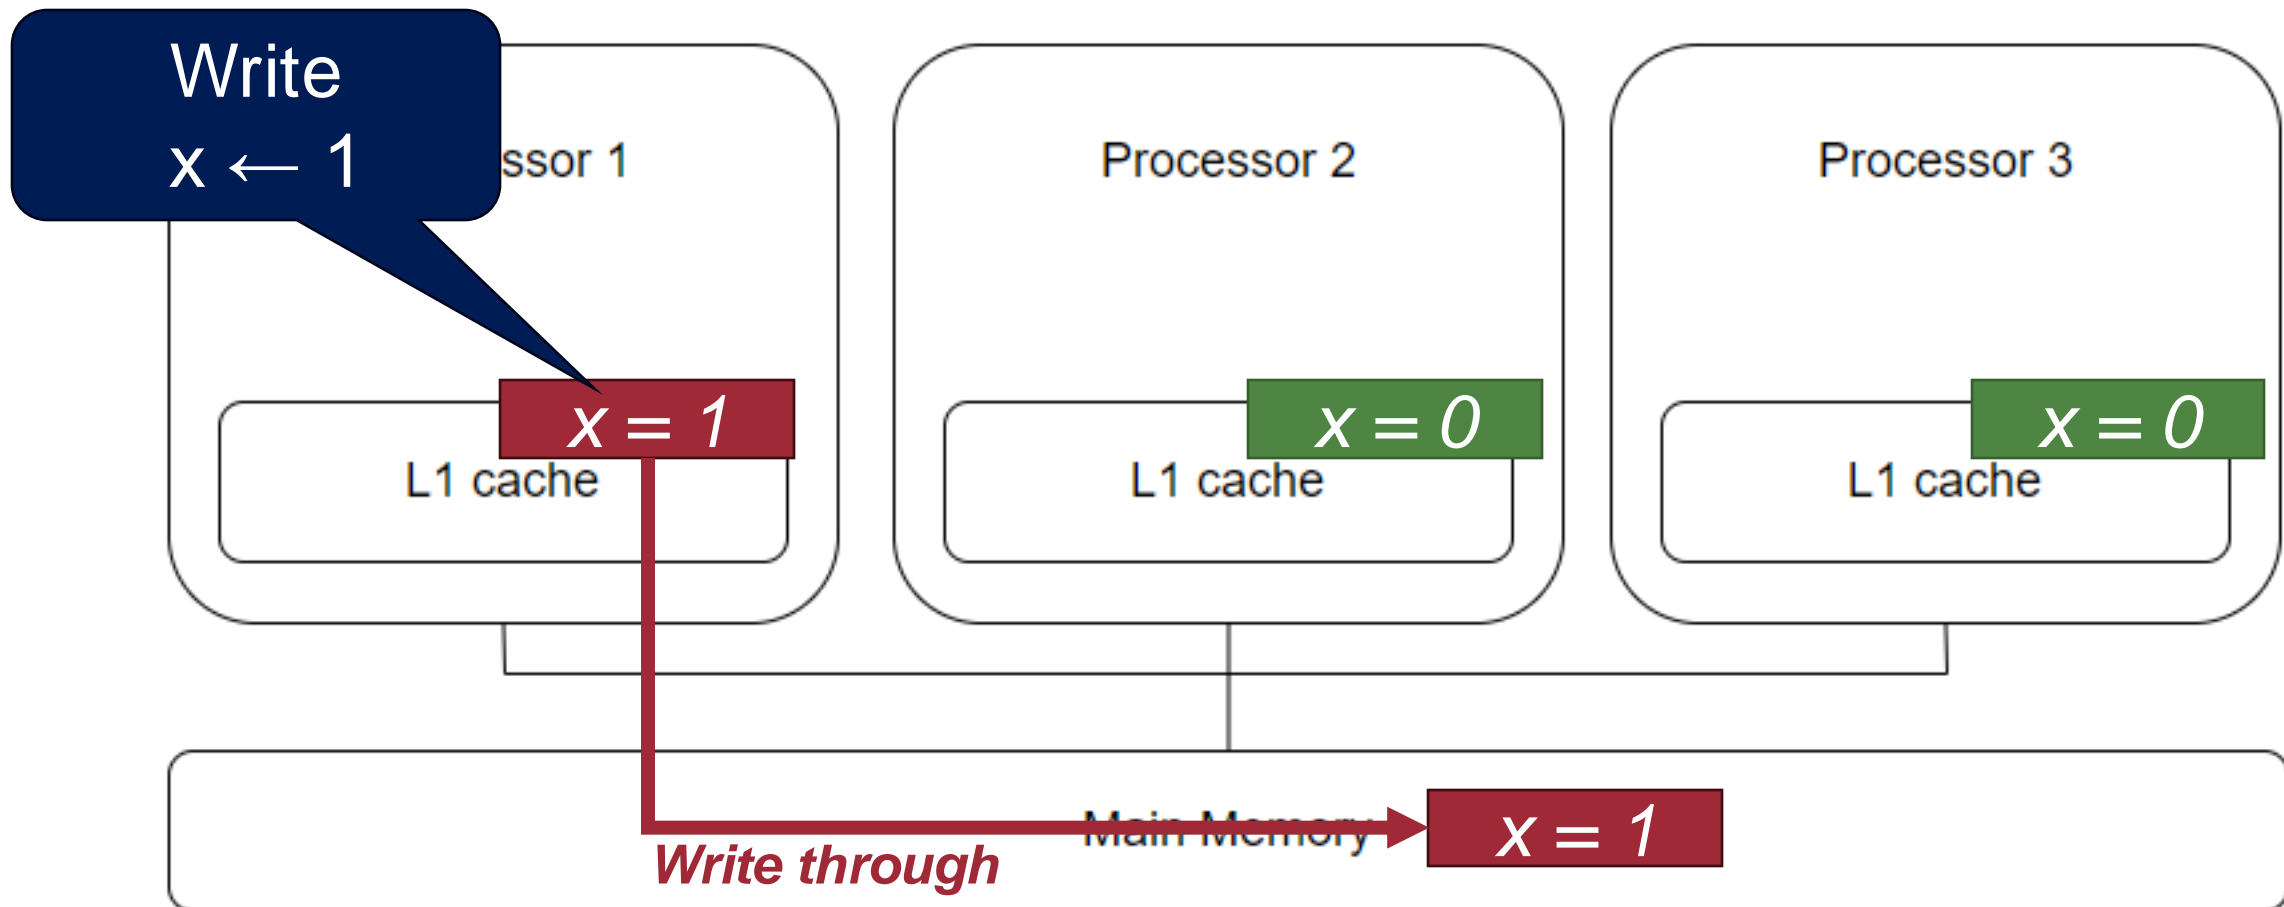
Cache Coherence Problems

- Multiple copies of same data may become inconsistent



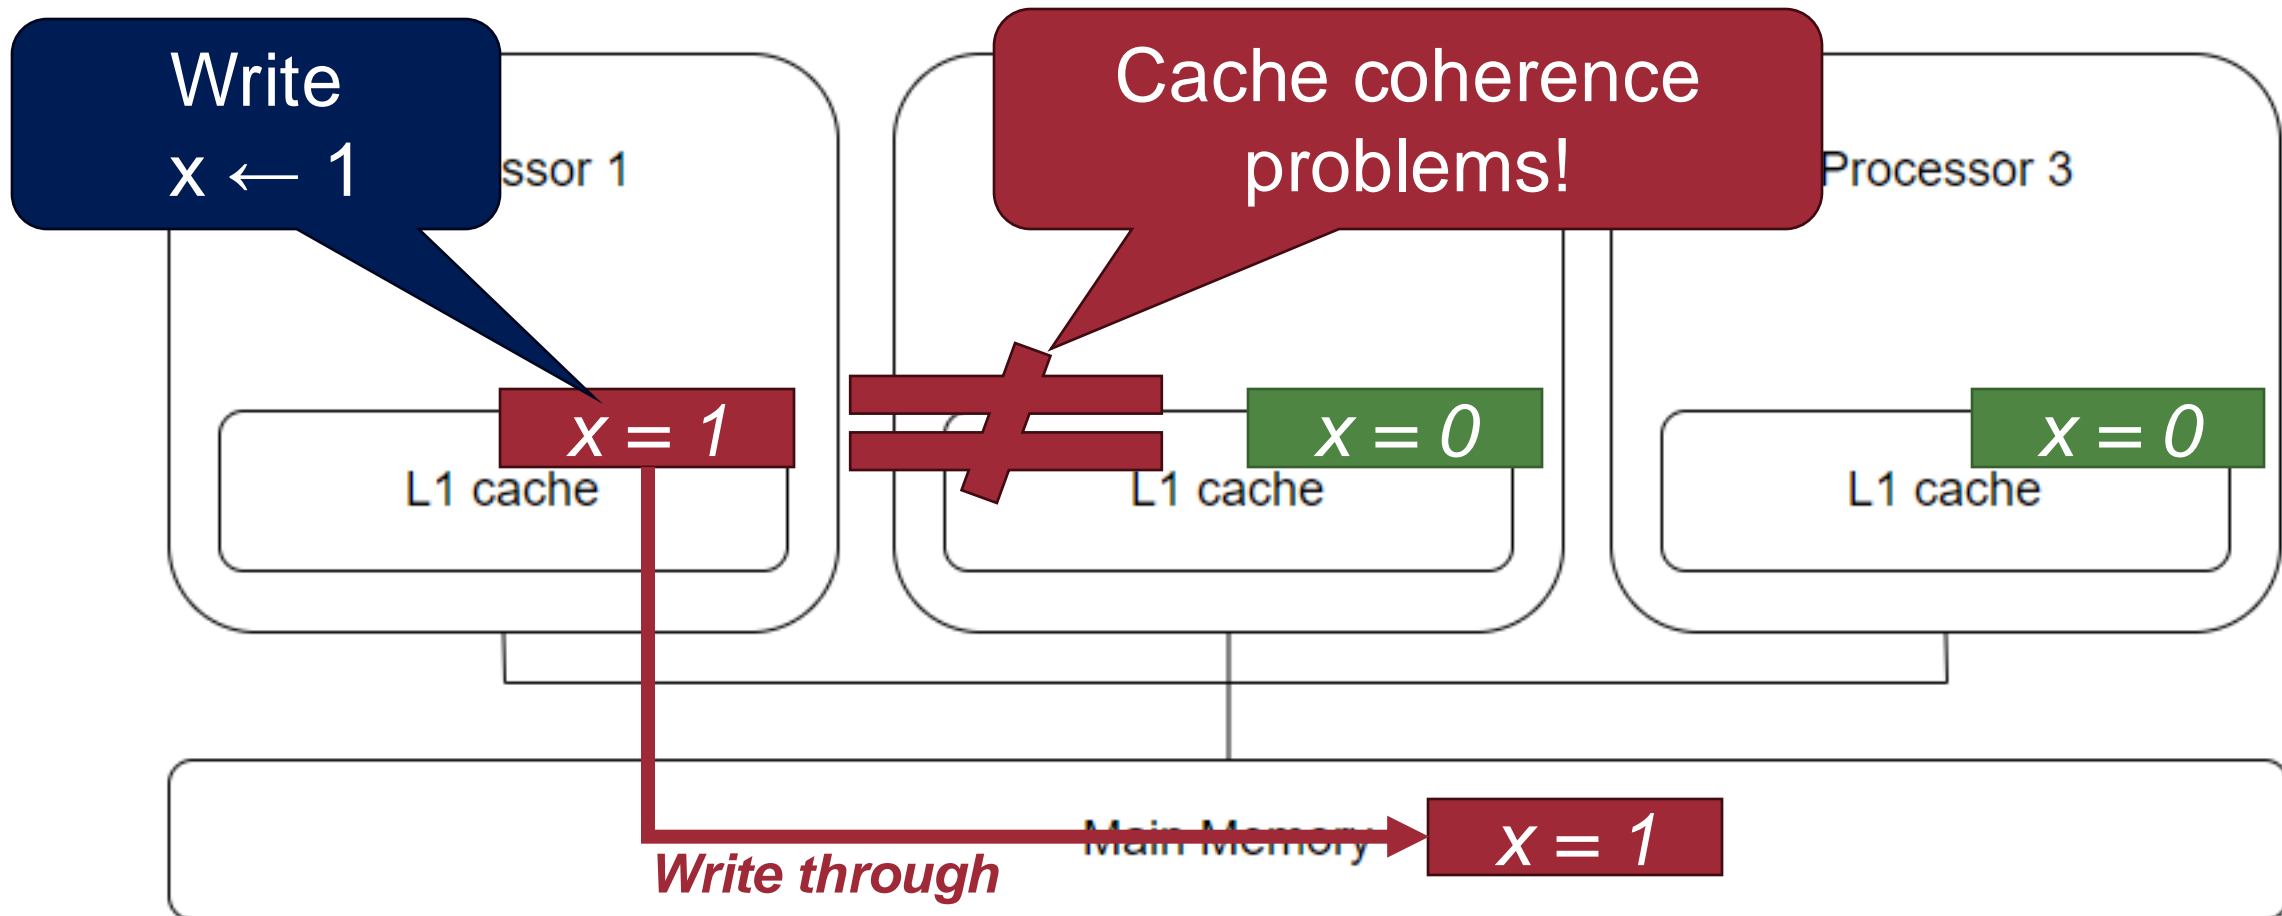
Cache Coherence Problems

- Multiple copies of same data may become inconsistent



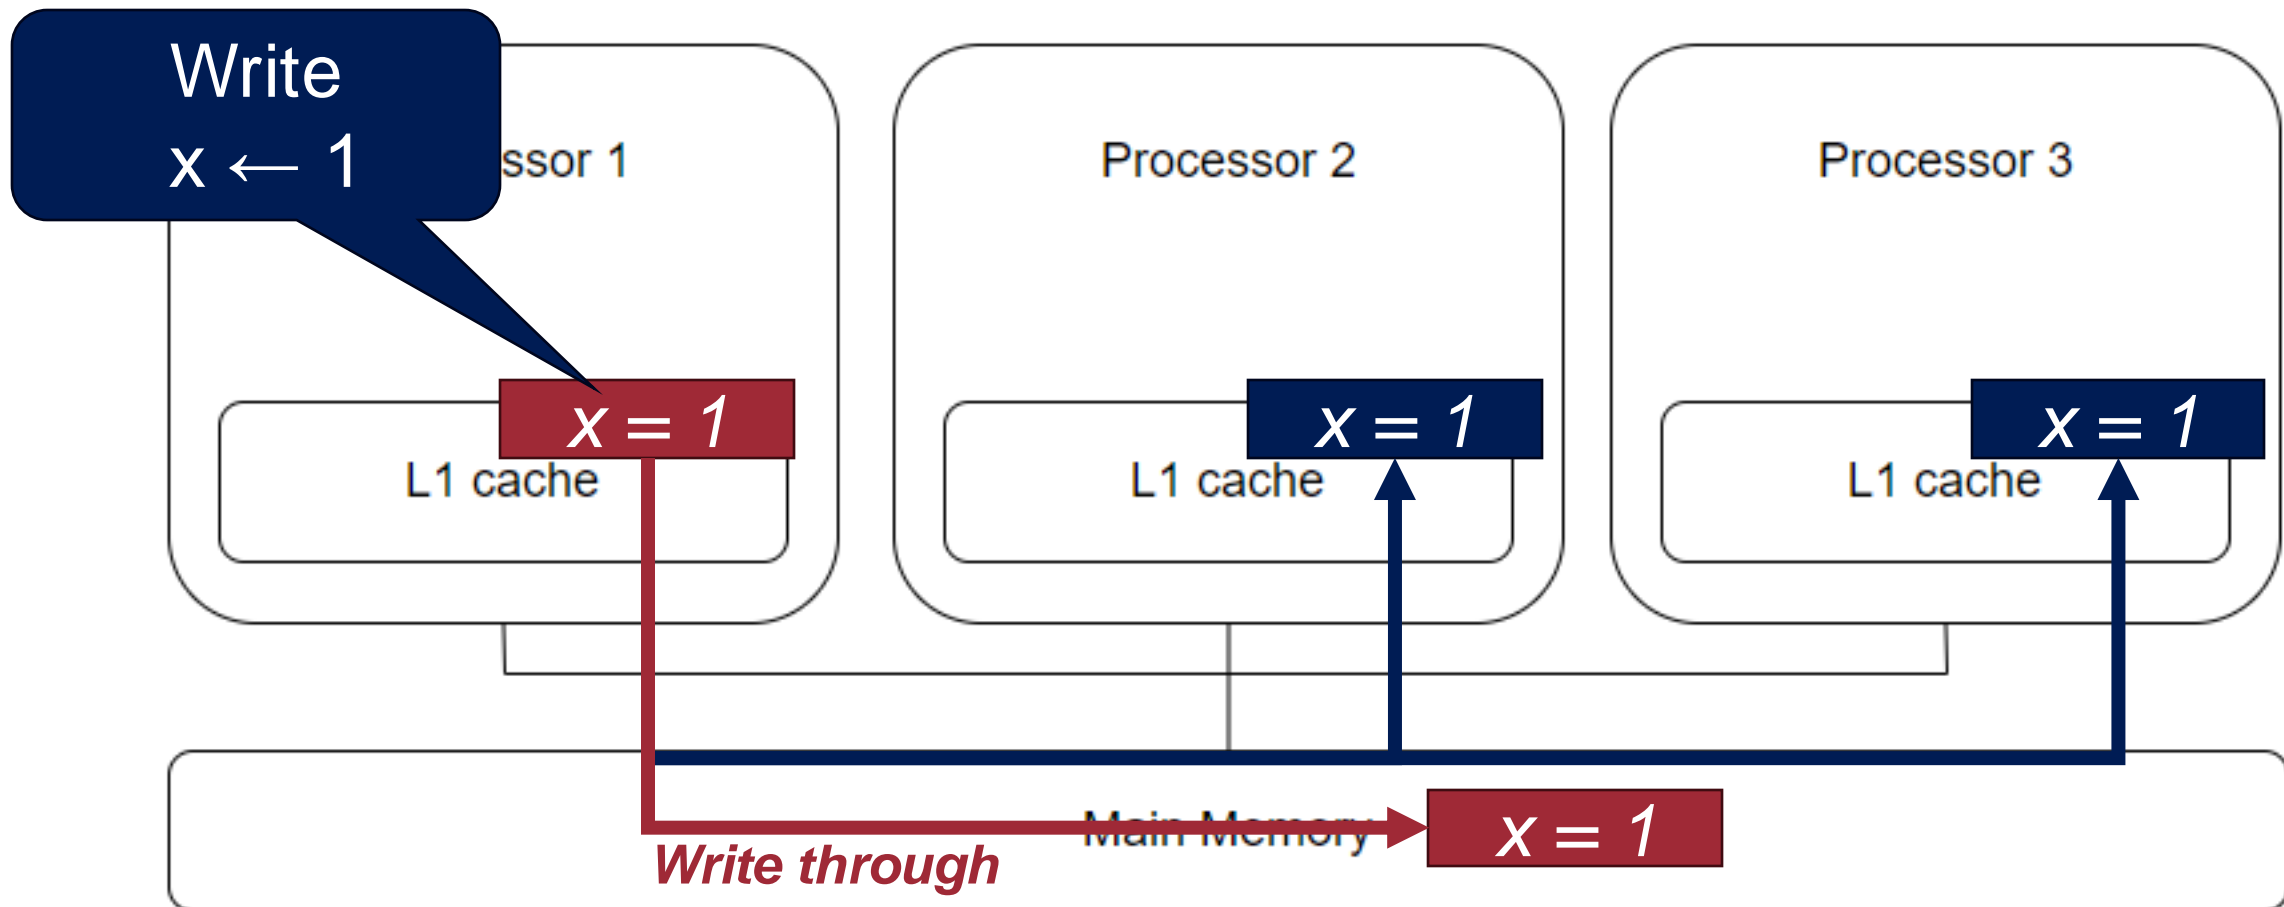
Cache Coherence Problems

- Multiple copies of same data may become inconsistent



Multicore Cache Coherence: Solution

- Use the cache coherence protocol (e.g., MSI, MESI protocol)

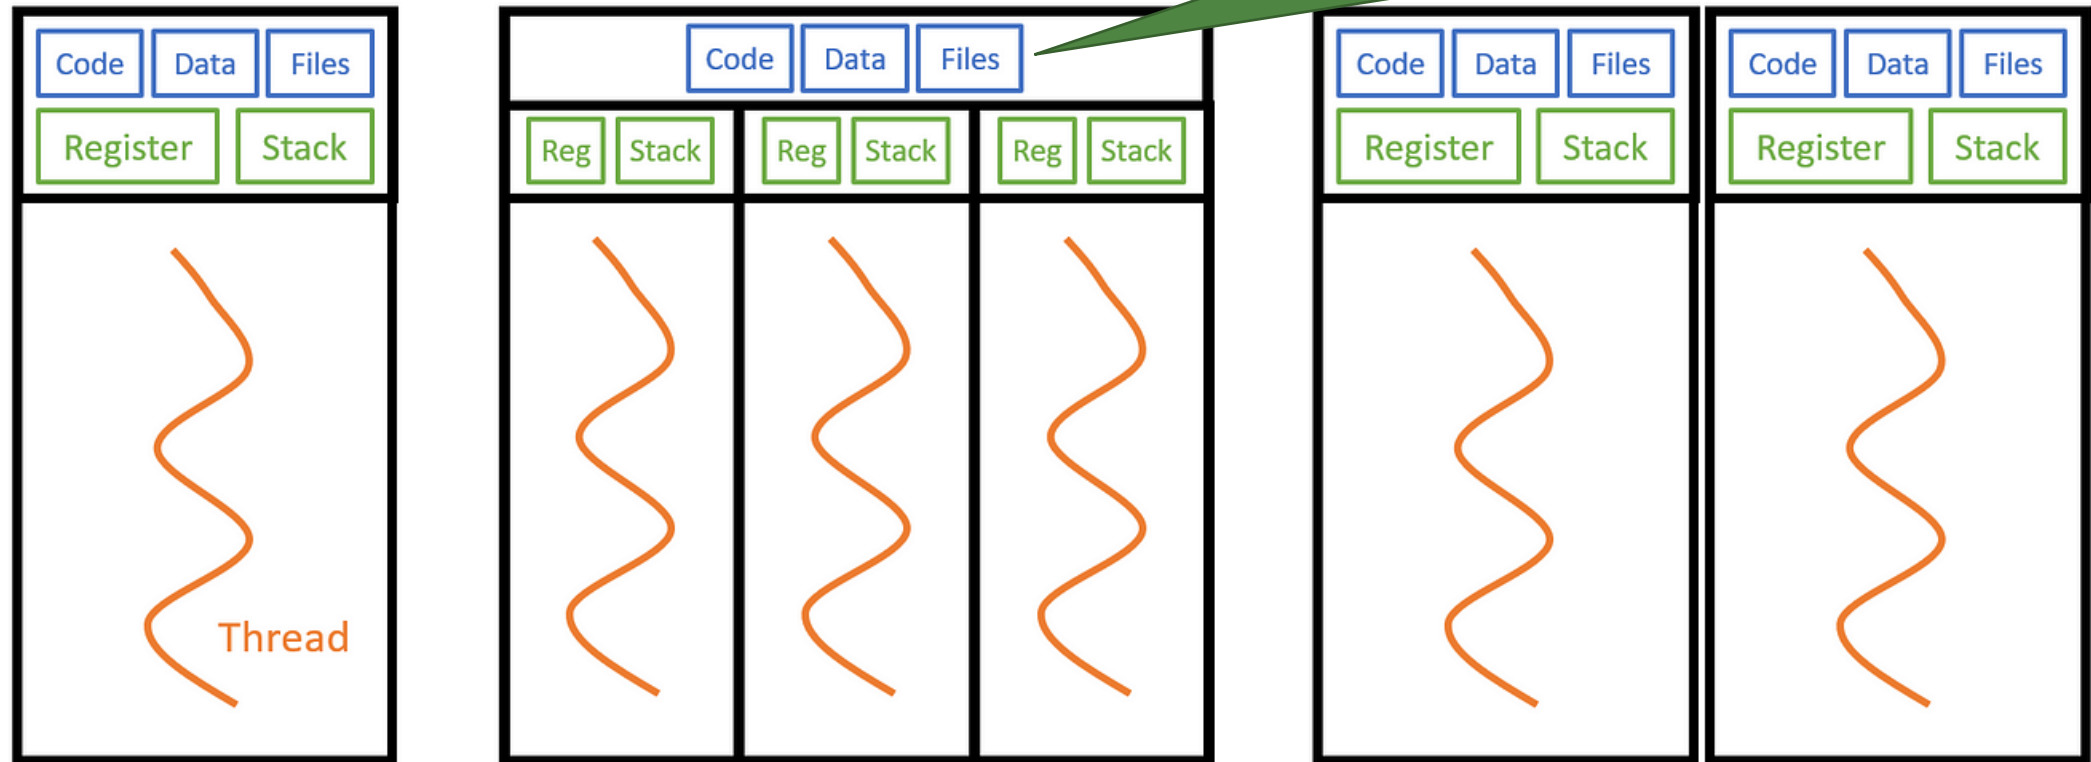


Multithreading



- Thread: the unit of execution within a process

Share the memory space



Single Processor Single Thread

Single Processor Multithread

Multiprocessing

Hardware Multithreading



- Run multiple threads of execution in parallel
 - Fast switching between threads
- Fine-grain multithreading
 - Switch threads after each cycle
 - Interleave instruction execution
 - If one thread stalls, others are executed
- Coarse-grain multithreading
 - Only switch on long stall (e.g., L2-cache miss)
 - Simplifies hardware, but doesn't hide short stalls (e.g., data hazards)

Instruction and Data Streams



- According to the Flynn's taxonomy of computers, ...

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345

Question?