

CSE551:

Advanced Computer Security

9. ROP & ASLR

Seongil Wi

Project Checkpoint



- Submit a single PDF file consisting of multiple presentation slides!
- For this checkpoint, there will be no presentation session; only the checkpoint submission is required
- Due: 10/23, 11:59PM

Project Checkpoint

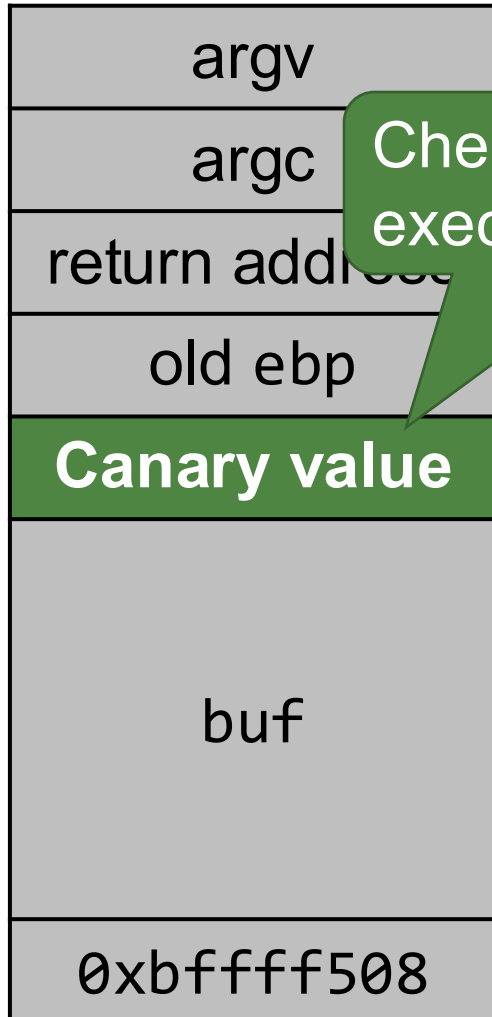


- You should upload a single PDF file on BlackBored.
- The name of the PDF file should have the following format: [your ID-last name.pdf]
 - If your name is Gil-dong Hong, and your ID is 20231234, then you should submit a file named “20231234-Hong.pdf”
 - If your team consists of two people, each member must submit a PDF file
- Submit a single PDF file consisting of multiple presentation slides. The PDF should include the following topics and contents (must be written in English!):
 - Introduction
 - Background
 - Motivation
 - Approach
 - Your Progress

Recap: Mitigating Memory Corruption Bugs

4

Mitigation #1: Canary



Check value before
executing return!

Mitigation #2: NX (No eXcute)

Corrupted
memory

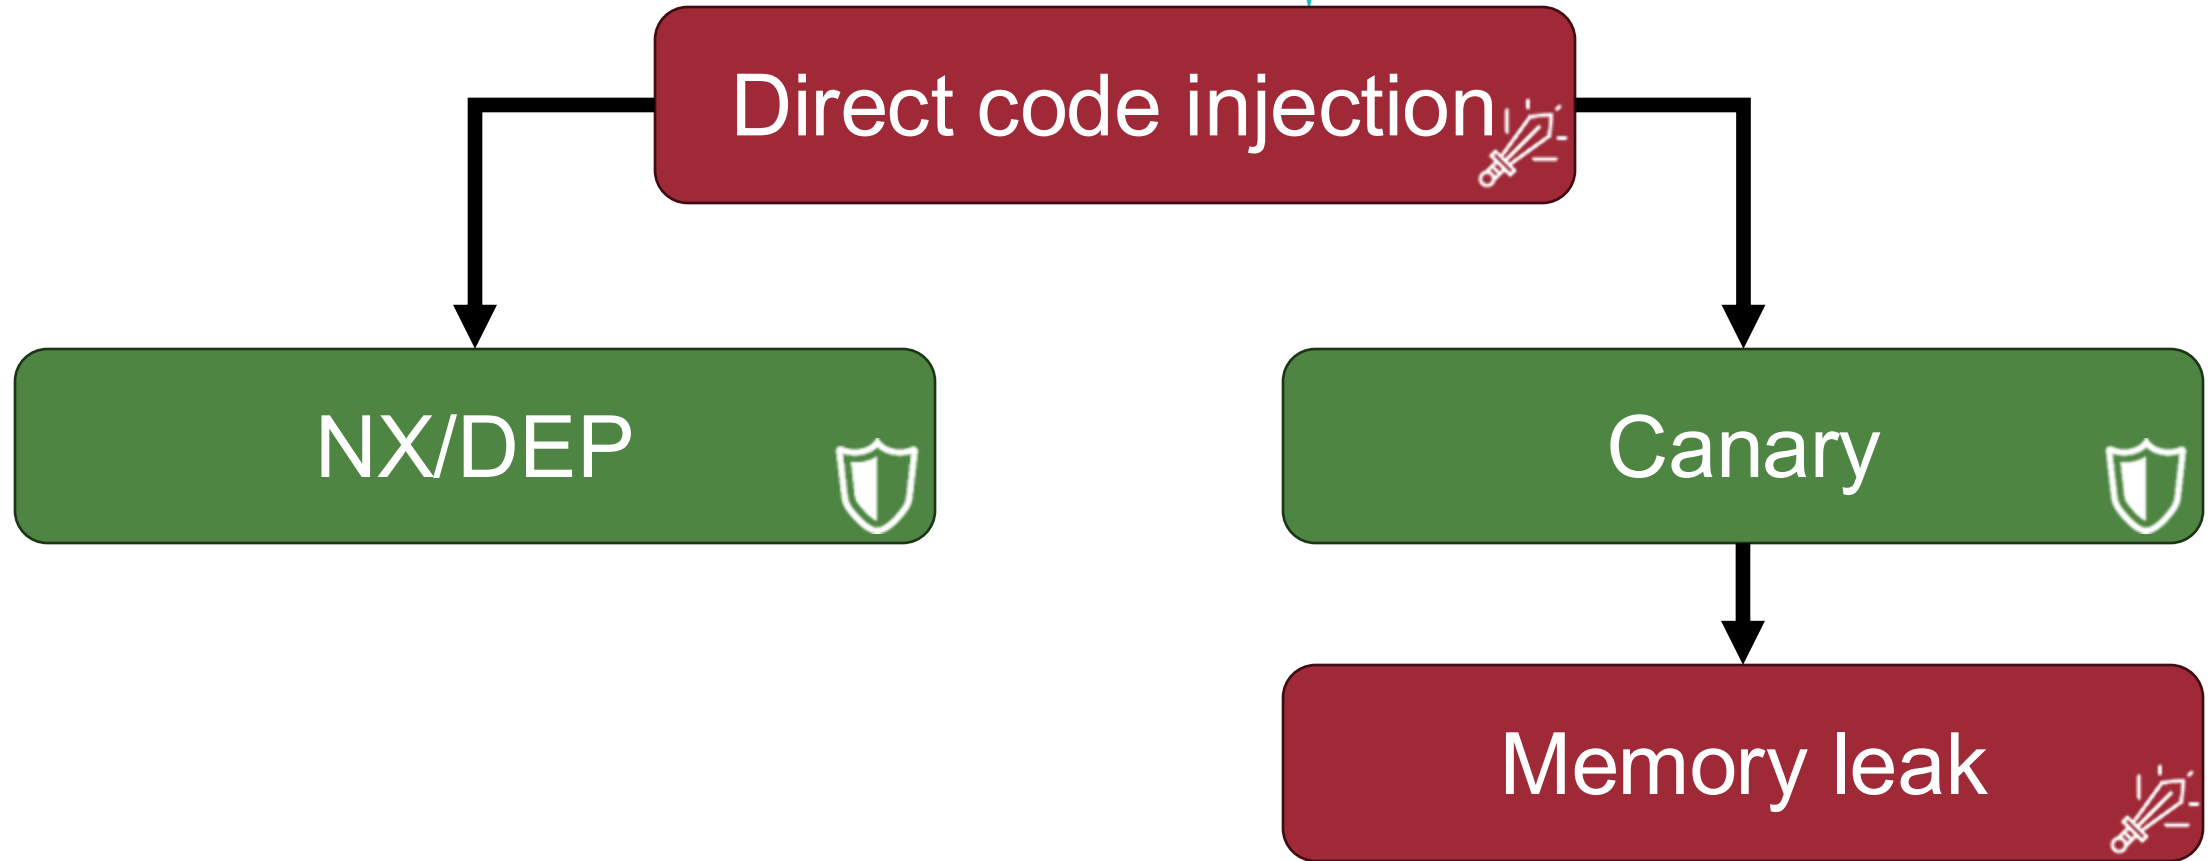
Attacker's code
(Shellcode)

Hijacked
control flow

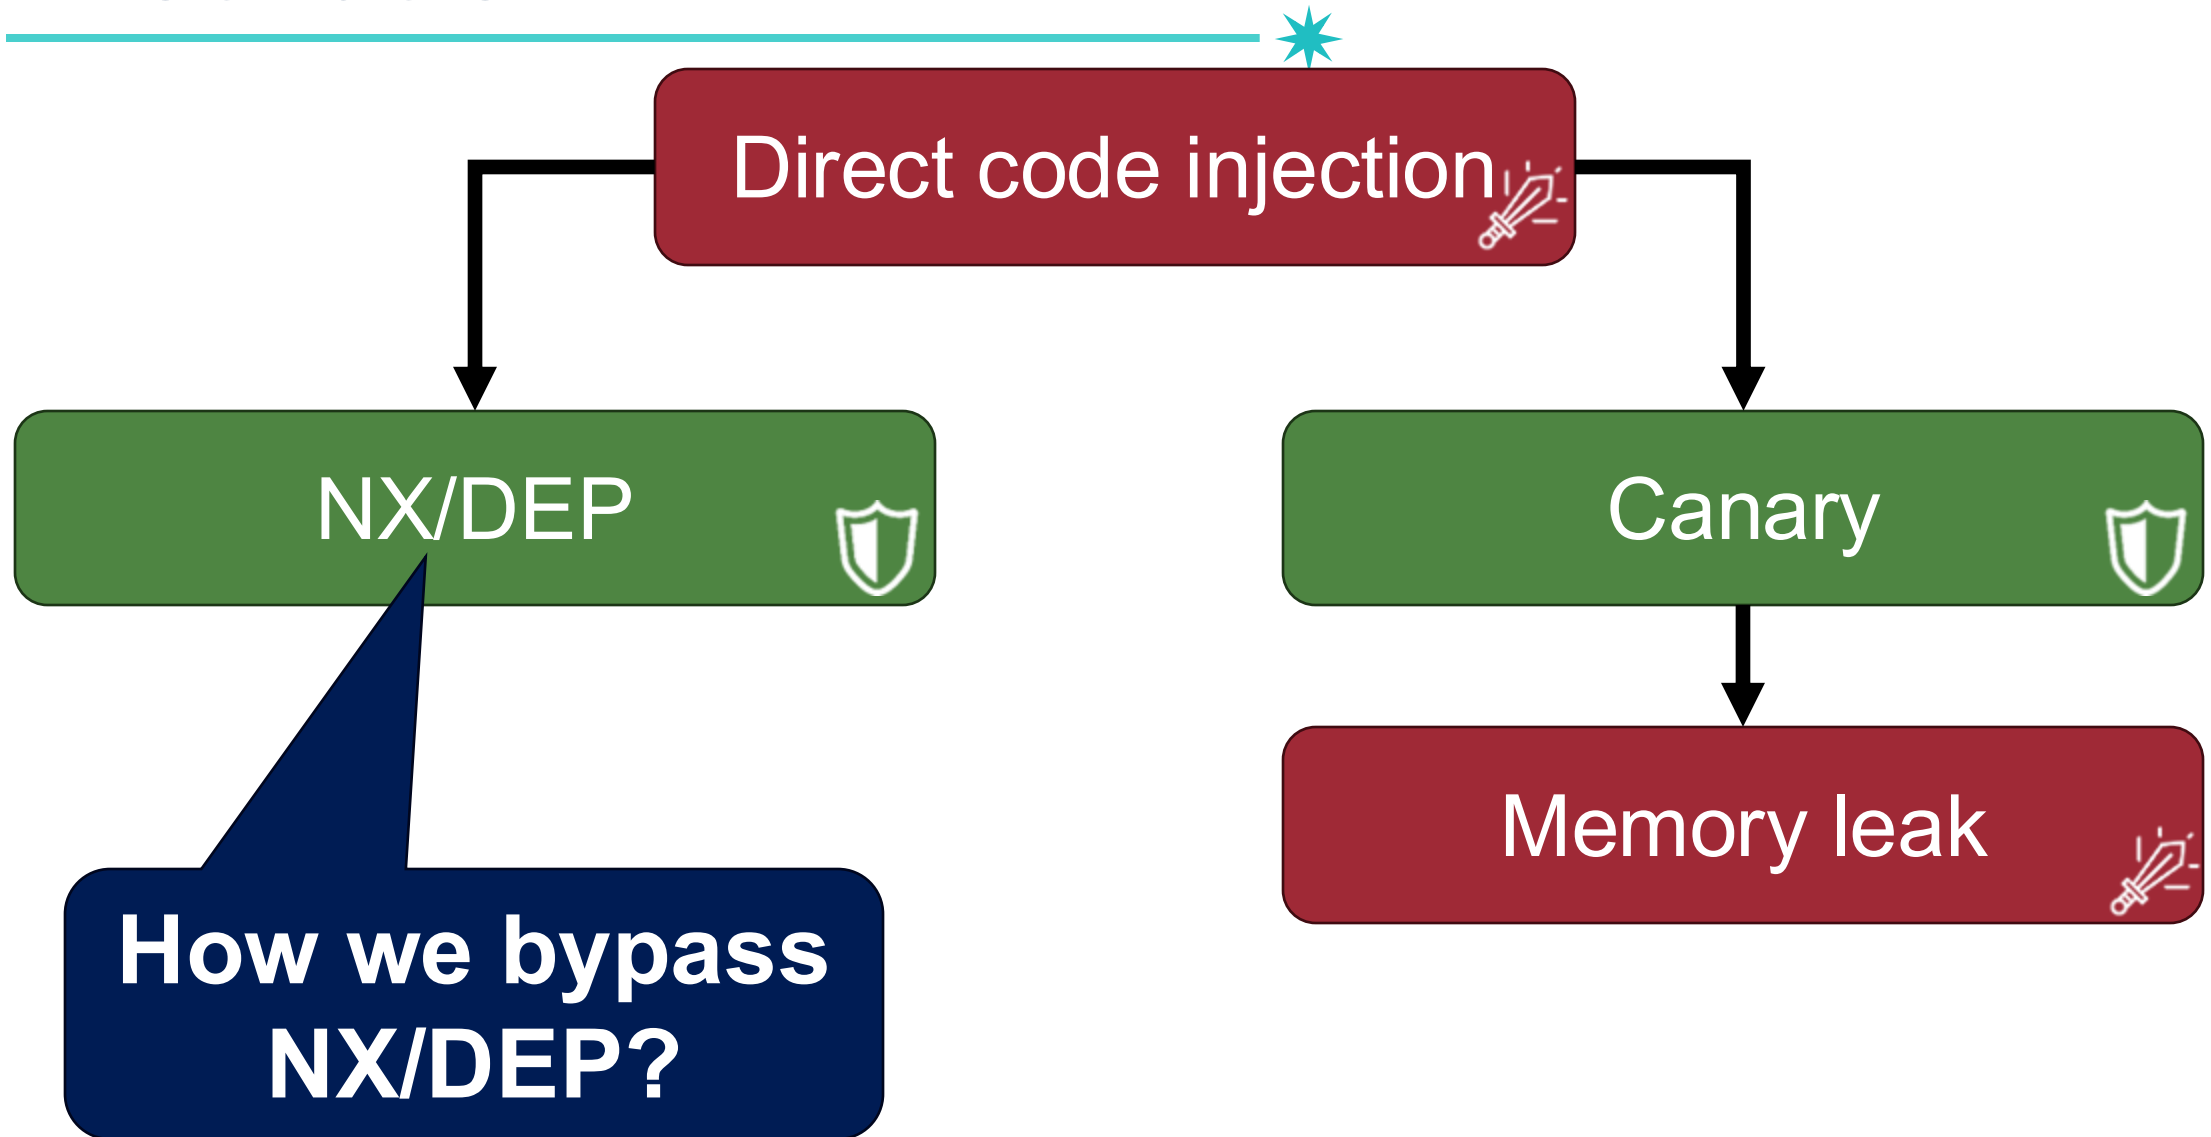
Make this region non-
executable! (e.g., stack
should be non-executable)

Control Hijack Attack / Defense So Far

5



Motivation



Code-Reuse Attacks

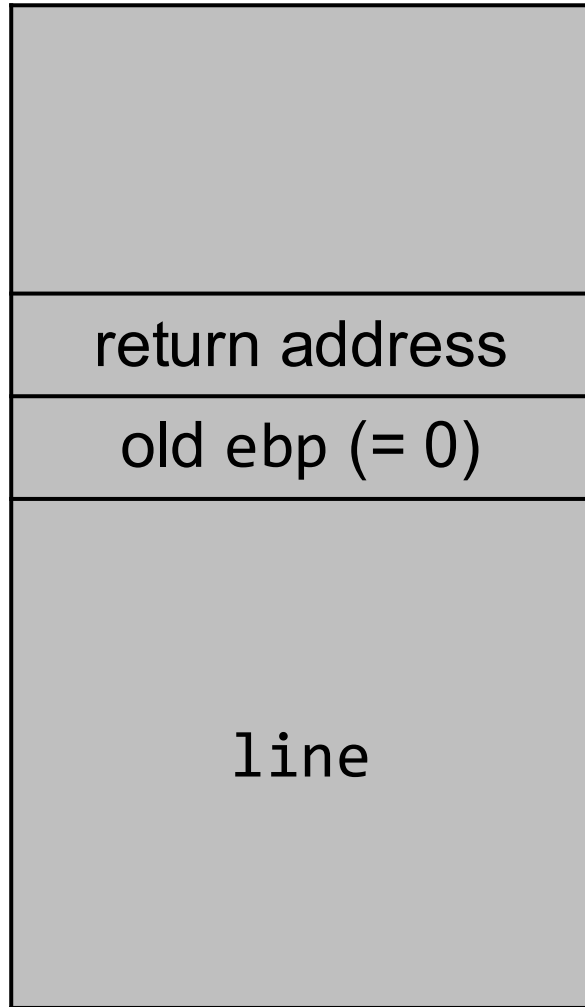
Bypassing DEP



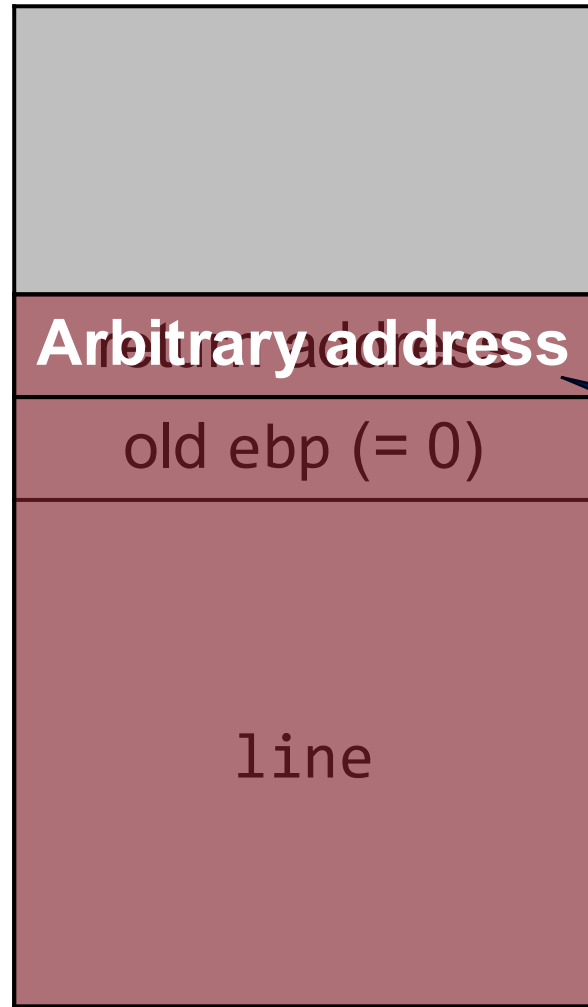
- Return-to-stack exploit is disabled
- But, we can still jump to an arbitrary address of ***existing code*** (= ***Code Reuse Attack***)

Main Idea: Jump to Existing Code

9



Main Idea: Jump to Existing Code



Jump to the *existing code space*, not to the stack

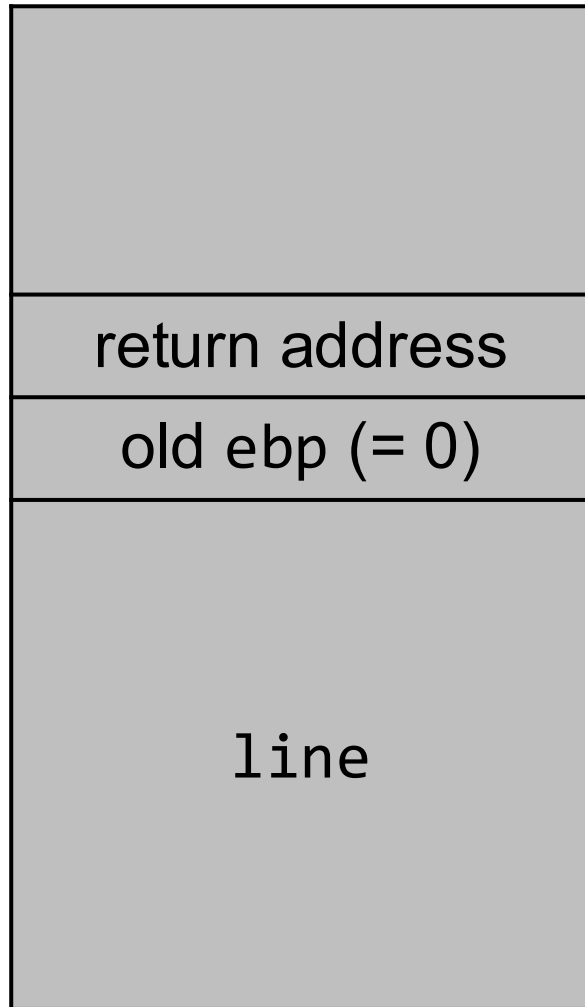
Code Reuse Attack #1: Return-to-Libc



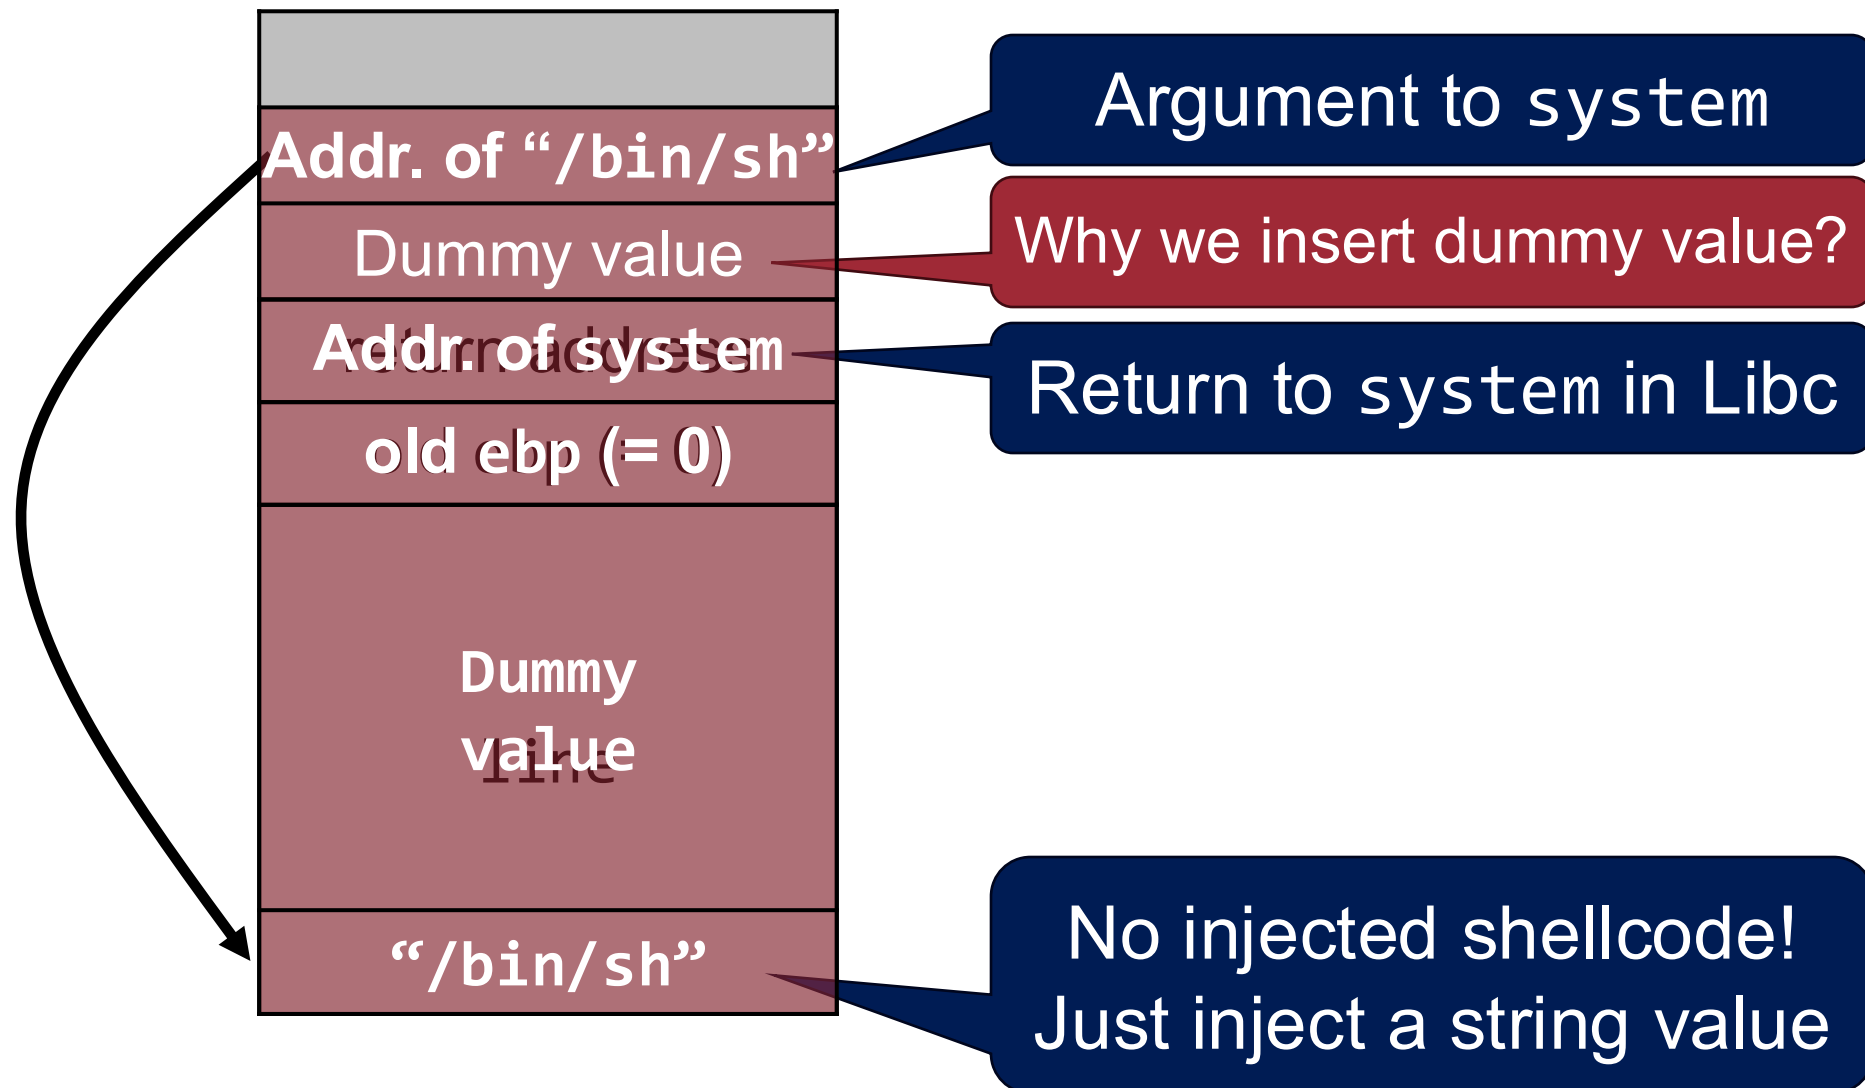
- LIBC (LIBrary C) is a standard library that most programs commonly use
 - For example, printf is in LIBC
- Many useful functions in LIBC to execute
 - exec family: execl, execlp, execl, ...
 - system
 - mprotect
 - mmap

Code Reuse Attack #1: Return-to-Libc

12



Code Reuse Attack #1: Return-to-Libc




Recap: Function Call (call)

14

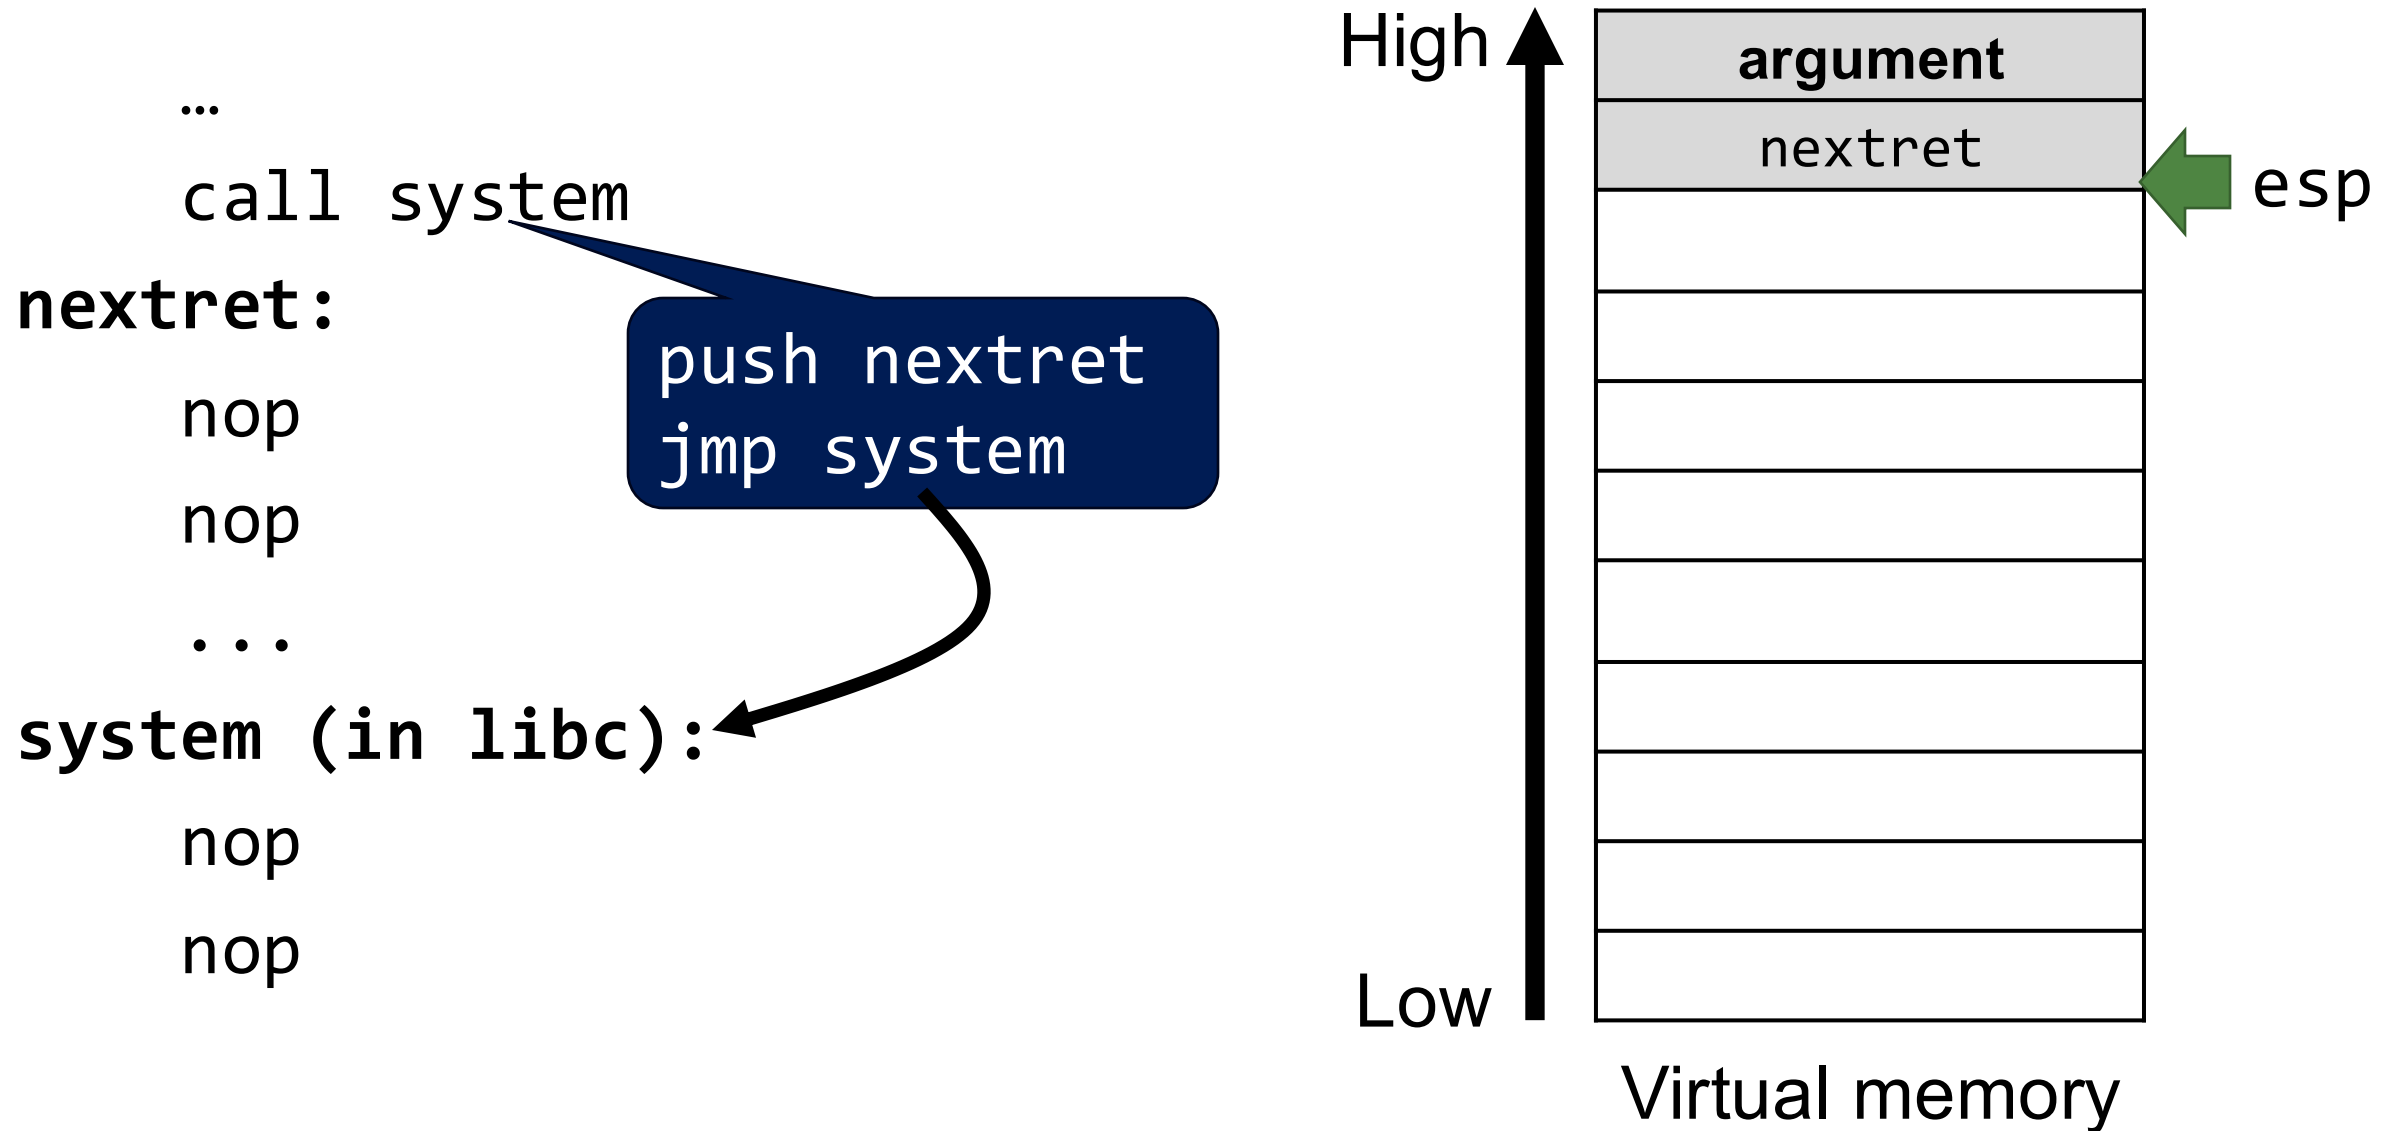
```
...  
call system  
nextret:  
  nop  
  nop  
  ...  
system (in libc):  
  nop  
  nop
```

push nextret
jmp system



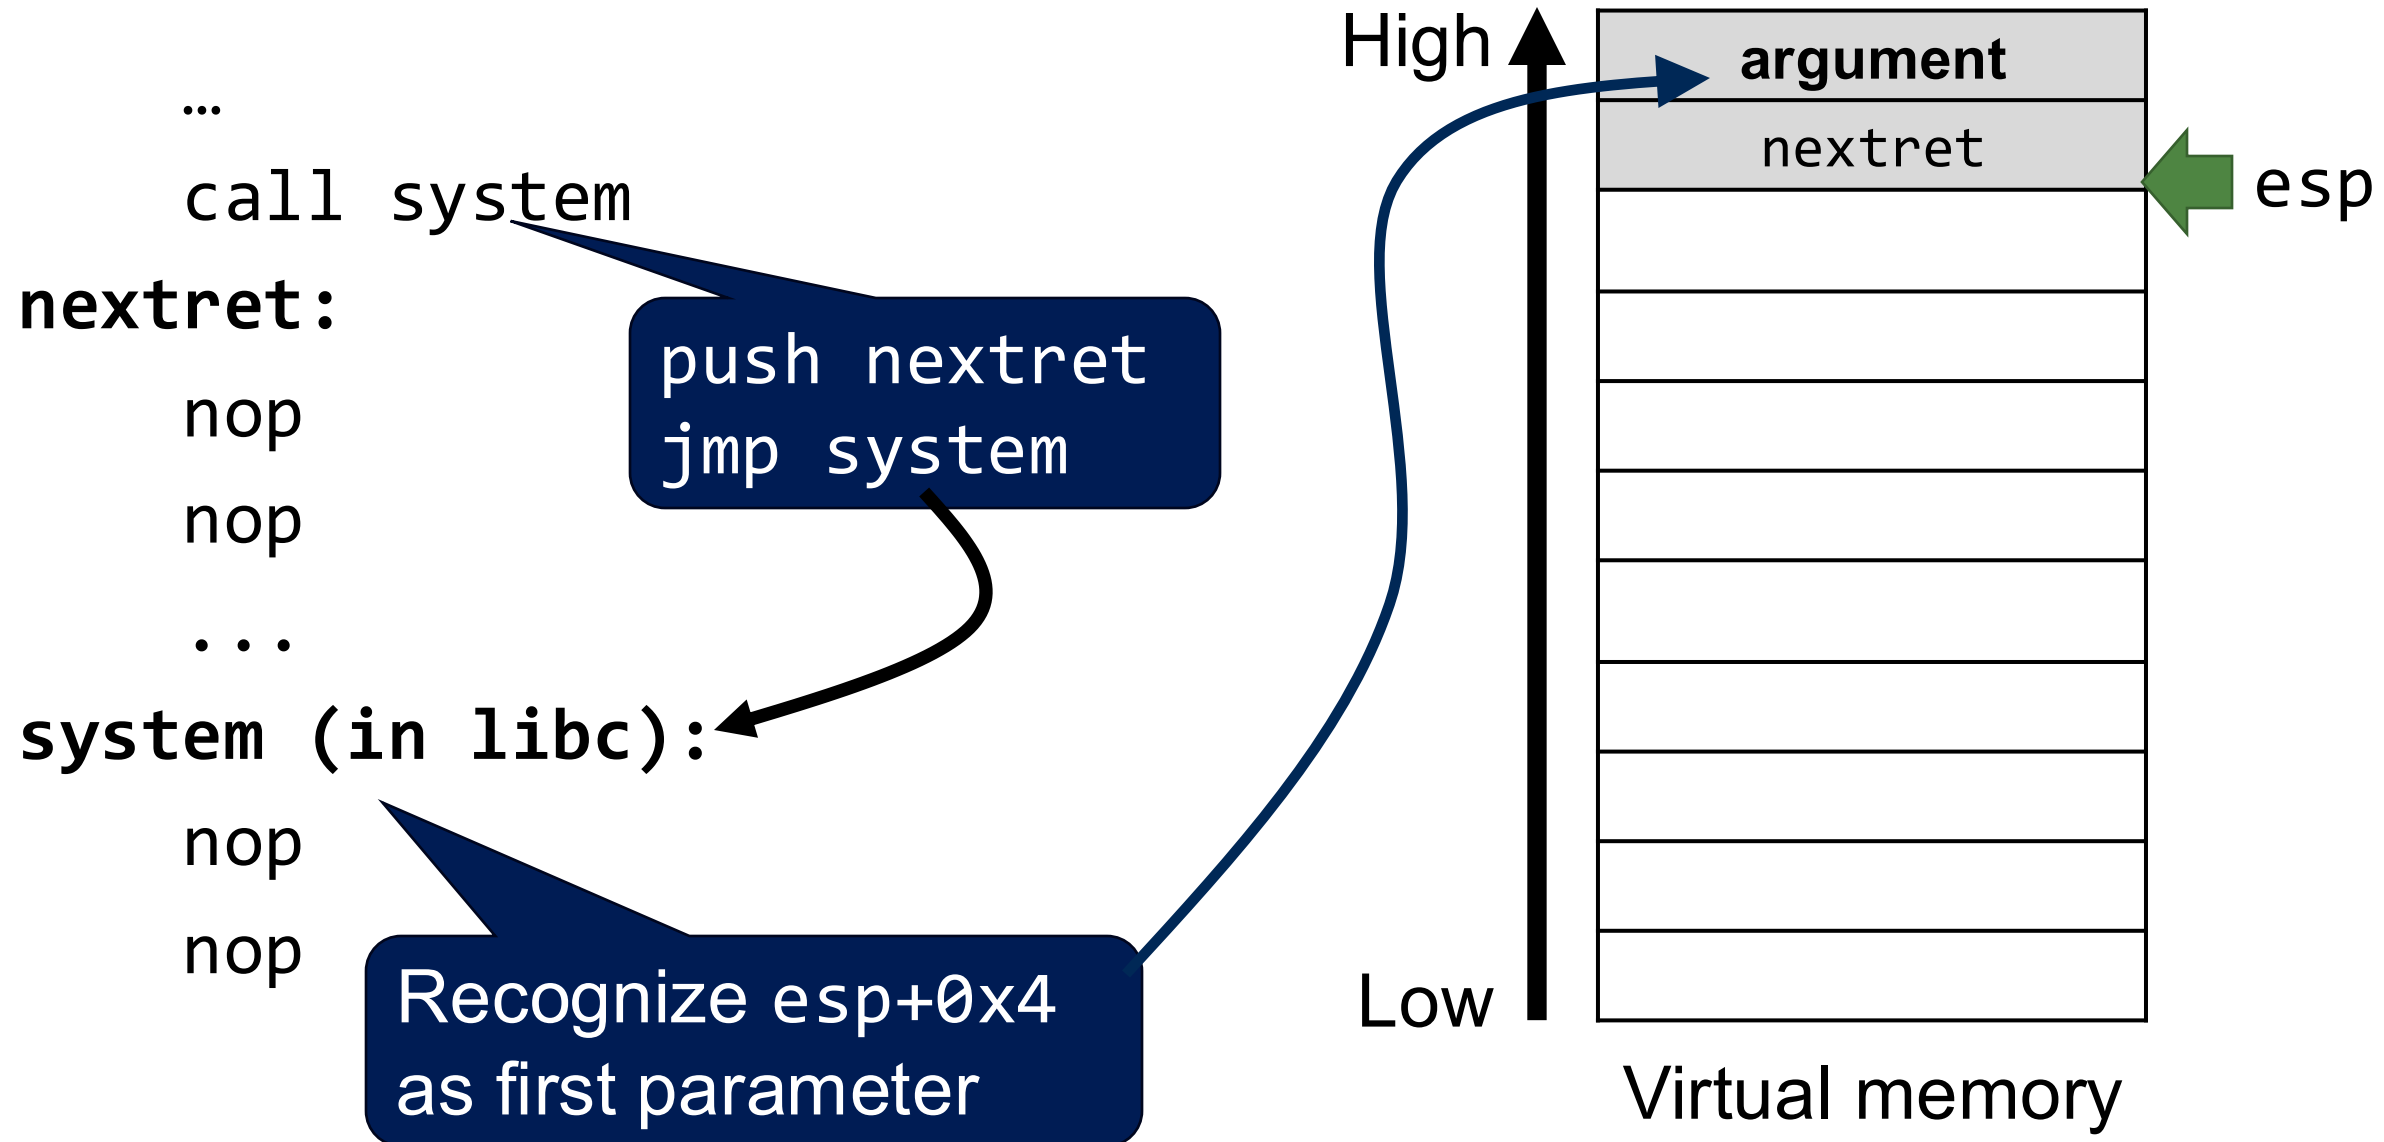
Recap: Function Call (call)

15



Recap: Function Call (call)

16



libc provides System Call Wrapper

08049162 <main>:

8049162:	55	push	ebp
8049163:	89 e5	mov	ebp, esp
8049165:	83 ec 08	sub	esp, 0x8
8049168:	c7 45 f8 08 a0 04 08	mov	DWORD PTR [ebp-0x8], 0x804a008
804916f:	c7 45 fc 00 00 00 00	mov	DWORD PTR [ebp-0x4], 0x0
8049176:	6a 00	push	0x0
8049178:	8d 45 f8	lea	eax, [ebp-0x8]
804917b:	50	push	eax
804917c:	68 08 a0 04 08	push	0x804a008
8049181:	e8 c7 29 02 00	call	806c4b0 <__execve>

First argument of execve

You are actually calling a wrapper function around the syscall

libc provides System Call Wrapper

libc code

0806c4b0 <__execve>:

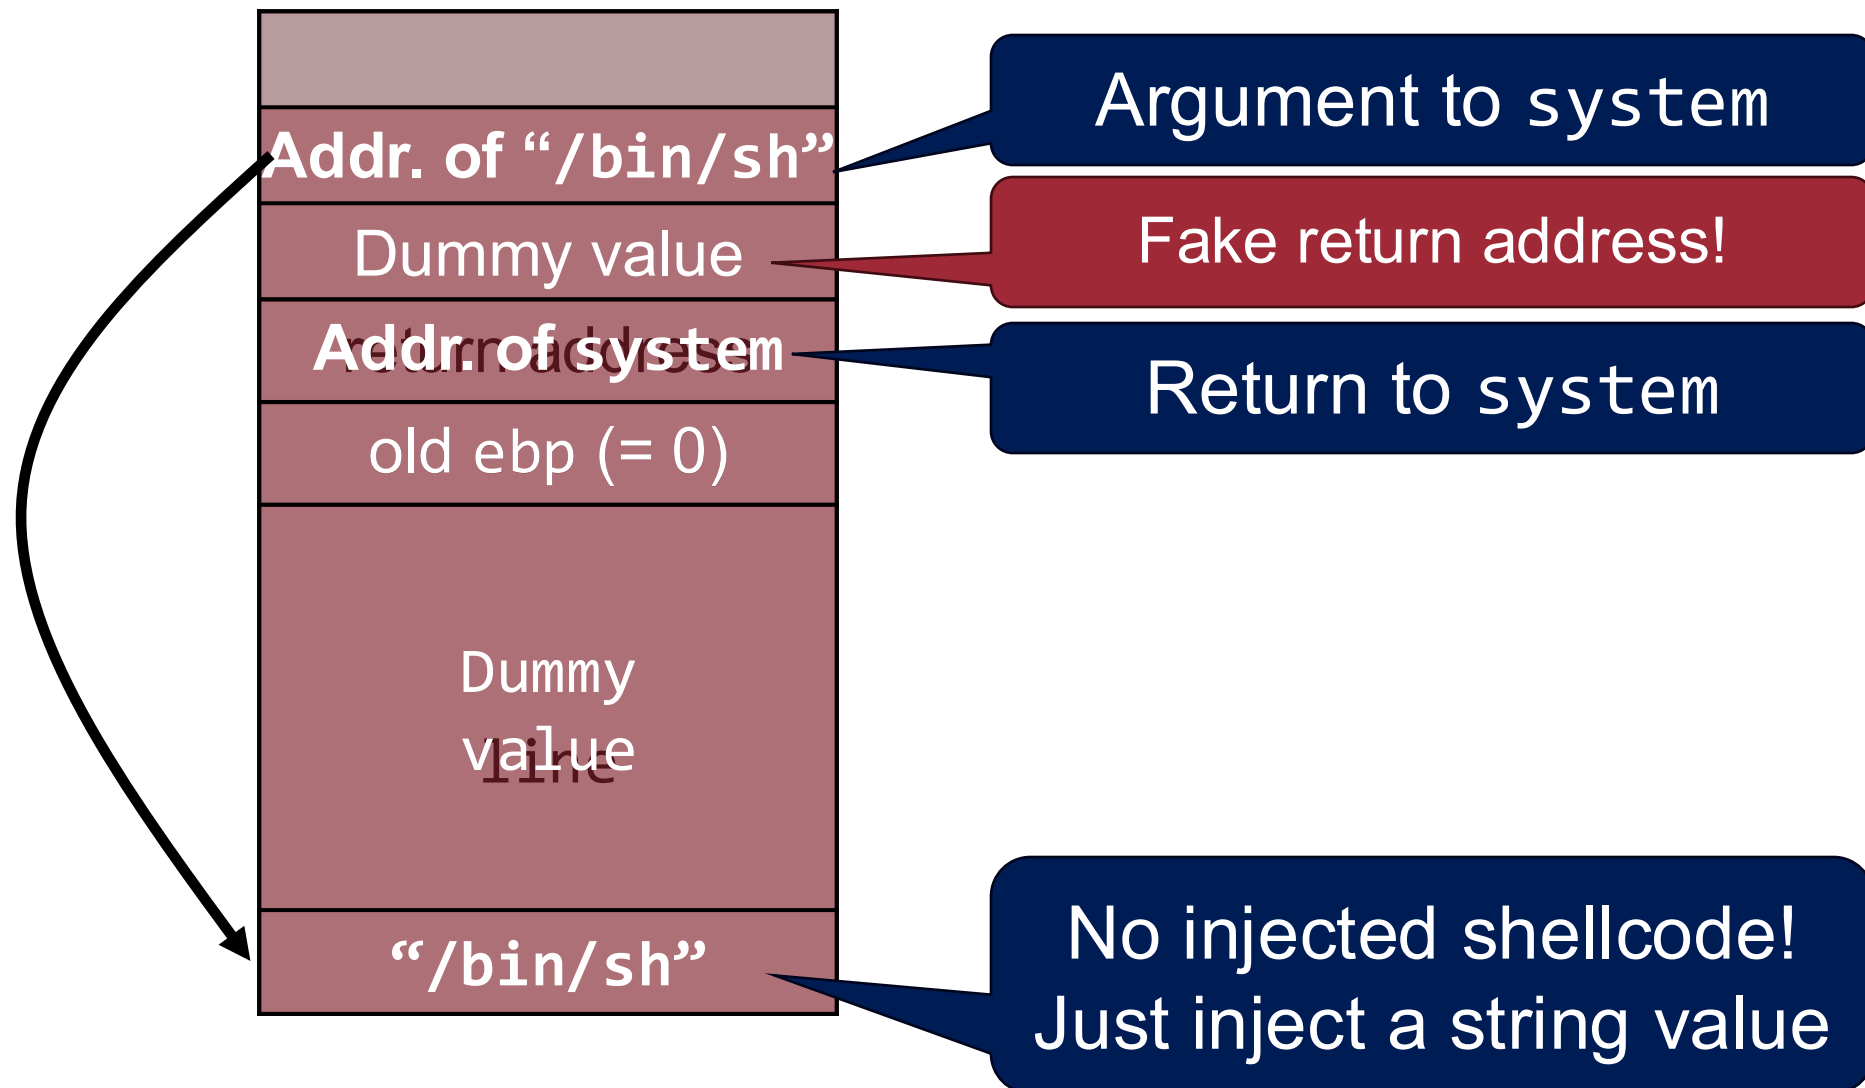
```
806c4b0: 53
806c4b1: 8b 54 24 10
806c4b5: 8b 4c 24 0c
806c4b9: 8b 5c 24 08
806c4bd: b8 0b 00 00 00
806c4c2: cd 80
```

```
push    ebx
mov     edx,DWORD PTR [esp+0x10]
mov     ecx,DWORD PTR [esp+0xc]
mov     ebx,DWORD PTR [esp+0x8]
mov     eax,0xb
int     0x80
```

System Call!

Get first
argument

Code Reuse Attack #1: Return-to-Libc



Motivation of Return-oriented Programming 20



Return-to-LIBC requires LIBC function calls, but ...☹

- Different versions of LIBC

```
attacker_local@environment:~$ ldd --version  
ldd (Ubuntu GLIBC 2.31-0ubuntu9.17) 2.31
```



```
victim@environment:/# ldd --version  
ldd (Ubuntu GLIBC 2.27-3ubuntu1) 2.27
```

Motivation of Return-oriented Programming ²¹



Return-to-LIBC requires LIBC function calls, but ...☹

- Different versions of LIBC
- LIBC may not be used at all
- Some functions in LIBC can be excluded

```
attacker_local@environment:~$ ldd --version  
ldd (Ubuntu GLIBC 2.31-0ubuntu9.17) 2.31
```



```
victim@environment:/# ldd --version  
ldd (Ubuntu GLIBC 2.27-3ubuntu1) 2.27
```

Motivation of Return-oriented Programming 22

Return-to-LIBC requires LIBC function calls, but ...☹

- Different versions of LIBC
- LIBC may not be used at all
- Some functions in LIBC can be excluded

```
attacker_local@environment:~$ ldd --version  
ldd (Ubuntu GLIBC 2.31-0ubuntu9.17) 2.31
```

***Can we spawn a shell
without the use of LIBC functions?***

Return-oriented Programming (ROP)

Code Reuse Attack #2: ROP



Generalized Code Reuse Attack

Formally introduced by Hovav in CCS 2007

“The Geometry of Innocent Flesh on the Bone: Return-to-libc without Function Calls (on the x86)”

The Geometry of Innocent Flesh on the Bone:
Return-into-libc without Function Calls (on the x86)

Hovav Shacham*
hovav@cs.ucsd.edu

Abstract

We present new techniques that allow a return-into-libc attack to be mounted on x86 executables that calls *no functions at all*. Our attack combines a large number of short instruction sequences to build *gadgets* that allow arbitrary computation. We show how to discover such instruction sequences by means of static analysis. We make use, in an essential way, of the properties of the x86 instruction set.

1 Introduction

We present new techniques that allow a return-into-libc attack to be mounted on x86 executables that is every bit as powerful as code injection. We thus demonstrate that the widely deployed “W⊕X” defense, which rules out code injection but allows return-into-libc attacks, is much less useful than previously thought.

Attacks using our technique call no functions whatsoever. In fact, the use instruction sequences from libc that weren’t placed there by the assembler. This makes our attack resilient to defenses that remove certain functions from libc or change the assembler’s code generation choices.

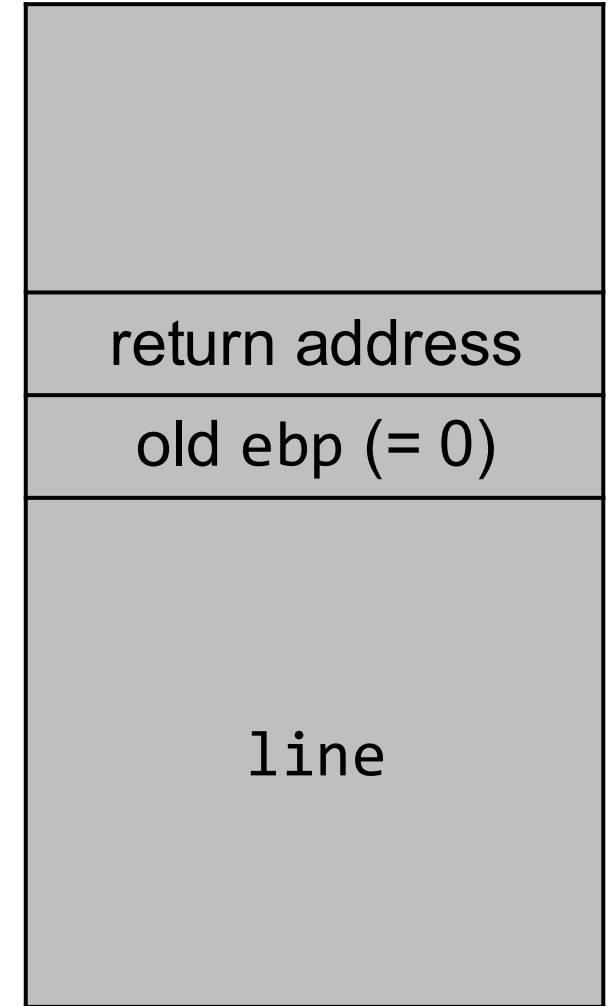
Unlike previous attacks, ours combines a large number of short instruction sequences to build

Main Idea: Return (ret) Chaining

Attacker's goal:

execute following instructions

```
add eax, ebx
mov ecx, eax
inc ecx
mov edx, 42
```

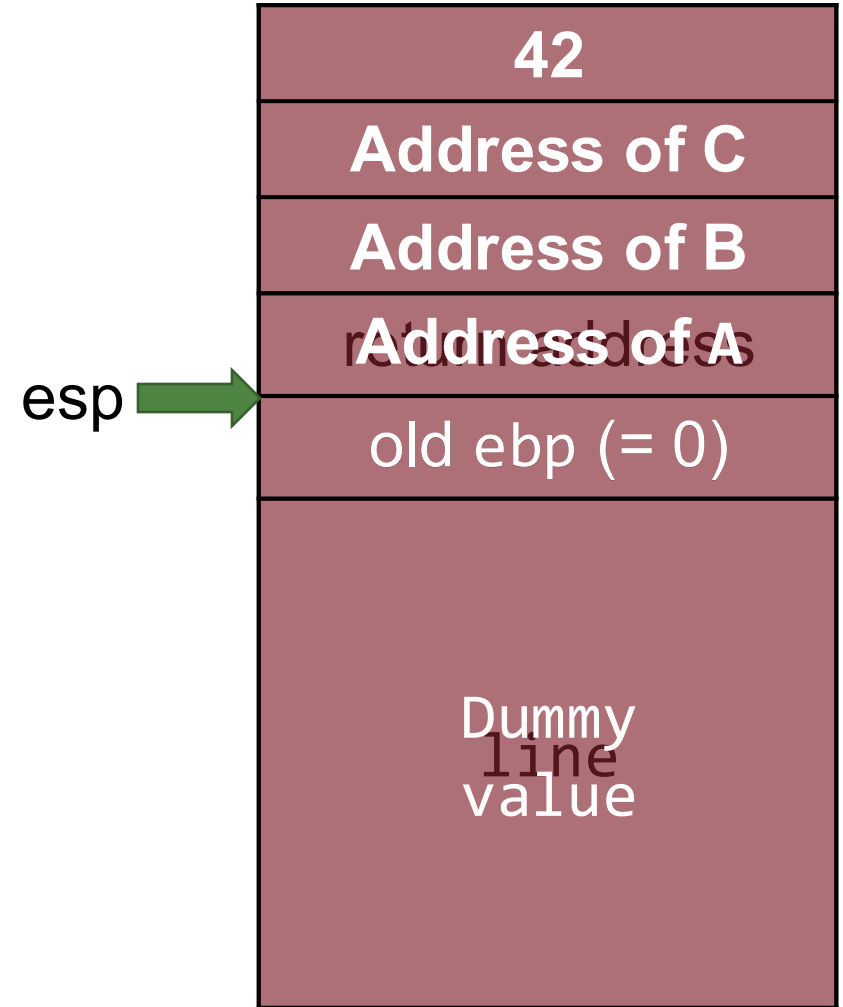


Main Idea: Return (ret) Chaining

Attacker's goal:

execute following instructions

```
add eax, ebx
mov ecx, eax
inc ecx
mov edx, 42
```



Main Idea: Return (ret) Chaining

Attacker's goal:

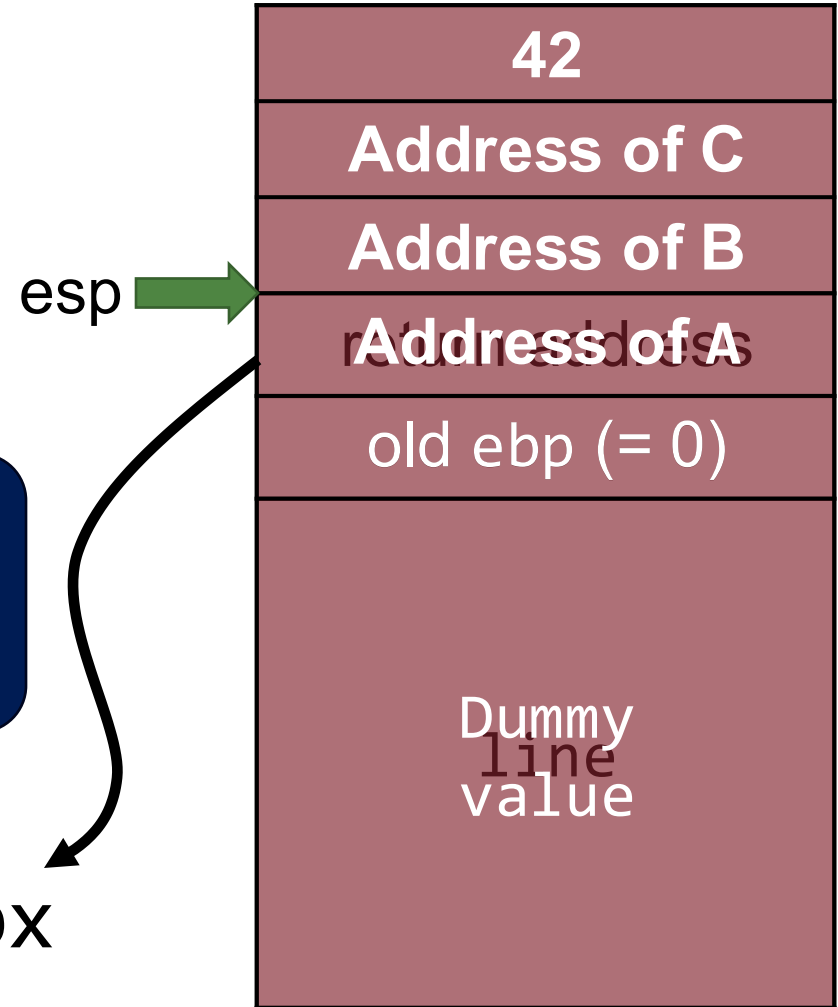
execute following instructions

```
add eax, ebx  
mov ecx, eax  
inc ecx  
mov edx, 42
```

Somewhere in the
binary code

A

```
add eax, ebx  
ret
```



Main Idea: Return (ret) Chaining


Attacker's goal:

execute following instructions

```
add eax, ebx
mov ecx, eax
inc ecx
mov edx, 42
```

ROP Gadget:
Instruction sequence
that ends with ret

A | add eax, ebx
ret



Main Idea: Return (ret) Chaining

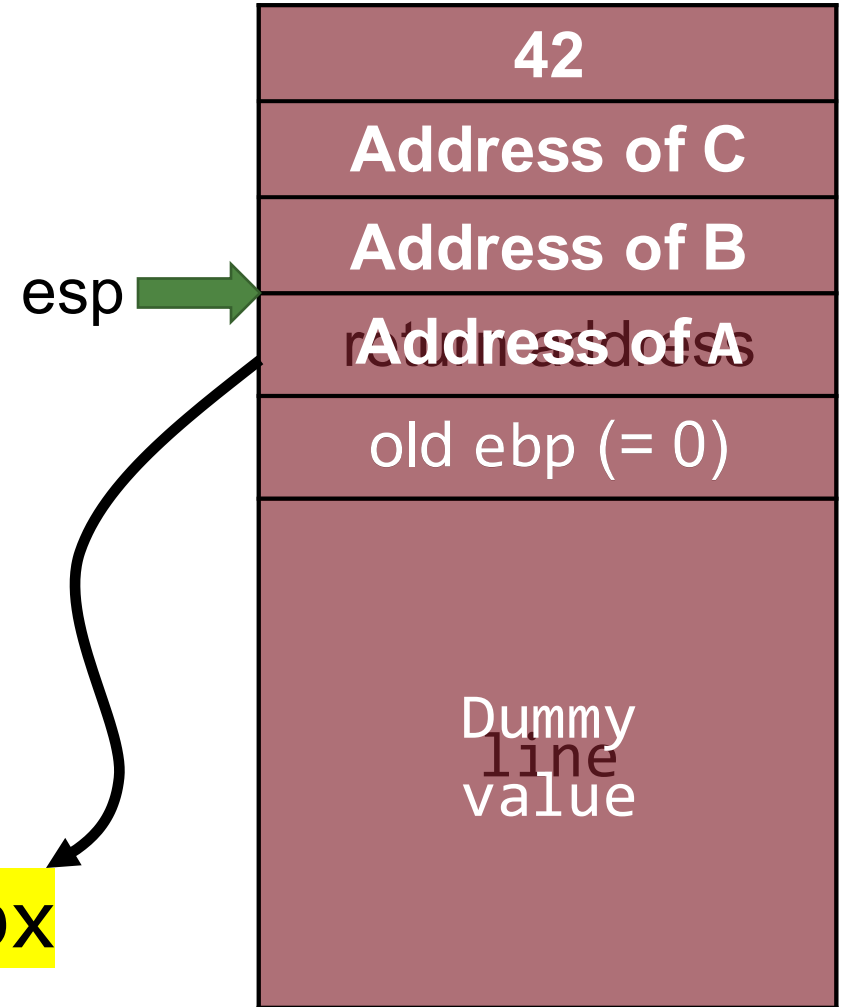
Attacker's goal:

execute following instructions

```
add eax, ebx  
mov ecx, eax  
inc ecx  
mov edx, 42
```

A

```
add eax, ebx  
ret
```



Main Idea: Return (ret) Chaining


Attacker's goal:

execute following instructions

```
add eax, ebx  
mov ecx, eax  
inc ecx  
mov edx, 42
```

pop eip
= jump to another
gadget

A | add eax, ebx
| ret



Main Idea: Return (ret) Chaining

Attacker's goal:

execute following instructions

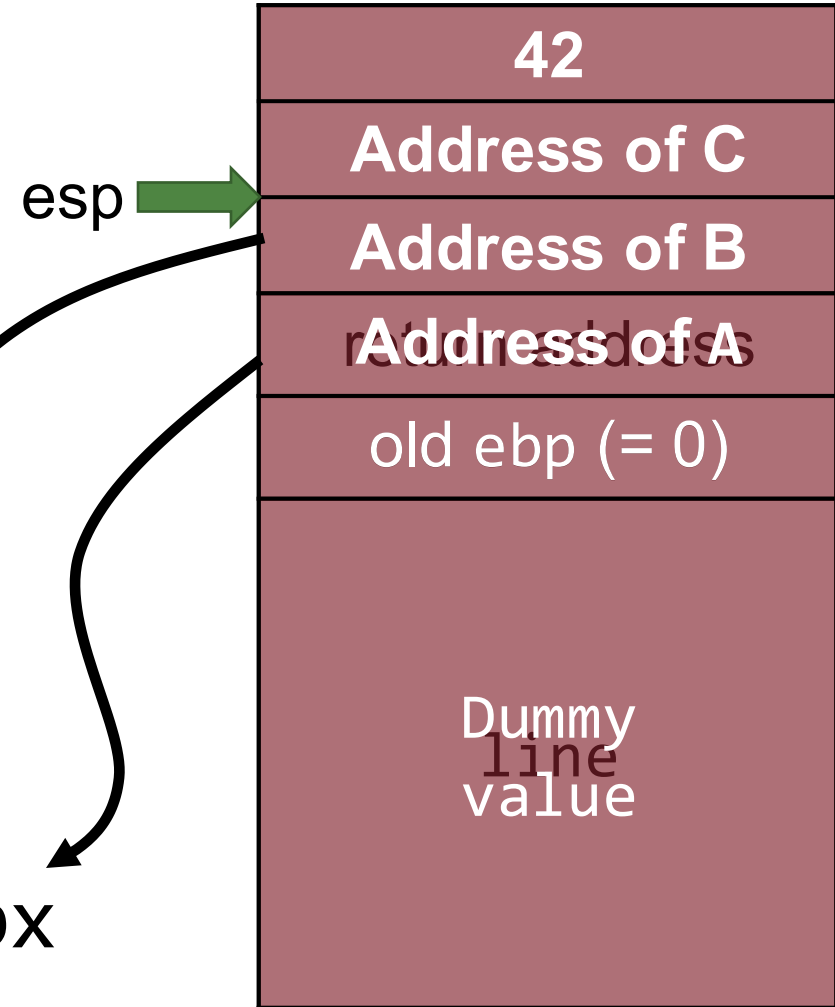
```
add eax, ebx  
mov ecx, eax  
inc ecx  
mov edx, 42
```

B

```
mov ecx, eax  
ret
```

A

```
add eax, ebx  
ret
```



Main Idea: Return (ret) Chaining


Attacker's goal:

execute following instructions

```
add eax, ebx  
mov ecx, eax  
inc ecx  
mov edx, 42
```

B | mov ecx, eax
| ret

A | add eax, ebx
| ret



Main Idea: Return (ret) Chaining

Attacker's goal:


execute following instructions

```
add eax, ebx
mov ecx, eax
inc ecx
mov edx, 42
```

C | **inc ecx**
pop edx
ret

B | mov ecx, eax
ret

A | add eax, ebx
ret



Main Idea: Return (ret) Chaining

Attacker's goal:


execute following instructions

```
add eax, ebx
mov ecx, eax
inc ecx
mov edx, 42
```

C | `inc ecx`
`pop edx`
`ret`

B | `mov ecx, eax`
`ret`

A | `add eax, ebx`
`ret`



Main Idea: Return (ret) Chaining

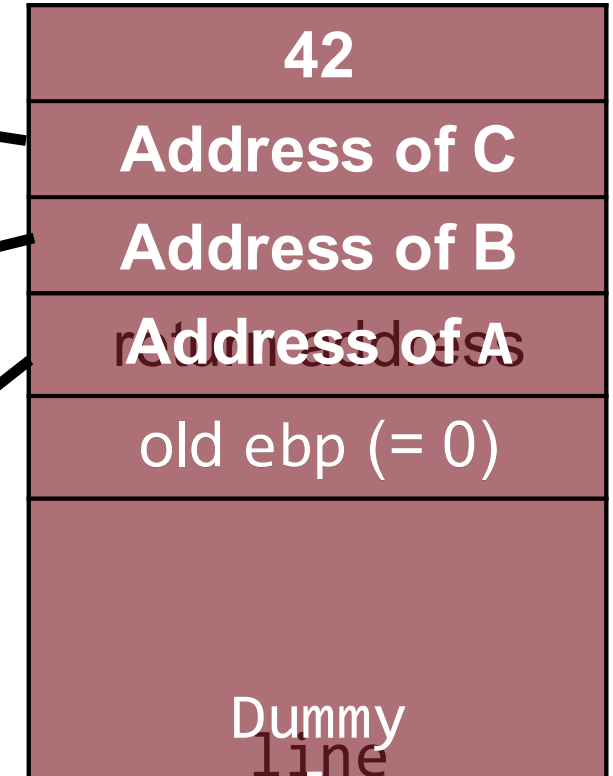
Attacker's goal:

execute following instructions

```
add eax, ebx
mov ecx, eax
inc ecx
mov edx, 42
```

C | `inc ecx`
 | **`pop edx`**
 | `ret`

B | `mov ecx, eax`
 | `ret`



Return chaining with ROP gadgets
allows arbitrary computation!

ROP Practice



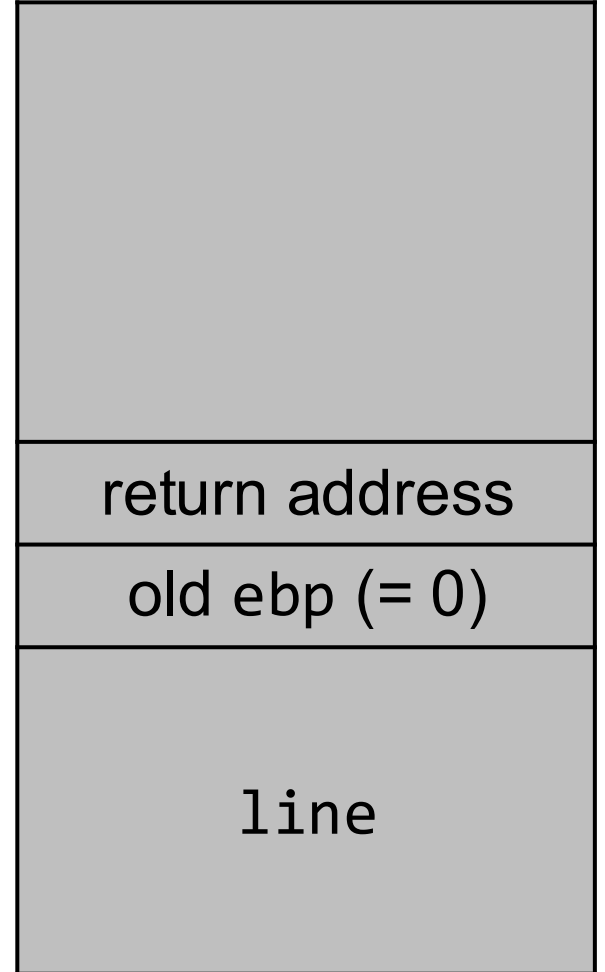
Goal: Modify ptr to be 0x42424242 with ROP

```
mov [ptr], 0x42424242
```

Gadget A		pop eax
		ret

Gadget B		pop ebx
		ret

Gadget C		mov [eax], ebx
		ret



ROP Workflow



1. Disassemble binary
2. Identify useful instruction sequences (i.e., gadgets)
 - E.g., an instruction sequence that ends with `ret` is useful
 - E.g., an instruction sequence that ends with `jmp reg` can be useful
(`pop eax; jmp eax`)
3. Assemble gadgets to perform some computation
 - E.g., spawning a shell

Challenge: Gathering as many gadgets as possible

Many Gadgets in Regular Binaries?

x86 instructions have their lengths ranging from 1 byte to 18 bytes, i.e., it uses ***variable-length encoding***

x86 instructions have
variable lengths

08048aac <main>:

```

8048aac: 8d 4c 24 04
8048ab0: 83 e4 f0
8048ab3: ff 71 fc
8048ab6: 55
8048ab7: 89 e5
8048ab9: 51
8048aba: 83 ec 14
8048abd: c7 45 f0 88 ad 0a 08
8048ac4: c7 45 f4 00 00 00 00
8048acb: 83 ec 04
8048ace: 6a 00
8048ad0: 8d 45 f0
8048ad3: 50
8048ad4: 68 88 ad 0a 08
8048ad9: e8 02 39 01 00

```

...

```

lea    ecx,[esp+0x4]
and     esp,0xffffffff0
push   DWORD PTR [ecx-0x4]
push   ebp
mov     ebp,esp
push   ecx
sub     esp,0x14
mov     DWORD PTR [ebp-0x10],0x80aad88
mov     DWORD PTR [ebp-0xc],0x0
sub     esp,0x4
push   0x0
lea     eax,[ebp-0x10]
push   eax
push   0x80aad88
call   805c3e0 <__execve>

```

Many Gadgets in Regular Binaries?



x86 instructions have their lengths ranging from 1 byte to 18 bytes, i.e., it uses ***variable-length encoding***

Therefore, there can be both **intended** and **unintended gadgets** in x86 binaries

Disassembling x86



eip



e8 05 ff ff ff

81 c3 59 12 00 00

call 8048330

add ebx,0x1259

What if we disassemble the code
from the second byte (05)?

Unintended ret Instruction



eip



e8 05 ff ff ff

81 c3 59 12 00 00

add eax, 0x81ffffff

ret

Totally different, but still valid instructions!

Unintended ret Instruction



eip



e8 05 ff ff ff
81 c3 59 12 00 00

add eax, 0x81ffffff
ret

Unintended ret Instruction



eip



e8	05	ff	ff	ff	
81	c3	59	12	00	00

add	eax,	0x81ffffff
ret		

Disassemble from Any Addresses in Memory Pages

45



- This is perfectly legal
- We can find lots of ***unintended*** ret instructions

Finding Unintended Gadgets

Algorithm GALILEO:

```
create a node, root, representing the ret instruction;  
place root in the trie;  
for pos from 1 to textseg_len do:  
    if the byte at pos is c3, i.e., a ret instruction, then:  
        call BUILDFROM(pos, root).
```

Find *c3* (*ret*)
instruction

Procedure BUILDFROM(index *pos*, instruction *parent_insn*):

```
for step from 1 to max_insn_len do:  
    if bytes  $[(pos - step) \dots (pos - 1)]$  decode as a valid instruction insn then:  
        ensure insn is in the trie as a child of parent_insn;  
        if insn isn't boring then:  
            call BUILDFROM(pos - step, insn).
```

Finding Unintended Gadgets

Algorithm GALILEO:

create a node, *root*, representing the `ret` instruction;
place *root* in the trie;

for *pos* **from** 1 **to** *textseg_len* **do**:

if the byte at *pos* is `c3`, i.e., a `ret` instruction, **then**:

call `BUILDFROM(pos, root)`.

Find `c3` (`ret`)
instruction

Procedure `BUILDFROM(index pos, instruction`

for *step* **from** 1 **to** *max_insn_len* **do**:

if bytes $[(pos - step) \dots (pos - 1)]$ **do**:

ensure *insn* is in the trie as a child of *parent_insn*;

if *insn* isn't boring **then**:

call `BUILDFROM(pos - step, insn)`.

"Boring" Instructions

1. The *insn* is a leave instruction
2. The *insn* is `pop ebp`
3. The *insn* is unconditional jump

Many Gadgets in Regular Binaries?



Also, program size may matter!

Larger code \Rightarrow More chance to get useful gadgets

Exploit Hardening Made Easy,
USENIX Security 2011

Q: Exploit Hardening Made Easy

Edward J. Schwartz, Thanassis Avgerinos and David Brumley
Carnegie Mellon University, Pittsburgh, PA
{edmcman, thanassis, dbrumley}@cmu.edu

Abstract

Prior work has shown that return oriented programming (ROP) can be used to bypass $W\oplus X$, a software defense that stops shellcode, by reusing instructions from large libraries such as `libc`. Modern operating systems have since enabled address randomization (ASLR), which randomizes the location of `libc`, making these techniques unusable in practice. However, modern ASLR implementations leave smaller amounts of executable code unrandomized and it has been unclear whether an attacker can use these small code fragments to construct payloads in the general case.

In this paper, we show defenses as currently deployed can be bypassed with new techniques for automatically

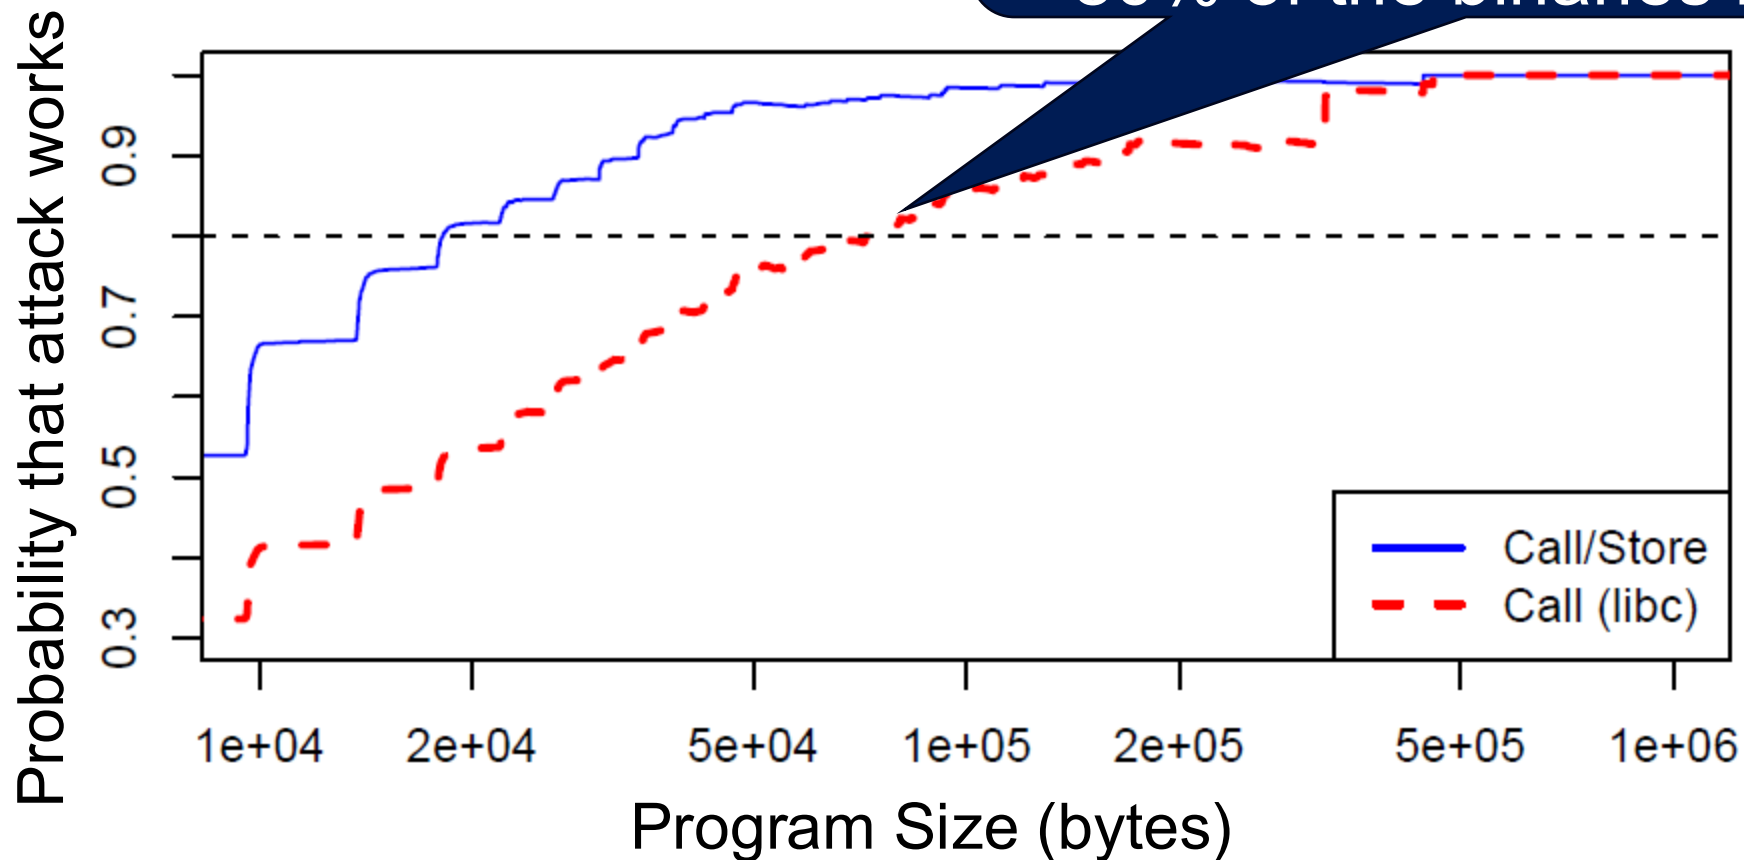
could be to spawn a remote shell to control the program, to install malware, or to exfiltrate sensitive information stored by the program.

Luckily, modern OSes now employ $W\oplus X$ and ASLR together — two defenses intended to thwart control flow hijacks. Write xor eXecute ($W\oplus X$, also known as DEP) prevents an attacker's payload itself from being directly executed. Address space layout randomization (ASLR) prevents an attacker from utilizing structures within the application itself as a payload by randomizing the addresses of program segments. These two defenses, when used together, make control flow hijack vulnerabilities difficult to exploit.

However, ASLR and $W\oplus X$ are not enforced com-

For All /usr/bin Programs

Show that 100KB was enough to successfully create exploits for 80% of the binaries in /usr/bin



Question



How can we mitigate code reuse attacks (ROP)?

Defenses against Code Reuse Attacks

- Detection
 - ROP Payload Detection Using Speculative Code Execution, **MALWARE 2011**
 - Transparent ROP Exploit Mitigation Using Indirect Branch Tracing, **USENIX Security 2013**
 - ROPecker: A Generic and Practical Approach for Defending Against ROP Attacks, **NDSS 2014**



New Attack

Size Does Matter: Why Using Gadget-Chain Length to Prevent Code-Reuse Attacks is Hard, **USENIX Security 2014**


Defenses against Code Reuse Attacks



- Code Modification & Randomization
 - Defeating Return-Oriented Rootkits With “Return-less” Kernels, ***EuroSys 2010***
 - Binary stirring: Self-randomizing Instruction Addresses of Legacy x86 Binary Code, ***CCS 2012***
 - Smashing the Gadgets: Hindering Return-Oriented Programming using in-Place Code Randomization, ***Oakland 2012***
- Enforcing Safety Policy
 - Control-Flow Integrity, ***CCS 2005***
 - Securing Software by Enforcing Data-Flow Integrity, ***NDSS 2014***
 - Code-Pointer Integrity, ***OSDI 2014***

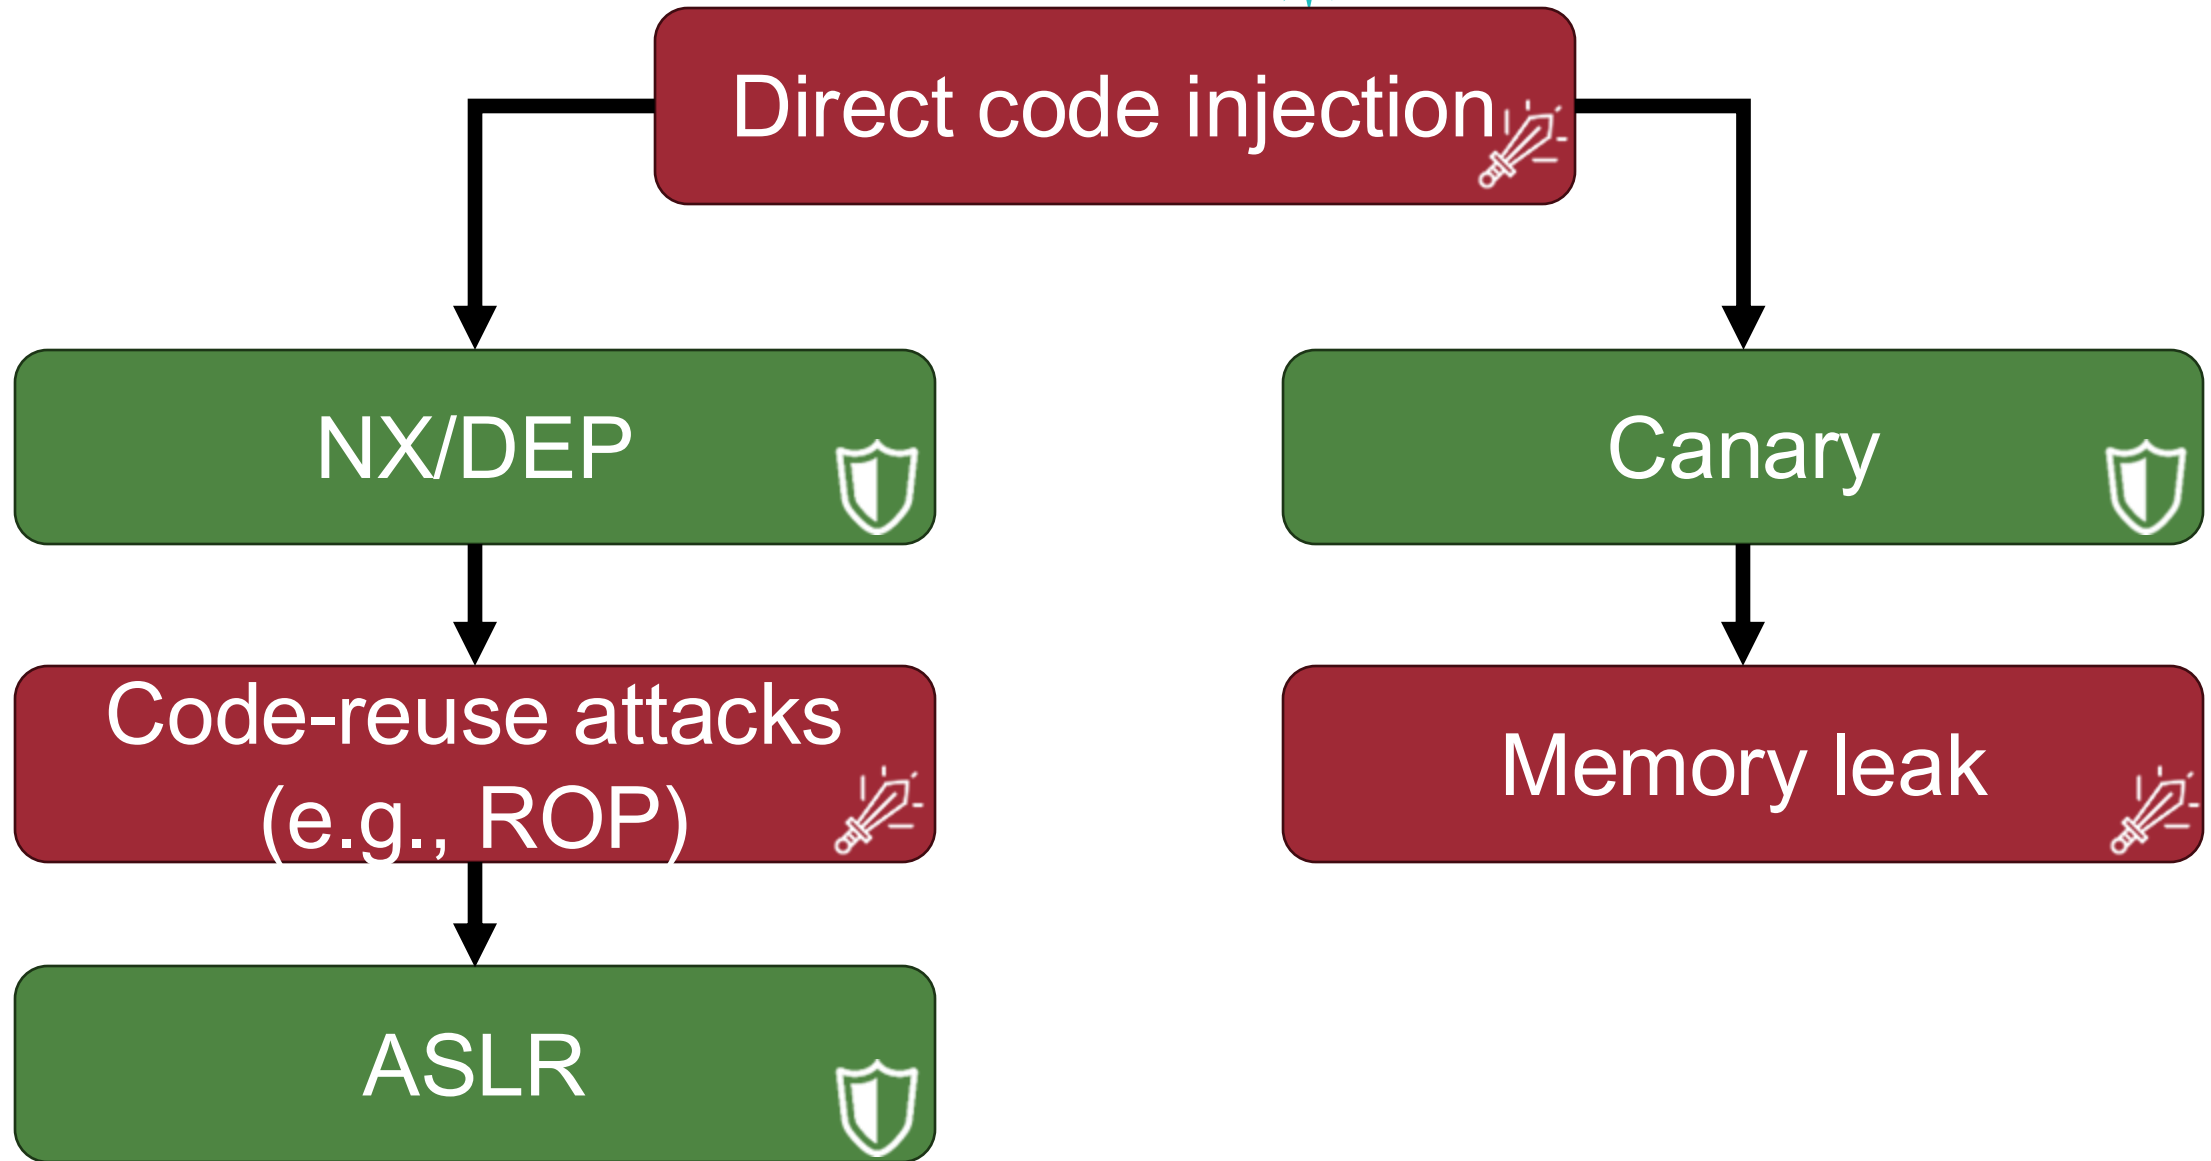
DEP and Code Reuse Attacks

53




Control Hijack Attack / Defense So Far

54



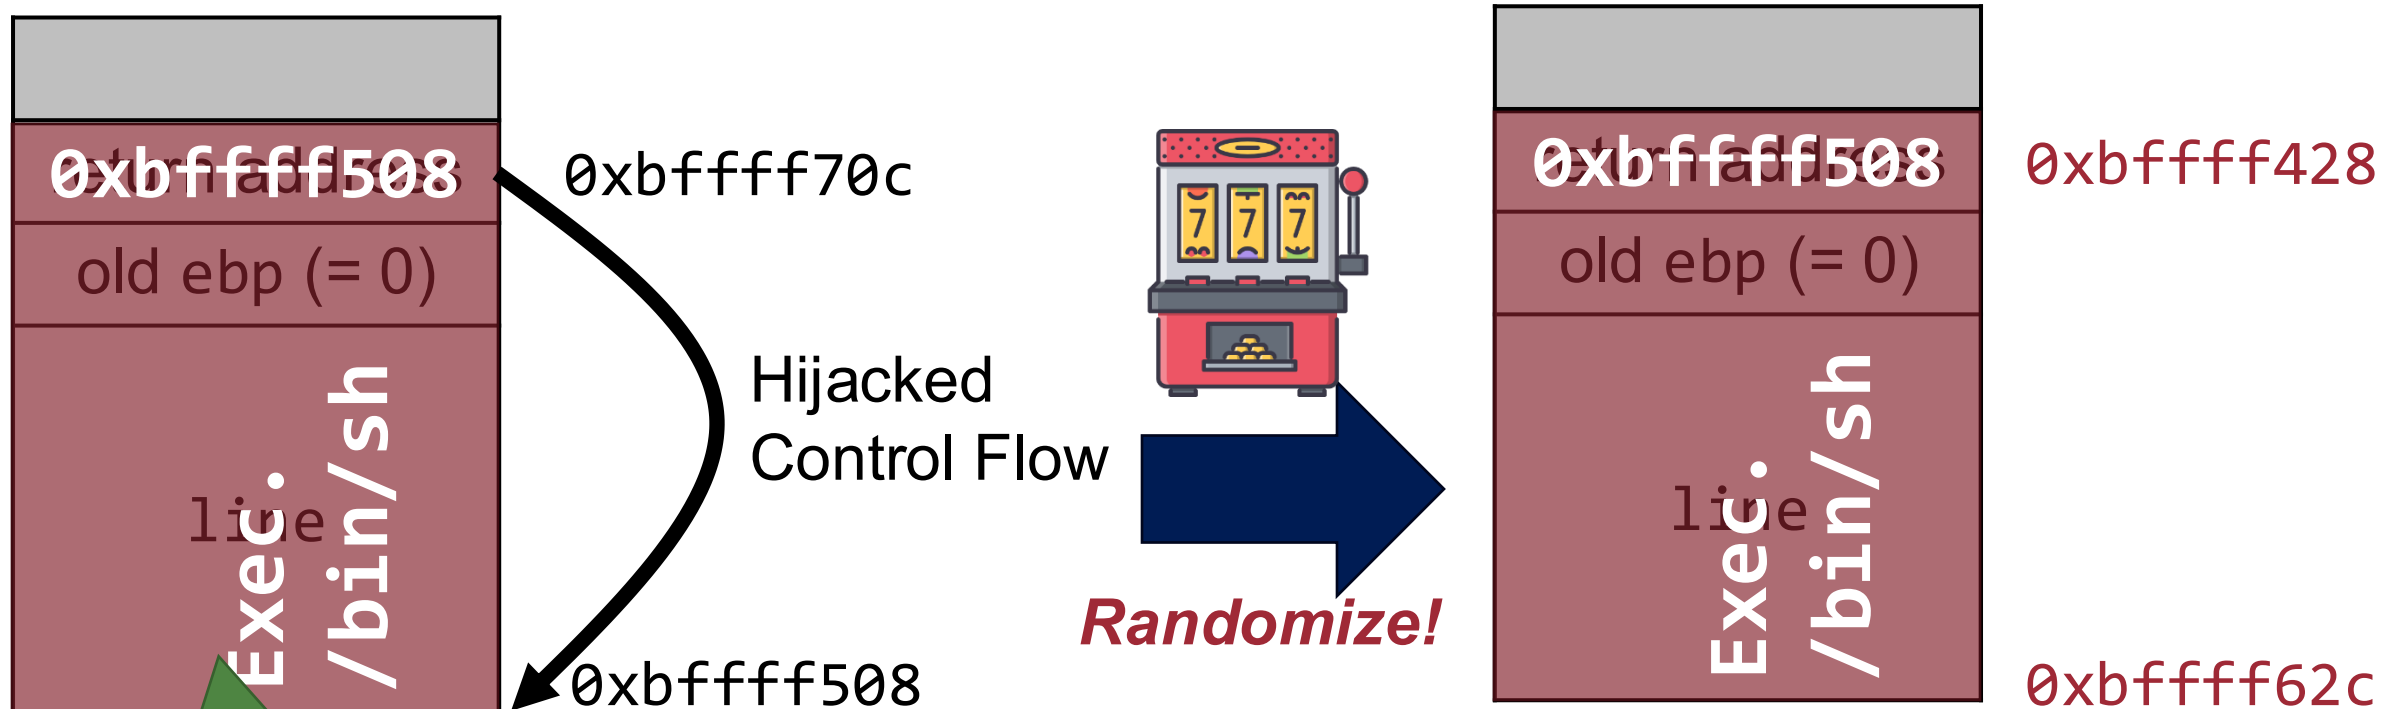
Address Space Layout Randomization (ASLR)

Control Flow Hijack Attack



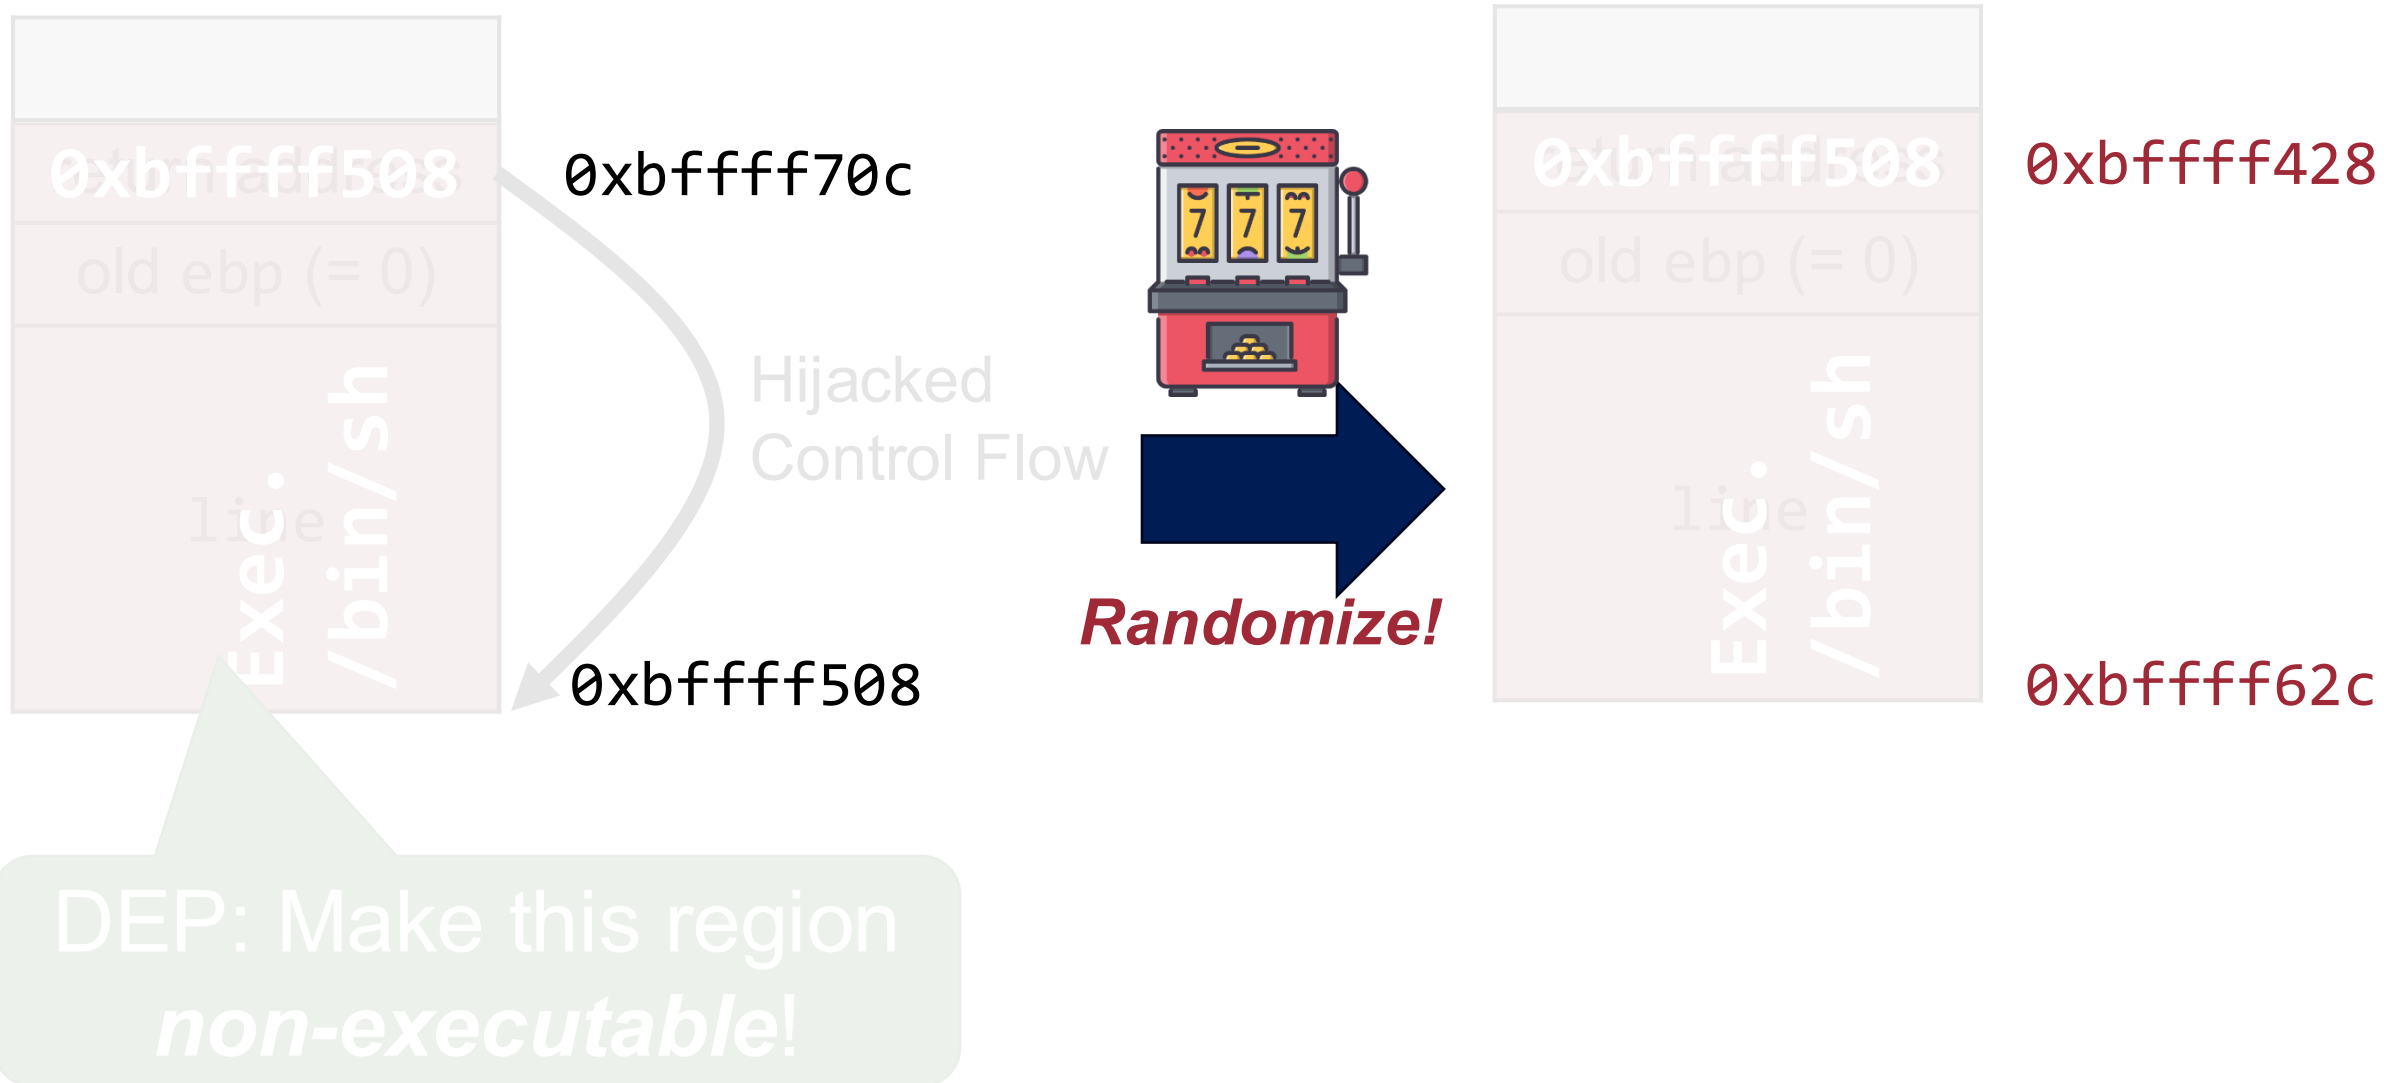
DEP: Make this region
non-executable!

Different Perspective: ASLR

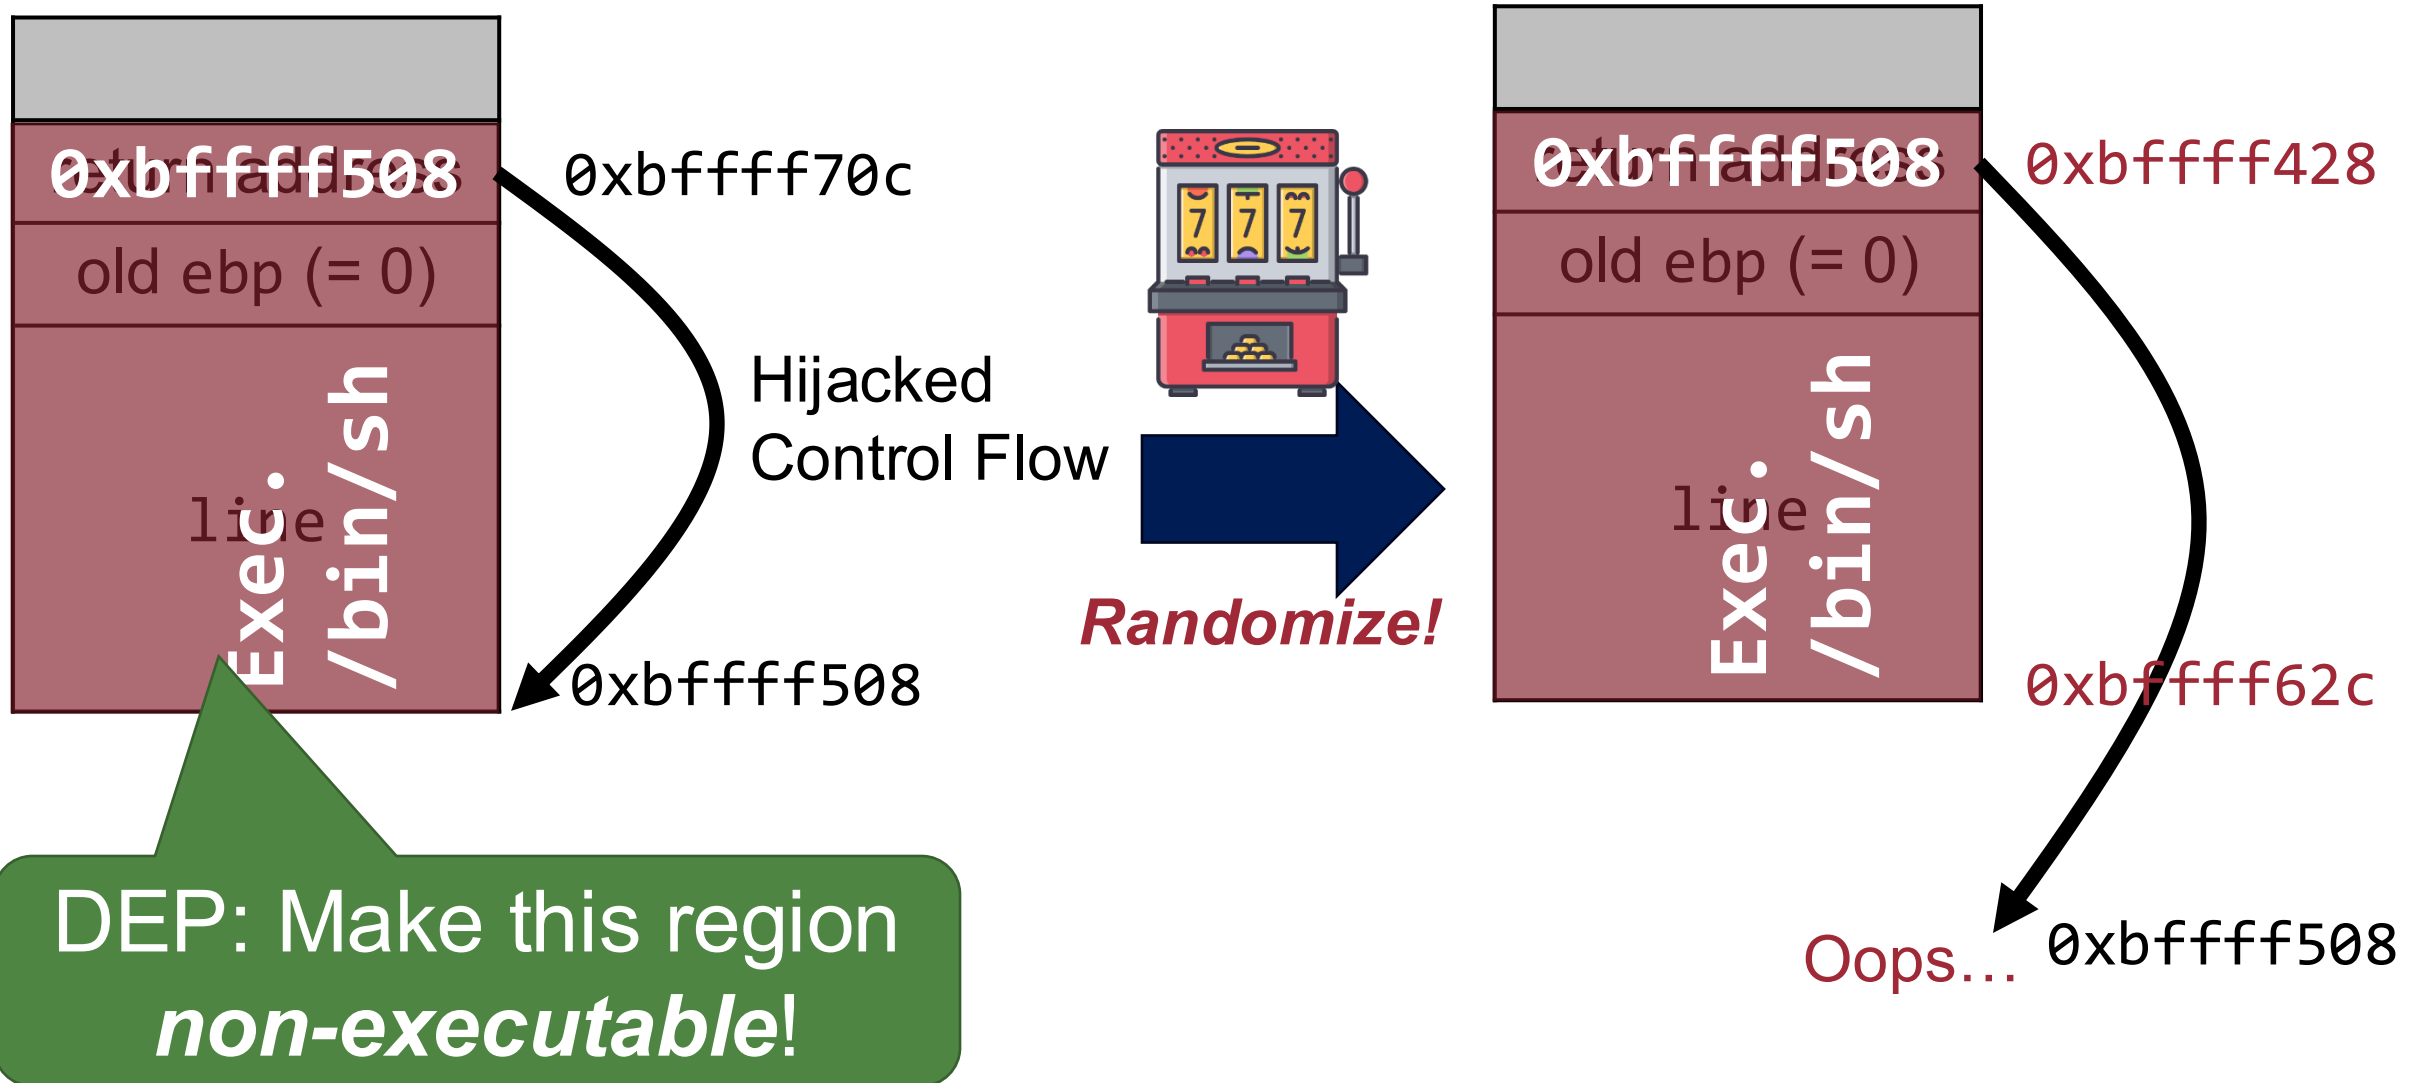


DEP: Make this region
non-executable!

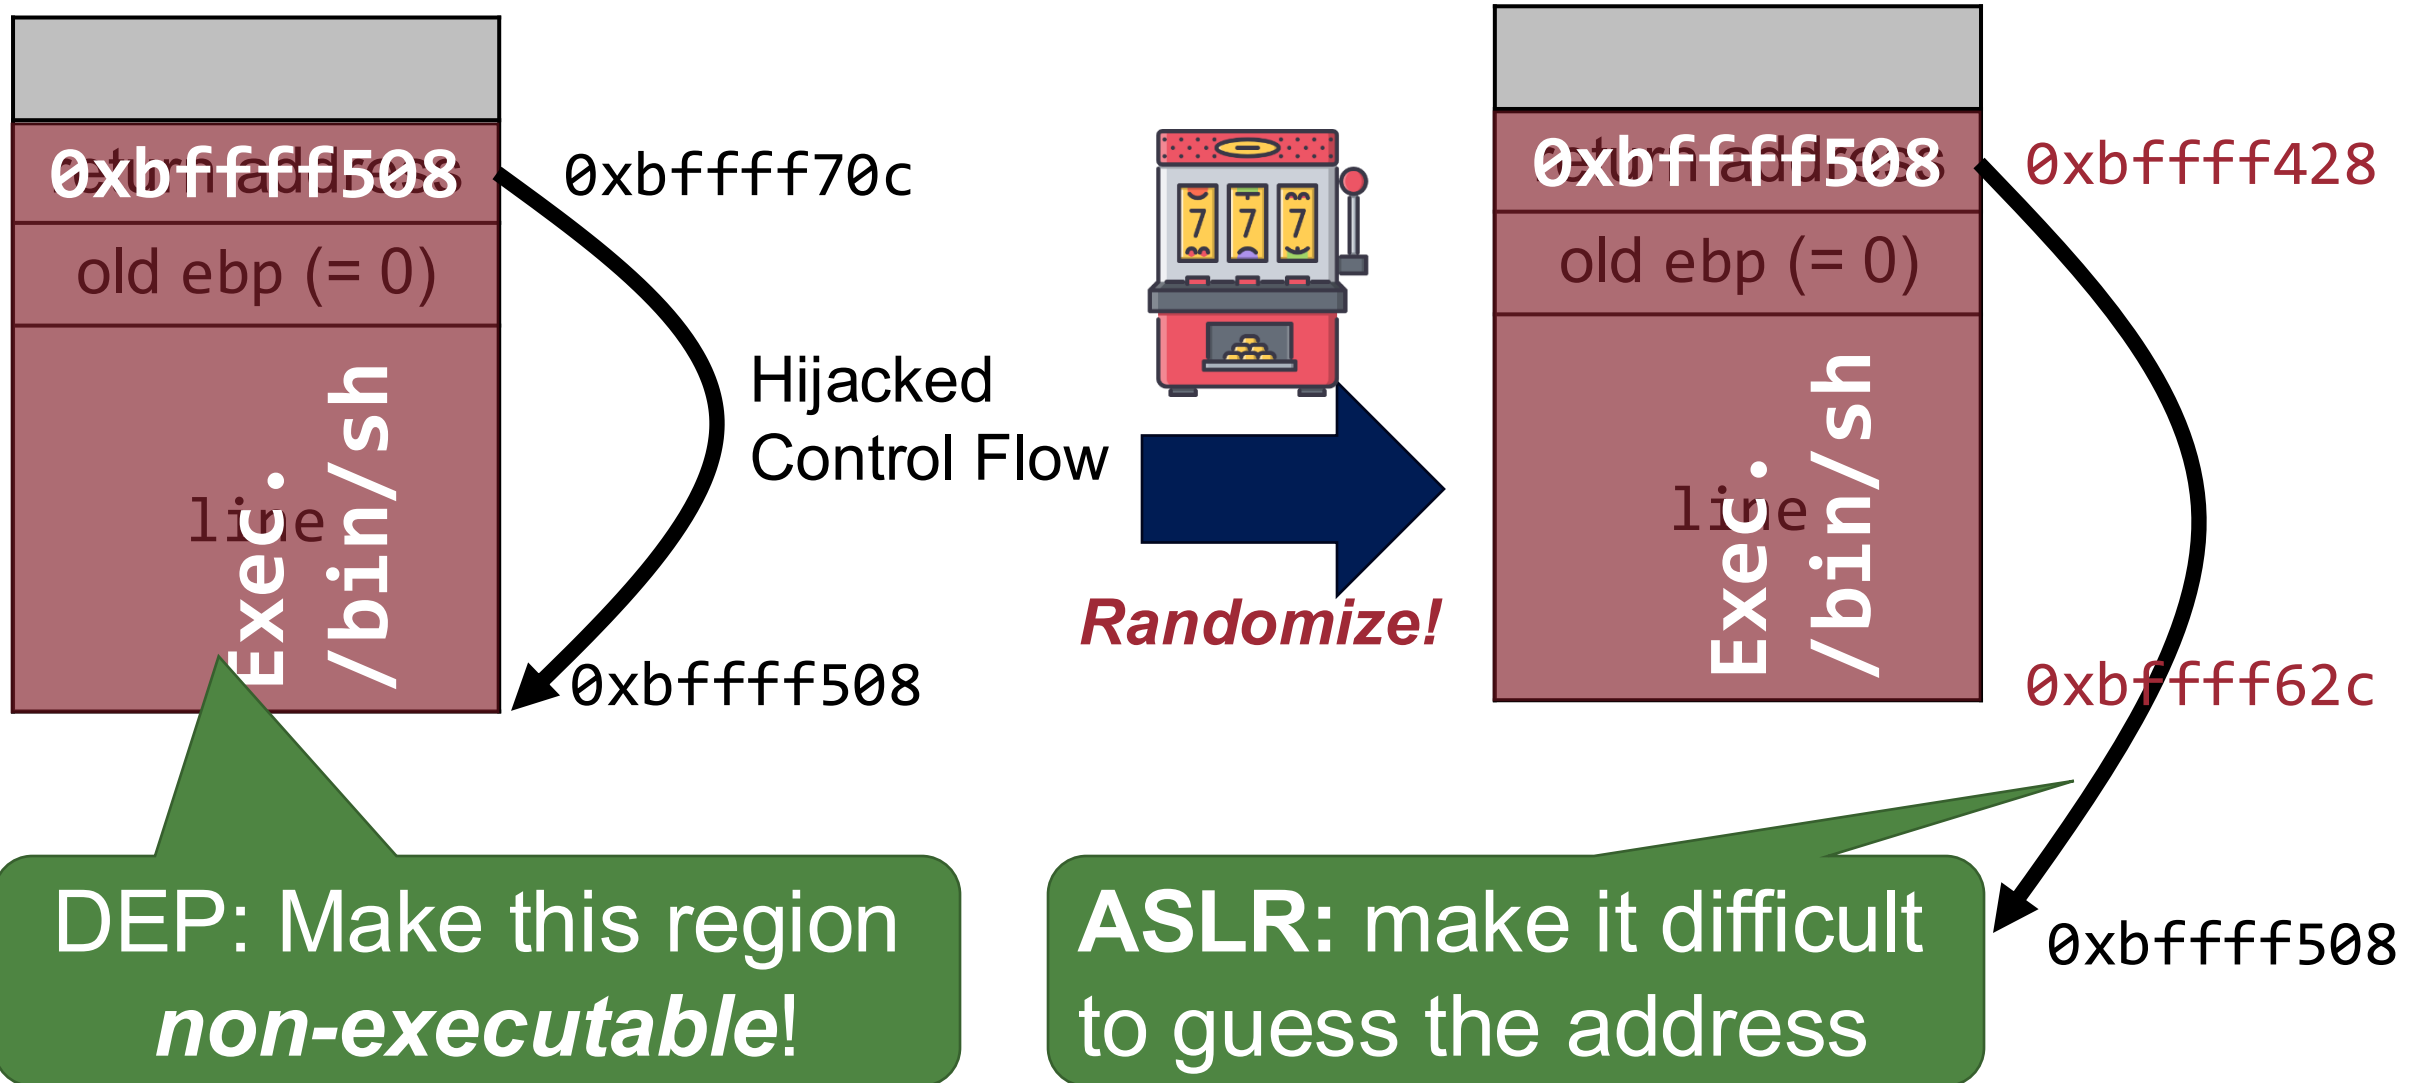
Different Perspective: ASLR



Different Perspective: ASLR



Different Perspective: ASLR



World without ASLR



- Use the same address space over and over again!

Printing out ESP



```
#include <stdio.h>
```

```
int main (void) {
```

```
    int x = 42;
```

```
    return printf("%08p\n", &x);    // printing out esp
```

```
}
```

World with ASLR



- ASLR is ON by default [Ubuntu-Security]

You can enable ASLR by:

```
$ echo 2 | sudo tee /proc/sys/kernel/randomize_va_space
```

World with ASLR



- ASLR is ON by default [Ubuntu-Security]

You can enable ASLR by:

```
$ echo 2 | sudo tee /proc/sys/kernel/randomize_va_space
```

Why 2?

Manual Says



Value	Description
-------	-------------

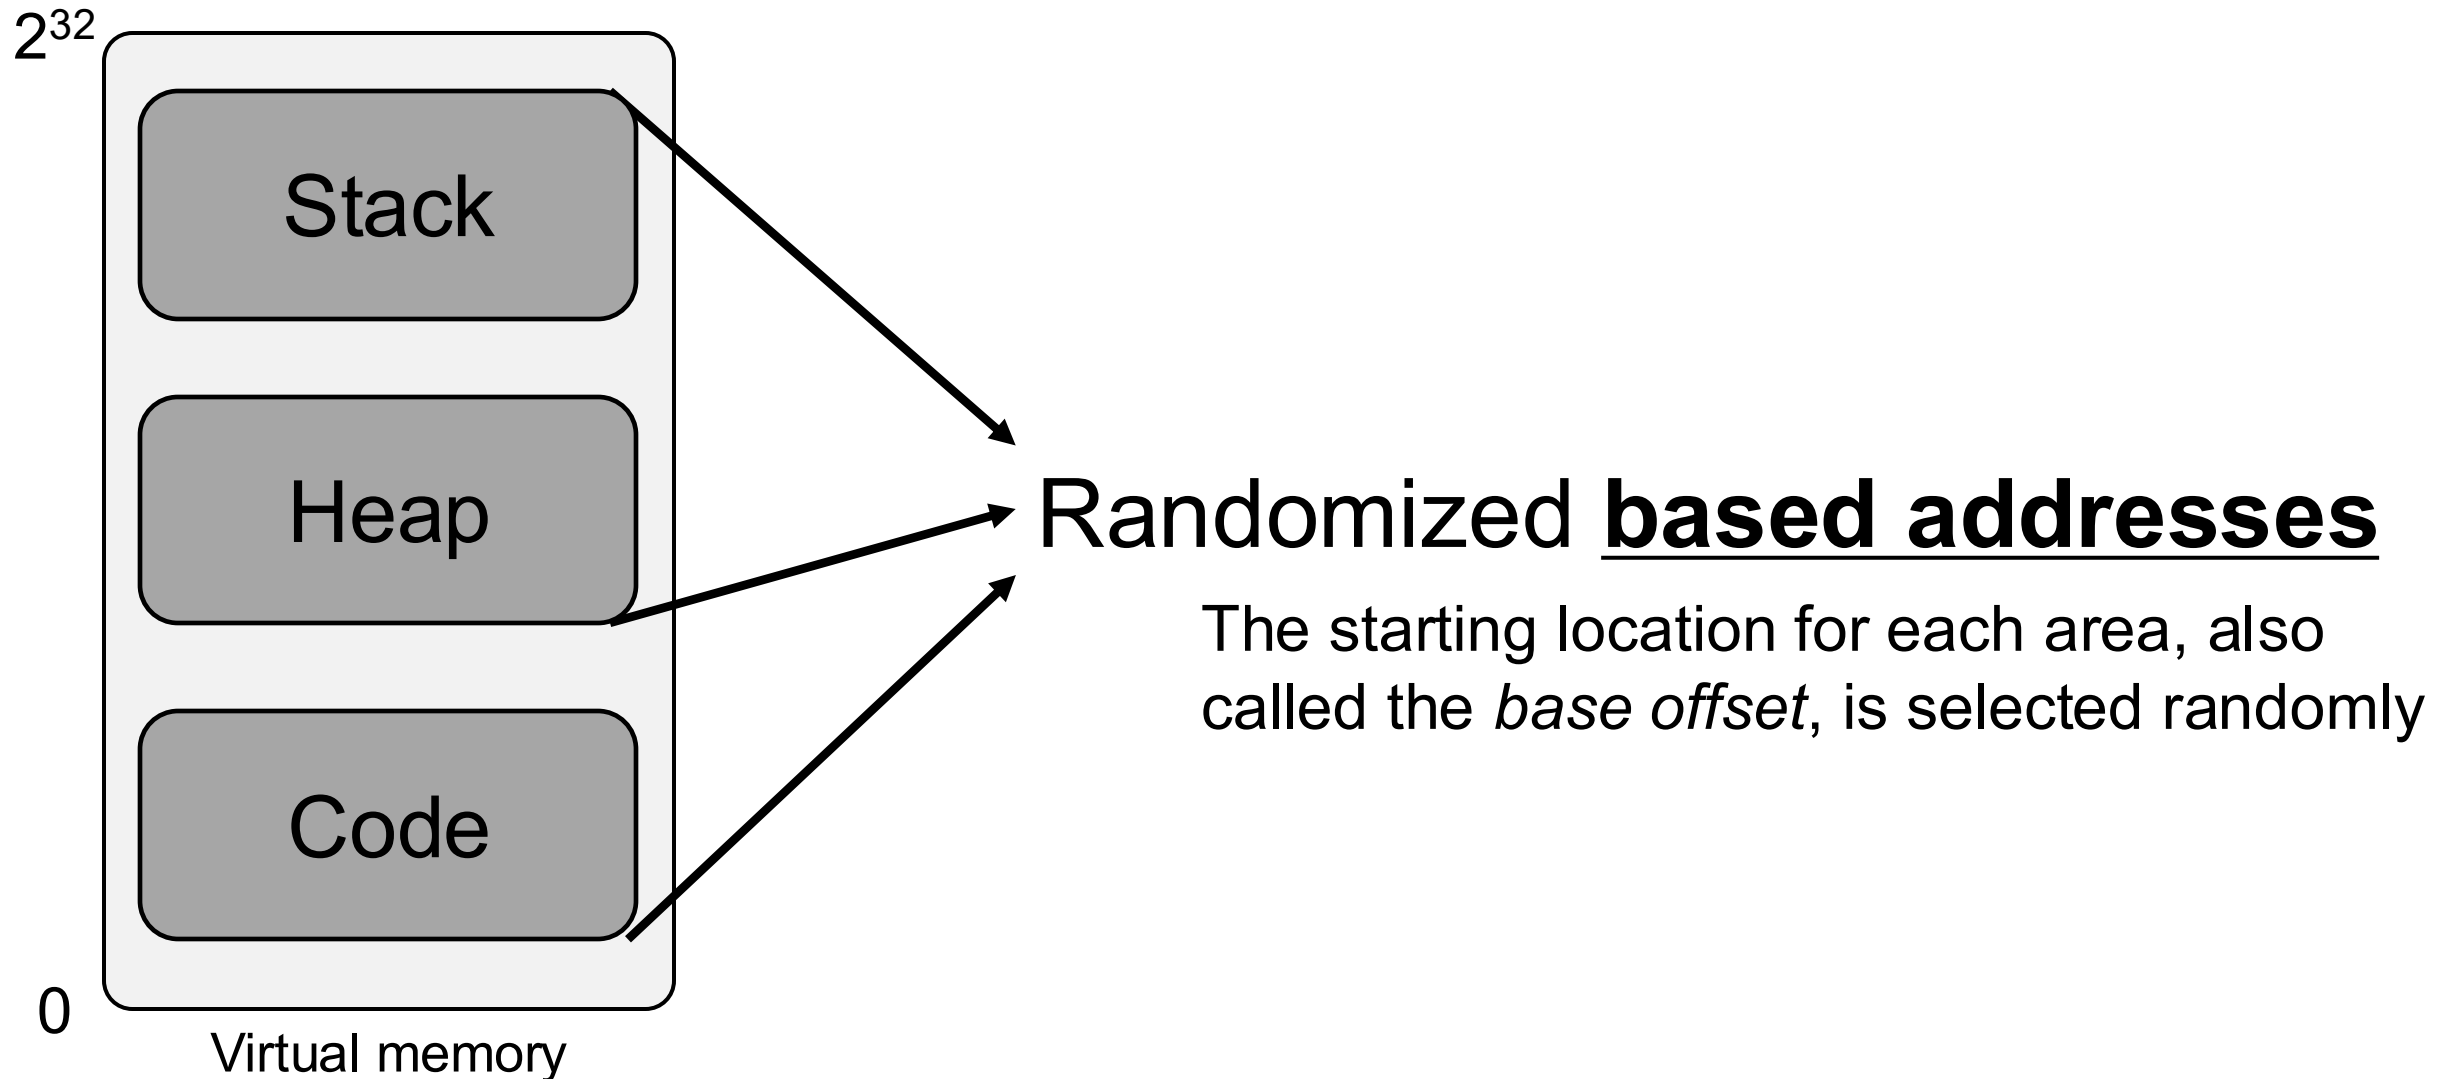
0	Turn ASLR off
---	---------------

1	Make the address the <u>stack</u> and the <u>library space</u> randomized
---	---

2	Also, support <u>heap randomization</u>
---	---


ASLR Randomizes Virtual Memory Areas

66



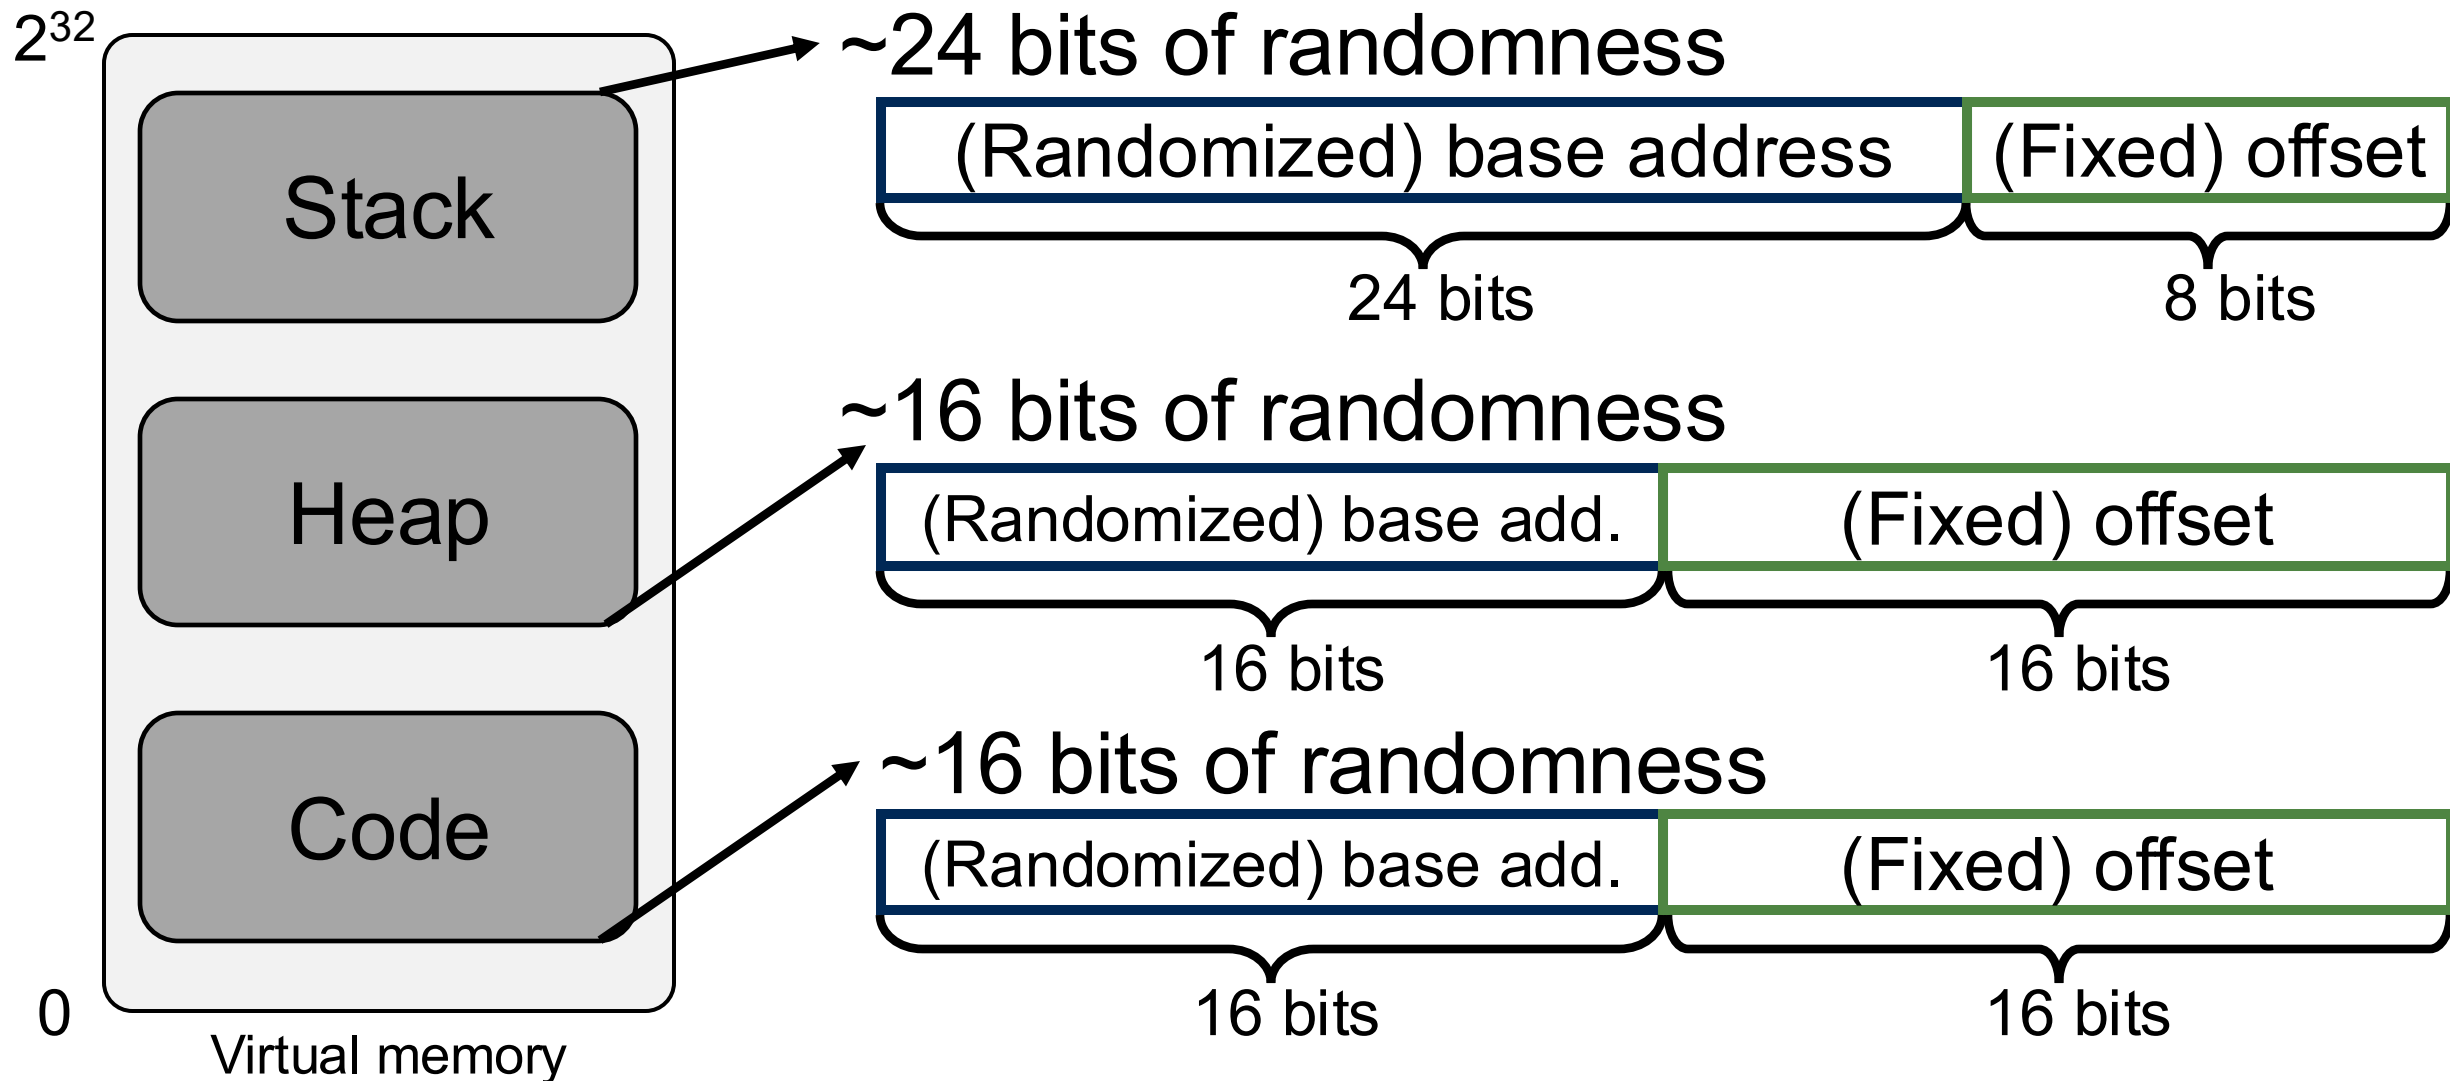
ASLR Randomizes Virtual Memory Areas

67



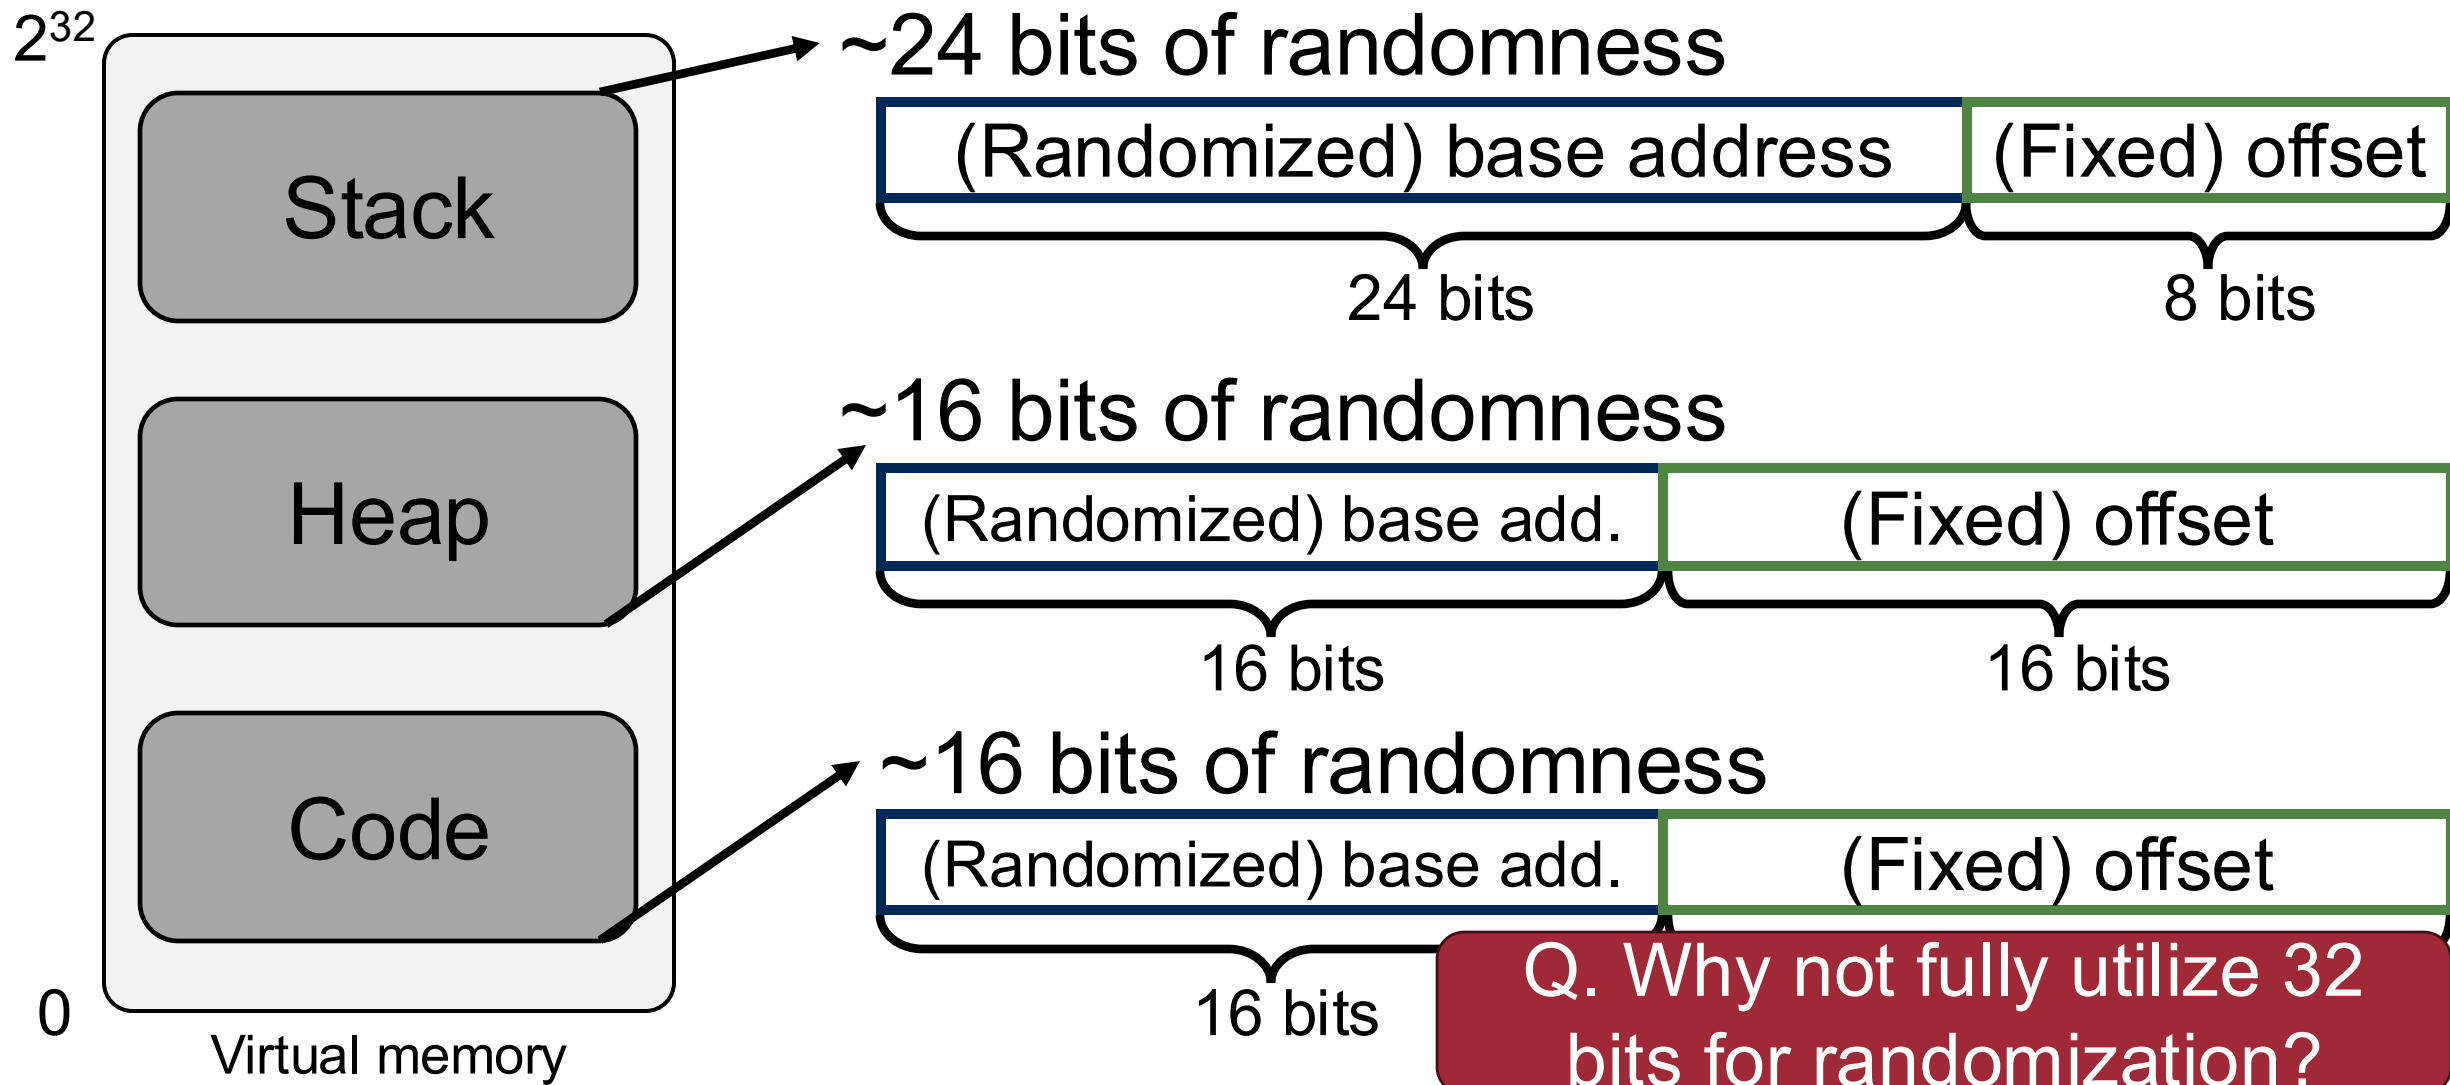
ASLR Randomizes Virtual Memory Areas

68



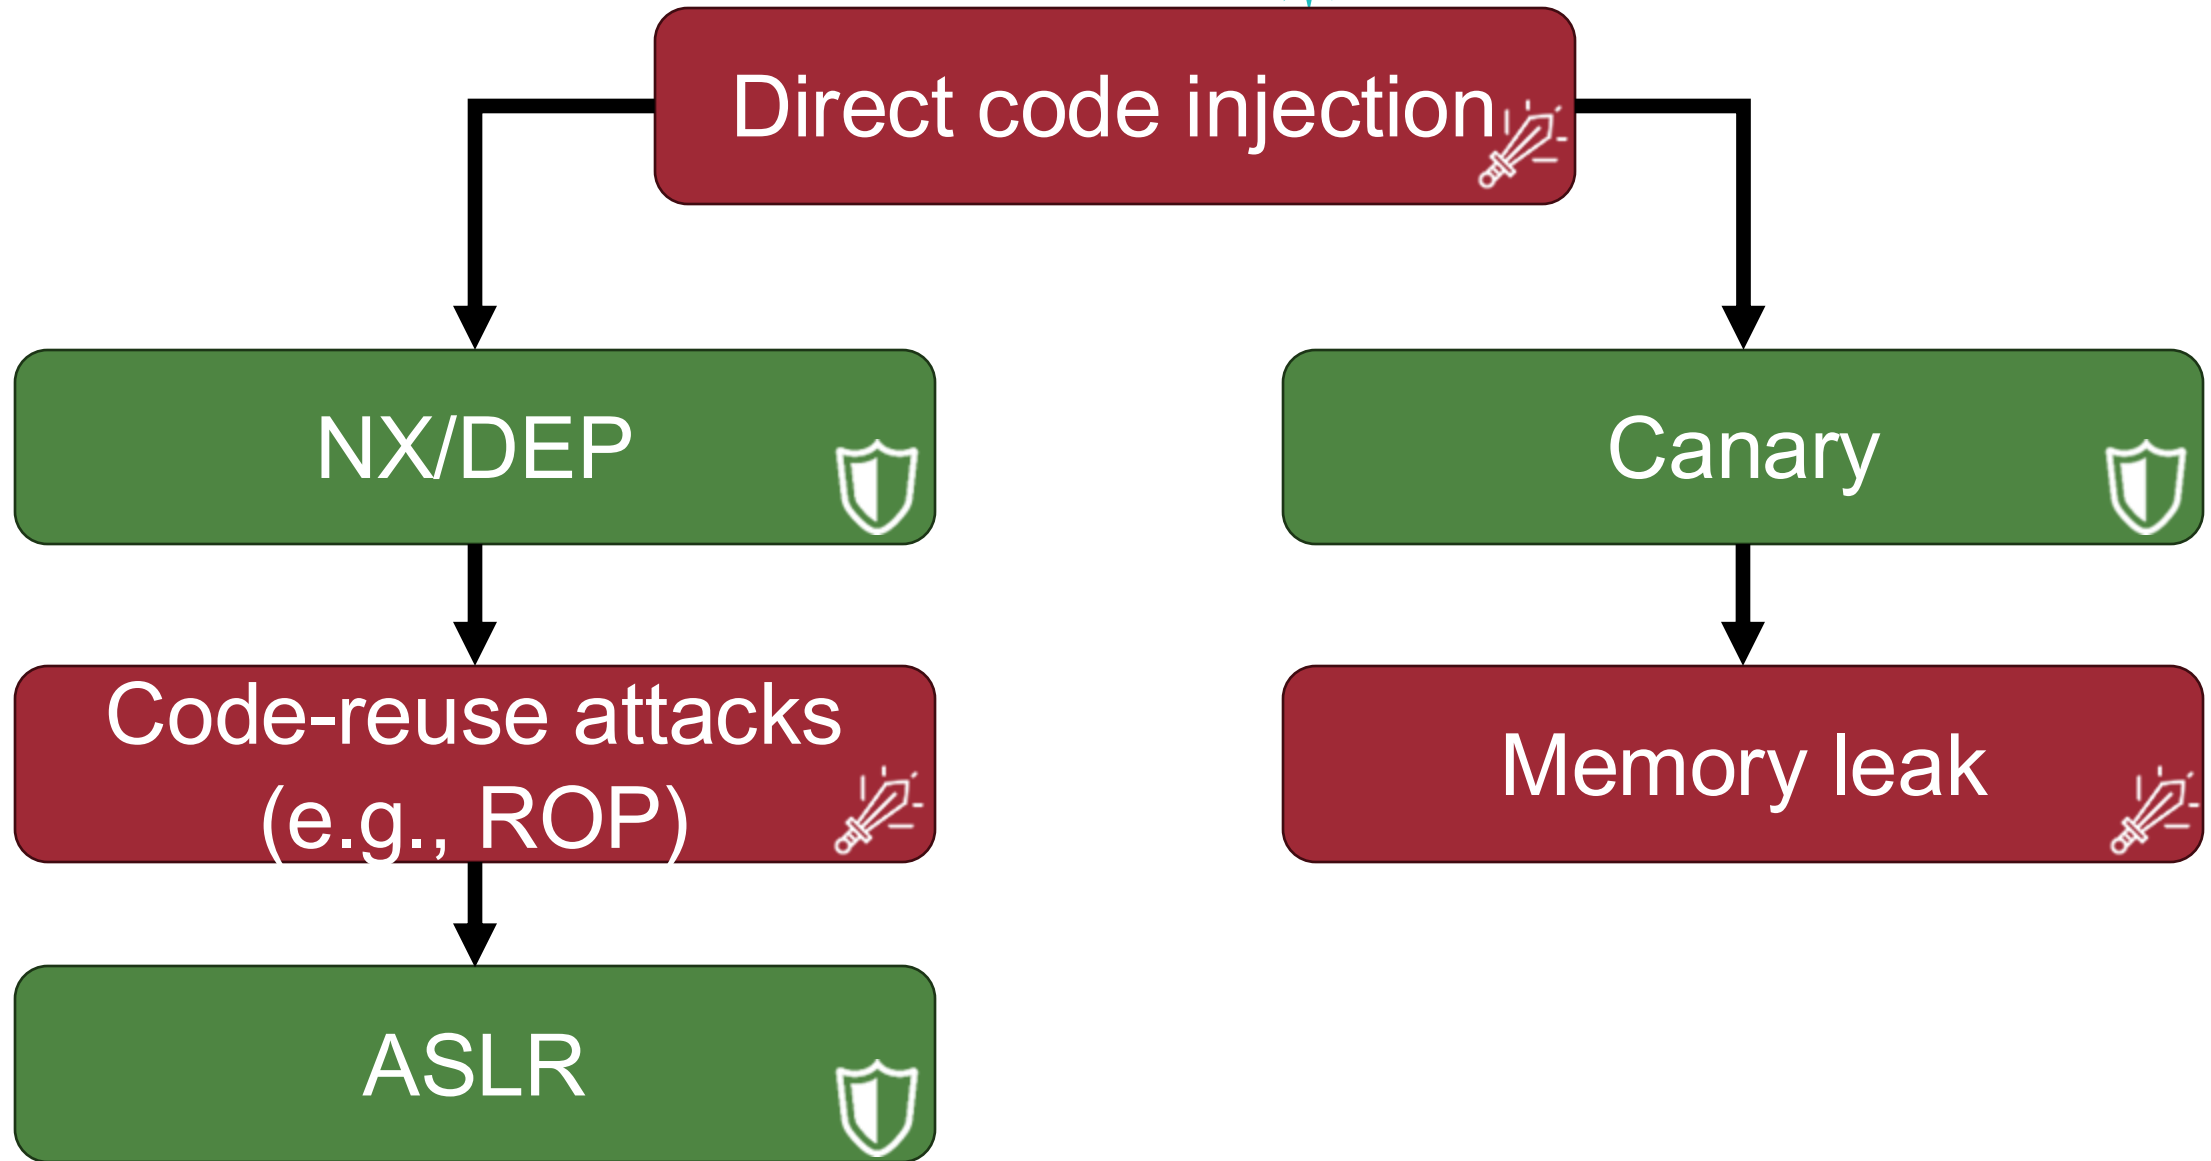
ASLR Randomizes Virtual Memory Areas

69



Control Hijack Attack / Defense So Far

70



Previous Exploits will *NOT* Work w/ ASLR

71

- ASLR will randomize the **base addresses** of the stack, heap, and code segments
- We cannot know the address of our shellcode nor library functions
 - Thus, no return-to-stack nor return-to-LIBC

Are we safe now?

Attacking ASLR

Part 1. Entropy

Attack #1: Entropy is Small on x86

- Just 16 bits are used for heap and libraries on x86 (Therefore, entropy is small on x86)
- **Brute-forcing** is possible for server applications that use *forking*

On the Effectiveness of Address-Space Randomization, **CCS 2004**

On the Effectiveness of Address-Space Randomization

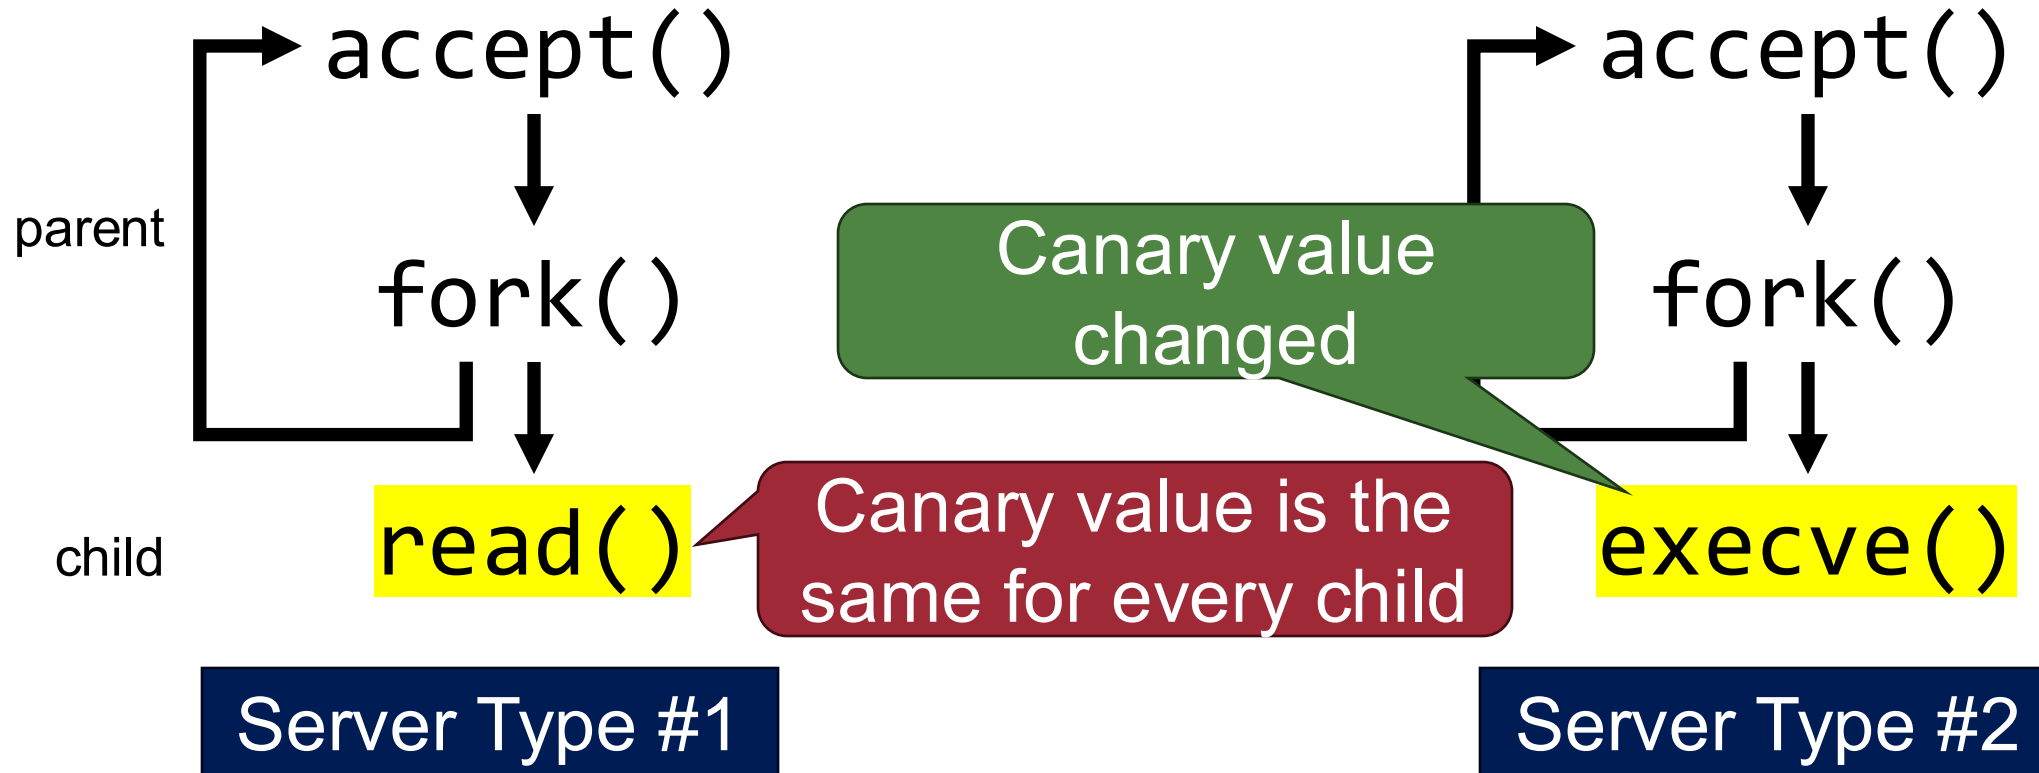
Hovav Shacham hovav@cs.stanford.edu	Matthew Page mpage@stanford.edu	Ben Pfaff blp@cs.stanford.edu
Eu-Jin Goh eujin@cs.stanford.edu	Nagendra Modadugu nagendra@cs.stanford.edu	Dan Boneh dabo@cs.stanford.edu

Abstract

Address-space randomization is a technique used to fortify systems against buffer overflow attacks. The idea is to introduce artificial diversity by randomizing the memory location of certain system components. This mechanism is available for both Linux (via PaX ASLR) and OpenBSD. We study the effectiveness of address-space randomization and find that its utility on 32-bit architectures is limited by the number of bits available for address randomization. In particular, we demonstrate a *derandomization attack* that will convert any standard buffer-overflow exploit into an exploit that works against systems protected by address-space randomization. The resulting exploit is as effective as the original, albeit somewhat slower: on average 216 sec-

Recap: Reused Canary Value

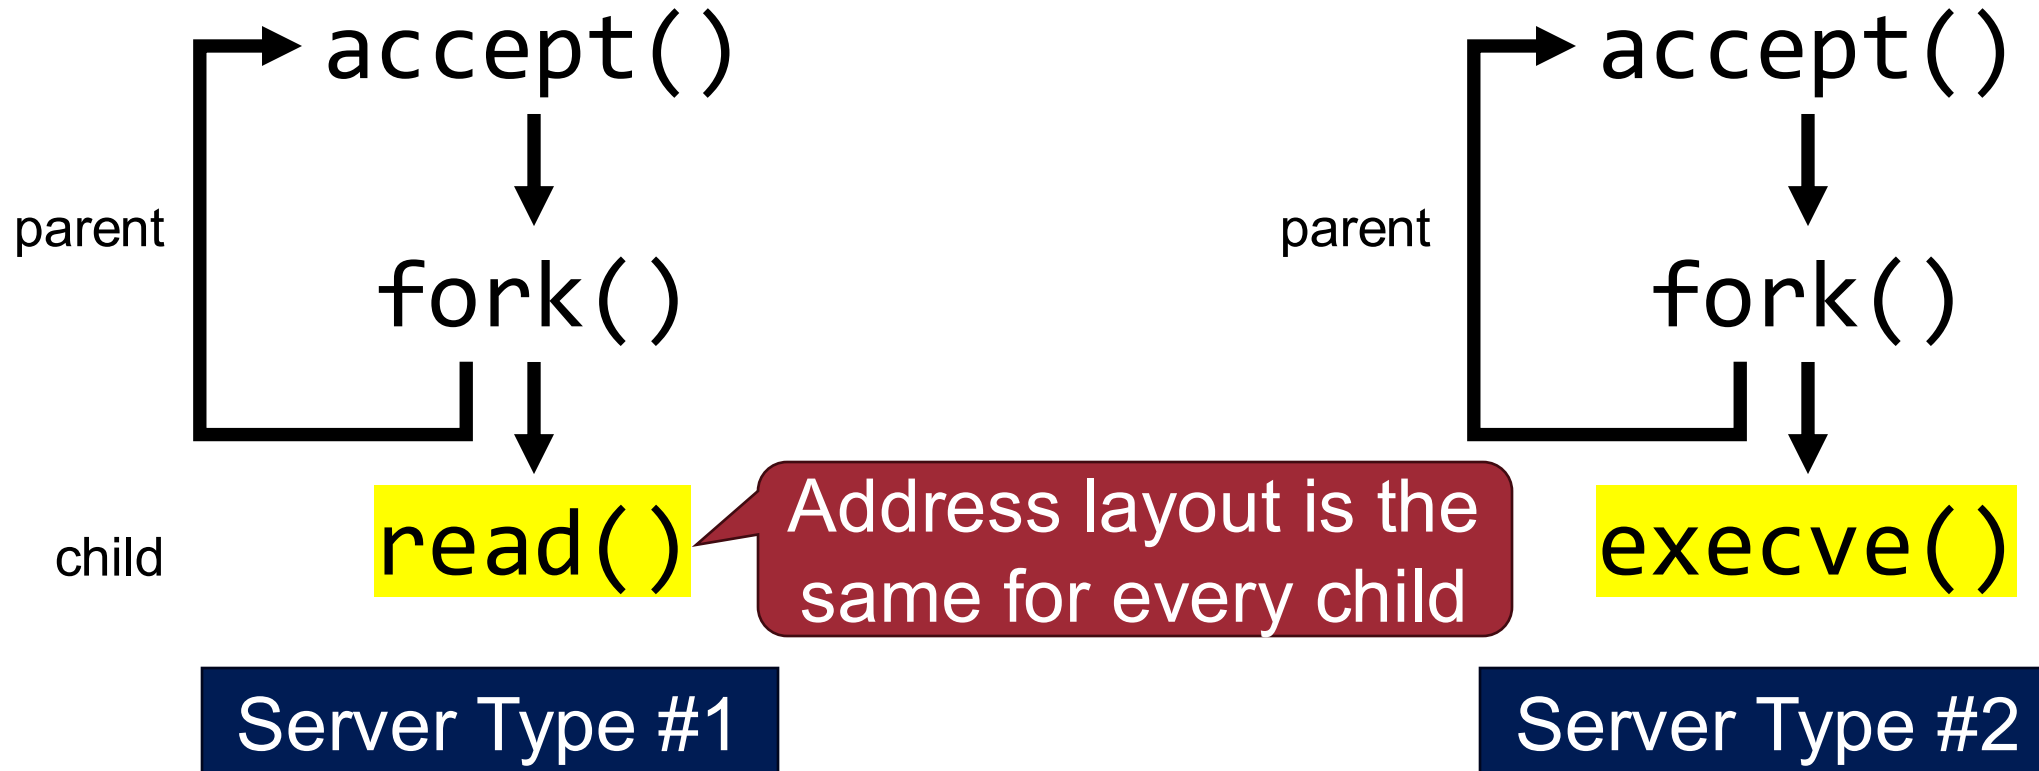
- Uses a random canary value for **every process creation**



e.g., OpenSSH does this

Remained Address Space

- Uses a random canary value for **every process creation**



Attack #1: Entropy is Small on x86



- Just 16 bits are used for heap and libraries on x86 (Therefore, entropy is small on x86)
- **Brute-forcing** is possible for server applications that use *forking*
 - Forked process has the same address space layout as its parent
 - Once we know the address of *a function in LIBC*, we can deduce the addresses of *all functions in LIBC*!

Key point: relative offsets between LIBC functions are the same regardless of ASLR

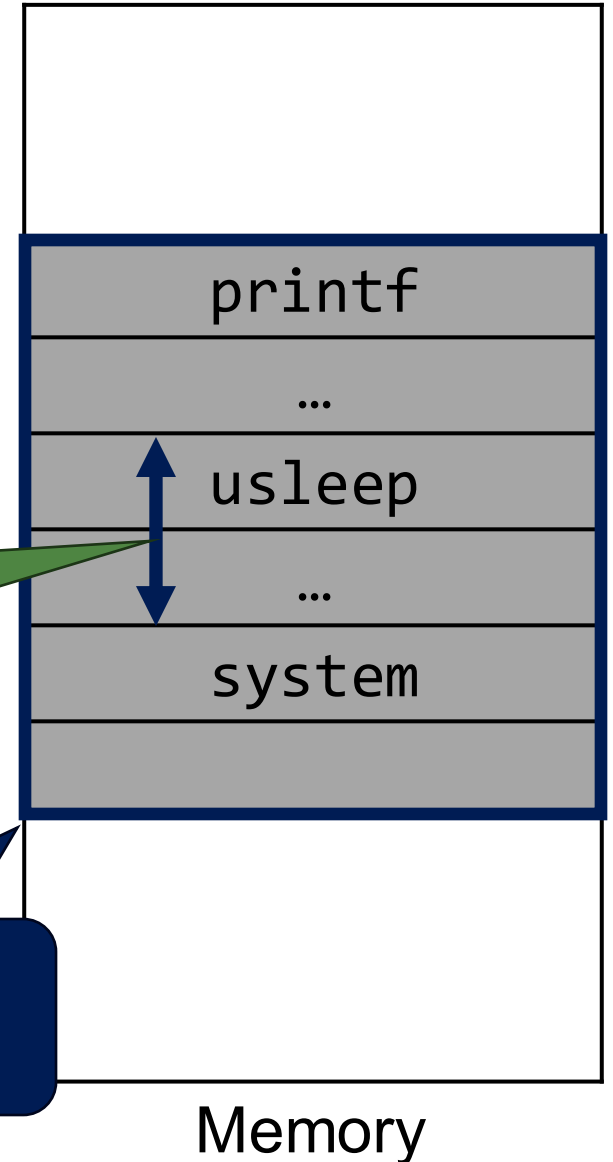
Observation: Relative Offsets are Same!

77

Relative offset between
usleep and system is fixed!

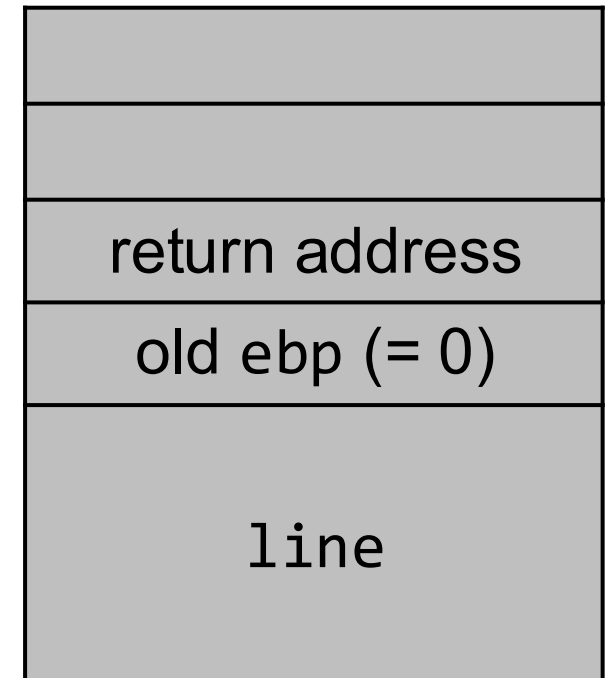
Publicly known!

LIBC
base address



Brute-forcing Attack Example

- Target: Apache web server
 - Forks children on requests
- Vulnerability: Buffer overflow vulnerability



Brute-forcing Attack Example

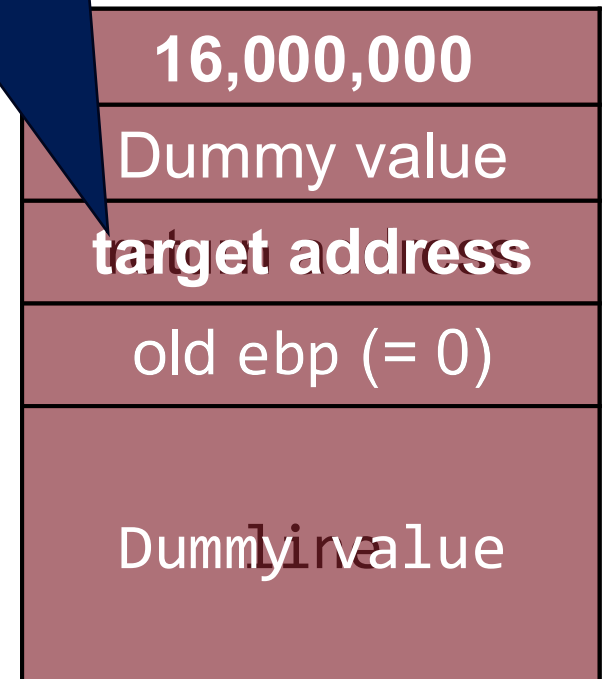
- Target: Apache web server
 - Forks children on requests
- Vulnerability: Buffer overflow vulnerability

16,000,000
Dummy value
target address
old ebp (= 0)
Dummy value

Brute-forcing Attack Example

- Target: Apache web server
 - Forks children on requests
- Vulnerability: Buffer overflow
- Method: Return-to-LIBC (usleep)
 - Try to brute-force the address of usleep with a fake parameter of 16,000,000 (waiting for 16 seconds)

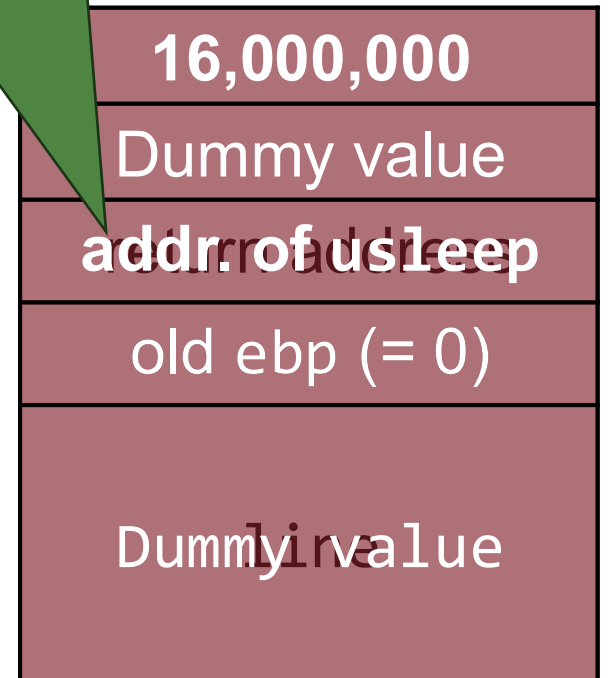
Brute-force on 16 bits to find the address of usleep



Brute-forcing Attack Example


- Target: Apache web server
 - Forks children on requests
- Vulnerability: Buffer overflow
- Method: Return-to-LIBC (usleep)
 - Try to brute-force the address of usleep with a fake parameter of 16,000,000 (waiting for 16 seconds)

If correct, the server will wait 16 seconds

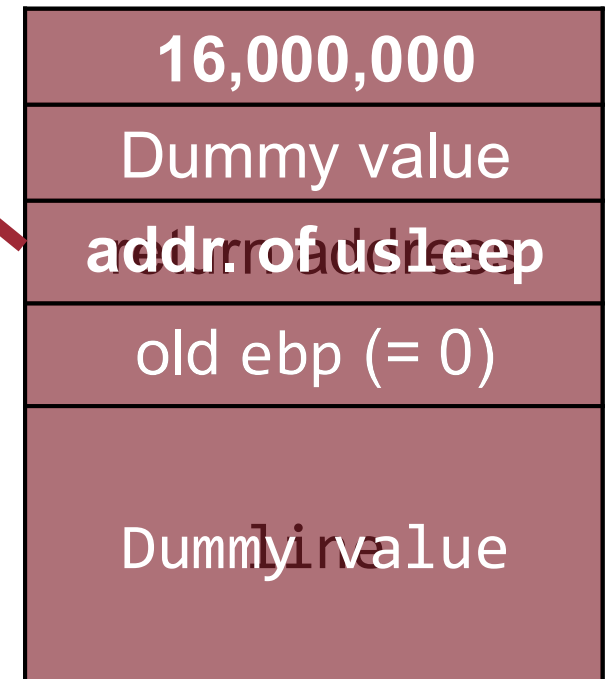


Brute-forcing Attack Example

- Target: Apache web server
 - Forks children on requests
- Vulnerability: Buffer overflow vulnerability
- Method: Return-to-LIBC (usleep)
 - Try to brute-force the address of `usleep` with a fake parameter of 16,000,000 (waiting for 16 seconds)
 - Once we know the address of `usleep`, we can determine the address of `exec` or `system`



LIBC




Brute-forcing Attack Example


- Target: Apache web server
 - Forks children on requests
- Vulnerability: Buffer overflow vulnerability
- Method: Return-to-LIBC (usleep)
 - Try to brute-force the address of usleep with a fake parameter of 16,000,000 (waiting for 16 seconds)
 - Once we know the address of usleep, we can determine the address of exec or system

Publicly known

offset



LIBC




Brute-forcing Attack Example

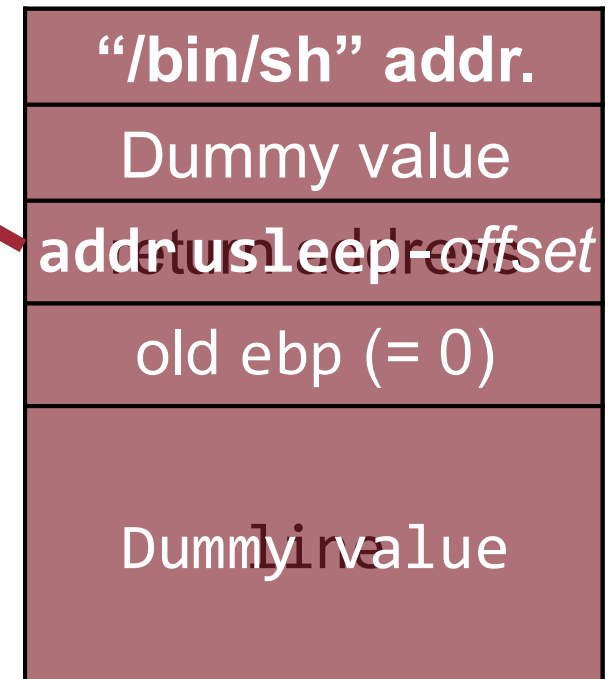
- Target: Apache web server
 - Forks children on requests
- Vulnerability: Buffer overflow vulnerability
- Method: Return-to-LIBC (usleep)
 - Try to brute-force the address of usleep with a fake parameter of 16,000,000 (waiting for 16 seconds)
 - Once we know the address of usleep, we can determine the address of exec or system

Publicly known

offset



LIBC



Randomization Frequency on Two Major OSes⁸⁵



- On **Windows**: every time the machine starts
 - Each module will get a random address once per boot (but, stack and heap will be randomized per execution)
- On **Linux**: every time a process loads
 - Each module will get a random address for every execution

Which one is better?

Performance: Which One is Better?

- On **Windows**: every time the machine starts
 - Each module will get a random address once per boot (but, stack and heap will be randomized per execution)

Faster: relocation once at boot time

- On **Linux**: every time a process loads
 - Each module will get a random address for every execution

Slower: relocation fixups for every execution

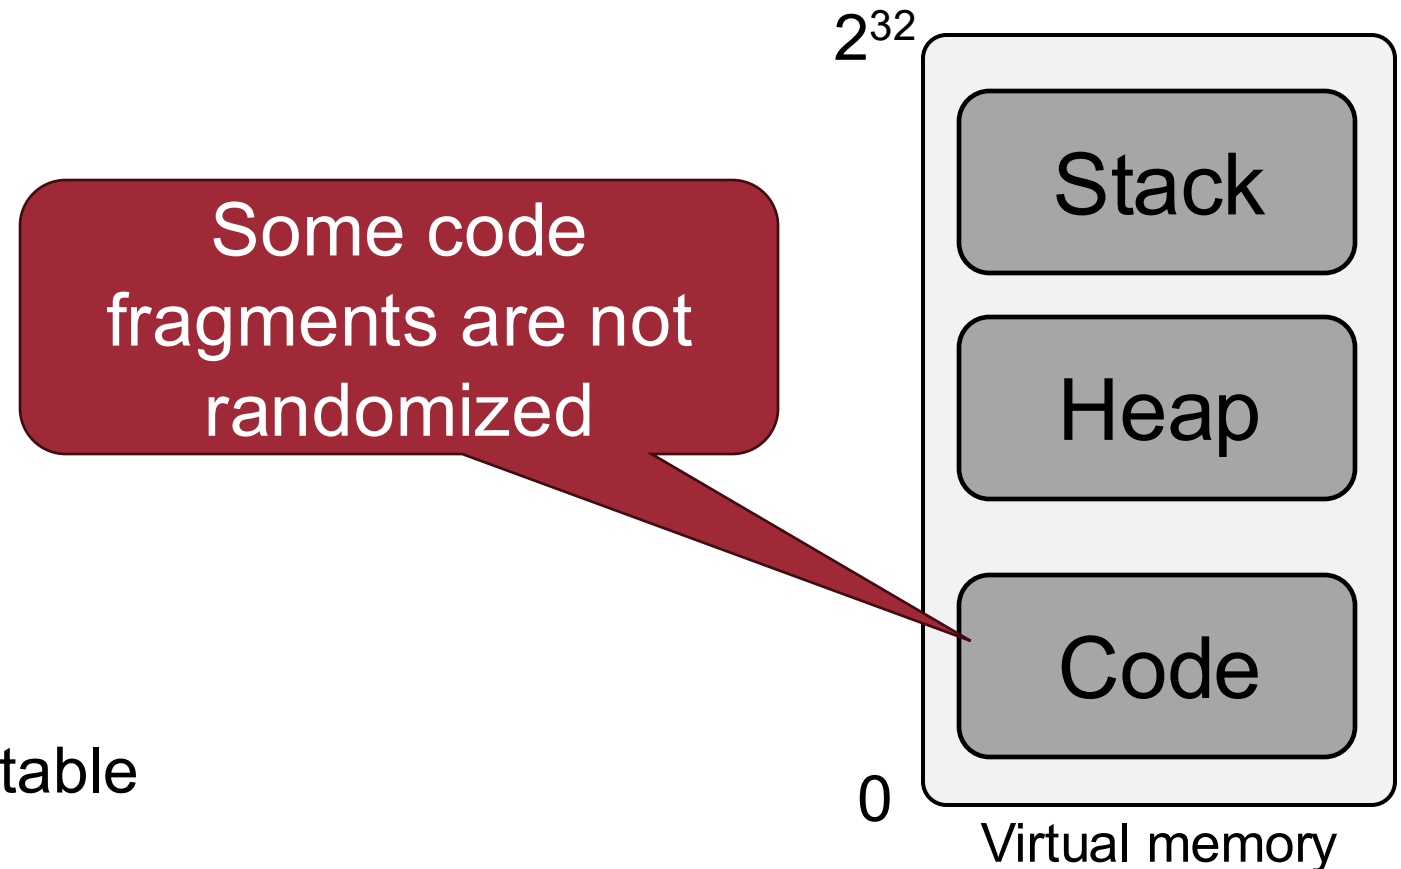
How about security?

Attacking ASLR

Part 2. Exploiting Fixed Addresses

Attack 2: Exploiting Fixed Addresses

- Most binaries (before 2016) had non-randomized segments
 - Before 2016, compilers created ***non-PIE***¹ executables by default



¹Non Position-Independent Executable

Position-Independent Executable (PIE)

- Position-Independent Code (PIC) or PIE is code that runs regardless of its location (e.g., shellcode)
 - “gcc” will produce a PIE by default
 - “gcc -fno-pic -no-pie” will produce a non-PIE

PIE vs. non-PIE



- Position-Independent Code (PIC) or PIE is code that runs regardless of its location (e.g., shellcode)
 - “gcc” will produce a PIE by default
 - “gcc -fno-pic -no-pie” will produce a non-PIE

080491ba <main>:

```

80491ba: lea    ecx,[esp+0x4]
80491be: and    esp,0xfffffffff0
80491c1: push   DWORD PTR [ecx-0x4]
80491c4: push   ebp
80491c5: mov    ebp,esp
80491c7: push   ecx
80491c8: sub    esp,0x14
80491cb: mov    eax,gs:0x14

```

\$ gcc -fno-pic -no-pie

000011f1 <main>:

```

11f1: lea    ecx,[esp+0x4]
11f5: and    esp,0xfffffffff0
11f8: push   DWORD PTR [ecx-0x4]
11fb: push   ebp
11fc: mov    ebp,esp
11fe: push   ebx
11ff: push   ecx
1200: sub    esp,0x10

```

\$ gcc (Produce a PIE)

PIE vs. non-PIE



- **Non-randomized segments even when ASLR is turned on**
- **Relative addresses – randomized when ASLR is turned on**

```

080491ba <main>:
80491ba: lea     ecx,[esp+0x4]
80491be: and     esp,0xfffffffff0
80491c1: push   DWORD PTR [ecx-0x4]
80491c4: push   ebp
80491c5: mov     ebp,esp
80491c7: push   ecx
80491c8: sub     esp,0x14
80491cb: mov     eax,gs:0x14
  
```

\$ gcc -fno-pic -no-pie

```

000011f1 <main>:
11f1: lea     ecx,[esp+0x4]
11f5: and     esp,0xfffffffff0
11f8: push   DWORD PTR [ecx-0x4]
11fb: push   ebp
11fc: mov     ebp,esp
11fe: push   ebx
11ff: push   ecx
1200: sub     esp,0x10
  
```

\$ gcc (Produce a PIE)

Legacy Binaries Are Not a PIE



- 93% of Linux binaries were not a PIE (in 2009)
- Thus, the code sections were not randomized



But, why?

Security vs. Performance



- Relative-addressing instructions are slower than absolute-addressing instructions
- Performance overhead of PIE on x86 is 10% on average
(Too much PIE is bad for performance, ETH Techreport, 2012)
- Most applications on current x86 are still not PIEs

ROP-based Attack on Legacy Binaries

- Some code fragments are not randomized!
- Why not use ROP on them

Fact: relative offsets between LIBC functions are the same regardless of ASLR

Exploitation Idea



- If a LIBC function has been invoked at least once, GOT should contain a concrete address of the function in LIBC
- Therefore, we will read the GOT entry using ROP and compute the address of `system` by using the relative offset between the LIBC function and `system`

Suppose we can get the address of `open` function from the GOT

$$\begin{aligned} (\text{addr of system}) &= (\text{addr of open}) \\ &+ (\text{offset from open to system in LIBC}) \end{aligned}$$

Example ROP

$$(\text{addr of system}) = (\text{addr of open}) + (\text{offset from open to system in LIBC})$$

Gadget **C**

jmp [eax]

Gadget **B**

pop eax
 add eax, edi
 ret

Gadget **A**

pop edi
 ret

Address of C

y

Address of B

x

Address of A

old ebp (= 0)

line

x

y

Possible Defenses?



- Use PIEs
- Use 64-bit CPU: lots of entropy
- Detect brute-forcing attacks
 - Many crashes in a short amount of time
- Use non-forking servers
- Code randomization (a.k.a. fine-grained ASLR)


Code Randomization

Motivation





- ASLR only changes base addresses of VMAs
 - A single pointer leak can reveal the entire memory layout of a VMA



Fine-Grained ASLR (Not Used in Practice) ¹⁰⁰



Classic ASLR



Function-level ASLR



Instruction-level ASLR

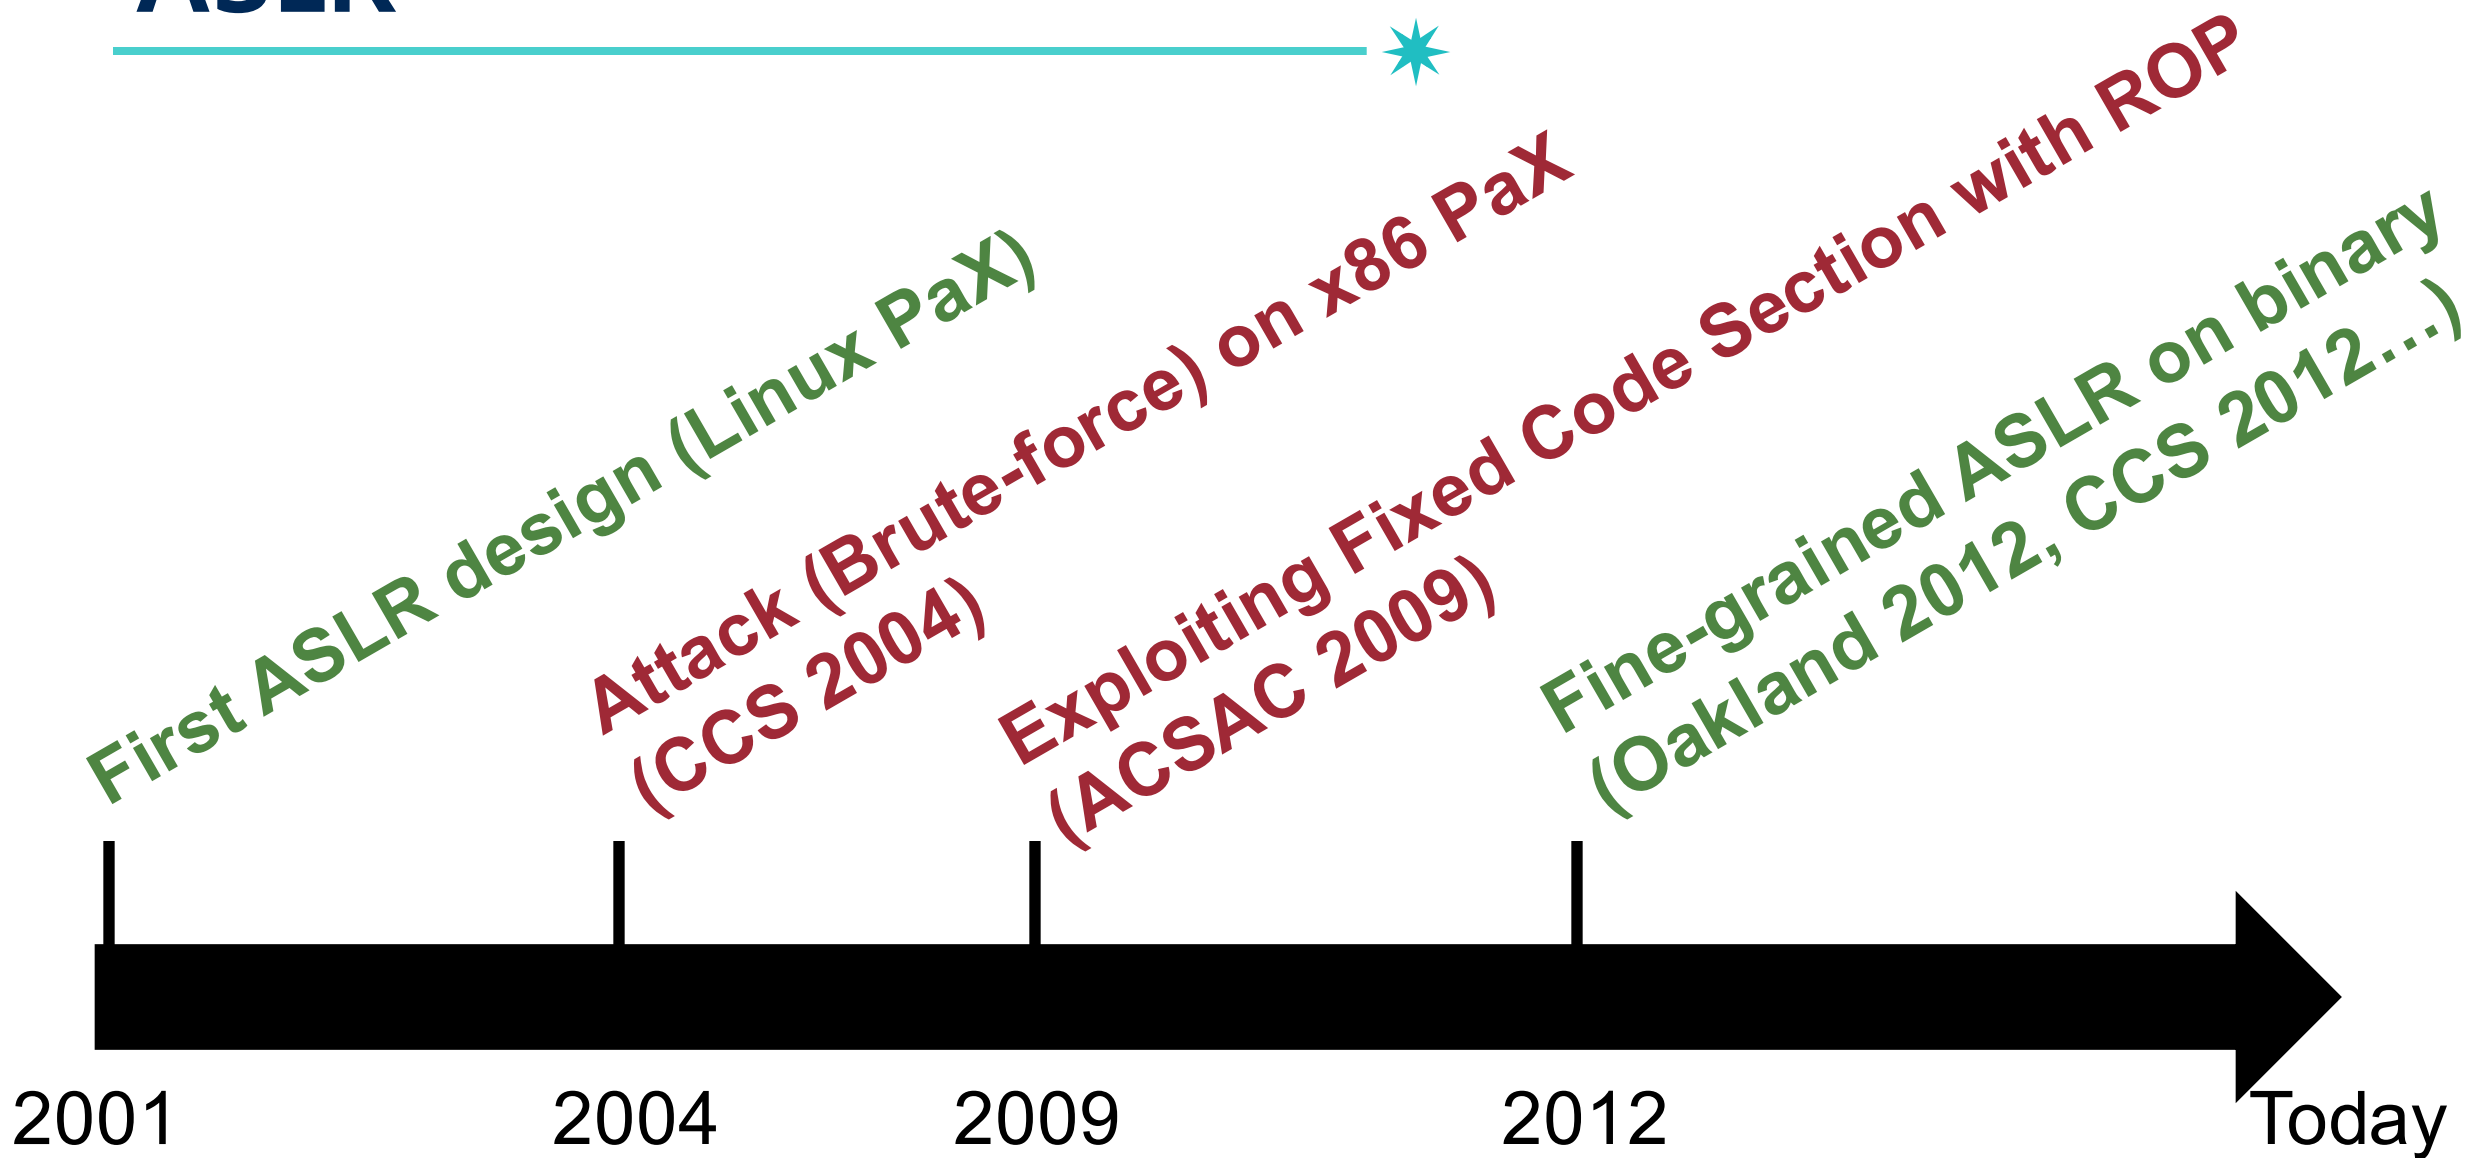
Fine-Grained ASLR (Not Used in Practice) ¹⁰



- Source-based approaches:
 - Efficient Techniques for Comprehensive Protection from Memory Error Exploits, **USENIX Security 2005**
 - Enhanced Operating System Security through Efficient and Fine-grained Address Space Randomization, **USENIX Security 2012**
- Binary-based approaches:
 - Smashing the Gadgets: Hindering Return-Oriented Programming Using In-Place Code Randomization, **S&P 2012**
 - ILR: Where'd My Gadgets Go?, **S&P 2012**
 - Binary Stirring: Self-randomizing Instruction Addresses of Legacy x86 Binary Code, **CCS 2012**

ASLR

102



Memory Disclosure

Memory Disclosure \neq Memory Corruption

104

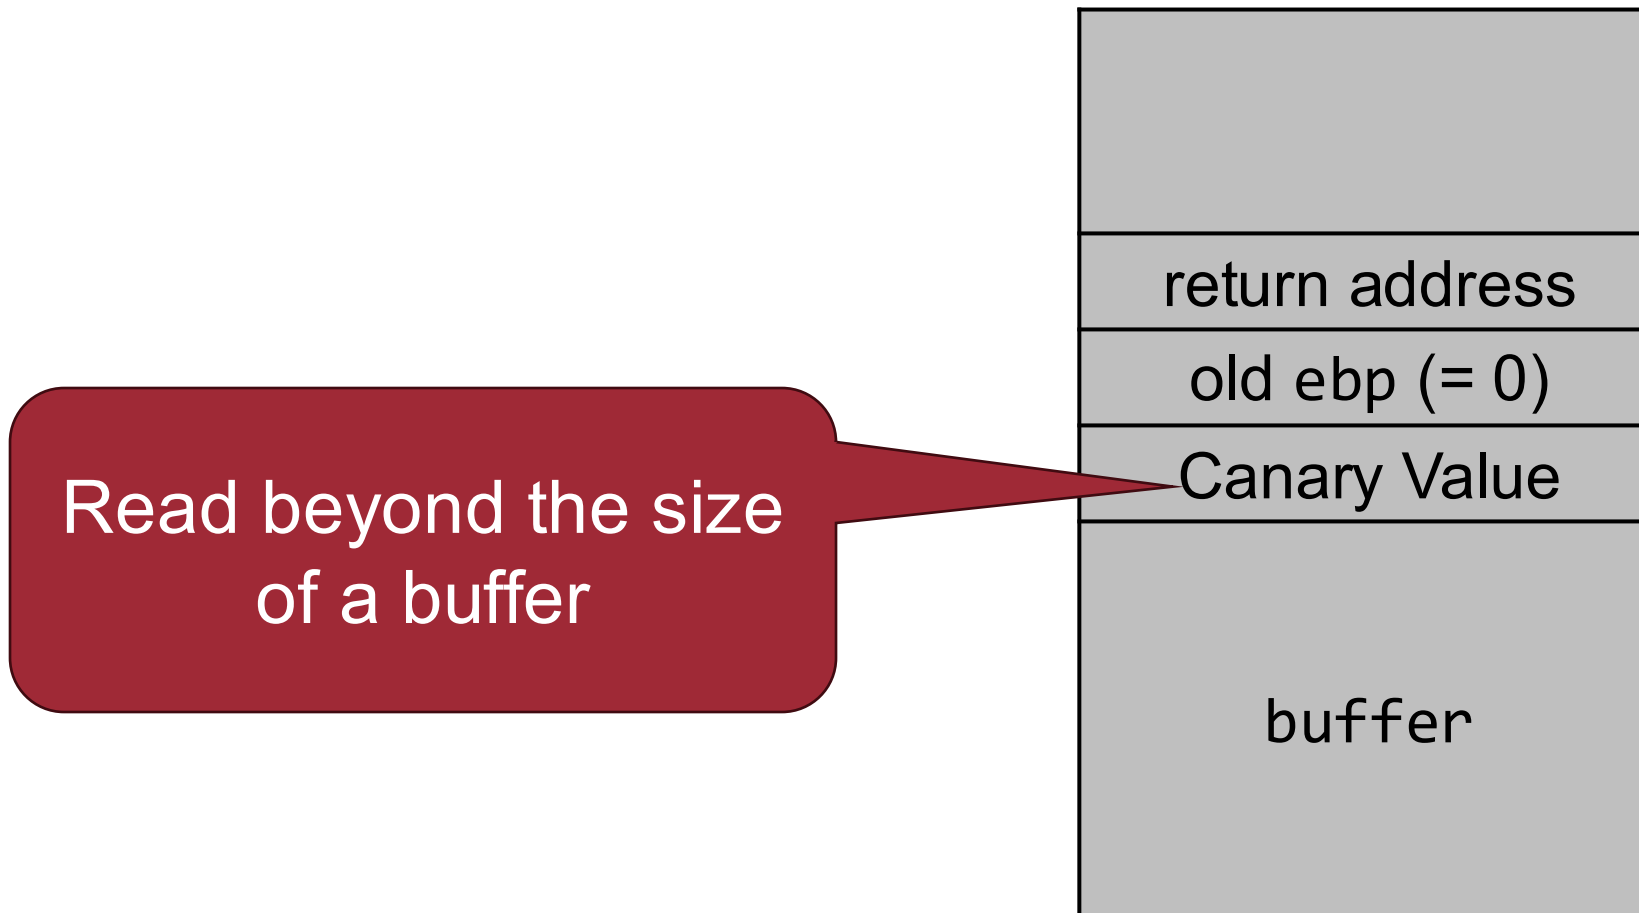


Memory disclosure does not necessarily involve memory corruption

Buffer Over-Read



Buffer over-read is a bug that allows an attacker to read beyond the size of a buffer

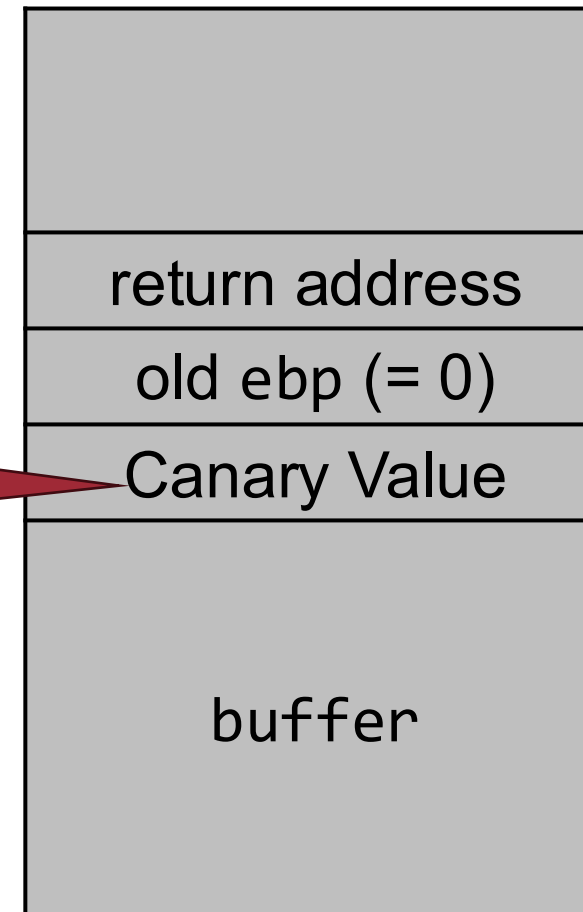


Buffer Over-Read



Buffer over-read is a bug that allows an attacker to read beyond the size of a buffer

Does ***not*** necessarily
involve memory
corruption!



Example: Heartbleed Bug (in 2014)

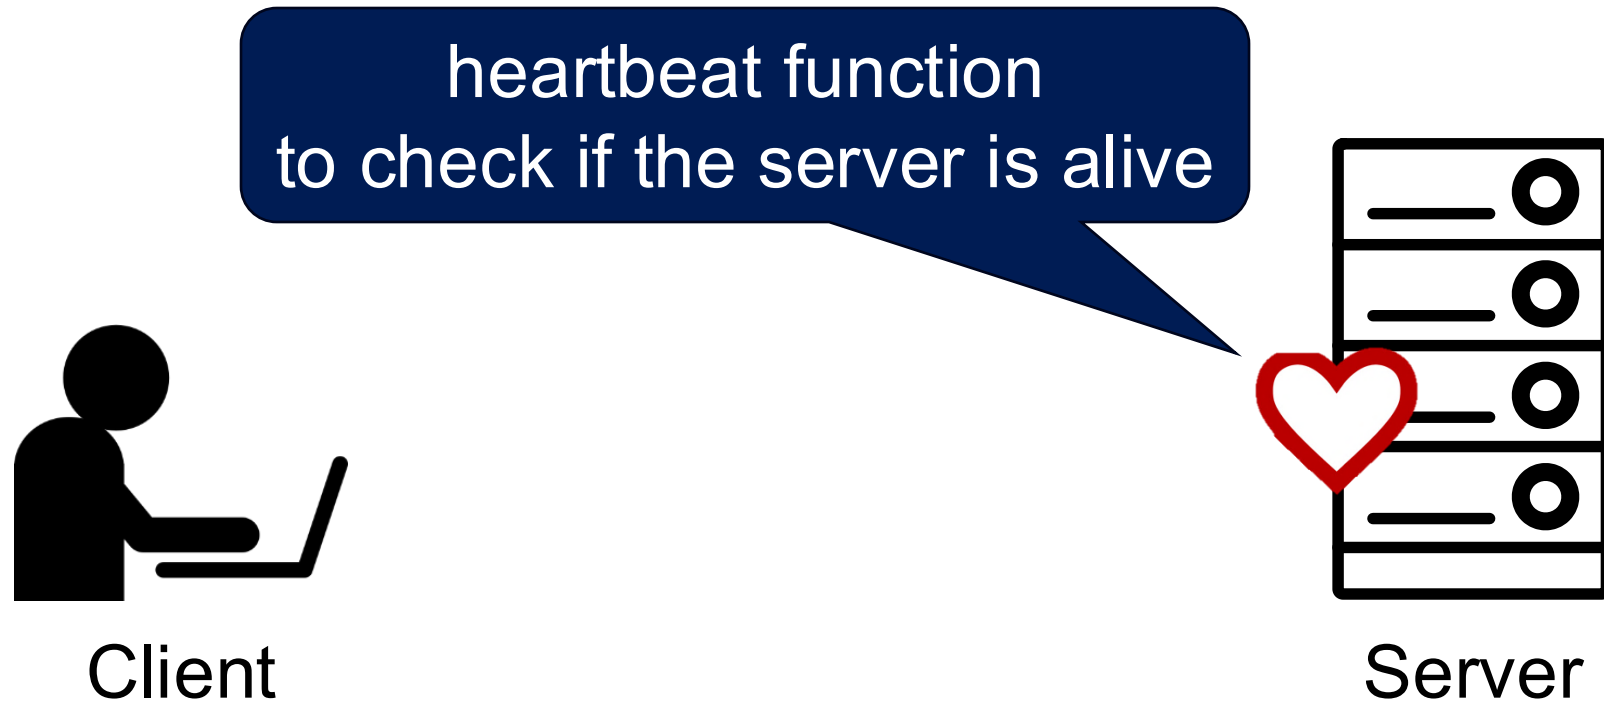
107

- Famous bug in OpenSSL (in TLS *heartbeat*)
- An attacker can steal private keys



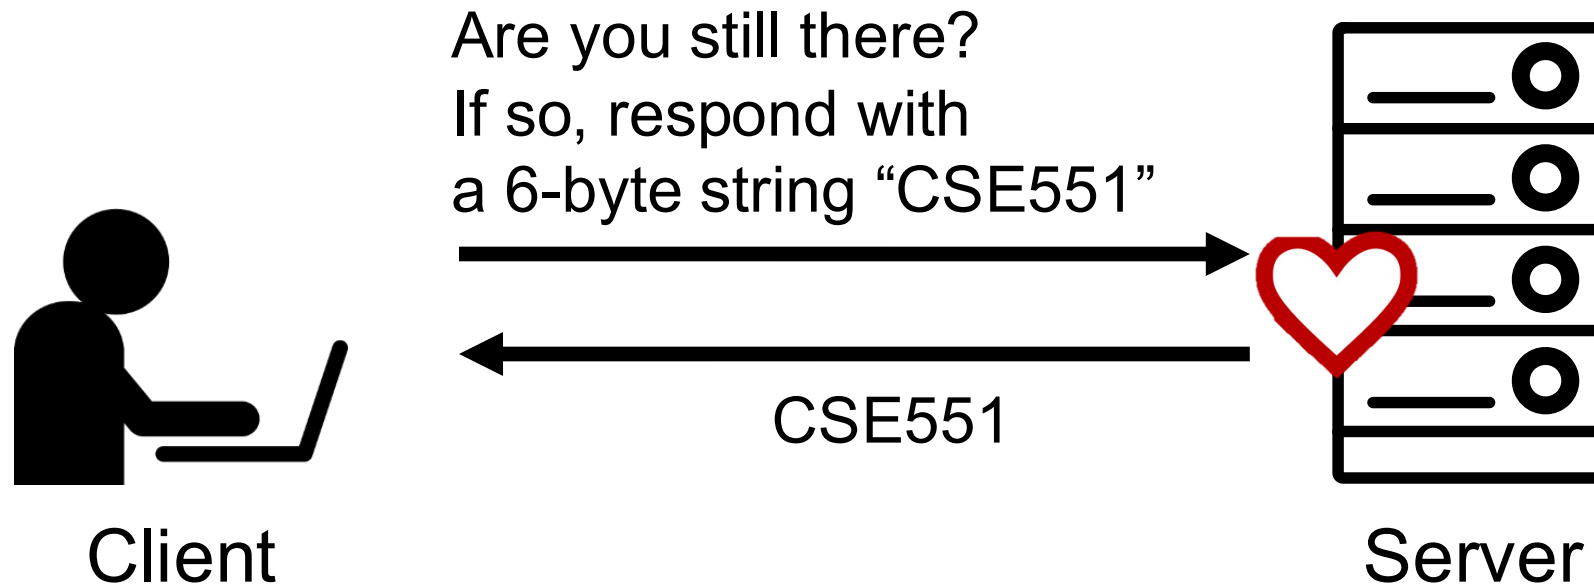
Heartbleed Bug: High-level Workflow

108



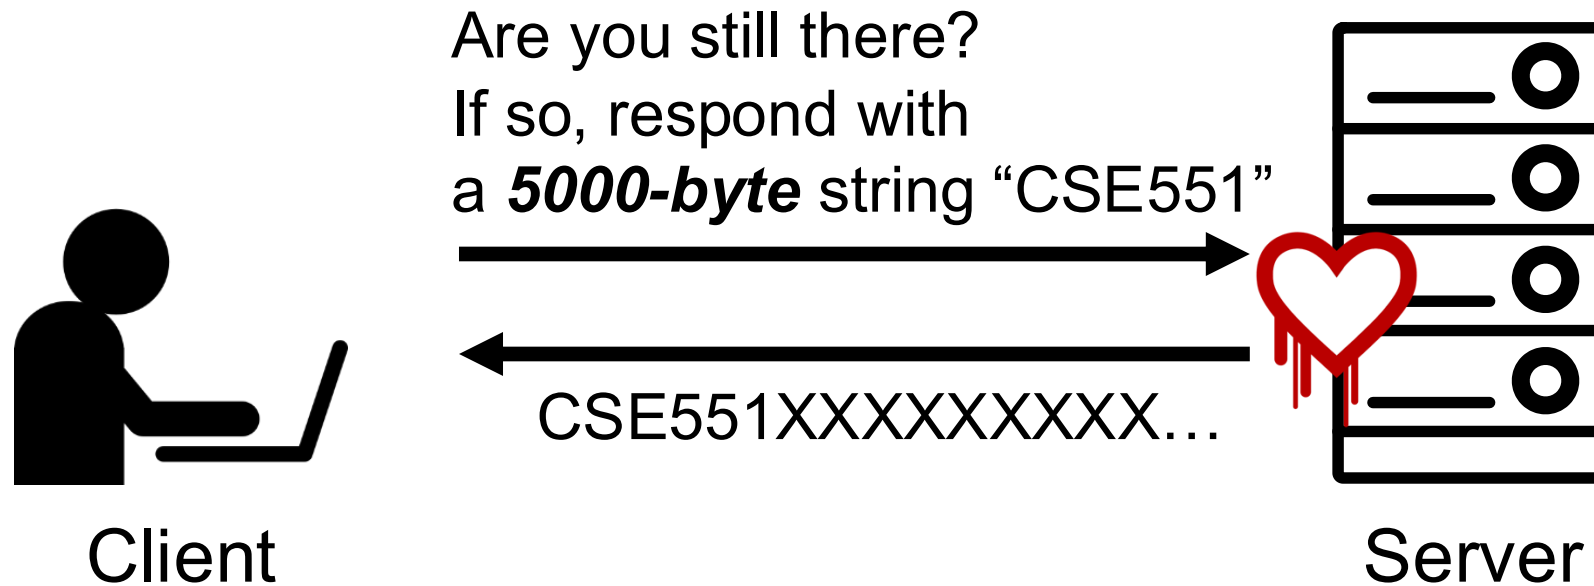
Heartbleed Bug: High-level Workflow

109



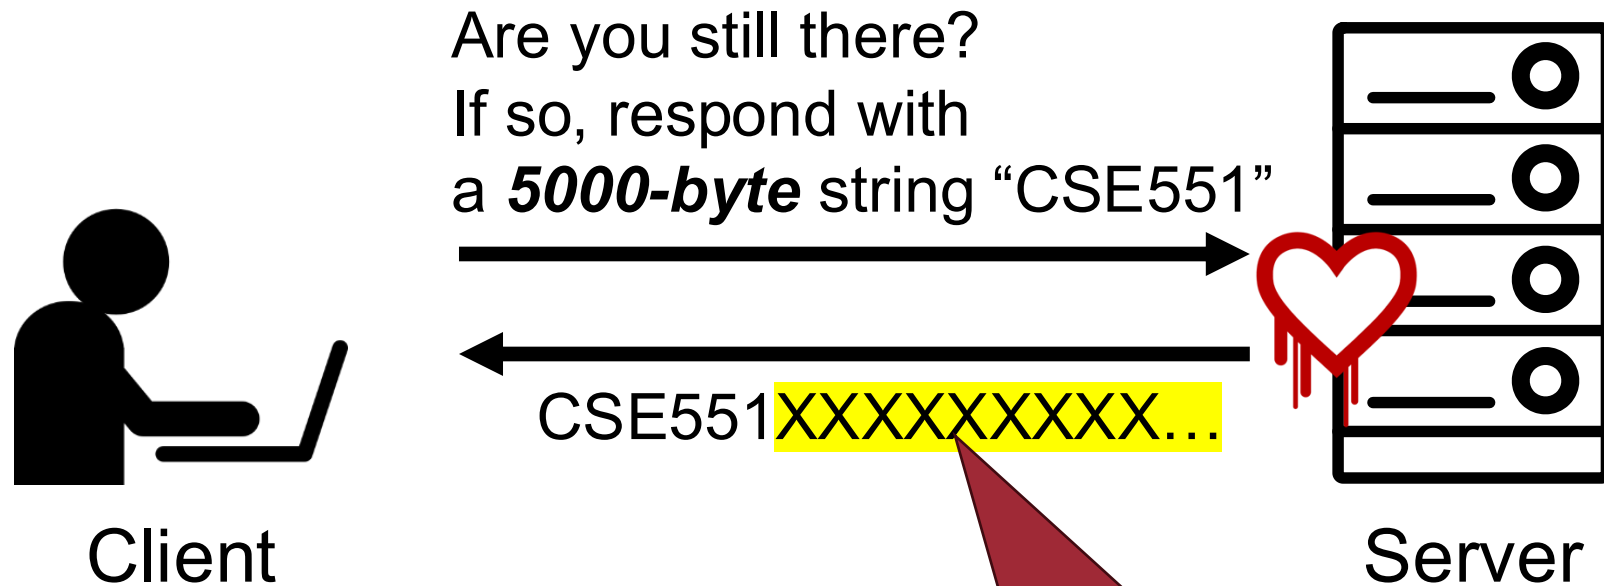
Heartbleed Bug: High-level Workflow

110



Heartbleed Bug: High-level Workflow

111



Memory disclosure!
(leak private keys)


The Bug



```
struct {  
    HeartbeatMessageType type;  
    uint16 payload_length;  
    opaque payload[HeartbeatMessage.payload_length];  
    opaque padding[padding_length];  
} HeartbeatMessage;
```

```
struct {  
    unsigned int length;  
    unsigned char *data;  
    ...  
} SSL3_RECORD;
```

The Bug

```
struct {  
    HeartbeatMessageType type;  
    uint16 payload_length;  
    opaque payload[HeartbeatMessage.payload_length];  
    opaque padding[HeartbeatMessage.padding_length];  
} HeartbeatMessage;  
  
struct {  
    unsigned int length;  
    unsigned char *data;  
    ...  
} SSL3_RECORD;  
  
memcpy(bp, pl, length); // vulnerable spot! 
```

Calculated from
the user's payload (i.e., 6)

Payload obtained from
HeartbeatMessage (i.e., CSE551)

Obtained from
the user's input (i.e., 5000)

Copy arbitrary memory contents of a
server! TLS private key may be available

The Bug

```
struct {  
    HeartbeatMessageType type;  
    uint16 payload_length;  
    opaque payload[HeartbeatMessage.payload_length];  
    opaque padding[HeartbeatMessage.payload_length];  
} HeartbeatMessage;
```

Calculated from
the user's payload (i.e., 6)

Payload obtained from
HeartbeatMessage (i.e., CSE551)

```
struct {  
    unsigned int length;  
    unsigned char *data;  
    ...  
} SSL3_RECORD;
```

Obtained from
the user's input (i.e., 5000)

```
memcpy(bp, pl, length); // vulnerable spot! 🐛
```

Copy arbitrary memory contents of a
server! TLS private key may be available

Root cause:

Did not check the
consistency of the values
of the two variables!

Other Memory Disclosure



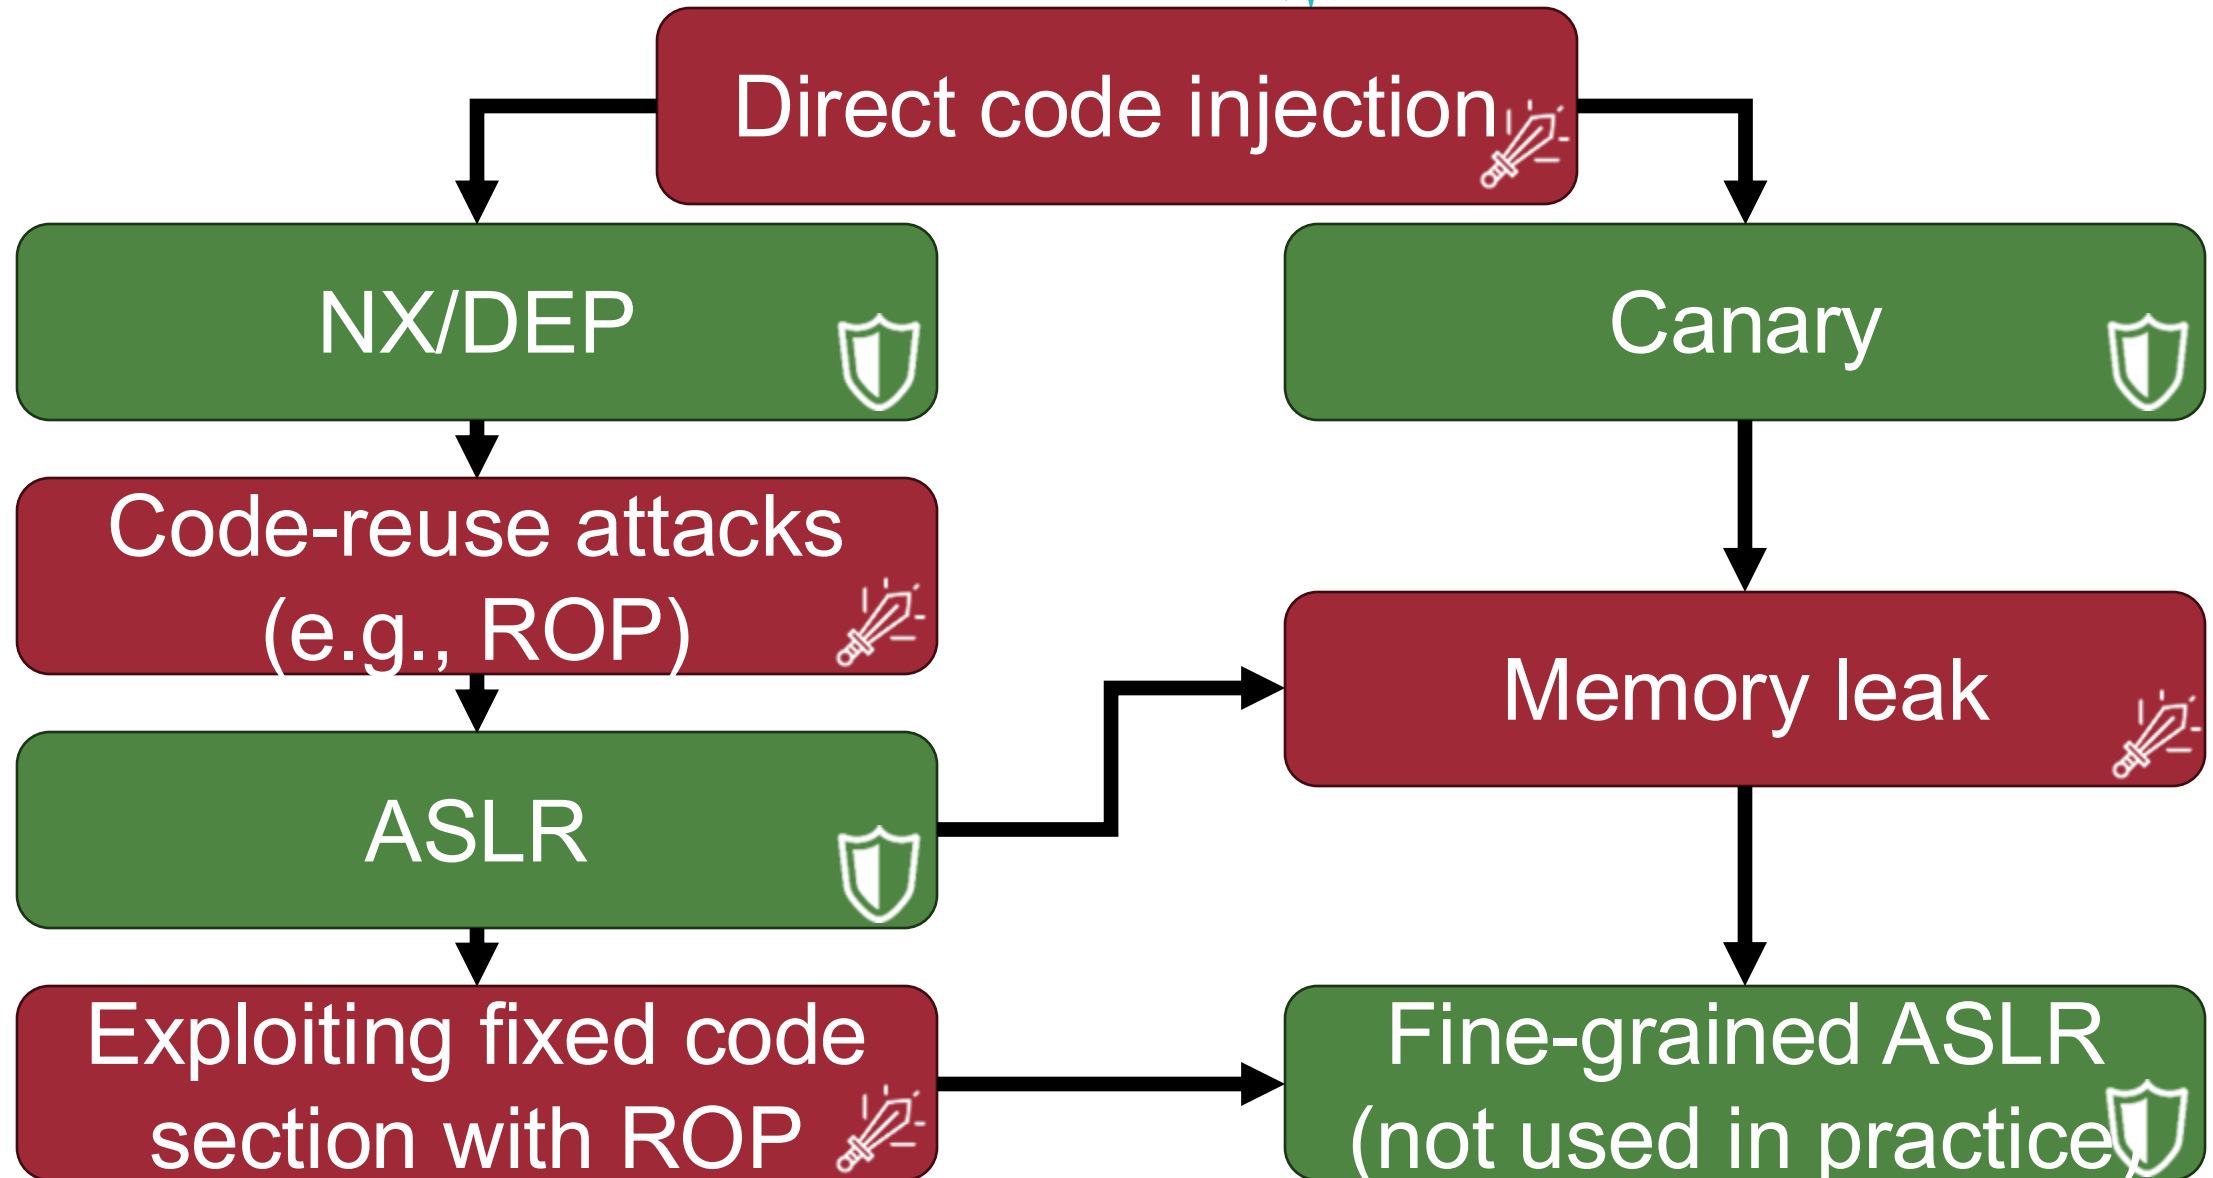
- Format string vulnerability also leaks memory info
 - “%08x.%08x.%08x...”
- Memory corruption bugs may allow memory leak
 - E.g., overwriting the length field of a string object

Memory Disclosure and Exploit

- It is possible that a program may have more than a single vulnerability
 - For example, one memory corruption and one memory disclosure
- In such a case, we can bypass existing defenses
 - **Canary bypass**: canary value could be leaked
 - **ASLR bypass**: code/stack pointers could be leaked


Caveat: we should be able to leak memory contents and trigger the memory corruption **within the same process**

Attack / Defense So Far



Arms Race in Security

118



Summary



- Code reuse attacks allow an attacker to bypass DEP
- Many mitigation techniques are proposed for code reuse attacks, which will be covered next
- ASLR: one of the mitigation techniques against code-reuse attacks
 - Brute-forcing attacks and ROP with fixed code section allow an attacker to bypass ASLR
- Memory disclosure (\neq Memory Corruption)
- Security vs. Performance

Question?