

CSE467: Computer Security

13. Control Flow Hijack

Seongil Wi

Software Bug




- Software bug is an ***error*** in a program

Q: If you only have time for fixing one bug out of hundred, which bug will you fix first?

Exploitable Bugs



We often call an *exploitable bug* as a vulnerability




Exploitable Bugs



We often call an **exploitable bug** as a vulnerability


Exploitation is an act of taking advantage of a bug to cause unintended behavior of the target program



Some vulnerabilities allow an attacker to run any **arbitrary code** on victim's machines without their consent

Control Flow Hijack

Control Flow Hijack



The Classic Exploitation

- The first computer worm (called Morris Worm) was born



Robert Tappan Morris

- Creator of the worm
- Cornell graduate
- Professor at MIT now

Morris Worm



- Exploited a **buffer overflow** vulnerability

```
int main(int argc, char* argv[]) {  
    char line[512];  
    /* omitted ... */  
    gets(line); /* Buffer Overflow! */  
    /* omitted ... */  
}
```

This simple line allowed the Morris Worm to infect 10% of the internet computers in 1988

Replicating Historic Exploitation



```
int main(int argc, char* argv[]) {  
    char line[512];  
    gets(line);  
    return 0;  
}
```

Compile this program with:
\$ gcc -mpreferred-stack-boundary=2
-O0 -fno-stack-protector -fno-pic
-no-pie -z execstack -o morris
morris.c -m32

Compiler Warning (ignore this for now):

morris.c:(.text+0x11): warning: the `gets' function is dangerous and should not be used.

gets(char *s)

11



Reads a line from STDIN into the buffer pointed to by s until a terminating new line or EOF, which it replaces with a NULL byte ('\0')


Disassembled Code for the Morris Worm 12

```
$ objdump -M intel -d morris
```

```
08049162 <main>:
```

8049162:	55	push	ebp
8049163:	89 e5	mov	ebp, esp
8049165:	81 ec 00 02 00 00	sub	esp, 0x200
804916b:	8d 85 00 fe ff ff	lea	eax, [ebp-0x200]
8049171:	50	push	eax
8049172:	e8 b9 fe ff ff	call	8049030 <gets@plt>
8049177:	83 c4 04	add	esp, 0x4
804917a:	b8 00 00 00 00	mov	eax, 0x0
804917f:	c9	leave	
8049180:	c3	ret	

Analyzing the Vulnerability



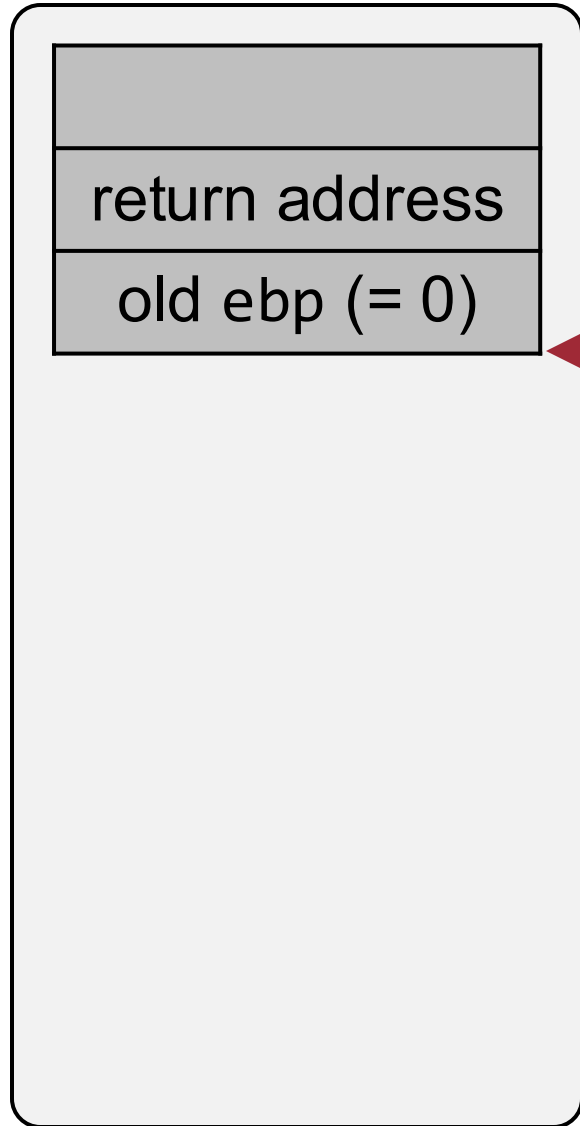
Virtual memory

eip: 0x8049162
ebp: 0x0
esp: 0xbffff70c

Execution context

08049162 <main>:
→ 8049162: push ebp
8049163: mov ebp, esp
8049165: sub esp, 0x200
804916b: lea eax, [ebp-0x200]
8049171: push eax
8049172: call 8049030 ; gets
8049177: add esp, 0x4
804917a: mov eax, 0x0
804917f: leave
8049180: ret

Analyzing the Vulnerability



Virtual memory

eip: 0x8049163
ebp: 0x0
esp: 0xbffff708

Execution context

08049162 <main>:

8049162: push ebp

8049163: mov ebp, esp

8049165: sub esp, 0x200

804916b: lea eax, [ebp-0x200]

8049171: push eax

8049172: call 8049030 ; gets

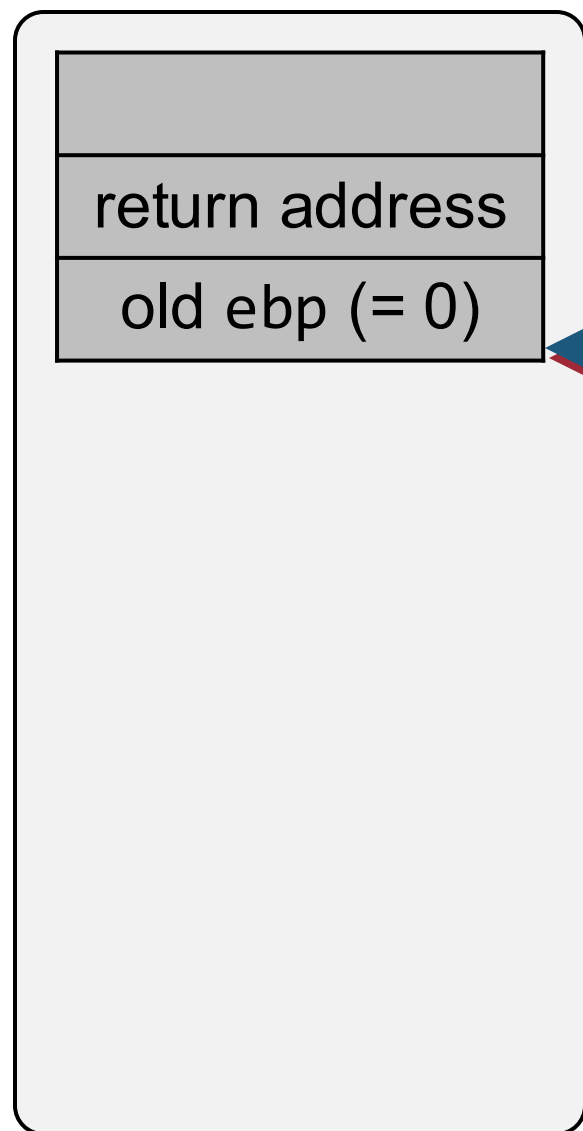
8049177: add esp, 0x4

804917a: mov eax, 0x0

804917f: leave

8049180: ret

Analyzing the Vulnerability



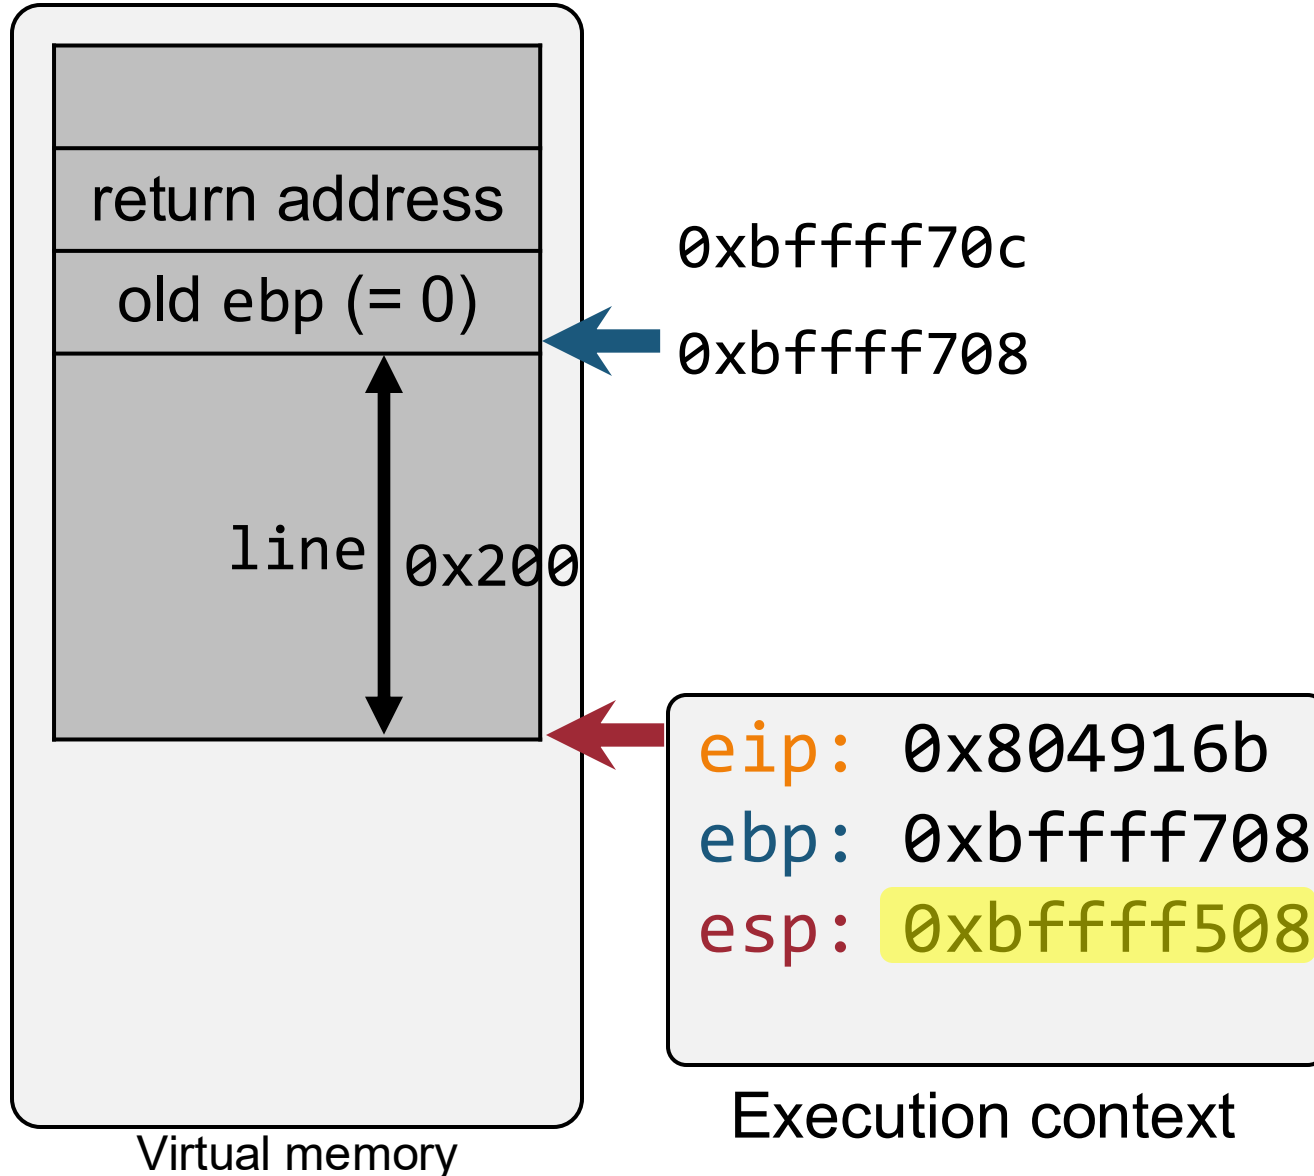
Virtual memory

eip: 0x8049165
ebp: 0xbffffff708
esp: 0xbffffff708

Execution context

08049162 <main>:
8049162: push ebp
8049163: mov ebp, esp
8049165: sub esp, 0x200
804916b: lea eax, [ebp-0x200]
8049171: push eax
8049172: call 8049030 ; gets
8049177: add esp, 0x4
804917a: mov eax, 0x0
804917f: leave
8049180: ret

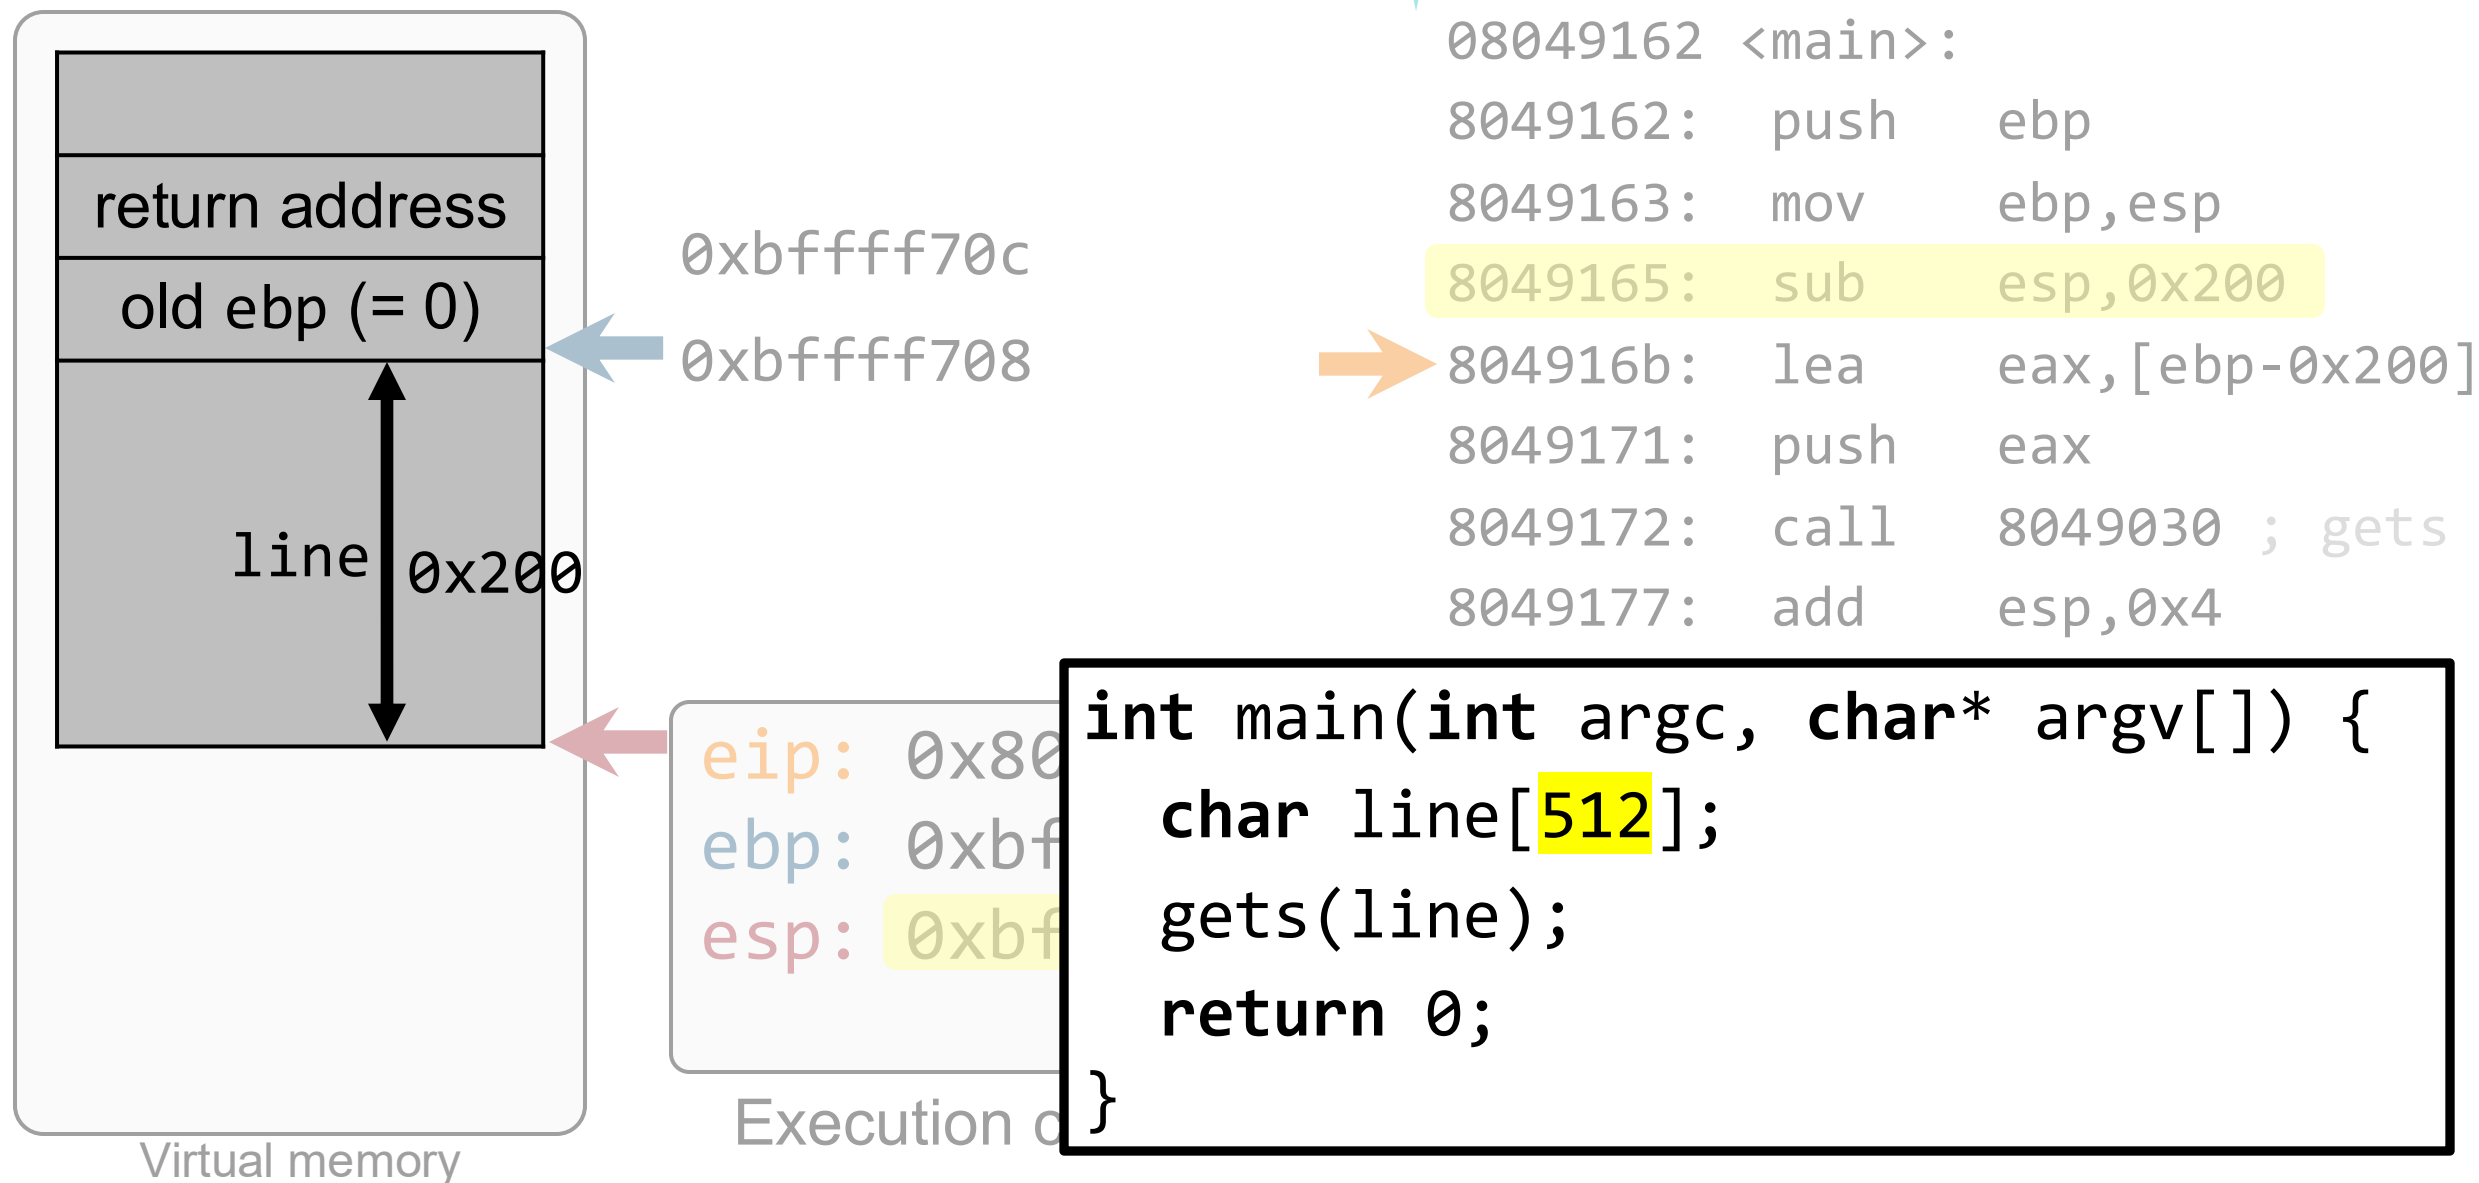
Analyzing the Vulnerability




```

08049162 <main>:
8049162:  push    ebp
8049163:  mov     ebp,esp
8049165:  sub     esp,0x200
804916b:  lea     eax,[ebp-0x200]
8049171:  push    eax
8049172:  call    8049030 ; gets
8049177:  add     esp,0x4
804917a:  mov     eax,0x0
804917f:  leave
8049180:  ret
  
```


Analyzing the Vulnerability



Analyzing the Vulnerability



0xbffffff70c

0xbffffff708

line

eip: 0x8049171
ebp: 0xbffffff708
esp: 0xbffffff508
eax: 0xbffffff508

Execution context

08049162 <main>:

8049162: push ebp

8049163: mov ebp, esp

8049165: sub esp, 0x200

804916b: lea eax, [ebp-0x200]

8049171: push eax

8049172: call 8049030 ; gets

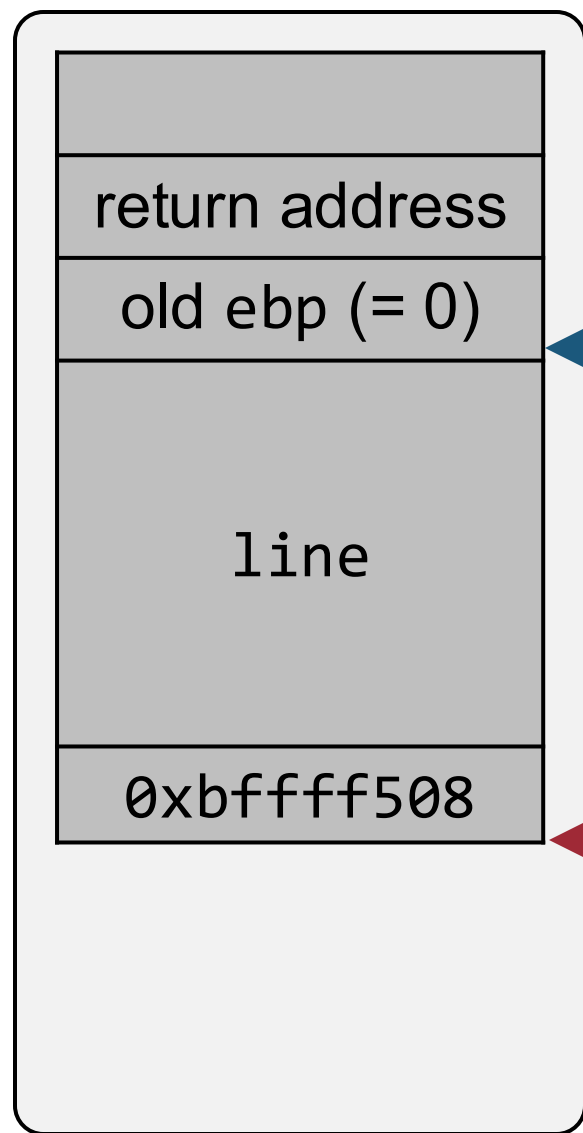
8049177: add esp, 0x4

804917a: mov eax, 0x0

804917f: leave

8049180: ret

Analyzing the Vulnerability



Virtual memory

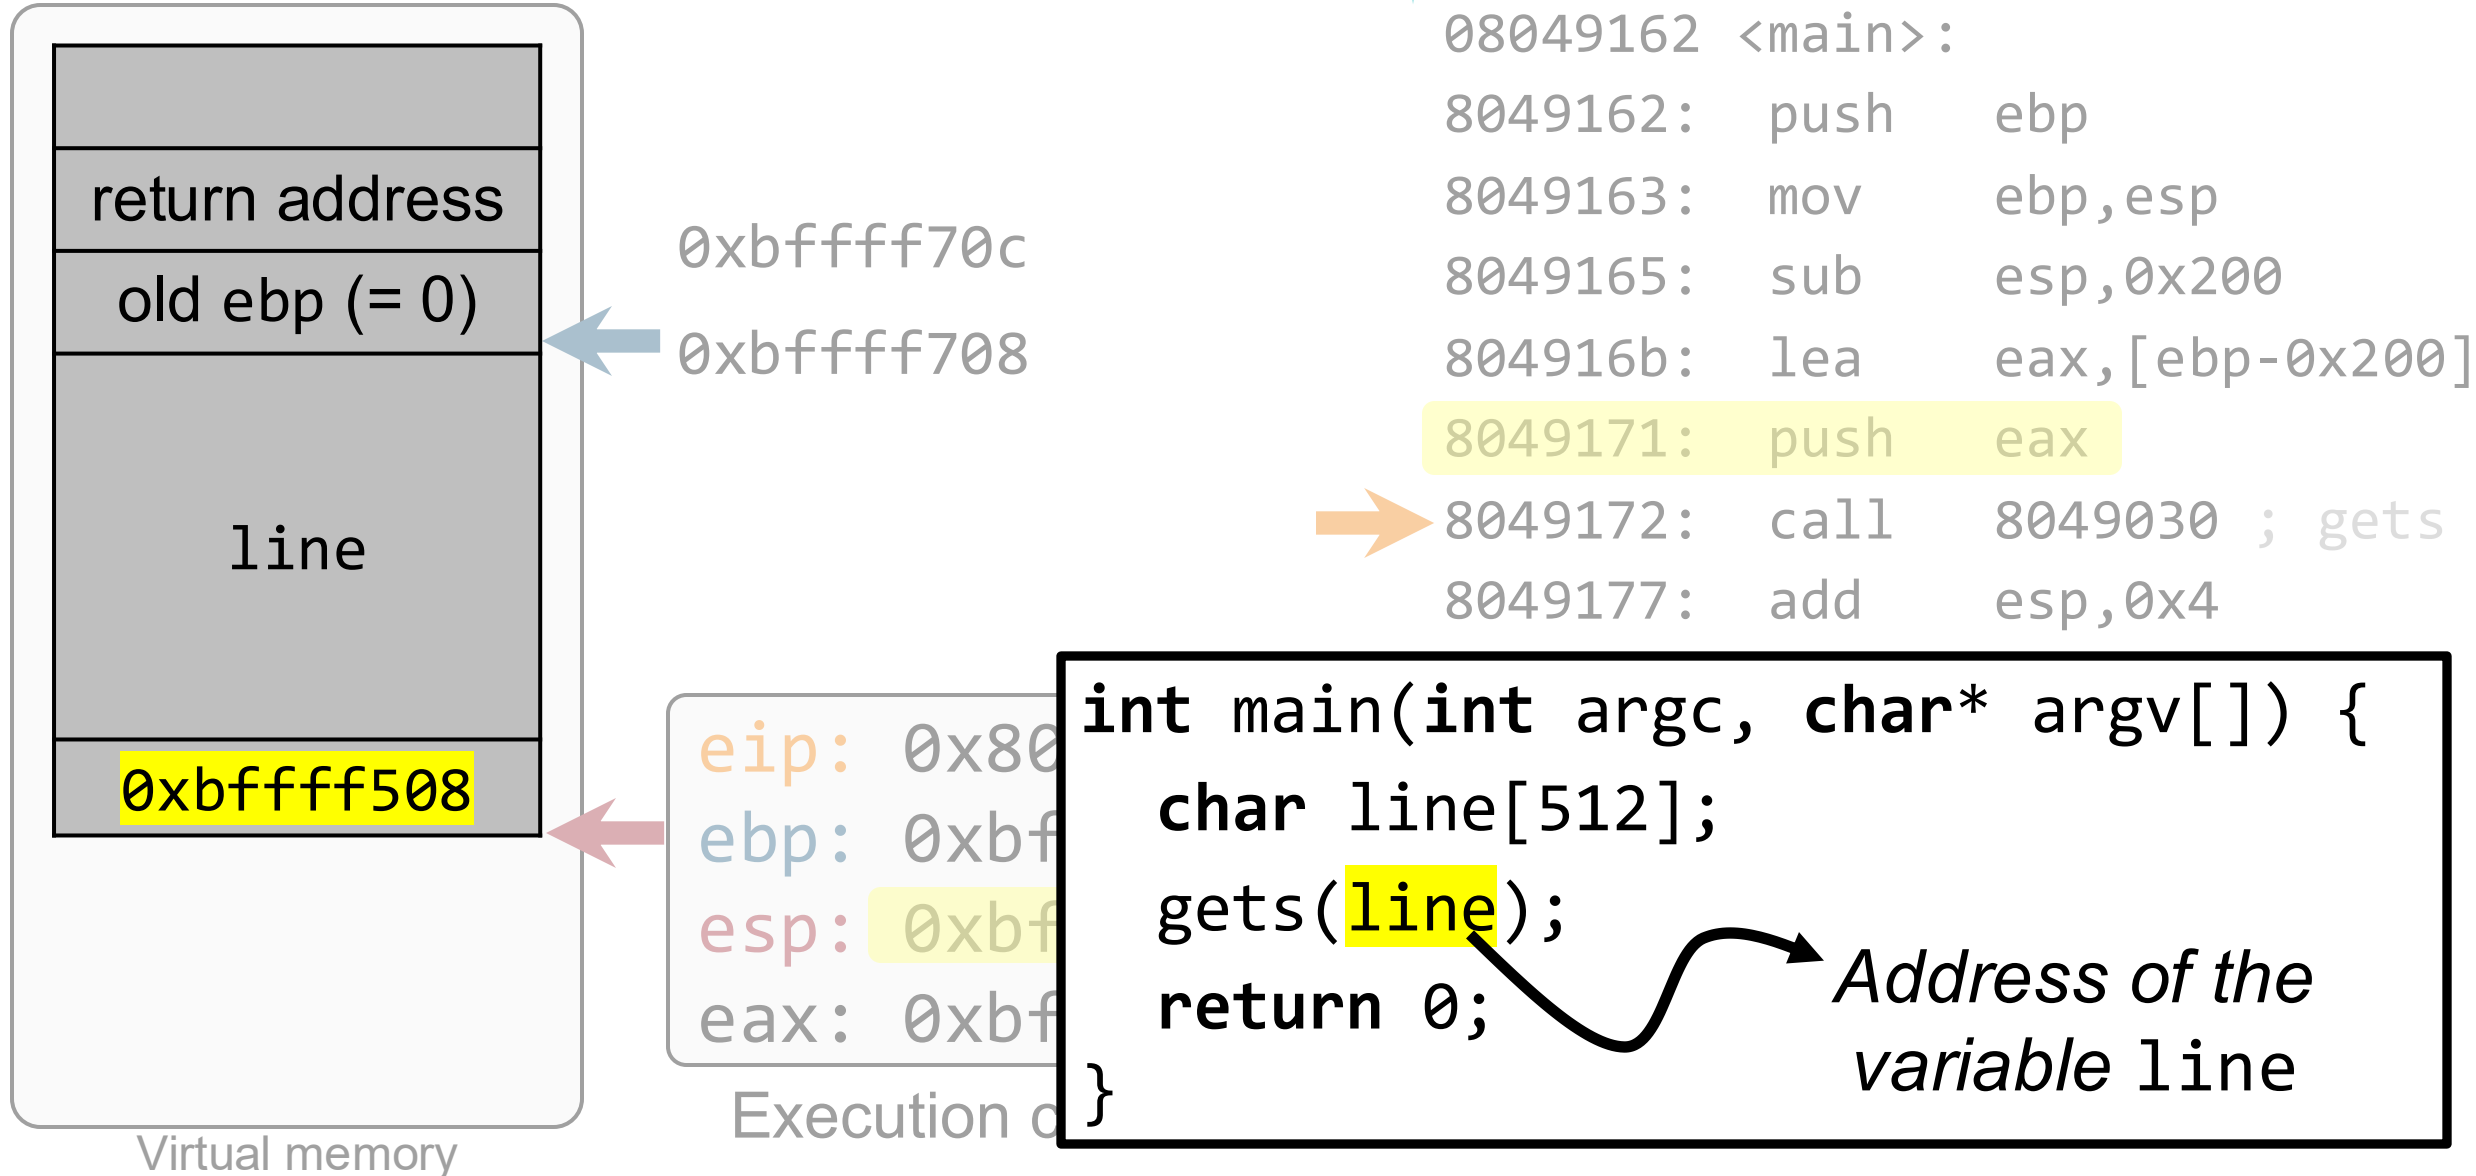
0xbffffff70c
0xbffffff708

eip: 0x8049172
ebp: 0xbffffff708
esp: 0xbffffff504
eax: 0xbffffff508

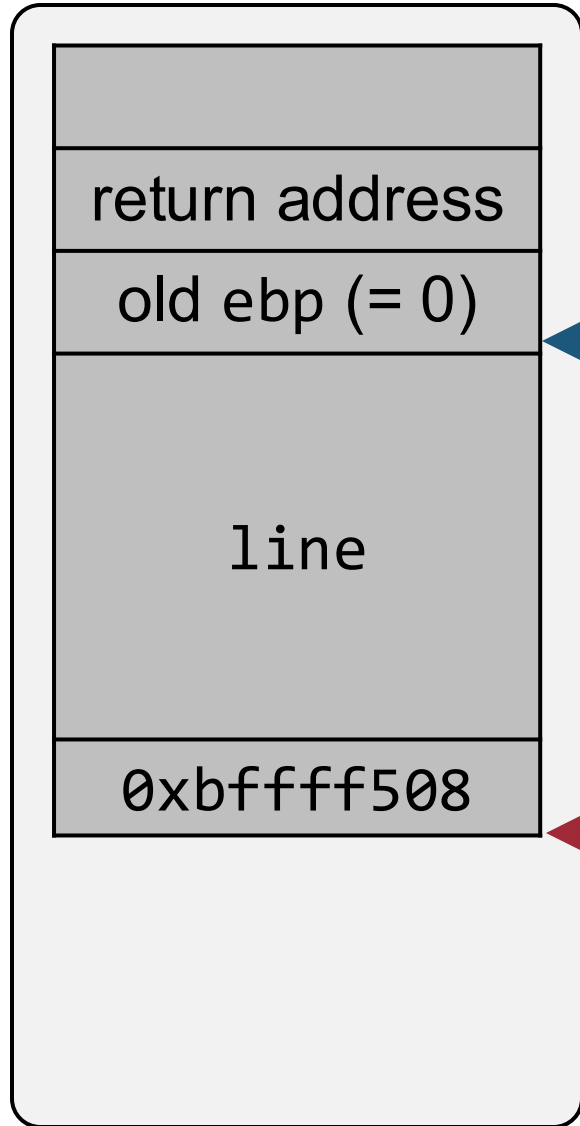
Execution context

```
08049162 <main>:  
8049162:  push    ebp  
8049163:  mov     ebp,esp  
8049165:  sub     esp,0x200  
804916b:  lea     eax,[ebp-0x200]  
8049171:  push    eax  
8049172:  call    8049030 ; gets  
8049177:  add     esp,0x4  
804917a:  mov     eax,0x0  
804917f:  leave  
8049180:  ret
```

Analyzing the Vulnerability



Analyzing the Vulnerability



What if user input is 520 consecutive 'A's?

0xbffff708

eip: 0x8049177
 ebp: 0xbffff708
 esp: 0xbffff504
 eax: 0xbffff508

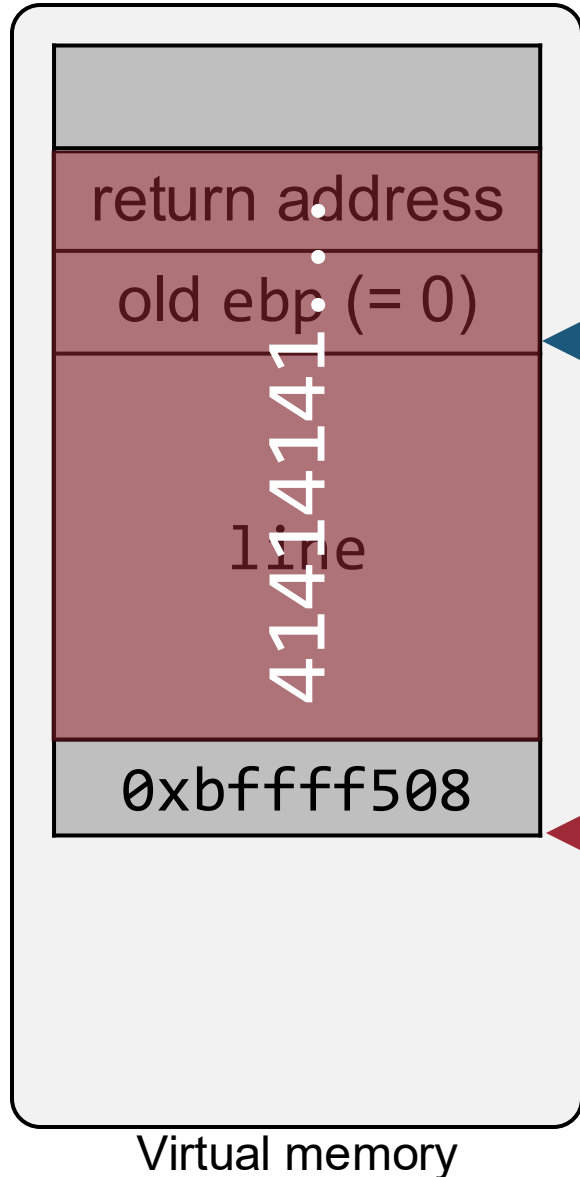
Execution context

08049162 <main>:

```

08049162:  push    ebp
08049163:  mov     ebp, esp
08049165:  sub     esp, 0x200
0804916b:  lea     eax, [ebp-0x200]
08049171:  push    eax
08049172:  call    8049030 ; gets
08049177:  add     esp, 0x4
0804917a:  mov     eax, 0x0
0804917f:  leave
08049180:  ret
  
```

Analyzing the Vulnerability



What if user input is 520 consecutive 'A's?

0xbffff708

eip: 0x8049177
 ebp: 0xbffff708
 esp: 0xbffff504
 eax: 0xbffff508

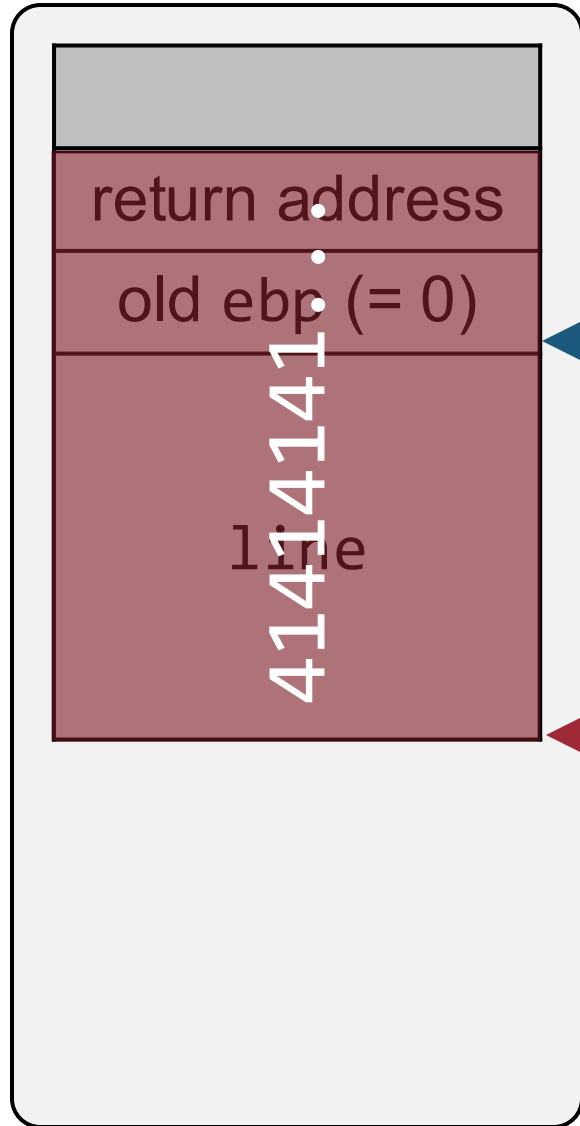
Execution context

08049162 <main>:

```

08049162:  push    ebp
08049163:  mov     ebp, esp
08049165:  sub     esp, 0x200
0804916b:  lea     eax, [ebp-0x200]
08049171:  push    eax
08049172:  call    8049030 ; gets
08049177:  add     esp, 0x4
0804917a:  mov     eax, 0x0
0804917f:  leave
08049180:  ret
  
```

Analyzing the Vulnerability



0xbffff70c

0xbffff708


eip: 0x804917a
 ebp: 0xbffff708
 esp: 0xbffff508
 eax: 0xbffff508

Execution context

```

08049162 <main>:
8049162:  push    ebp
8049163:  mov     ebp,esp
8049165:  sub     esp,0x200
804916b:  lea     eax,[ebp-0x200]
8049171:  push    eax
8049172:  call    8049030 ; gets
8049177:  add     esp,0x4
804917a:  mov     eax,0x0
804917f:  leave
8049180:  ret
  
```

Analyzing the Vulnerability



0xbffff70c

0xbffff708

eip: 0x804917f
 ebp: 0xbffff708
 esp: 0xbffff508
 eax: 0x0

Execution context

08049162 <main>:

8049162: push ebp

8049163: mov ebp, esp

8049165: sub esp, 0x200

804916b: lea eax, [ebp-0x200]

8049171: push eax

8049172: call 8049030 ; gets

8049177: add esp, 0x4


804917a: mov eax, 0x0

804917f: leave

8049180: ret

mov esp, ebp
 pop ebp

Analyzing the Vulnerability




Virtual memory

eip: 0x8049180
ebp: 0x41414141
esp: 0xbffff70c
eax: 0x0

Execution context

08049162 <main>:
8049162: push ebp
8049163: mov ebp, esp
8049165: sub esp, 0x200
804916b: lea eax, [ebp-0x200]
8049171: push eax
8049172: call 8049030 ; gets
8049177: add esp, 0x4
804917a: mov eax, 0x0
804917f: leave
8049180: ret

Analyzing the Vulnerability



Virtual memory

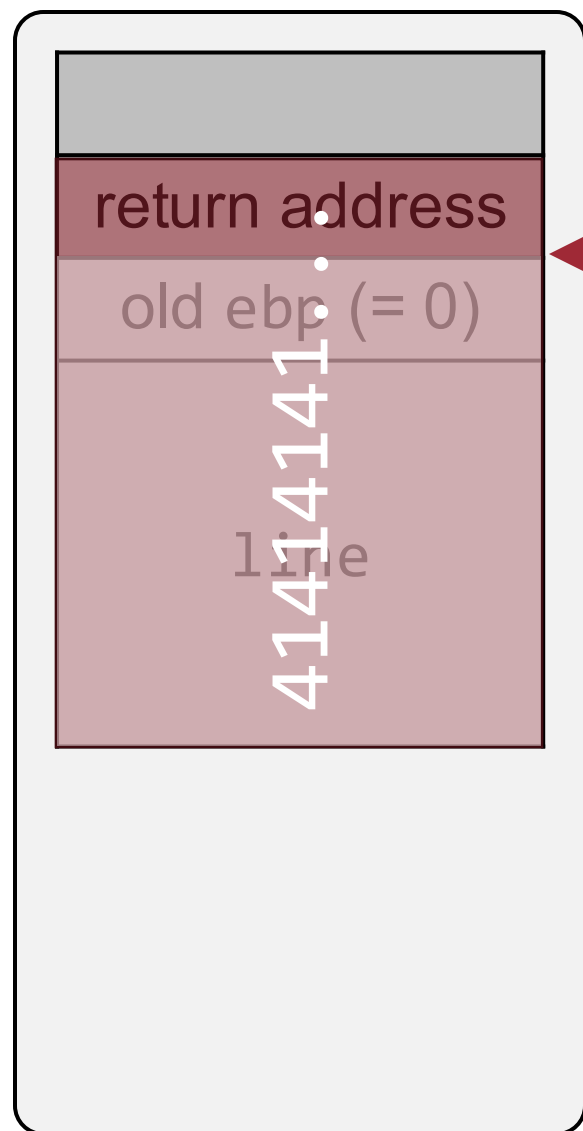
```
eip: 0x8049180
ebp: 0x41414141
esp: 0xbffff70c
eax: 0x0
```

Execution context

```
08049162 <main>:
8049162:  push    ebp
8049163:  mov     ebp,esp
8049165:  sub     esp,0x200
804916b:  lea     eax,[ebp-0x200]
8049171:  push    eax
8049172:  call    8049030 ; gets
8049177:  add     esp,0x4
804917a:  mov     eax,0x0
804917f:  leave
8049180:  ret
```

pop eip

Analyzing the Vulnerability



Virtual memory

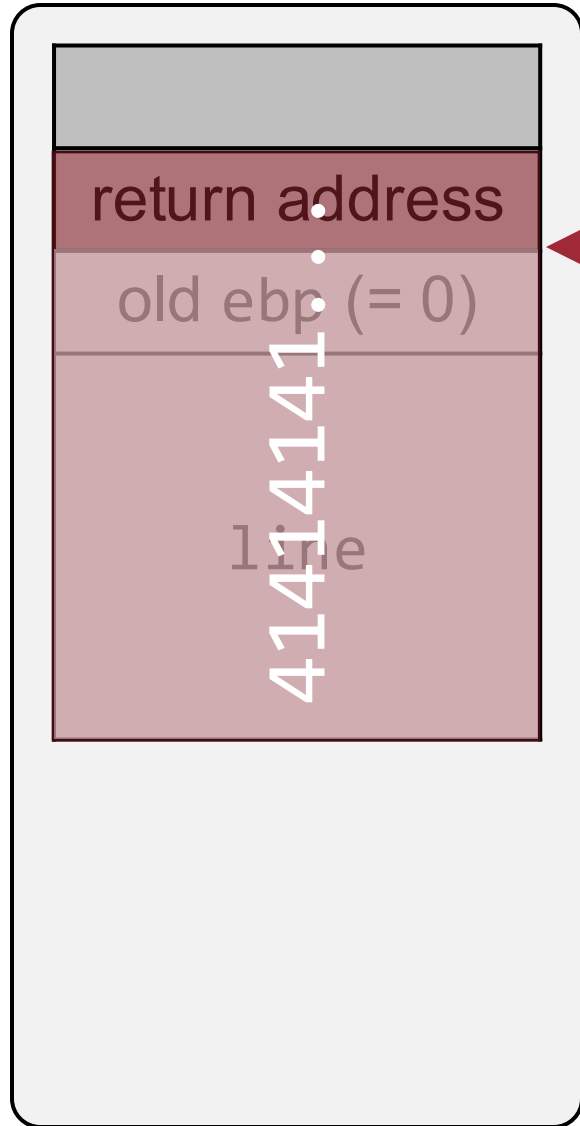
eip: 0x41414141
ebp: 0x41414141
esp: 0xbffff70c
eax: 0x0

Execution context

```
08049162 <main>:  
8049162:  push    ebp  
8049163:  mov     ebp,esp  
8049165:  sub     esp,0x200  
804916b:  lea     eax,[ebp-0x200]  
8049171:  push    eax  
8049172:  call    8049030 ; gets  
8049177:  add     esp,0x4  
804917a:  mov     eax,0x0  
804917f:  leave  
8049180:  ret
```

pop eip

Analyzing the Vulnerability



Virtual memory

eip: 0x41414141
ebp: 0x41414141
esp: 0xbffff70c
eax: 0x0

Execution context

Control flow hijacked!

```
08049162 <main>:  
8049162:  push    ebp  
8049163:  mov     ebp,esp  
8049165:  sub     esp,0x200  
804916b:  lea     eax,[ebp-0x200]  
8049171:  push    eax  
8049172:  call    8049030 ; gets  
8049177:  add     esp,0x4  
804917a:  mov     eax,0x0  
804917f:  leave  
8049180:  ret
```

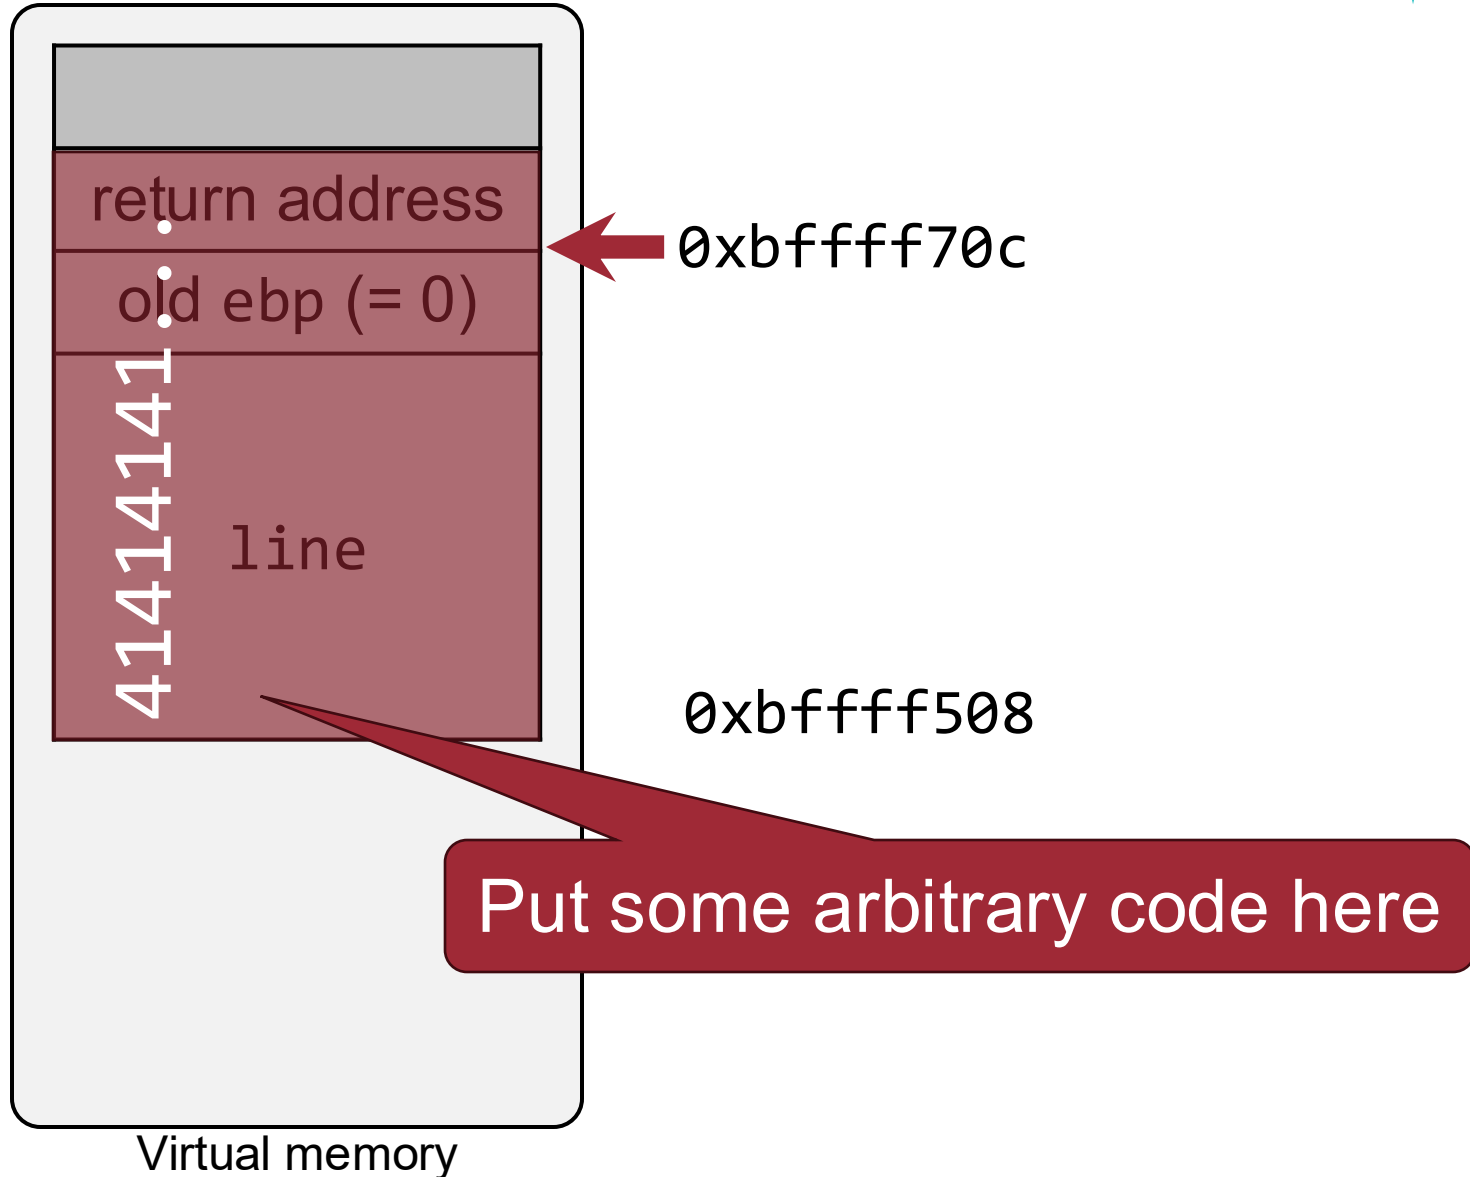
pop eip

So Far ...

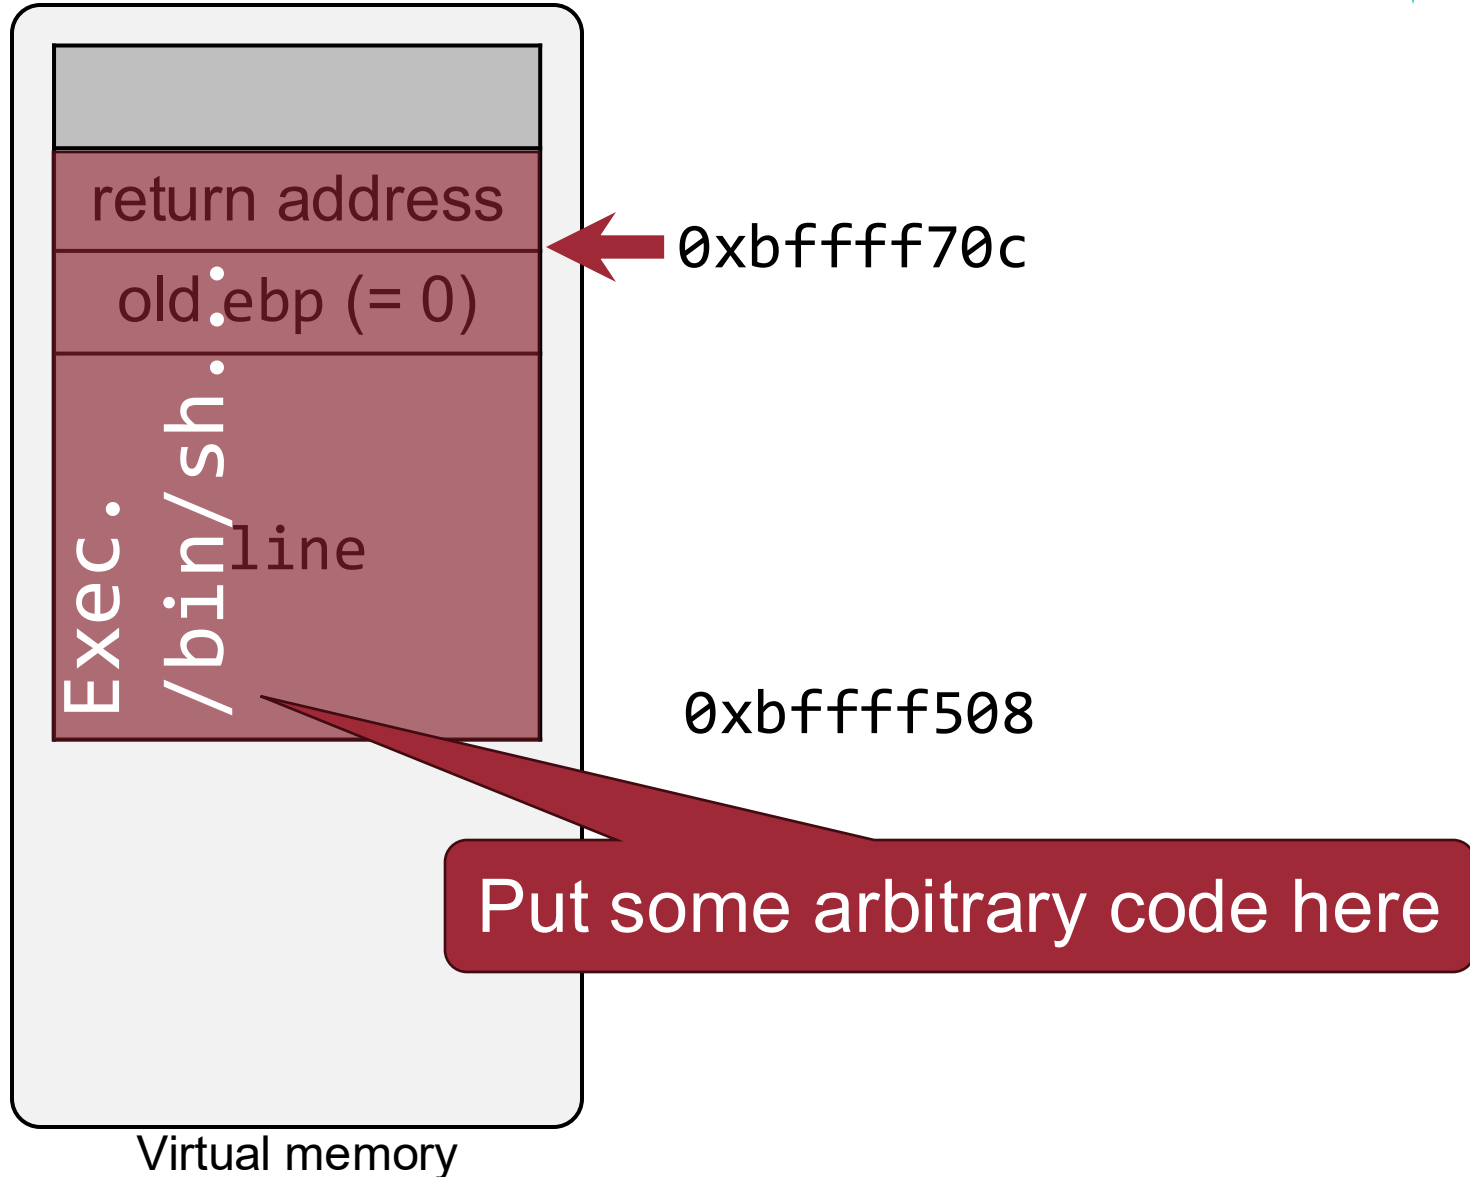


- We hijacked the control flow of the program, i.e., we can jump to any where!
- But, where do we jump to?
- We want to inject some ***arbitrary code*** to run!

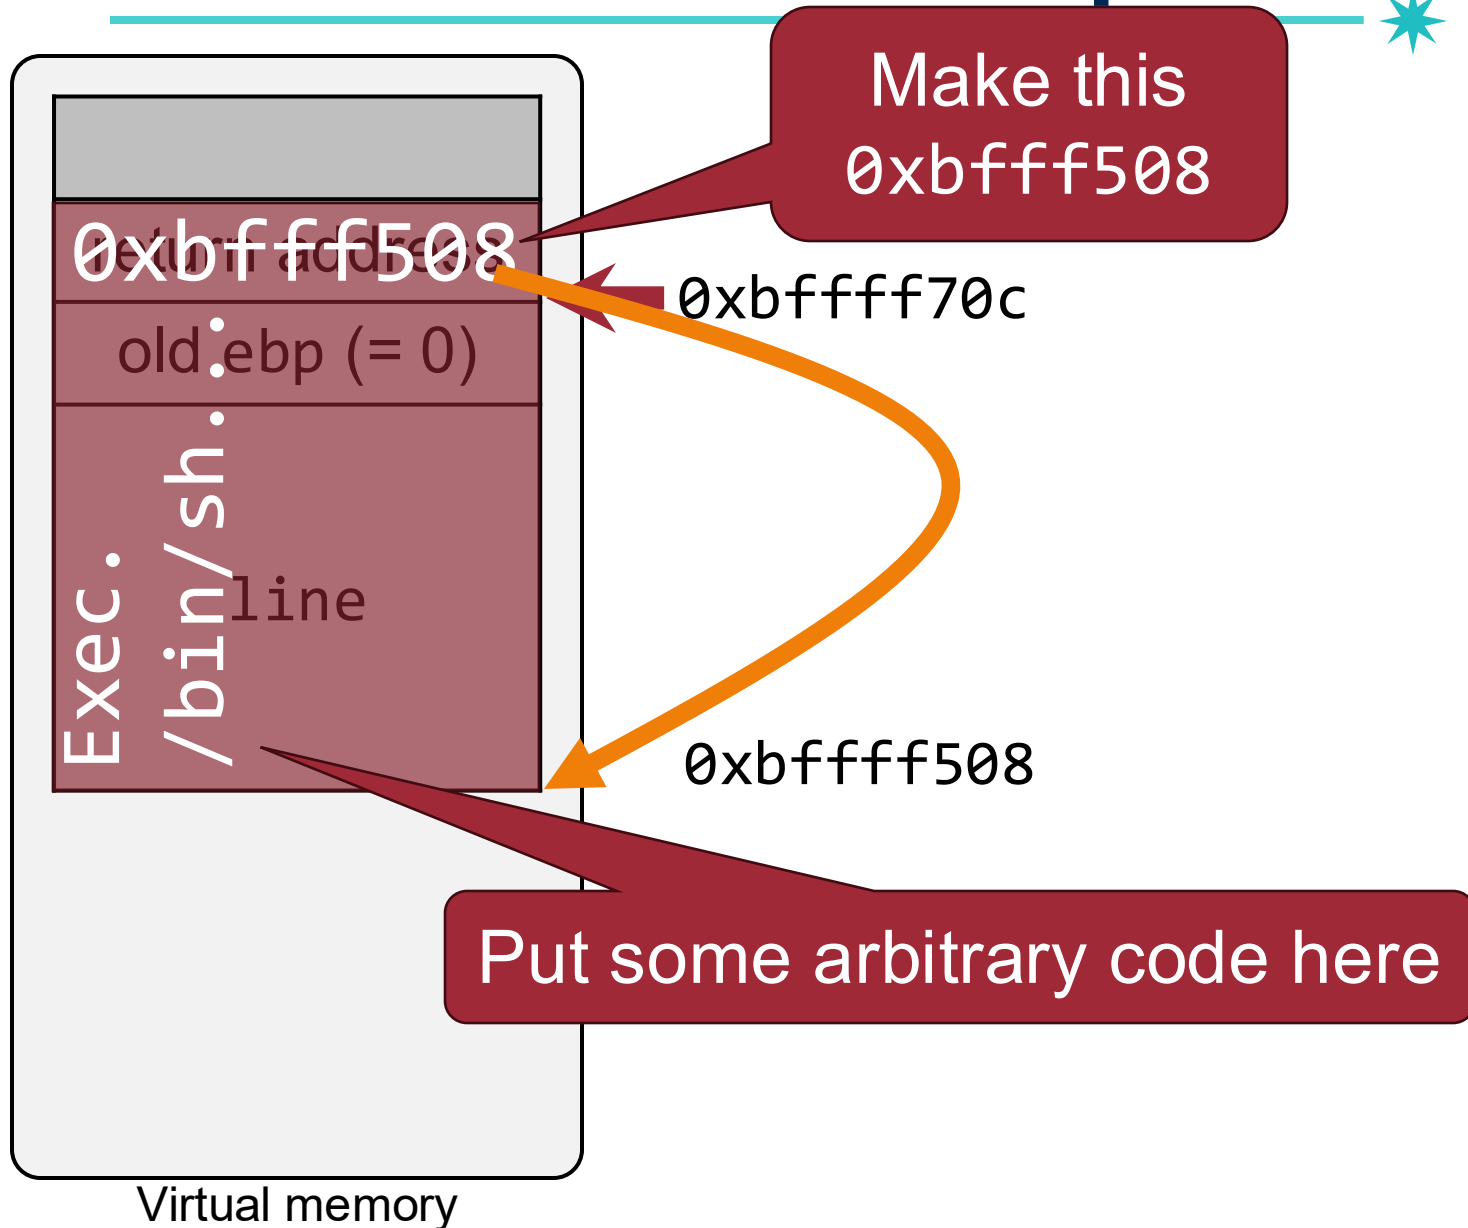
Return-to-Stack Exploit



Return-to-Stack Exploit



Return-to-Stack Exploit



Executing *Shellcode*



- Shellcode can run any arbitrary logic
 - Download /etc/passwd
 - Install malicious software (malware)
 - ...
- Typically, executing `/bin/sh` is enough
 - This is the most powerful attack: we can run arbitrary commands
 - You can also achieve this with relatively ***small piece of code***
 - This is the reason why we call it as ***shellcode*** (code that typically runs shell)

Shellcoding



How to write code that executes `/bin/sh`?

execve() Function in C Library

\$ man execve

```
EXECVE(2)                Linux Programmer's Manual                EXECVE(2)
NAME
    execve - execute program

SYNOPSIS
    #include <unistd.h>

    int execve(const char *filename, char *const argv[],
               char *const envp[]);
```

Executable path

Environment variables

Command line arguments

Shellcode in C



```
/*  
    int execve(const char *filename, char *const argv[],  
               char *const envp[]);  
*/  
  
#include <stdio.h>  
  
void main(void) {  
    char* argv[] = { "/bin/sh", NULL };  
    execve("/bin/sh", argv, NULL);  
}
```

Compile with (-static) Option

08049162 <main>:

8049162:	55	push	ebp
8049163:	89 e5	mov	ebp, esp
8049165:	83 ec 08	sub	esp, 0x8
8049168:	c7 45 f8 08 a0 04 08	mov	DWORD PTR [ebp-0x8], 0x804a008
804916f:	c7 45 fc 00 00 00 00	mov	DWORD PTR [ebp-0x4], 0x0
8049176:	6a 00	push	0x0
8049178:	8d 45 f8	lea	eax, [ebp-0x8]
804917b:	50	push	eax
804917c:	68 08 a0 04 08	push	0x804a008
8049181:	e8 c7 29 02 00	call	806c4b0 <__execve>

Is it possible to use this assembly for our exploitation?

Challenge #1: Null Bytes



08049162 <main>:

8049162:	55	push	ebp
8049163:	89 e5	mov	ebp, esp
8049165:	83 ec 08	sub	esp, 0x8
8049168:	c7 45 f8 08 a0 04 08	mov	DWORD PTR [ebp-0x8], 0x804a008
804916f:	c7 45 fc 00 00 00 00	mov	DWORD PTR [ebp-0x4], 0x0
8049176:	6a 00	push	0x0
8049178:	8d 45 f8	lea	eax, [ebp-0x8]
804917b:	50	push	eax
804917c:	68 08 a0 04 08	push	0x804a008
8049181:	e8 c7 29 02 00	call	806c4b0 <__execve>

gets(char *s)



Reads a line from STDIN into the buffer pointed to by s until a terminating new line or EOF, which it replaces with a NULL byte ('\0')

Challenge #1: Null Bytes



08049162 <main>:

```

8049162: 55                push    ebp
8049163: 89 e5            mov     ebp, esp
8049165: 83 ec 08         sub     esp, 0x8
8049168: c7 45 f8 08 a0 04 08 mov     DWORD PTR [ebp-0x8], 0x804a008
804916f: c7 45 fc 00 00 00 00 mov     DWORD PTR [ebp-0x4], 0x0
8049176: 6a 00           push    0x0
8049178: 8d 45 f8        lea     eax, [ebp-0x8]
804917b: 50             push    eax
804917c: 68 08 a0 04 08  push    0x804a008
8049181: c7 45 f8 00 00 00 00 mov     DWORD PTR [ebp-0x8], 0x0

```

Solution:

Write your own assembly code (shellcode) that does not contain any zero (NULL) byte

Challenge #2: String Pointer



08049162 <main>:

8049162:	55	push	ebp
8049163:	89 e5	mov	ebp, esp
8049165:	83 ec 08	sub	esp, 0x8
8049168:	c7 45 f8 08 a0 04 08	mov	DWORD PTR [ebp-0x8], 0x804a008
804916f:	c7 45 fc 00 00 00 00	mov	DWORD PTR [ebp-0x4], 0x0
8049176:	6a 00	push	0x0
8049178:	8d 45 f8	lea	eax, [ebp-0x8]
804917b:	50	push	eax
804917c:	68 08 a0 04 08	push	0x804a008
8049181:	e8 c7 29 02 00	call	806c4b0 <__execve>

Pointer to "/bin/sh"

Solution:

Push "/bin/sh" string to stack and get the pointer from esp

Challenge #3: External Call (execve)

08049162 <main>:

8049162:	55	push	ebp
8049163:	89 e5	mov	ebp, esp
8049165:	83 ec 08	sub	esp, 0x8
8049168:	c7 45 f8 08 a0 04 08	mov	DWORD PTR [ebp-0x8], 0x804a008
804916f:	c7 45 fc 00 00 00 00	mov	DWORD PTR [ebp-0x4], 0x0
8049176:	6a 00	push	0x0
8049178:	8d 45 f8	lea	eax, [ebp-0x8]
804917b:	50	push	eax
804917c:	68 08 a0 04 08	push	0x804a008
8049181:	e8 c7 29 02 00	call	806c4b0 <__execve>

Just a wrapper function
in the C library (libc)

Solution:

We can just inline this function

Challenge #3: External Call (execve)



0806c4b0 <__execve>:

806c4b0:	53	push	ebx
806c4b1:	8b 54 24 10	mov	edx,DWORD PTR [esp+0x10]
806c4b5:	8b 4c 24 0c	mov	ecx,DWORD PTR [esp+0xc]
806c4b9:	8b 5c 24 08	mov	ebx,DWORD PTR [esp+0x8]
806c4bd:	b8 0b 00 00 00	mov	eax,0xb
806c4c2:	cd 80	int	0x80

System Call!

Solution:

We can just inline this function

System Calls

allow a program to interface with OS



0806c4b0 <__execve>:

```

806c4b0:  53                push    ebx
806c4b1:  8b 54 24 10       mov     edx,DWORD PTR [esp+0x10]
806c4b5:  8b 4c 24 0c       mov     ecx,DWORD PTR [esp+0xc]
806c4b9:  8b 5c 24 08       mov     ebx,DWORD PTR [esp+0x8]
806c4bd:  b8 0b 00 00 00    mov     eax,0xb
806c4c2:  cd 80            int     0x80

```

Register	Meaning
eax	System call number
ebx	1 st argument
ecx	2 nd argument
edx	3 rd argument

Register	Meaning
esi	4 th argument
edi	5 th argument
ebp	6 th argument
eax	Return value

List of System Calls for x86

- See: /usr/include/i386-linux-gnu/asm/unistd_32.h

```
#define __NR_restart_syscall 0
#define __NR_exit 1
#define __NR_fork 2
#define __NR_read 3
#define __NR_write 4
#define __NR_open 5
#define __NR_close 6
#define __NR_waitpid 7
#define __NR_creat 8
#define __NR_link 9
#define __NR_unlink 10
#define __NR_execve 11
#define __NR_chdir 12 ...
```



0xb

Writing a Shellcode



- Shellcode should run regardless of the address it is loaded. In other words, it should be ***position independent***.

Writing a Shellcode




```
.intel_syntax noprefix
```

```
; This is a comment
```

Assemble this code
to see the binary

Final Exploitation



- Fill the buffer with our shellcode (Let's assume that it is 31 bytes)
- The rest of the buffer (481 bytes = 512-31) can be filled with any characters
- The old ebp can be filled with any characters (4 bytes)
- The return address should point to the shellcode (`0xbffff508`)¹

¹The buffer address should differ from machine to machine. Thus, it is necessary to obtain the right address from a debugger (e.g., GDB)

Caveat



We assume that we know the exact address of the buffer

This is very difficult even without modern defenses such as ASLR

Using GDB



- GDB reference: <http://www.yolinux.com/TUTORIALS/GDB-Commands.html>
- It is recommended to always turn on the Intel syntax by using this command:

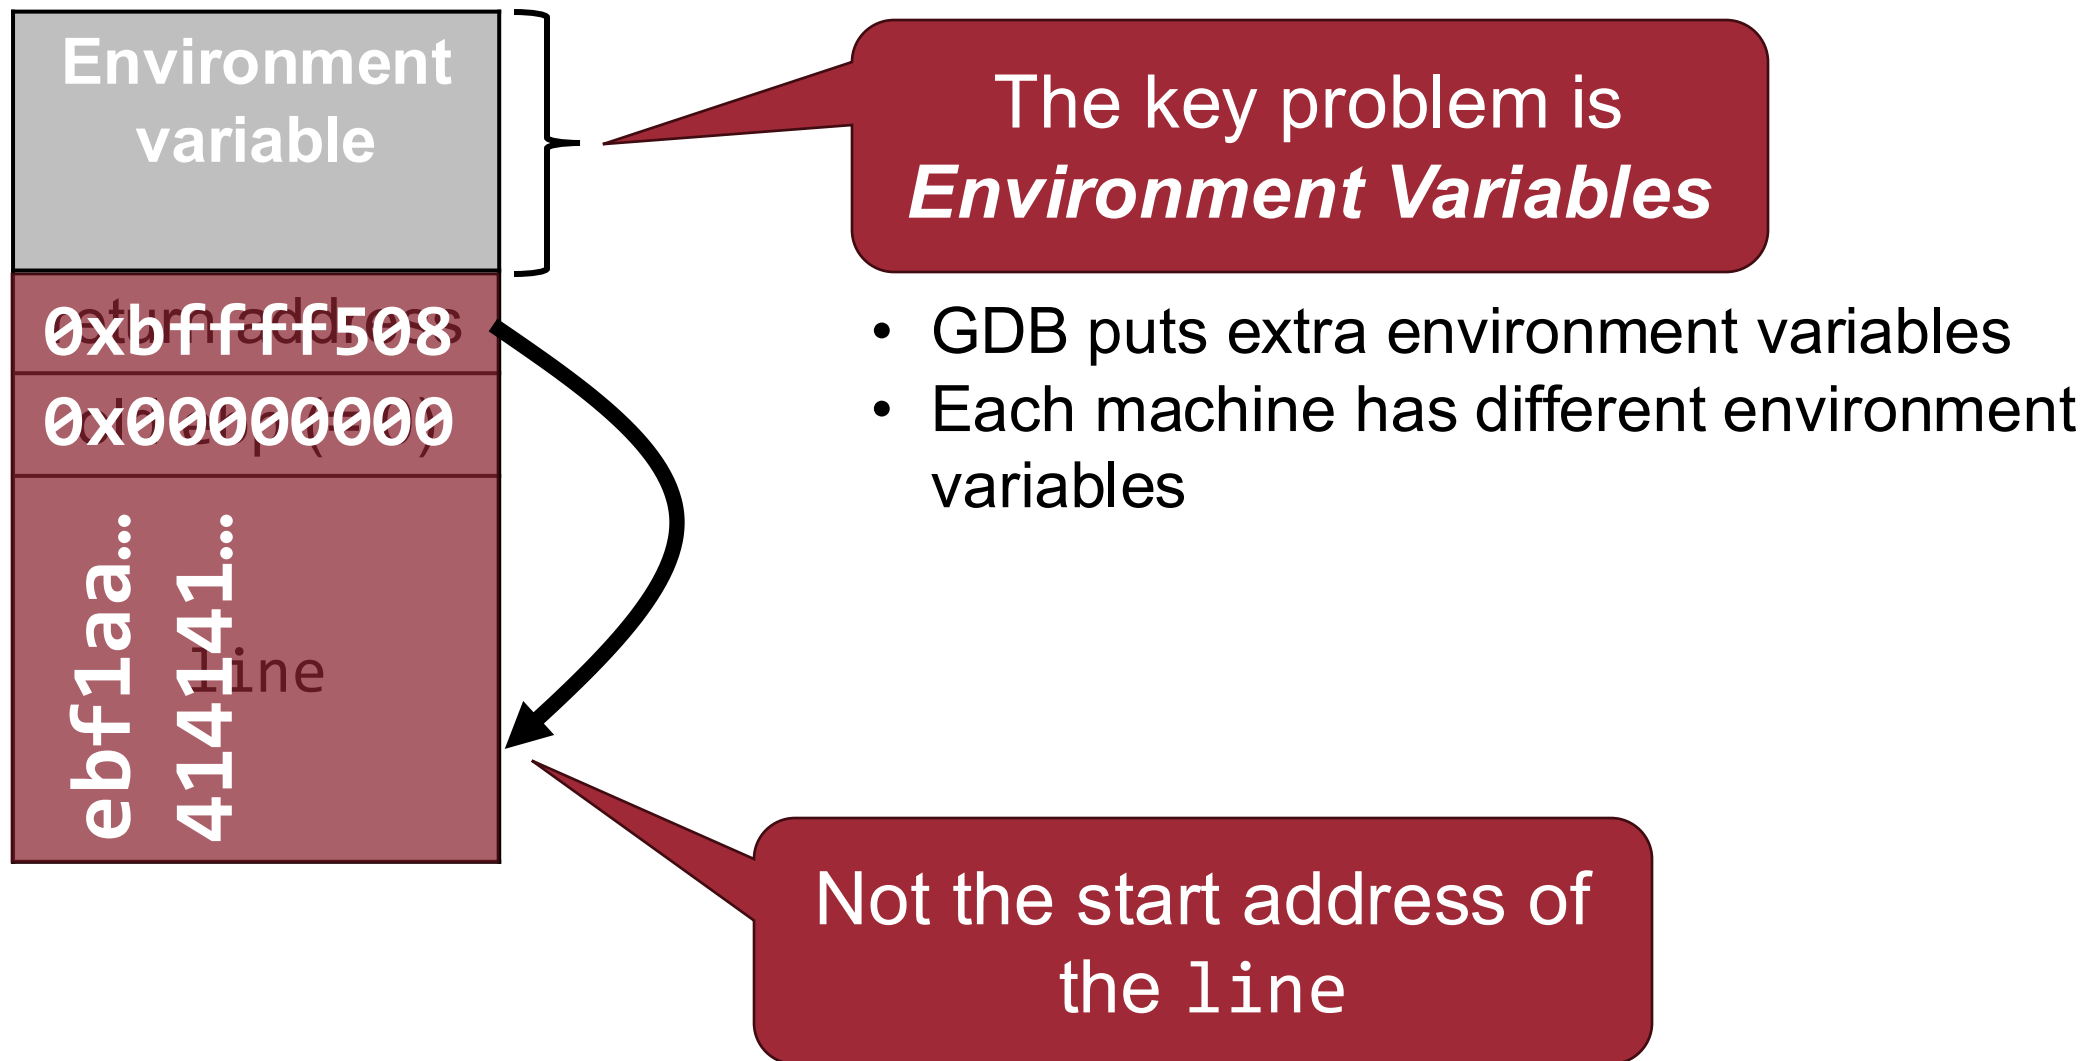
```
$ echo "set disassembly-flavor intel" > ~/.gdbinit
```

Exploit w/ or w/o GDB




- The buffer address identified through GDB is not the same as it without GDB
- Thus, our exploit doesn't work outside GDB!

Why Different?




Making Exploit Robust: NOP Sled

- NOP sled (a.k.a., NOP slide) is used to make the exploit robust against different buffer addresses
 - One-byte NO-OP (NOP) instruction is equivalent to `xchg eax, eax`
 - `0x90` represents the NOP instruction



Making Exploit Robust: NOP Sled

- NOP sled (a.k.a., NOP slide) is used to make the exploit robust against different buffer addresses
 - One-byte NO-OP (NOP) instruction is equivalent to `xchg eax, eax`
 - `0x90` represents the NOP instruction



Making Exploit Robust: NOP Sled

- NOP sled (a.k.a., NOP slide) is used to make the exploit robust against different buffer addresses
 - One-byte NO-OP (NOP) instruction is equivalent to `xchg eax, eax`
 - `0x90` represents the NOP instruction



Making Exploit Robust: NOP Sled

- NOP sled (a.k.a., NOP slide) is used to make the exploit robust against different buffer addresses
 - One-byte NO-OP (NOP) instruction is equivalent to `xchg eax, eax`
 - `0x90` represents the NOP instruction



Off-by-One Error

Subtle Error



```
#include <stdio.h>
#include <string.h>
#define BUFSIZE (512)

void printer(char* str) {
    char buf[BUFSIZE];
    strcpy(buf, str);
    printf("%s\n", buf);
}

int main(int argc, char* argv[]) {
    if ( argc < 2 || strlen(argv[1]) > BUFSIZE ) return -1;
    printer(argv[1]);
    return 0;
}
```

Subtle Error



```
#include <stdio.h>
#include <string.h>
#define BUFSIZE (512)
```

```
void printer(char* str) {
    char buf[BUFSIZE];
    strcpy(buf, str);
    printf("%s\n", buf);
}
```

```
int main(int argc, char* argv[]) {
    if ( argc < 2 || strlen(argv[1]) > BUFSIZE ) return -1;
    printer(argv[1]);
    return 0;
}
```

We can just overwrite 1 byte NULL beyond the size of the buffer (buf)

But, some off-by-one bugs are exploitable!

Subtle Error



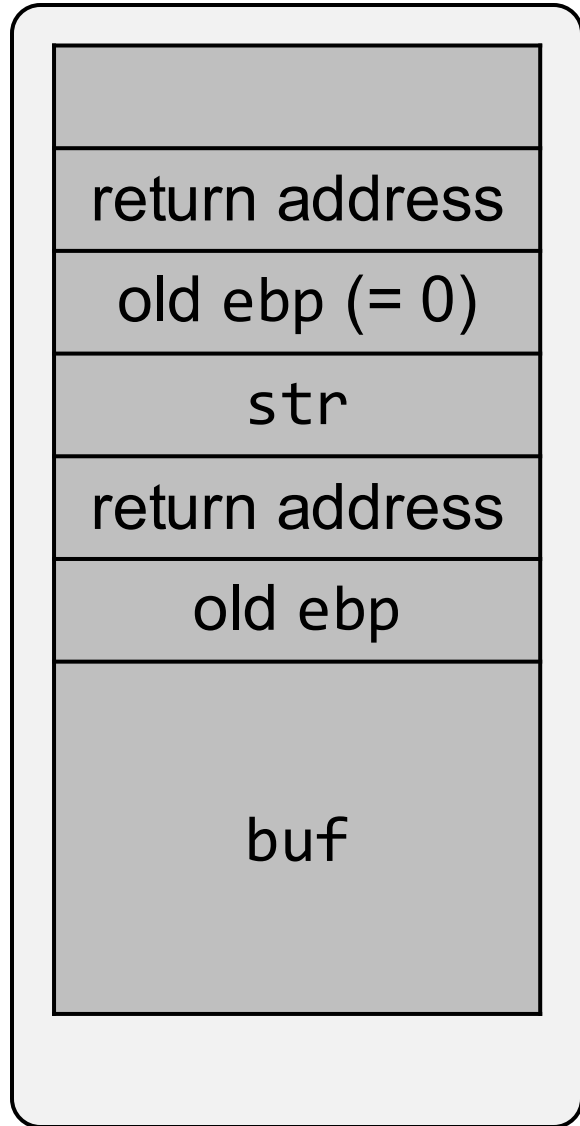
```
#include <stdio.h>
#include <string.h>
#define BUFSIZE (512)

void printer(char* str) {
    char buf[BUFSIZE];
    strcpy(buf, str);
    printf("%s\n", buf);
}
```

```
int main(int argc, char* argv[]) {
    if ( argc < 2 || strlen(argv[1]) > BUFSIZE ) return -1;
    printer(argv[1]);
    return 0;
}
```

Exercise: Can you draw the stack diagram?

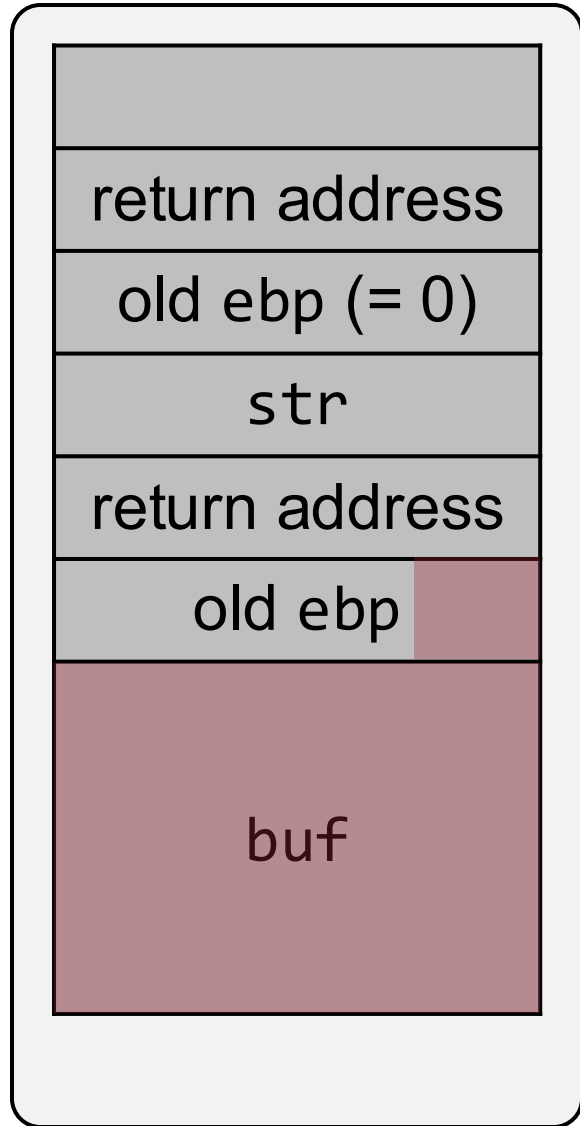
Off-by-one Bugs Can be Exploitable



Virtual memory

```
leave = mov esp, ebp;  
pop ebp
```

Off-by-one Bugs Can be Exploitable



Virtual memory

```
leave = mov esp, ebp;  
pop ebp
```

GDB Usage



- Start: `$ gdb <your binary>`
- Disassemble:
`(gdb) disass <func name>`
- Breakpoint setting:
`(gdb) b *<address>`
- Run:
`(gdb) r`
- Step:
`(gdb) step # Go to next instruction, diving into function`
`(gdb) next # Go to next instruction but don't dive into function`
`(gdb) continue # Continue normal execution`

GDB Usage



- Register information:

```
(gdb) i r <register_name>
(gdb) info register
```

- Memory information:

```
(gdb) x/16w <address>
(gdb) x/4w <register_name>
...
```

x/nfu <address>

Print memory.

n: How many units to print (default 1).

f: Format character (like „print“).

u: Unit.

Unit is one of:

b: Byte,

h: Half-word (two bytes)

w: Word (four bytes)

g: Giant word (eight bytes)).

Recommended Reading

GDB Cheatsheet

– <https://darkdust.net/files/GDB%20Cheat%20Sheet.pdf>

Smashing the Stack for Fun and Profit, Phrack 1996, by Alphe One

– <http://phrack.org/issues/49/14.html>

Summary



- Only some bugs are exploitable
- Some exploits allow an attacker to hijack the control flow of the program and to run any arbitrary code
- Return-to-stack exploit puts a shellcode into a stack buffer and jumps to it by overwriting the return address
- We can make return-to-stack exploit robust by using NOP sleds
- Off-by-one errors can sometimes be exploitable

Question?